

Building user-based recommendation model for Amazon

Please Note that:

In the code below

- The project requirement description titles are in yellow
- code is in light grey
- help on `melt()` output is in turquoise
- Output of the code is in plain text with **bold** headings in them

DESCRIPTION

The dataset provided contains movie reviews given by Amazon customers. Reviews were given between May 1996 and July 2014.

Data Dictionary

UserID – 4848 customers who provided a rating for each movie

Movie 1 to Movie 206 – 206 movies for which ratings are provided by 4848 distinct users

Data Considerations

- All the users have not watched all the movies and therefore, all movies are not rated. These missing values are represented by NA.
- Ratings are on a scale of -1 to 10 where -1 is the least rating and 10 is the best.

Analysis Task

- Exploratory Data Analysis:

Which movies have maximum views/ratings?

What is the average rating for each movie? Define the top 5 movies with the maximum ratings.

Define the top 5 movies with the least audience.

- Recommendation Model: Some of the movies hadn't been watched and therefore, are not rated by the users.

Netflix would like to take this as an opportunity and build a machine learning recommendation algorithm which provides the ratings for each of the users.

Divide the data into training and test data

Build a recommendation model on training data

Make predictions on the test data

Code for: Building user-based recommendation model for Amazon

```
import pandas as pd
```

```
df = pd.read_csv('C:\\Users\\SandipG\\Desktop\\Python Programs\\Amazon - Movies and TV Ratings.csv')
```

```
df.head()
```

	user_id	Mov1	Mov2	Mov3	Mov4	Mov5	Mov6	Mov7	Mov8	Mov9	Mov10	Mov11	Mov12	Mov13	Mov14	Mov15	Mov16	Mov17	Mov18	Mov19
0	A3R5O BKS7 OM2IR	5.0	5.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	AH3Q C2PC1 VTGP	NaN	NaN	2.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	A3LKP 6WPM P9UK X	NaN	NaN	NaN	5.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	AVIY6 8KEPQ 5ZD	NaN	NaN	NaN	5.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	A1CV1 WROP 5KTT W	NaN	NaN	NaN	NaN	5.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

5 rows × 207 columns

```
df.describe()
```

	Mov1	Mov2	Mov3	Mov4	Mov5	Mov6	Mov7	Mov8	Mov9	Mov10	Mov11	Mov12	Mov13	Mov14	Mov15	Mov16	Mov17	Mov18	Mov19	Mov20
count	1.0	1.0	1.0	2.0	29.00	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

	M ov ie 1	M ov ie 2	M ov ie 3	M ov ie 4	Mo vie 5	M ov ie 6	M ov ie 7	M ov ie 8	M ov ie 9	M ov ie1 0	.	M ovi e1 97	M ovi e1 98	M ovi e1 99	M ovi e2 00	M ovi e2 01	M ovi e2 02	M ovi e2 03	M ovi e2 04	Mo vie 20 5	Mo vie 20 6
u n t					00 00						.	00 0			00 0	00 0	00 0		00 0	00 00	00 00
m e a n	5. 0	5. 0	2. 0	5. 0	4.1 03 44 8	4. 0	5. 0	5. 0	5. 0	5.0	.	3.8 00 00 0	5.0	5.0	4.6 25 00 0	4.3 33 33 3	4.3 33 33 3	3.0	4.3 75 00 0	4.6 28 57 1	4.9 23 07 7
st d	N a N	N a N	N a N	0. 0	1.4 96 30 1	N a N	N a N	N a N	N a N	Na N	.	1.6 43 16 8	0.0	Na N	0.5 17 54 9	1.1 54 70 1	1.6 32 99 3	Na N	1.4 07 88 6	0.9 10 25 9	0.2 77 35 0
m i n	5. 0	5. 0	2. 0	5. 0	1.0 00 00 0	4. 0	5. 0	5. 0	5. 0	5.0	.	1.0 00 00 0	5.0	5.0	4.0 00 00 0	3.0 00 00 0	1.0 00 00 0	3.0	1.0 00 00 0	1.0 00 00 0	4.0 00 00 0
2 5 %	5. 0	5. 0	2. 0	5. 0	4.0 00 00 0	4. 0	5. 0	5. 0	5. 0	5.0	.	4.0 00 00 0	5.0	5.0	4.0 00 00 0	4.0 00 00 0	5.0 00 00 0	3.0	4.7 50 00 0	5.0 00 00 0	5.0 00 00 0
5 0 %	5. 0	5. 0	2. 0	5. 0	5.0 00 00 0	4. 0	5. 0	5. 0	5. 0	5.0	.	4.0 00 00 0	5.0	5.0	5.0 00 00 0	5.0 00 00 0	5.0 00 00 0	3.0	5.0 00 00 0	5.0 00 00 0	5.0 00 00 0
7 5 %	5. 0	5. 0	2. 0	5. 0	5.0 00 00 0	4. 0	5. 0	5. 0	5. 0	5.0	.	5.0 00 00 0	5.0	5.0	5.0 00 00 0	5.0 00 00 0	5.0 00 00 0	3.0	5.0 00 00 0	5.0 00 00 0	5.0 00 00 0
m a x	5. 0	5. 0	2. 0	5. 0	5.0 00 00 0	4. 0	5. 0	5. 0	5. 0	5.0	.	5.0 00 00 0	5.0	5.0	5.0 00 00 0	5.0 00 00 0	5.0 00 00 0	3.0	5.0 00 00 0	5.0 00 00 0	5.0 00 00 0

8 rows × 206 columns

```
df_main = df.copy()
```

```
df.describe().T["count"].sort_values(ascending = False)[:10].to_frame()
```

	count
Movie127	2313.0
Movie140	578.0
Movie16	320.0
Movie103	272.0
Movie29	243.0
Movie91	128.0
Movie92	101.0
Movie89	83.0
Movie158	66.0
Movie108	54.0

```
df.drop('user_id', axis=1).sum().sort_values(ascending=False)[:10].to_frame()
```

	0
Movie127	9511.0
Movie140	2794.0
Movie16	1446.0
Movie103	1241.0
Movie29	1168.0
Movie91	586.0

0

Movie92 482.0

Movie89 380.0

Movie158 318.0

Movie108 252.0

```
!pip install surprise
```

```
Requirement already satisfied: surprise in e:\programs\anaconda3\lib\site-packages (0.1)
```

```
Requirement already satisfied: scikit-surprise in e:\programs\anaconda3\lib\site-packages (from surprise) (1.1.0)
```

```
Requirement already satisfied: joblib>=0.11 in e:\programs\anaconda3\lib\site-packages (from scikit-surprise->surprise) (0.13.2)
```

```
Requirement already satisfied: numpy>=1.11.2 in e:\programs\anaconda3\lib\site-packages (from scikit-surprise->surprise) (1.16.5)
```

```
Requirement already satisfied: scipy>=1.0.0 in e:\programs\anaconda3\lib\site-packages (from scikit-surprise->surprise) (1.4.1)
```

```
Requirement already satisfied: six>=1.10.0 in e:\programs\anaconda3\lib\site-packages (from scikit-surprise->surprise) (1.12.0)
```

```
from surprise import Reader
```

```
from surprise import accuracy
```

```
from surprise.model_selection import train_test_split
```

```
help(df.melt())
```

```
Help on DataFrame in module pandas.core.frame object:
```

```
class DataFrame(pandas.core.generic.NDFrame)
```

```
| DataFrame(data=None, index=None, columns=None, dtype=None, copy=False)
```

```
|
```

```
| Two-dimensional size-mutable, potentially heterogeneous tabular data  
| structure with labeled axes (rows and columns). Arithmetic operations  
| align on both row and column labels. Can be thought of as a dict-like  
| container for Series objects. The primary pandas data structure.
```

```
|
```

```

| Parameters
| -----
| data : ndarray (structured or homogeneous), Iterable, dict, or DataFrame
|       Dict can contain Series, arrays, constants, or list-like objects
|
|       .. versionchanged :: 0.23.0
|           If data is a dict, column order follows insertion-order for
|           Python 3.6 and later.
|
|       .. versionchanged :: 0.25.0
|           If data is a list of dicts, column order follows insertion-order
|           Python 3.6 and later.
|
| index : Index or array-like
|       Index to use for resulting frame. Will default to RangeIndex if
|       no indexing information part of input data and no index provided
| columns : Index or array-like
|       Column labels to use for resulting frame. Will default to
|       RangeIndex (0, 1, 2, ..., n) if no column labels are provided
| dtype : dtype, default None
|       Data type to force. Only a single dtype is allowed. If None, infer
| copy : boolean, default False
|       Copy data from inputs. Only affects DataFrame / 2d ndarray input
|
| See Also
| -----
| DataFrame.from_records : Constructor from tuples, also record arrays.
| DataFrame.from_dict : From dicts of Series, arrays, or dicts.
| DataFrame.from_items : From sequence of (key, value) pairs
|       read_csv, pandas.read_table, pandas.read_clipboard.
|
| Examples
| -----
| Constructing DataFrame from a dictionary.
|
| >>> d = {'col1': [1, 2], 'col2': [3, 4]}
| >>> df = pd.DataFrame(data=d)
| >>> df
|      col1  col2
| 0        1     3
| 1        2     4
|
| Notice that the inferred dtype is int64.
|

```

```

| >>> df.dtypes
| col1    int64
| col2    int64
| dtype: object
|
| To enforce a single dtype:
|
| >>> df = pd.DataFrame(data=d, dtype=np.int8)
| >>> df.dtypes
| col1    int8
| col2    int8
| dtype: object
|
| Constructing DataFrame from numpy ndarray:
|
| >>> df2 = pd.DataFrame(np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]]),
| ...                      columns=['a', 'b', 'c'])
| >>> df2
|    a  b  c
| 0  1  2  3
| 1  4  5  6
| 2  7  8  9
|
| Method resolution order:
|    DataFrame
|    pandas.core.generic.NDFrame
|    pandas.core.base.PandasObject
|    pandas.core.accessor.DirNamesMixin
|    pandas.core.base.SelectionMixin
|    builtins.object
|
| Methods defined here:
|
| __add__(self, other, axis=None, level=None, fill_value=None)
|    Binary operator __add__ with support to substitute a fill_value for
missing data in
|    one of the inputs
|
| Parameters
| -----
|    other : Series, DataFrame, or constant
|    axis : {0, 1, 'index', 'columns'}
|           For Series input, axis to match Series index on
|    fill_value : None or float value, default None

```

```

|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|         level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __and__(self, other, axis='columns', level=None, fill_value=None)
|     Binary operator __and__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|         level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes

```



```

|         -----
|         Mismatched indices will be unioned together
|
|     __div__ = __truediv__(self, other, axis=None, level=None,
fill_value=None)
|
|     __eq__(self, other)
|         Wrapper for comparison method __eq__
|
|     __floordiv__(self, other, axis=None, level=None, fill_value=None)
|         Binary operator __floordiv__ with support to substitute a fill_value
for missing data in
|         one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|     level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|         Mismatched indices will be unioned together
|
|     __ge__(self, other)
|         Wrapper for comparison method __ge__
|
|     __getitem__(self, key)
|
|     __gt__(self, other)

```

```

|     Wrapper for comparison method __gt__
|
|     __iadd__(self, other)
|
|     __iand__(self, other)
|
|     __ifloordiv__(self, other)
|
|     __imod__(self, other)
|
|     __imul__(self, other)
|
|     __init__(self, data=None, index=None, columns=None, dtype=None,
copy=False)
|         Initialize self.  See help(type(self)) for accurate signature.
|
|     __ior__(self, other)
|
|     __ipow__(self, other)
|
|     __isub__(self, other)
|
|     __itruediv__(self, other)
|
|     __ixor__(self, other)
|
|     __le__(self, other)
|         Wrapper for comparison method __le__
|
|     __len__(self)
|         Returns length of info axis, but here we use the index.
|
|     __lt__(self, other)
|         Wrapper for comparison method __lt__
|
|     __matmul__(self, other)
|         Matrix multiplication using binary `@` operator in Python>=3.5.
|
|     __mod__(self, other, axis=None, level=None, fill_value=None)
|         Binary operator __mod__ with support to substitute a fill_value for
missing data in
|         one of the inputs
|
|         Parameters

```

```

|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|     level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __mul__(self, other, axis=None, level=None, fill_value=None)
|     Binary operator __mul__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|     level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|

```

```

|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __ne__(self, other)
|         Wrapper for comparison method __ne__
|
|     __or__(self, other, axis='columns', level=None, fill_value=None)
|         Binary operator __or__ with support to substitute a fill_value for
missing data in
|         one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|     level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __pow__(self, other, axis=None, level=None, fill_value=None)
|         Binary operator __pow__ with support to substitute a fill_value for
missing data in
|         one of the inputs

```

```

|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|           For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|           Fill existing missing (NaN) values, and any new element needed
for
|           successful DataFrame alignment, with this value before
computation.
|           If data in both corresponding DataFrame locations is missing
|           the result will be missing
|     level : int or name
|           Broadcast across a level, matching Index values on the
|           passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __radd__(self, other, axis=None, level=None, fill_value=None)
|     Binary operator __radd__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|           For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|           Fill existing missing (NaN) values, and any new element needed
for
|           successful DataFrame alignment, with this value before
computation.
|           If data in both corresponding DataFrame locations is missing
|           the result will be missing
|     level : int or name
|           Broadcast across a level, matching Index values on the

```

```

|         passed MultiIndex level
|
|
|     Returns
|     -----
|     result : DataFrame
|
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|
|     __rand__(self, other, axis='columns', level=None, fill_value=None)
|     Binary operator __rand__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|     level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|
|     Returns
|     -----
|     result : DataFrame
|
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|
|     __rdiv__ = __rtruediv__(self, other, axis=None, level=None,
fill_value=None)
|
|
|     __repr__(self)
|     Return a string representation for a particular DataFrame.

```

```

|
|  __rfloordiv__(self, other, axis=None, level=None, fill_value=None)
|      Binary operator __rfloordiv__ with support to substitute a fill_value
for missing data in
|      one of the inputs
|
|      Parameters
|      -----
|      other : Series, DataFrame, or constant
|      axis : {0, 1, 'index', 'columns'}
|          For Series input, axis to match Series index on
|      fill_value : None or float value, default None
|          Fill existing missing (NaN) values, and any new element needed
for
|          successful DataFrame alignment, with this value before
computation.
|          If data in both corresponding DataFrame locations is missing
|          the result will be missing
|      level : int or name
|          Broadcast across a level, matching Index values on the
|          passed MultiIndex level
|
|      Returns
|      -----
|      result : DataFrame
|
|      Notes
|      -----
|      Mismatched indices will be unioned together
|
|  __rmatmul__(self, other)
|      Matrix multiplication using binary `@` operator in Python>=3.5.
|
|  __rmod__(self, other, axis=None, level=None, fill_value=None)
|      Binary operator __rmod__ with support to substitute a fill_value for
missing data in
|      one of the inputs
|
|      Parameters
|      -----
|      other : Series, DataFrame, or constant
|      axis : {0, 1, 'index', 'columns'}
|          For Series input, axis to match Series index on
|      fill_value : None or float value, default None

```

```

|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|         level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __rmul__(self, other, axis=None, level=None, fill_value=None)
|     Binary operator __rmul__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|         level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes

```



```

|     -----
|     Mismatched indices will be unioned together
|
|     __ror__(self, other, axis='columns', level=None, fill_value=None)
|     Binary operator __ror__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|     level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __rpow__(self, other, axis=None, level=None, fill_value=None)
|     Binary operator __rpow__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None

```

```

|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|         level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __rsub__(self, other, axis=None, level=None, fill_value=None)
|     Binary operator __rsub__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|         level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes

```

```

|     -----
|     Mismatched indices will be unioned together
|
|     __rtruediv__(self, other, axis=None, level=None, fill_value=None)
|     Binary operator __rtruediv__ with support to substitute a fill_value
for missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|     level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __rxor__(self, other, axis='columns', level=None, fill_value=None)
|     Binary operator __rxor__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None

```

```

|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|         level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __setitem__(self, key, value)
|
|     __sub__(self, other, axis=None, level=None, fill_value=None)
|     Binary operator __sub__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|         level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame

```

```

|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __truediv__(self, other, axis=None, level=None, fill_value=None)
|     Binary operator __truediv__ with support to substitute a fill_value
for missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on
|     fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|     level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     __xor__(self, other, axis='columns', level=None, fill_value=None)
|     Binary operator __xor__ with support to substitute a fill_value for
missing data in
|     one of the inputs
|
|     Parameters
|     -----
|     other : Series, DataFrame, or constant
|     axis : {0, 1, 'index', 'columns'}
|         For Series input, axis to match Series index on

```

```

|         fill_value : None or float value, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing
|         level : int or name
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level
|
|     Returns
|     -----
|     result : DataFrame
|
|     Notes
|     -----
|     Mismatched indices will be unioned together
|
|     add(self, other, axis='columns', level=None, fill_value=None)
|     Get Addition of dataframe and other, element-wise (binary operator
`add`).
|
|     Equivalent to ``dataframe + other``, but with support to substitute a
fill_value
|     for missing data in one of the inputs. With reverse version, `radd`.
|
|     Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
|     arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
|     Parameters
|     -----
|     other : scalar, sequence, Series, or DataFrame
|     Any single or multiple element data structure, or list-like
object.
|
|     axis : {0 or 'index', 1 or 'columns'}
|     Whether to compare by the index (0 or 'index') or columns
|     (1 or 'columns'). For Series input, axis to match Series index
on.
|
|     level : int or label
|     Broadcast across a level, matching Index values on the
|     passed MultiIndex level.
|
|     fill_value : float or None, default None

```

```

|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing.
|
|     Returns
|     -----
|     DataFrame
|         Result of the arithmetic operation.
|
|     See Also
|     -----
|     DataFrame.add : Add DataFrames.
|     DataFrame.sub : Subtract DataFrames.
|     DataFrame.mul : Multiply DataFrames.
|     DataFrame.div : Divide DataFrames (float division).
|     DataFrame.truediv : Divide DataFrames (float division).
|     DataFrame.floordiv : Divide DataFrames (integer division).
|     DataFrame.mod : Calculate modulo (remainder after division).
|     DataFrame.pow : Calculate exponential power.
|
|     Notes
|     -----
|     Mismatched indices will be unioned together.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'angles': [0, 3, 4],
|     ...                    'degrees': [360, 180, 360]},
|     ...                    index=['circle', 'triangle', 'rectangle'])
|     >>> df
|
|           angles  degrees
|     circle         0     360
|     triangle        3     180
|     rectangle        4     360
|
|     Add a scalar with operator version which return the same
|     results.
|
|     >>> df + 1
|
|           angles  degrees
|     circle         1     361

```

```

|     triangle      4      181
|     rectangle     5      361
|
|     >>> df.add(1)
|           angles  degrees
|     circle      1      361
|     triangle     4      181
|     rectangle     5      361
|
|     Divide by constant with reverse version.
|
|     >>> df.div(10)
|           angles  degrees
|     circle      0.0     36.0
|     triangle     0.3     18.0
|     rectangle     0.4     36.0
|
|     >>> df.rdiv(10)
|           angles  degrees
|     circle      inf  0.027778
|     triangle  3.333333  0.055556
|     rectangle  2.500000  0.027778
|
|     Subtract a list and Series by axis with operator version.
|
|     >>> df - [1, 2]
|           angles  degrees
|     circle      -1     358
|     triangle      2     178
|     rectangle      3     358
|
|     >>> df.sub([1, 2], axis='columns')
|           angles  degrees
|     circle      -1     358
|     triangle      2     178
|     rectangle      3     358
|
|     >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
|rectangle']),
|     ...         axis='index')
|           angles  degrees
|     circle      -1     359
|     triangle      2     179
|     rectangle      3     359

```



```

|
|     Multiply a DataFrame of different shape with operator version.
|
|     >>> other = pd.DataFrame({'angles': [0, 3, 4]},
|     ...                       index=['circle', 'triangle', 'rectangle'])
|     >>> other
|
|           angles
|     circle      0
|     triangle    3
|     rectangle   4
|
|     >>> df * other
|
|           angles  degrees
|     circle      0      NaN
|     triangle     9      NaN
|     rectangle   16      NaN
|
|     >>> df.mul(other, fill_value=0)
|
|           angles  degrees
|     circle      0      0.0
|     triangle     9      0.0
|     rectangle   16      0.0
|
|     Divide by a MultiIndex by level.
|
|     >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
|     ...                              'degrees': [360, 180, 360, 360, 540,
720]}),
|     ...                              index=[['A', 'A', 'A', 'B', 'B',
'B'],
|     ...                              ['circle', 'triangle',
'rectangle',
|     ...                              'square', 'pentagon',
'hexagon']])
|     >>> df_multindex
|
|           angles  degrees
|     A circle      0      360
|       triangle     3      180
|       rectangle     4      360
|     B square      4      360
|       pentagon     5      540
|       hexagon      6      720
|
|     >>> df.div(df_multindex, level=1, fill_value=0)

```

```

|           angles  degrees
|   A circle      NaN     1.0
|   triangle      1.0     1.0
|   rectangle     1.0     1.0
|   B square      0.0     0.0
|   pentagon      0.0     0.0
|   hexagon       0.0     0.0
|
|  agg = aggregate(self, func, axis=0, *args, **kwargs)
|
|  aggregate(self, func, axis=0, *args, **kwargs)
|      Aggregate using one or more operations over the specified axis.
|
|  .. versionadded:: 0.20.0
|
|  Parameters
|  -----
|  func : function, str, list or dict
|      Function to use for aggregating the data. If a function, must
either
|      work when passed a DataFrame or when passed to DataFrame.apply.
|
|      Accepted combinations are:
|
|      - function
|      - string function name
|      - list of functions and/or function names, e.g. ``[np.sum,
'mean']``
|      - dict of axis labels -> functions, function names or list of
such.
|
|      axis : {0 or 'index', 1 or 'columns'}, default 0
|          If 0 or 'index': apply function to each column.
|          If 1 or 'columns': apply function to each row.
|
|  *args
|      Positional arguments to pass to `func`.
|
|  **kwargs
|      Keyword arguments to pass to `func`.
|
|  Returns
|  -----
|  scalar, Series or DataFrame
|
|  The return can be:
|

```

```

|         * scalar : when Series.agg is called with single function
|         * Series : when DataFrame.agg is called with a single function
|         * DataFrame : when DataFrame.agg is called with several functions
|
|         Return scalar, Series or DataFrame.
|
|         The aggregation operations are always performed over an axis, either
the
|         index (default) or the column axis. This behavior is different from
|         `numpy` aggregation functions (`mean`, `median`, `prod`, `sum`,
`std`,
|         `var`), where the default is to compute the aggregation of the
flattened
|         array, e.g., ``numpy.mean(arr_2d)`` as opposed to
|         ``numpy.mean(arr_2d, axis=0)``.
|
|         `agg` is an alias for `aggregate`. Use the alias.
|
|         See Also
|         -----
|         DataFrame.apply : Perform any type of operations.
|         DataFrame.transform : Perform transformation type operations.
|         core.groupby.GroupBy : Perform operations over groups.
|         core.resample.Resampler : Perform operations over resampled bins.
|         core.window.Rolling : Perform operations over rolling window.
|         core.window.Expanding : Perform operations over expanding window.
|         core.window.EWM : Perform operation over exponential weighted
|         window.
|
|         Notes
|         -----
|         `agg` is an alias for `aggregate`. Use the alias.
|
|         A passed user-defined-function will be passed a Series for
evaluation.
|
|         Examples
|         -----
|         >>> df = pd.DataFrame([[1, 2, 3],
|         ...                    [4, 5, 6],
|         ...                    [7, 8, 9],
|         ...                    [np.nan, np.nan, np.nan]],
|         ...                    columns=['A', 'B', 'C'])
|

```

```

|     Aggregate these functions over the rows.
|
|     >>> df.agg(['sum', 'min'])
|           A      B      C
|     sum  12.0  15.0  18.0
|     min   1.0   2.0   3.0
|
|     Different aggregations per column.
|
|     >>> df.agg({'A' : ['sum', 'min'], 'B' : ['min', 'max']})
|           A      B
|     max   NaN   8.0
|     min    1.0   2.0
|     sum  12.0  NaN
|
|     Aggregate over the columns.
|
|     >>> df.agg("mean", axis="columns")
|     0      2.0
|     1      5.0
|     2      8.0
|     3      NaN
|     dtype: float64
|
|     align(self, other, join='outer', axis=None, level=None, copy=True,
fill_value=None, method=None, limit=None, fill_axis=0, broadcast_axis=None)
|
|     Align two objects on their axes with the
|     specified join method for each axis Index.
|
|     Parameters
|     -----
|     other : DataFrame or Series
|     join : {'outer', 'inner', 'left', 'right'}, default 'outer'
|     axis : allowed axis of the other object, default None
|           Align on index (0), columns (1), or both (None)
|     level : int or level name, default None
|           Broadcast across a level, matching Index values on the
|           passed MultiIndex level
|     copy : boolean, default True
|           Always returns new objects. If copy=False and no reindexing is
|           required then original objects are returned.
|     fill_value : scalar, default np.NaN
|           Value to use for missing values. Defaults to NaN, but can be any
|           "compatible" value

```

```

|     method : {'backfill', 'bfill', 'pad', 'ffill', None}, default None
|         Method to use for filling holes in reindexed Series
|     pad / ffill: propagate last valid observation forward to next
valid
|         backfill / bfill: use NEXT valid observation to fill gap
|     limit : int, default None
|         If method is specified, this is the maximum number of consecutive
|         NaN values to forward/backward fill. In other words, if there is
|         a gap with more than this number of consecutive NaNs, it will
only
|         be partially filled. If method is not specified, this is the
|         maximum number of entries along the entire axis where NaNs will
be
|         filled. Must be greater than 0 if not None.
|     fill_axis : {0 or 'index', 1 or 'columns'}, default 0
|         Filling axis, method and limit
|     broadcast_axis : {0 or 'index', 1 or 'columns'}, default None
|         Broadcast values along this axis, if aligning two objects of
|         different dimensions
|
|     Returns
|     -----
|     (left, right) : (DataFrame, type of other)
|         Aligned objects.
|
|     all(self, axis=0, bool_only=None, skipna=True, level=None, **kwargs)
|         Return whether all elements are True, potentially over an axis.
|
|     Returns True unless there at least one element within a series or
|     along a Dataframe axis that is False or equivalent (e.g. zero or
|     empty).
|
|     Parameters
|     -----
|     axis : {0 or 'index', 1 or 'columns', None}, default 0
|         Indicate which axis or axes should be reduced.
|
|         * 0 / 'index' : reduce the index, return a Series whose index is
the
|         original column labels.
|         * 1 / 'columns' : reduce the columns, return a Series whose index
is the
|         original index.
|         * None : reduce all axes, return a scalar.

```

```

|
|     bool_only : bool, default None
|         Include only boolean columns. If None, will attempt to use
everything,
|         then use only boolean data. Not implemented for Series.
|     skipna : bool, default True
|         Exclude NA/null values. If the entire row/column is NA and skipna
is
|         True, then the result will be True, as for an empty row/column.
|         If skipna is False, then NA are treated as True, because these
are not
|         equal to zero.
|     level : int or level name, default None
|         If the axis is a MultiIndex (hierarchical), count along a
|         particular level, collapsing into a Series.
|     **kwargs : any, default None
|         Additional keywords have no effect but might be accepted for
|         compatibility with NumPy.
|
|     Returns
|     -----
|     Series or DataFrame
|         If level is specified, then, DataFrame is returned; otherwise,
Series
|         is returned.
|
|     See Also
|     -----
|     Series.all : Return True if all elements are True.
|     DataFrame.any : Return True if one (or more) elements are True.
|
|     Examples
|     -----
|     **Series**
|
|     >>> pd.Series([True, True]).all()
|     True
|     >>> pd.Series([True, False]).all()
|     False
|     >>> pd.Series([]).all()
|     True
|     >>> pd.Series([np.nan]).all()
|     True
|     >>> pd.Series([np.nan]).all(skipna=False)

```

```

|     True
|
|     **DataFrames**
|
|     Create a dataframe from a dictionary.
|
|     >>> df = pd.DataFrame({'col1': [True, True], 'col2': [True, False]})
|     >>> df
|           col1  col2
|     0  True   True
|     1  True  False
|
|     Default behaviour checks if column-wise values all return True.
|
|     >>> df.all()
|           col1    True
|           col2   False
|           dtype: bool
|
|     Specify ``axis='columns'`` to check if row-wise values all return
True.
|
|     >>> df.all(axis='columns')
|           0     True
|           1    False
|           dtype: bool
|
|     Or ``axis=None`` for whether every value is True.
|
|     >>> df.all(axis=None)
|           False
|
|     any(self, axis=0, bool_only=None, skipna=True, level=None, **kwargs)
|     Return whether any element is True, potentially over an axis.
|
|     Returns False unless there at least one element within a series or
|     along a Dataframe axis that is True or equivalent (e.g. non-zero or
|     non-empty).
|
|     Parameters
|     -----
|     axis : {0 or 'index', 1 or 'columns', None}, default 0
|           Indicate which axis or axes should be reduced.
|

```

```

|         * 0 / 'index' : reduce the index, return a Series whose index is
the
|         original column labels.
|         * 1 / 'columns' : reduce the columns, return a Series whose index
is the
|         original index.
|         * None : reduce all axes, return a scalar.
|
|         bool_only : bool, default None
|         Include only boolean columns. If None, will attempt to use
everything,
|         then use only boolean data. Not implemented for Series.
|         skipna : bool, default True
|         Exclude NA/null values. If the entire row/column is NA and skipna
is
|         True, then the result will be False, as for an empty row/column.
|         If skipna is False, then NA are treated as True, because these
are not
|         equal to zero.
|         level : int or level name, default None
|         If the axis is a MultiIndex (hierarchical), count along a
|         particular level, collapsing into a Series.
|         **kwargs : any, default None
|         Additional keywords have no effect but might be accepted for
|         compatibility with NumPy.
|
|     Returns
|     -----
|     Series or DataFrame
|         If level is specified, then, DataFrame is returned; otherwise,
Series
|         is returned.
|
|     See Also
|     -----
|     numpy.any : Numpy version of this method.
|     Series.any : Return whether any element is True.
|     Series.all : Return whether all elements are True.
|     DataFrame.any : Return whether any element is True over requested
axis.
|     DataFrame.all : Return whether all elements are True over requested
axis.
|
|     Examples

```



```

|     -----
|     **Series**
|
|     For Series input, the output is a scalar indicating whether any
element
|     is True.
|
|     >>> pd.Series([False, False]).any()
|     False
|     >>> pd.Series([True, False]).any()
|     True
|     >>> pd.Series([]).any()
|     False
|     >>> pd.Series([np.nan]).any()
|     False
|     >>> pd.Series([np.nan]).any(skipna=False)
|     True
|
|     **DataFrame**
|
|     Whether each column contains at least one True element (the default).
|
|     >>> df = pd.DataFrame({"A": [1, 2], "B": [0, 2], "C": [0, 0]})
|     >>> df
|           A  B  C
|     0    1  0  0
|     1    2  2  0
|
|     >>> df.any()
|     A      True
|     B      True
|     C     False
|     dtype: bool
|
|     Aggregating over the columns.
|
|     >>> df = pd.DataFrame({"A": [True, False], "B": [1, 2]})
|     >>> df
|           A  B
|     0   True  1
|     1  False  2
|
|     >>> df.any(axis='columns')
|     0      True

```

```

|     1     True
|     dtype: bool
|
|     >>> df = pd.DataFrame({"A": [True, False], "B": [1, 0]})
|     >>> df
|           A  B
|     0  True  1
|     1 False  0
|
|     >>> df.any(axis='columns')
|     0     True
|     1    False
|     dtype: bool
|
|     Aggregating over the entire DataFrame with ``axis=None``.
|
|     >>> df.any(axis=None)
|     True
|
|     `any` for an empty DataFrame is an empty Series.
|
|     >>> pd.DataFrame([]).any()
|     Series([], dtype: bool)
|
|     append(self, other, ignore_index=False, verify_integrity=False,
sort=None)
|         Append rows of `other` to the end of caller, returning a new object.
|
|         Columns in `other` that are not in the caller are added as new
columns.
|
|         Parameters
|         -----
|         other : DataFrame or Series/dict-like object, or list of these
|             The data to append.
|         ignore_index : boolean, default False
|             If True, do not use the index labels.
|         verify_integrity : boolean, default False
|             If True, raise ValueError on creating index with duplicates.
|         sort : boolean, default None
|             Sort columns if the columns of `self` and `other` are not
aligned.
|
|             The default sorting is deprecated and will change to not-sorting
|             in a future version of pandas. Explicitly pass ``sort=True`` to

```

```

|         silence the warning and sort. Explicitly pass ``sort=False`` to
|         silence the warning and not sort.
|
|         .. versionadded:: 0.23.0
|
|     Returns
|     -----
|     DataFrame
|
|     See Also
|     -----
|     concat : General function to concatenate DataFrame or Series objects.
|
|     Notes
|     -----
|     If a list of dict/series is passed and the keys are all contained in
|     the DataFrame's index, the order of the columns in the resulting
|     DataFrame will be unchanged.
|
|     Iteratively appending rows to a DataFrame can be more computationally
|     intensive than a single concatenate. A better solution is to append
|     those rows to a list and then concatenate the list with the original
|     DataFrame all at once.
|
|     Examples
|     -----
|
|     >>> df = pd.DataFrame([[1, 2], [3, 4]], columns=list('AB'))
|     >>> df
|          A  B
|     0  1  2
|     1  3  4
|     >>> df2 = pd.DataFrame([[5, 6], [7, 8]], columns=list('AB'))
|     >>> df.append(df2)
|          A  B
|     0  1  2
|     1  3  4
|     0  5  6
|     1  7  8
|
|     With `ignore_index` set to True:
|
|     >>> df.append(df2, ignore_index=True)
|          A  B

```

```

|      0  1  2
|      1  3  4
|      2  5  6
|      3  7  8
|
|      The following, while not recommended methods for generating
DataFrames,
|      show two ways to generate a DataFrame from multiple data sources.
|
|      Less efficient:
|
|      >>> df = pd.DataFrame(columns=['A'])
|      >>> for i in range(5):
|      ...     df = df.append({'A': i}, ignore_index=True)
|      >>> df
|
|      A
|      0  0
|      1  1
|      2  2
|      3  3
|      4  4
|
|      More efficient:
|
|      >>> pd.concat([pd.DataFrame([i], columns=['A']) for i in range(5)],
|      ...           ignore_index=True)
|
|      A
|      0  0
|      1  1
|      2  2
|      3  3
|      4  4
|
|      apply(self, func, axis=0, broadcast=None, raw=False, reduce=None,
result_type=None, args=(), **kws)
|      Apply a function along an axis of the DataFrame.
|
|      Objects passed to the function are Series objects whose index is
|      either the DataFrame's index (`axis=0`) or the DataFrame's columns
|      (`axis=1`). By default (`result_type=None`), the final return
type
|      is inferred from the return type of the applied function. Otherwise,
|      it depends on the `result_type` argument.
|

```

```

| Parameters
| -----
| func : function
|     Function to apply to each column or row.
| axis : {0 or 'index', 1 or 'columns'}, default 0
|     Axis along which the function is applied:
|
|     * 0 or 'index': apply function to each column.
|     * 1 or 'columns': apply function to each row.
| broadcast : bool, optional
|     Only relevant for aggregation functions:
|
|     * ``False`` or ``None`` : returns a Series whose length is the
|       length of the index or the number of columns (based on the
|       `axis` parameter)
|     * ``True`` : results will be broadcast to the original shape
|       of the frame, the original index and columns will be retained.
|
| .. deprecated:: 0.23.0
|     This argument will be removed in a future version, replaced
|     by result_type='broadcast'.
|
| raw : bool, default False
|     * ``False`` : passes each row or column as a Series to the
|       function.
|     * ``True`` : the passed function will receive ndarray objects
|       instead.
|
|     If you are just applying a NumPy reduction function this will
|     achieve much better performance.
| reduce : bool or None, default None
|     Try to apply reduction procedures. If the DataFrame is empty,
|     `apply` will use `reduce` to determine whether the result
|     should be a Series or a DataFrame. If ``reduce=None`` (the
|     default), `apply`'s return value will be guessed by calling
|     `func` on an empty Series
|     (note: while guessing, exceptions raised by `func` will be
|     ignored).
|
|     If ``reduce=True`` a Series will always be returned, and if
|     ``reduce=False`` a DataFrame will always be returned.
|
| .. deprecated:: 0.23.0
|     This argument will be removed in a future version, replaced
|     by ``result_type='reduce'``.
|

```

```

|     result_type : {'expand', 'reduce', 'broadcast', None}, default None
|     These only act when ``axis=1`` (columns):
|
|     * 'expand' : list-like results will be turned into columns.
|     * 'reduce' : returns a Series if possible rather than expanding
|                   list-like results. This is the opposite of 'expand'.
|     * 'broadcast' : results will be broadcast to the original shape
|                     of the DataFrame, the original index and columns will be
|                     retained.
|
|     The default behaviour (None) depends on the return value of the
|     applied function: list-like results will be returned as a Series
|     of those. However if the apply function returns a Series these
|     are expanded to columns.
|
|     .. versionadded:: 0.23.0
|
|     args : tuple
|         Positional arguments to pass to `func` in addition to the
|         array/series.
|     **kwargs
|         Additional keyword arguments to pass as keywords arguments to
|         `func`.
|
|     Returns
|     -----
|     Series or DataFrame
|         Result of applying ``func`` along the given axis of the
|         DataFrame.
|
|     See Also
|     -----
|     DataFrame.applymap: For elementwise operations.
|     DataFrame.aggregate: Only perform aggregating type operations.
|     DataFrame.transform: Only perform transforming type operations.
|
|     Notes
|     -----
|     In the current implementation apply calls `func` twice on the
|     first column/row to decide whether it can take a fast or slow
|     code path. This can lead to unexpected behavior if `func` has
|     side-effects, as they will take effect twice for the first
|     column/row.
|

```

| Examples

| -----

| >>> df = pd.DataFrame([[4, 9]] * 3, columns=['A', 'B'])

| >>> df

| A B
| 0 4 9
| 1 4 9
| 2 4 9

| Using a numpy universal function (in this case the same as
| `np.sqrt(df)`):

| >>> df.apply(np.sqrt)

| A B
| 0 2.0 3.0
| 1 2.0 3.0
| 2 2.0 3.0

| Using a reducing function on either axis

| >>> df.apply(np.sum, axis=0)

| A 12
| B 27
| dtype: int64

| >>> df.apply(np.sum, axis=1)

| 0 13
| 1 13
| 2 13
| dtype: int64

| Returning a list-like will result in a Series

| >>> df.apply(lambda x: [1, 2], axis=1)

| 0 [1, 2]
| 1 [1, 2]
| 2 [1, 2]
| dtype: object

| Passing result_type='expand' will expand list-like results
| to columns of a Dataframe

| >>> df.apply(lambda x: [1, 2], axis=1, result_type='expand')

```

|         0  1
|         0  1  2
|         1  1  2
|         2  1  2
|
|         Returning a Series inside the function is similar to passing
|         ``result_type='expand'``. The resulting column names
|         will be the Series index.
|
|         >>> df.apply(lambda x: pd.Series([1, 2], index=['foo', 'bar']),
axis=1)
|         foo  bar
|         0    1    2
|         1    1    2
|         2    1    2
|
|         Passing ``result_type='broadcast'`` will ensure the same shape
|         result, whether list-like or scalar is returned by the function,
|         and broadcast it along the axis. The resulting column names will
|         be the originals.
|
|         >>> df.apply(lambda x: [1, 2], axis=1, result_type='broadcast')
|         A  B
|         0  1  2
|         1  1  2
|         2  1  2
|
|         applymap(self, func)
|         Apply a function to a Dataframe elementwise.
|
|         This method applies a function that accepts and returns a scalar
|         to every element of a DataFrame.
|
|         Parameters
|         -----
|         func : callable
|             Python function, returns a single value from a single value.
|
|         Returns
|         -----
|         DataFrame
|             Transformed DataFrame.
|
|         See Also

```



```

| -----
| DataFrame.apply : Apply a function along input axis of DataFrame.
|
| Notes
| -----
| In the current implementation applymap calls `func` twice on the
| first column/row to decide whether it can take a fast or slow
| code path. This can lead to unexpected behavior if `func` has
| side-effects, as they will take effect twice for the first
| column/row.
|
| Examples
| -----
| >>> df = pd.DataFrame([[1, 2.12], [3.356, 4.567]])
| >>> df
|      0      1
| 0  1.000  2.120
| 1  3.356  4.567
|
| >>> df.applymap(lambda x: len(str(x)))
|      0  1
| 0  3  4
| 1  5  5
|
| Note that a vectorized version of `func` often exists, which will
| be much faster. You could square each number elementwise.
|
| >>> df.applymap(lambda x: x**2)
|      0      1
| 0  1.000000  4.494400
| 1 11.262736 20.857489
|
| But it's better to avoid applymap in that case.
|
| >>> df ** 2
|      0      1
| 0  1.000000  4.494400
| 1 11.262736 20.857489
|
| assign(self, **kwargs)
| Assign new columns to a DataFrame.
|
| Returns a new object with all original columns in addition to new
| ones.

```

```

|     Existing columns that are re-assigned will be overwritten.
|
| Parameters
| -----
|     **kwargs : dict of {str: callable or Series}
|         The column names are keywords. If the values are
|         callable, they are computed on the DataFrame and
|         assigned to the new columns. The callable must not
|         change input DataFrame (though pandas doesn't check it).
|         If the values are not callable, (e.g. a Series, scalar, or
array),
|         they are simply assigned.
|
| Returns
| -----
|     DataFrame
|         A new DataFrame with the new columns in addition to
|         all the existing columns.
|
| Notes
| -----
|     Assigning multiple columns within the same ``assign`` is possible.
|     For Python 3.6 and above, later items in ``*\kwargs`` may refer to
|     newly created or modified columns in 'df'; items are computed and
|     assigned into 'df' in order. For Python 3.5 and below, the order of
|     keyword arguments is not specified, you cannot refer to newly created
|     or modified columns. All items are computed first, and then assigned
|     in alphabetical order.
|
|     .. versionchanged :: 0.23.0
|
|     Keyword argument order is maintained for Python 3.6 and later.
|
| Examples
| -----
|     >>> df = pd.DataFrame({'temp_c': [17.0, 25.0]},
|     ...                    index=['Portland', 'Berkeley'])
|     >>> df
|
|          temp_c
|     Portland   17.0
|     Berkeley   25.0
|
|     Where the value is a callable, evaluated on `df`:
|

```

```
|     >>> df.assign(temp_f=lambda x: x.temp_c * 9 / 5 + 32)
```

	temp_c	temp_f
Portland	17.0	62.6
Berkeley	25.0	77.0

```
|
```

| Alternatively, the same behavior can be achieved by directly
| referencing an existing Series or sequence:

```
|
```

```
|     >>> df.assign(temp_f=df['temp_c'] * 9 / 5 + 32)
```

	temp_c	temp_f
Portland	17.0	62.6
Berkeley	25.0	77.0

```
|
```

| In Python 3.6+, you can create multiple columns within the same
assign
| where one of the columns depends on another one defined within the
same
| assign:

```
|
```

```
|     >>> df.assign(temp_f=lambda x: x['temp_c'] * 9 / 5 + 32,
```

```
|     ...         temp_k=lambda x: (x['temp_f'] + 459.67) * 5 / 9)
```

	temp_c	temp_f	temp_k
Portland	17.0	62.6	290.15
Berkeley	25.0	77.0	298.15

```
|
```

| boxplot = boxplot_frame(self, column=None, by=None, ax=None,
fontsize=None, rot=0, grid=True, figsize=None, layout=None, return_type=None,
**kwds)

| Make a box plot from DataFrame columns.

|

| Make a box-and-whisker plot from DataFrame columns, optionally
grouped
| by some other columns. A box plot is a method for graphically
depicting
| groups of numerical data through their quartiles.
| The box extends from the Q1 to Q3 quartile values of the data,
| with a line at the median (Q2). The whiskers extend from the edges
| of box to show the range of the data. The position of the whiskers
| is set by default to ``1.5 * IQR (IQR = Q3 - Q1)`` from the edges of
the box.
| Outlier points are those past the end of the whiskers.
|

| For further details see

```

|     Wikipedia's entry for `boxplot`
| <https://en.wikipedia.org/wiki/Box_plot>`.
|
|     Parameters
|     -----
|     column : str or list of str, optional
|         Column name or list of names, or vector.
|         Can be any valid input to :meth:`pandas.DataFrame.groupby`.
|     by : str or array-like, optional
|         Column in the DataFrame to :meth:`pandas.DataFrame.groupby`.
|         One box-plot will be done per value of columns in `by`.
|     ax : object of class matplotlib.axes.Axes, optional
|         The matplotlib axes to be used by boxplot.
|     fontsize : float or str
|         Tick label font size in points or as a string (e.g., `large`).
|     rot : int or float, default 0
|         The rotation angle of labels (in degrees)
|         with respect to the screen coordinate system.
|     grid : bool, default True
|         Setting this to True will show the grid.
|     figsize : A tuple (width, height) in inches
|         The size of the figure to create in matplotlib.
|     layout : tuple (rows, columns), optional
|         For example, (3, 5) will display the subplots
|         using 3 columns and 5 rows, starting from the top-left.
|     return_type : {'axes', 'dict', 'both'} or None, default 'axes'
|         The kind of object to return. The default is ``axes``.
|
|     * 'axes' returns the matplotlib axes the boxplot is drawn on.
|     * 'dict' returns a dictionary whose values are the matplotlib
|         Lines of the boxplot.
|     * 'both' returns a namedtuple with the axes and dict.
|     * when grouping with ``by``, a Series mapping columns to
|         ``return_type`` is returned.
|
|     If ``return_type`` is `None`, a NumPy array
|     of axes with the same shape as ``layout`` is returned.
|
|     **kwargs
|
|     All other plotting keyword arguments to be passed to
|     :func:`matplotlib.pyplot.boxplot`.
|
|     Returns
|     -----
|     result

```

```

|         See Notes.
|
|     See Also
|     -----
|     Series.plot.hist: Make a histogram.
|     matplotlib.pyplot.boxplot : Matplotlib equivalent plot.
|
|     Notes
|     -----
|     The return type depends on the `return_type` parameter:
|
|     * 'axes' : object of class matplotlib.axes.Axes
|     * 'dict' : dict of matplotlib.lines.Line2D objects
|     * 'both' : a namedtuple with structure (ax, lines)
|
|     For data grouped with ``by``, return a Series of the above or a numpy
|     array:
|
|     * :class:`~pandas.Series`
|     * :class:`~numpy.array` (for ``return_type = None``)
|
|     Use ``return_type='dict'`` when you want to tweak the appearance
|     of the lines after plotting. In this case a dict containing the Lines
|     making up the boxes, caps, fliers, medians, and whiskers is returned.
|
|     Examples
|     -----
|
|     Boxplots can be created for every column in the dataframe
|     by ``df.boxplot()`` or indicating the columns to be used:
|
|     .. plot::
|        :context: close-figs
|
|        >>> np.random.seed(1234)
|        >>> df = pd.DataFrame(np.random.randn(10,4),
|        ...                    columns=['Col1', 'Col2', 'Col3', 'Col4'])
|        >>> boxplot = df.boxplot(column=['Col1', 'Col2', 'Col3'])
|
|     Boxplots of variables distributions grouped by the values of a third
|     variable can be created using the option ``by``. For instance:
|
|     .. plot::
|        :context: close-figs

```

```

|
|     >>> df = pd.DataFrame(np.random.randn(10, 2),
|     ...                     columns=['Col1', 'Col2'])
|     >>> df['X'] = pd.Series(['A', 'A', 'A', 'A', 'A',
|     ...                     'B', 'B', 'B', 'B', 'B'])
|     >>> boxplot = df.boxplot(by='X')
|
|     A list of strings (i.e. ``['X', 'Y']``) can be passed to boxplot
|     in order to group the data by combination of the variables in the x-
axis:
|
|     .. plot::
|         :context: close-figs
|
|     >>> df = pd.DataFrame(np.random.randn(10,3),
|     ...                     columns=['Col1', 'Col2', 'Col3'])
|     >>> df['X'] = pd.Series(['A', 'A', 'A', 'A', 'A',
|     ...                     'B', 'B', 'B', 'B', 'B'])
|     >>> df['Y'] = pd.Series(['A', 'B', 'A', 'B', 'A',
|     ...                     'B', 'A', 'B', 'A', 'B'])
|     >>> boxplot = df.boxplot(column=['Col1', 'Col2'], by=['X', 'Y'])
|
|     The layout of boxplot can be adjusted giving a tuple to ``layout``:
|
|     .. plot::
|         :context: close-figs
|
|     >>> boxplot = df.boxplot(column=['Col1', 'Col2'], by='X',
|     ...                     layout=(2, 1))
|
|     Additional formatting can be done to the boxplot, like suppressing
the grid
|     (``grid=False``), rotating the labels in the x-axis (i.e. ``rot=45``)
|     or changing the fontsize (i.e. ``fontsize=15``):
|
|     .. plot::
|         :context: close-figs
|
|     >>> boxplot = df.boxplot(grid=False, rot=45, fontsize=15)
|
|     The parameter ``return_type`` can be used to select the type of
element
|     returned by `boxplot`. When ``return_type='axes'`` is selected,
|     the matplotlib axes on which the boxplot is drawn are returned:

```

```

|
|         >>> boxplot = df.boxplot(column=['Col1', 'Col2'],
return_type='axes')
|         >>> type(boxplot)
|         <class 'matplotlib.axes._subplots.AxesSubplot'>
|
|     When grouping with ``by``, a Series mapping columns to
``return_type``
|     is returned:
|
|         >>> boxplot = df.boxplot(column=['Col1', 'Col2'], by='X',
...                               return_type='axes')
|         >>> type(boxplot)
|         <class 'pandas.core.series.Series'>
|
|     If ``return_type`` is `None`, a NumPy array of axes with the same
shape
|     as ``layout`` is returned:
|
|         >>> boxplot = df.boxplot(column=['Col1', 'Col2'], by='X',
...                               return_type=None)
|         >>> type(boxplot)
|         <class 'numpy.ndarray'>
|
|     combine(self, other, func, fill_value=None, overwrite=True)
|         Perform column-wise combine with another DataFrame.
|
|     Combines a DataFrame with `other` DataFrame using `func`
|     to element-wise combine columns. The row and column indexes of the
|     resulting DataFrame will be the union of the two.
|
|     Parameters
|     -----
|     other : DataFrame
|         The DataFrame to merge column-wise.
|     func : function
|         Function that takes two series as inputs and return a Series or a
|         scalar. Used to merge the two dataframes column by columns.
|     fill_value : scalar value, default None
|         The value to fill NaNs with prior to passing any column to the
|         merge func.
|     overwrite : bool, default True
|         If True, columns in `self` that do not exist in `other` will be
|         overwritten with NaNs.

```

```

|
| Returns
| -----
| DataFrame
|     Combination of the provided DataFrames.
|
| See Also
| -----
| DataFrame.combine_first : Combine two DataFrame objects and default
to
|     non-null values in frame calling the method.
|
| Examples
| -----
| Combine using a simple function that chooses the smaller column.
|
| >>> df1 = pd.DataFrame({'A': [0, 0], 'B': [4, 4]})
| >>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
| >>> take_smaller = lambda s1, s2: s1 if s1.sum() < s2.sum() else s2
| >>> df1.combine(df2, take_smaller)
|      A  B
|    0  0  3
|    1  0  3
|
| Example using a true element-wise combine function.
|
| >>> df1 = pd.DataFrame({'A': [5, 0], 'B': [2, 4]})
| >>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
| >>> df1.combine(df2, np.minimum)
|      A  B
|    0  1  2
|    1  0  3
|
| Using `fill_value` fills Nones prior to passing the column to the
| merge function.
|
| >>> df1 = pd.DataFrame({'A': [0, 0], 'B': [None, 4]})
| >>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
| >>> df1.combine(df2, take_smaller, fill_value=-5)
|      A  B
|    0  0 -5.0
|    1  0  4.0
|
| However, if the same element in both dataframes is None, that None

```



```
|         is preserved
```

```
|         >>> df1 = pd.DataFrame({'A': [0, 0], 'B': [None, 4]})  
|         >>> df2 = pd.DataFrame({'A': [1, 1], 'B': [None, 3]})  
|         >>> df1.combine(df2, take_smaller, fill_value=-5)
```

```
|           A    B  
|         0  0 -5.0  
|         1  0  3.0
```

```
|         Example that demonstrates the use of `overwrite` and behavior when  
|         the axis differ between the dataframes.
```

```
|         >>> df1 = pd.DataFrame({'A': [0, 0], 'B': [4, 4]})  
|         >>> df2 = pd.DataFrame({'B': [3, 3], 'C': [-10, 1], }, index=[1, 2])  
|         >>> df1.combine(df2, take_smaller)
```

```
|           A    B    C  
|         0 NaN NaN NaN  
|         1 NaN 3.0 -10.0  
|         2 NaN 3.0  1.0
```

```
|         >>> df1.combine(df2, take_smaller, overwrite=False)
```

```
|           A    B    C  
|         0 0.0 NaN NaN  
|         1 0.0 3.0 -10.0  
|         2 NaN 3.0  1.0
```

```
|         Demonstrating the preference of the passed in dataframe.
```

```
|         >>> df2 = pd.DataFrame({'B': [3, 3], 'C': [1, 1], }, index=[1, 2])  
|         >>> df2.combine(df1, take_smaller)
```

```
|           A    B    C  
|         0 0.0 NaN NaN  
|         1 0.0 3.0 NaN  
|         2 NaN 3.0 NaN
```

```
|         >>> df2.combine(df1, take_smaller, overwrite=False)
```

```
|           A    B    C  
|         0 0.0 NaN NaN  
|         1 0.0 3.0 1.0  
|         2 NaN 3.0 1.0
```

```
|         combine_first(self, other)
```

```
|         Update null elements with value in the same location in `other`.
```

Combine two DataFrame objects by filling null values in one DataFrame with non-null values from other DataFrame. The row and column indexes of the resulting DataFrame will be the union of the two.

Parameters

other : DataFrame

Provided DataFrame to use to fill null values.

Returns

DataFrame

See Also

DataFrame.combine : Perform series-wise operation on two DataFrames using a given function.

Examples

```
>>> df1 = pd.DataFrame({'A': [None, 0], 'B': [None, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
>>> df1.combine_first(df2)
```

```
   A  B
0  1.0 3.0
1  0.0 4.0
```

Null values still persist if the location of that null value does not exist in `other`

```
>>> df1 = pd.DataFrame({'A': [None, 0], 'B': [4, None]})
>>> df2 = pd.DataFrame({'B': [3, 3], 'C': [1, 1]}, index=[1, 2])
>>> df1.combine_first(df2)
```

```
   A  B  C
0 NaN 4.0 NaN
1  0.0 3.0 1.0
2 NaN 3.0 1.0
```

```
compound(self, axis=None, skipna=None, level=None)
```

Return the compound percentage of the values for the requested axis.

.. deprecated:: 0.25.0

```

|     Parameters
|     -----
|     axis : {index (0), columns (1)}
|         Axis for the function to be applied on.
|     skipna : bool, default True
|         Exclude NA/null values when computing the result.
|     level : int or level name, default None
|         If the axis is a MultiIndex (hierarchical), count along a
|         particular level, collapsing into a Series.
|     numeric_only : bool, default None
|         Include only float, int, boolean columns. If None, will attempt
to use
|         everything, then use only numeric data. Not implemented for
Series.
|     **kwargs
|         Additional keyword arguments to be passed to the function.
|
|     Returns
|     -----
|     Series or DataFrame (if level specified)
|
|     corr(self, method='pearson', min_periods=1)
|         Compute pairwise correlation of columns, excluding NA/null values.
|
|     Parameters
|     -----
|     method : {'pearson', 'kendall', 'spearman'} or callable
|         * pearson : standard correlation coefficient
|         * kendall : Kendall Tau correlation coefficient
|         * spearman : Spearman rank correlation
|         * callable: callable with input two 1d ndarrays
|         and returning a float. Note that the returned matrix from
corr
|         will have 1 along the diagonals and will be symmetric
|         regardless of the callable's behavior
|         .. versionadded:: 0.24.0
|
|     min_periods : int, optional
|         Minimum number of observations required per pair of columns
|         to have a valid result. Currently only available for Pearson
|         and Spearman correlation.
|
|     Returns
|     -----

```

```

|     DataFrame
|         Correlation matrix.
|
|     See Also
|     -----
|     DataFrame.corrwith
|     Series.corr
|
|     Examples
|     -----
|     >>> def histogram_intersection(a, b):
|     ...     v = np.minimum(a, b).sum().round(decimals=1)
|     ...     return v
|     >>> df = pd.DataFrame([(0.2, 0.3), (0.0, 0.6), (0.6, 0.0), (0.2, 0.1)],
|     ...                     columns=['dogs', 'cats'])
|     >>> df.corr(method=histogram_intersection)
|
|         dogs  cats
|     dogs    1.0  0.3
|     cats    0.3  1.0
|
|     corrwith(self, other, axis=0, drop=False, method='pearson')
|         Compute pairwise correlation between rows or columns of DataFrame
|         with rows or columns of Series or DataFrame. DataFrames are first
|         aligned along both axes before computing the correlations.
|
|     Parameters
|     -----
|     other : DataFrame, Series
|         Object with which to compute correlations.
|     axis : {0 or 'index', 1 or 'columns'}, default 0
|         0 or 'index' to compute column-wise, 1 or 'columns' for row-wise.
|     drop : bool, default False
|         Drop missing indices from result.
|     method : {'pearson', 'kendall', 'spearman'} or callable
|         * pearson : standard correlation coefficient
|         * kendall : Kendall Tau correlation coefficient
|         * spearman : Spearman rank correlation
|         * callable: callable with input two 1d ndarrays
|           and returning a float
|
|     .. versionadded:: 0.24.0
|
|     Returns
|     -----

```

```

|     Series
|         Pairwise correlations.
|
|     See Also
|     -----
|     DataFrame.corr
|
|     count(self, axis=0, level=None, numeric_only=False)
|         Count non-NA cells for each column or row.
|
|         The values `None`, `NaN`, `NaT`, and optionally `numpy.inf`
| (depending
|     on `pandas.options.mode.use_inf_as_na`) are considered NA.
|
|     Parameters
|     -----
|     axis : {0 or 'index', 1 or 'columns'}, default 0
|         If 0 or 'index' counts are generated for each column.
|         If 1 or 'columns' counts are generated for each **row**.
|     level : int or str, optional
|         If the axis is a `MultiIndex` (hierarchical), count along a
|         particular `level`, collapsing into a `DataFrame`.
|         A `str` specifies the level name.
|     numeric_only : bool, default False
|         Include only `float`, `int` or `boolean` data.
|
|     Returns
|     -----
|     Series or DataFrame
|         For each column/row the number of non-NA/null entries.
|         If `level` is specified returns a `DataFrame`.
|
|     See Also
|     -----
|     Series.count: Number of non-NA elements in a Series.
|     DataFrame.shape: Number of DataFrame rows and columns (including NA
|         elements).
|     DataFrame.isna: Boolean same-sized DataFrame showing places of NA
|         elements.
|
|     Examples
|     -----
|     Constructing DataFrame from a dictionary:
|

```

```

|     >>> df = pd.DataFrame({"Person":
|     ...                     ["John", "Myla", "Lewis", "John", "Myla"],
|     ...                     "Age": [24., np.nan, 21., 33, 26],
|     ...                     "Single": [False, True, True, True, False]})
|
|     >>> df
|
|         Person  Age  Single
|     0   John  24.0   False
|     1   Myla   NaN    True
|     2  Lewis  21.0    True
|     3   John  33.0    True
|     4   Myla  26.0   False
|
|
|     Notice the uncounted NA values:
|
|
|     >>> df.count()
|
|     Person      5
|     Age         4
|     Single      5
|     dtype: int64
|
|
|     Counts for each **row**:
|
|
|     >>> df.count(axis='columns')
|
|     0      3
|     1      2
|     2      3
|     3      3
|     4      3
|     dtype: int64
|
|
|     Counts for one level of a `MultiIndex`:
|
|
|     >>> df.set_index(["Person", "Single"]).count(level="Person")
|
|         Age
|     Person
|     John      2
|     Lewis     1
|     Myla      1
|
|
|     cov(self, min_periods=None)
|
|     Compute pairwise covariance of columns, excluding NA/null values.
|
|
|     Compute the pairwise covariance among the series of a DataFrame.
|
|     The returned data frame is the `covariance matrix

```

```

|     <https://en.wikipedia.org/wiki/Covariance\_matrix>`__` of the columns
|     of the DataFrame.
|
|     Both NA and null values are automatically excluded from the
|     calculation. (See the note below about bias from missing values.)
|     A threshold can be set for the minimum number of
|     observations for each value created. Comparisons with observations
|     below this threshold will be returned as ``NaN``.
|
|     This method is generally used for the analysis of time series data to
|     understand the relationship between different measures
|     across time.
|
|     Parameters
|     -----
|     min_periods : int, optional
|         Minimum number of observations required per pair of columns
|         to have a valid result.
|
|     Returns
|     -----
|     DataFrame
|         The covariance matrix of the series of the DataFrame.
|
|     See Also
|     -----
|     Series.cov : Compute covariance with another Series.
|     core.window.EWM.cov: Exponential weighted sample covariance.
|     core.window.Expanding.cov : Expanding sample covariance.
|     core.window.Rolling.cov : Rolling sample covariance.
|
|     Notes
|     -----
|     Returns the covariance matrix of the DataFrame's time series.
|     The covariance is normalized by N-1.
|
|     For DataFrames that have Series that are missing data (assuming that
|     data is `missing at random`
|     <https://en.wikipedia.org/wiki/Missing\_data#Missing\_at\_random>`)
|     the returned covariance matrix will be an unbiased estimate
|     of the variance and covariance between the member Series.
|
|     However, for many applications this estimate may not be acceptable

```

```

|     because the estimate covariance matrix is not guaranteed to be
positive
|     semi-definite. This could lead to estimate correlations having
|     absolute values which are greater than one, and/or a non-invertible
|     covariance matrix. See `Estimation of covariance matrices
|     <http://en.wikipedia.org/w/index.php?title=Estimation\_of\_covariance\_matrices>`
|     for more details.
|
|     Examples
|     -----
|
|     >>> df = pd.DataFrame([(1, 2), (0, 3), (2, 0), (1, 1)],
|     ...                    columns=['dogs', 'cats'])
|     >>> df.cov()
|
|           dogs      cats
|     dogs  0.666667 -1.000000
|     cats -1.000000  1.666667
|
|     >>> np.random.seed(42)
|     >>> df = pd.DataFrame(np.random.randn(1000, 5),
|     ...                    columns=['a', 'b', 'c', 'd', 'e'])
|     >>> df.cov()
|
|           a         b         c         d         e
|     a  0.998438 -0.020161  0.059277 -0.008943  0.014144
|     b -0.020161  1.059352 -0.008543 -0.024738  0.009826
|     c  0.059277 -0.008543  1.010670 -0.001486 -0.000271
|     d -0.008943 -0.024738 -0.001486  0.921297 -0.013692
|     e  0.014144  0.009826 -0.000271 -0.013692  0.977795
|
|     **Minimum number of periods**
|
|     This method also supports an optional ``min_periods`` keyword
|     that specifies the required minimum number of non-NA observations for
|     each column pair in order to have a valid result:
|
|     >>> np.random.seed(42)
|     >>> df = pd.DataFrame(np.random.randn(20, 3),
|     ...                    columns=['a', 'b', 'c'])
|     >>> df.loc[df.index[:5], 'a'] = np.nan
|     >>> df.loc[df.index[5:10], 'b'] = np.nan
|     >>> df.cov(min_periods=12)
|
|           a         b         c
|     a  0.316741      NaN -0.150812
|     b      NaN  1.248003  0.191417
|     c -0.150812  0.191417  0.895202

```



```

|
| cummax(self, axis=None, skipna=True, *args, **kwargs)
|     Return cumulative maximum over a DataFrame or Series axis.
|
|     Returns a DataFrame or Series of the same size containing the
cumulative
|     maximum.
|
| Parameters
| -----
|     axis : {0 or 'index', 1 or 'columns'}, default 0
|         The index or the name of the axis. 0 is equivalent to None or
'index'.
|     skipna : boolean, default True
|         Exclude NA/null values. If an entire row/column is NA, the result
|         will be NA.
|     *args, **kwargs :
|         Additional keywords have no effect but might be accepted for
|         compatibility with NumPy.
|
| Returns
| -----
|     Series or DataFrame
|
| See Also
| -----
|     core.window.Expanding.max : Similar functionality
|         but ignores ``NaN`` values.
|     DataFrame.max : Return the maximum over
|         DataFrame axis.
|     DataFrame.cummax : Return cumulative maximum over DataFrame axis.
|     DataFrame.cummin : Return cumulative minimum over DataFrame axis.
|     DataFrame.cumsum : Return cumulative sum over DataFrame axis.
|     DataFrame.cumprod : Return cumulative product over DataFrame axis.
|
| Examples
| -----
|     **Series**
|
|     >>> s = pd.Series([2, np.nan, 5, -1, 0])
|     >>> s
|      0    2.0
|      1    NaN
|      2    5.0

```

```

|      3   -1.0
|      4    0.0
|      dtype: float64
|
|      By default, NA values are ignored.
|
|      >>> s.cummax()
|      0    2.0
|      1   NaN
|      2    5.0
|      3    5.0
|      4    5.0
|      dtype: float64
|
|      To include NA values in the operation, use ``skipna=False``
|
|      >>> s.cummax(skipna=False)
|      0    2.0
|      1   NaN
|      2   NaN
|      3   NaN
|      4   NaN
|      dtype: float64
|
|      **DataFrame**
|
|      >>> df = pd.DataFrame([[2.0, 1.0],
|      ...                    [3.0, np.nan],
|      ...                    [1.0, 0.0]],
|      ...                    columns=list('AB'))
|      >>> df
|           A    B
|      0  2.0  1.0
|      1  3.0  NaN
|      2  1.0  0.0
|
|      By default, iterates over rows and finds the maximum
|      in each column. This is equivalent to ``axis=None`` or
|      ``axis='index'``.
|
|      >>> df.cummax()
|           A    B
|      0  2.0  1.0
|      1  3.0  NaN

```

```

|      2  3.0  1.0
|
|      To iterate over columns and find the maximum in each row,
|      use ``axis=1``
|
|      >>> df.cummax(axis=1)
|
|           A    B
|      0  2.0  2.0
|      1  3.0  NaN
|      2  1.0  1.0
|
|      cummin(self, axis=None, skipna=True, *args, **kwargs)
|          Return cumulative minimum over a DataFrame or Series axis.
|
|          Returns a DataFrame or Series of the same size containing the
cumulative
|          minimum.
|
|          Parameters
|          -----
|          axis : {0 or 'index', 1 or 'columns'}, default 0
|              The index or the name of the axis. 0 is equivalent to None or
'index'.
|          skipna : boolean, default True
|              Exclude NA/null values. If an entire row/column is NA, the result
|              will be NA.
|          *args, **kwargs :
|              Additional keywords have no effect but might be accepted for
|              compatibility with NumPy.
|
|          Returns
|          -----
|          Series or DataFrame
|
|          See Also
|          -----
|          core.window.Expanding.min : Similar functionality
|              but ignores ``NaN`` values.
|          DataFrame.min : Return the minimum over
|              DataFrame axis.
|          DataFrame.cummax : Return cumulative maximum over DataFrame axis.
|          DataFrame.cummin : Return cumulative minimum over DataFrame axis.
|          DataFrame.cumsum : Return cumulative sum over DataFrame axis.
|          DataFrame.cumprod : Return cumulative product over DataFrame axis.

```

| Examples

| -----

| ****Series****

| >>> s = pd.Series([2, np.nan, 5, -1, 0])

| >>> s

| 0 2.0

| 1 NaN

| 2 5.0

| 3 -1.0

| 4 0.0

| dtype: float64

| By default, NA values are ignored.

| >>> s.cummin()

| 0 2.0

| 1 NaN

| 2 2.0

| 3 -1.0

| 4 -1.0

| dtype: float64

| To include NA values in the operation, use ``skipna=False``

| >>> s.cummin(skipna=False)

| 0 2.0

| 1 NaN

| 2 NaN

| 3 NaN

| 4 NaN

| dtype: float64

| ****DataFrame****

| >>> df = pd.DataFrame([[2.0, 1.0],

| ... [3.0, np.nan],

| ... [1.0, 0.0]],

| ... columns=list('AB'))

| >>> df

| A B

| 0 2.0 1.0

| 1 3.0 NaN

```

|      2  1.0  0.0
|
|      By default, iterates over rows and finds the minimum
|      in each column. This is equivalent to ``axis=None`` or
|      ``axis='index'``.
|
|      >>> df.cummin()
|           A    B
|      0  2.0  1.0
|      1  2.0  NaN
|      2  1.0  0.0
|
|      To iterate over columns and find the minimum in each row,
|      use ``axis=1``
|
|      >>> df.cummin(axis=1)
|           A    B
|      0  2.0  1.0
|      1  3.0  NaN
|      2  1.0  0.0
|
|      cumprod(self, axis=None, skipna=True, *args, **kwargs)
|      Return cumulative product over a DataFrame or Series axis.
|
|      Returns a DataFrame or Series of the same size containing the
cumulative
|      product.
|
|      Parameters
|      -----
|      axis : {0 or 'index', 1 or 'columns'}, default 0
|          The index or the name of the axis. 0 is equivalent to None or
'index'.
|      skipna : boolean, default True
|          Exclude NA/null values. If an entire row/column is NA, the result
|          will be NA.
|      *args, **kwargs :
|          Additional keywords have no effect but might be accepted for
|          compatibility with NumPy.
|
|      Returns
|      -----
|      Series or DataFrame
|

```

```

|     See Also
|     -----
|     core.window.Expanding.prod : Similar functionality
|         but ignores ``NaN`` values.
|     DataFrame.prod : Return the product over
|         DataFrame axis.
|     DataFrame.cummax : Return cumulative maximum over DataFrame axis.
|     DataFrame.cummin : Return cumulative minimum over DataFrame axis.
|     DataFrame.cumsum : Return cumulative sum over DataFrame axis.
|     DataFrame.cumprod : Return cumulative product over DataFrame axis.
|
|     Examples
|     -----
|     **Series**
|
|     >>> s = pd.Series([2, np.nan, 5, -1, 0])
|     >>> s
|     0    2.0
|     1    NaN
|     2    5.0
|     3   -1.0
|     4    0.0
|     dtype: float64
|
|     By default, NA values are ignored.
|
|     >>> s.cumprod()
|     0    2.0
|     1    NaN
|     2   10.0
|     3  -10.0
|     4   -0.0
|     dtype: float64
|
|     To include NA values in the operation, use ``skipna=False``
|
|     >>> s.cumprod(skipna=False)
|     0    2.0
|     1    NaN
|     2    NaN
|     3    NaN
|     4    NaN
|     dtype: float64
|

```

```

|     **DataFrame**
|
|     >>> df = pd.DataFrame([[2.0, 1.0],
|     ...                     [3.0, np.nan],
|     ...                     [1.0, 0.0]],
|     ...                     columns=list('AB'))
|     >>> df
|
|           A    B
|     0  2.0  1.0
|     1  3.0  NaN
|     2  1.0  0.0
|
|     By default, iterates over rows and finds the product
|     in each column. This is equivalent to ``axis=None`` or
|     ``axis='index'``.
|
|     >>> df.cumprod()
|
|           A    B
|     0  2.0  1.0
|     1  6.0  NaN
|     2  6.0  0.0
|
|     To iterate over columns and find the product in each row,
|     use ``axis=1``
|
|     >>> df.cumprod(axis=1)
|
|           A    B
|     0  2.0  2.0
|     1  3.0  NaN
|     2  1.0  0.0
|
|     cumsum(self, axis=None, skipna=True, *args, **kwargs)
|         Return cumulative sum over a DataFrame or Series axis.
|
|         Returns a DataFrame or Series of the same size containing the
cumulative
|         sum.
|
|         Parameters
|         -----
|         axis : {0 or 'index', 1 or 'columns'}, default 0
|             The index or the name of the axis. 0 is equivalent to None or
'index'.
|         skipna : boolean, default True

```

```

|         Exclude NA/null values. If an entire row/column is NA, the result
|         will be NA.
|     *args, **kwargs :
|         Additional keywords have no effect but might be accepted for
|         compatibility with NumPy.
|
|     Returns
|     -----
|     Series or DataFrame
|
|     See Also
|     -----
|     core.window.Expanding.sum : Similar functionality
|         but ignores ``NaN`` values.
|     DataFrame.sum : Return the sum over
|         DataFrame axis.
|     DataFrame.cummax : Return cumulative maximum over DataFrame axis.
|     DataFrame.cummin : Return cumulative minimum over DataFrame axis.
|     DataFrame.cumsum : Return cumulative sum over DataFrame axis.
|     DataFrame.cumprod : Return cumulative product over DataFrame axis.
|
|     Examples
|     -----
|     **Series**
|
|     >>> s = pd.Series([2, np.nan, 5, -1, 0])
|     >>> s
|     0    2.0
|     1    NaN
|     2    5.0
|     3   -1.0
|     4    0.0
|     dtype: float64
|
|     By default, NA values are ignored.
|
|     >>> s.cumsum()
|     0    2.0
|     1    NaN
|     2    7.0
|     3    6.0
|     4    6.0
|     dtype: float64
|

```



```

|         To include NA values in the operation, use ``skipna=False``
|
|         >>> s.cumsum(skipna=False)
|         0    2.0
|         1    NaN
|         2    NaN
|         3    NaN
|         4    NaN
|         dtype: float64
|
|         **DataFrame**
|
|         >>> df = pd.DataFrame([[2.0, 1.0],
|         ...                    [3.0, np.nan],
|         ...                    [1.0, 0.0]],
|         ...                    columns=list('AB'))
|         >>> df
|           A    B
|         0  2.0  1.0
|         1  3.0  NaN
|         2  1.0  0.0
|
|         By default, iterates over rows and finds the sum
|         in each column. This is equivalent to ``axis=None`` or
|         ``axis='index'``.
|
|         >>> df.cumsum()
|           A    B
|         0  2.0  1.0
|         1  5.0  NaN
|         2  6.0  1.0
|
|         To iterate over columns and find the sum in each row,
|         use ``axis=1``
|
|         >>> df.cumsum(axis=1)
|           A    B
|         0  2.0  3.0
|         1  3.0  NaN
|         2  1.0  1.0
|
|         diff(self, periods=1, axis=0)
|         First discrete difference of element.
|

```

```

|     Calculates the difference of a DataFrame element compared with
another
|     element in the DataFrame (default is the element in the same column
|     of the previous row).
|
|     Parameters
|     -----
|     periods : int, default 1
|         Periods to shift for calculating difference, accepts negative
|         values.
|     axis : {0 or 'index', 1 or 'columns'}, default 0
|         Take difference over rows (0) or columns (1).
|
|     .. versionadded:: 0.16.1.
|
|     Returns
|     -----
|     DataFrame
|
|     See Also
|     -----
|     Series.diff: First discrete difference for a Series.
|     DataFrame.pct_change: Percent change over given number of periods.
|     DataFrame.shift: Shift index by desired number of periods with an
|         optional time freq.
|
|     Examples
|     -----
|     Difference with previous row
|
|     >>> df = pd.DataFrame({'a': [1, 2, 3, 4, 5, 6],
|     ...                   'b': [1, 1, 2, 3, 5, 8],
|     ...                   'c': [1, 4, 9, 16, 25, 36]})
|     >>> df
|          a  b  c
|     0  1  1  1
|     1  2  1  4
|     2  3  2  9
|     3  4  3  16
|     4  5  5  25
|     5  6  8  36
|
|     >>> df.diff()
|          a  b  c

```

```

|      0  NaN  NaN  NaN
|      1  1.0  0.0  3.0
|      2  1.0  1.0  5.0
|      3  1.0  1.0  7.0
|      4  1.0  2.0  9.0
|      5  1.0  3.0 11.0
|
|      Difference with previous column
|
|      >>> df.diff(axis=1)
|           a      b      c
|      0 NaN  0.0  0.0
|      1 NaN -1.0  3.0
|      2 NaN -1.0  7.0
|      3 NaN -1.0 13.0
|      4 NaN  0.0 20.0
|      5 NaN  2.0 28.0
|
|      Difference with 3rd previous row
|
|      >>> df.diff(periods=3)
|           a      b      c
|      0 NaN  NaN  NaN
|      1 NaN  NaN  NaN
|      2 NaN  NaN  NaN
|      3  3.0  2.0 15.0
|      4  3.0  4.0 21.0
|      5  3.0  6.0 27.0
|
|      Difference with following row
|
|      >>> df.diff(periods=-1)
|           a      b      c
|      0 -1.0  0.0 -3.0
|      1 -1.0 -1.0 -5.0
|      2 -1.0 -1.0 -7.0
|      3 -1.0 -2.0 -9.0
|      4 -1.0 -3.0 -11.0
|      5 NaN  NaN  NaN
|
|      div = truediv(self, other, axis='columns', level=None, fill_value=None)
|
|      divide = truediv(self, other, axis='columns', level=None,
fill_value=None)

```

```

|
| dot(self, other)
|     Compute the matrix multiplication between the DataFrame and other.
|
|     This method computes the matrix product between the DataFrame and the
|     values of an other Series, DataFrame or a numpy array.
|
|     It can also be called using ``self @ other`` in Python >= 3.5.
|
| Parameters
| -----
| other : Series, DataFrame or array-like
|     The other object to compute the matrix product with.
|
| Returns
| -----
| Series or DataFrame
|     If other is a Series, return the matrix product between self and
|     other as a Serie. If other is a DataFrame or a numpy.array,
return
|     the matrix product of self and other in a DataFrame of a
np.array.
|
| See Also
| -----
| Series.dot: Similar method for Series.
|
| Notes
| -----
|     The dimensions of DataFrame and other must be compatible in order to
|     compute the matrix multiplication. In addition, the column names of
|     DataFrame and the index of other must contain the same values, as
they
|     will be aligned prior to the multiplication.
|
|     The dot method for Series computes the inner product, instead of the
|     matrix product here.
|
| Examples
| -----
|     Here we multiply a DataFrame with a Series.
|
|     >>> df = pd.DataFrame([[0, 1, -2, -1], [1, 1, 1, 1]])
|     >>> s = pd.Series([1, 1, 2, 1])

```

```

|     >>> df.dot(s)
|     0    -4
|     1     5
|     dtype: int64
|
|     Here we multiply a DataFrame with another DataFrame.
|
|     >>> other = pd.DataFrame([[0, 1], [1, 2], [-1, -1], [2, 0]])
|     >>> df.dot(other)
|           0    1
|     0    1    4
|     1    2    2
|
|     Note that the dot method give the same result as @
|
|     >>> df @ other
|           0    1
|     0    1    4
|     1    2    2
|
|     The dot method works also if other is an np.array.
|
|     >>> arr = np.array([[0, 1], [1, 2], [-1, -1], [2, 0]])
|     >>> df.dot(arr)
|           0    1
|     0    1    4
|     1    2    2
|
|     Note how shuffling of the objects does not change the result.
|
|     >>> s2 = s.reindex([1, 0, 2, 3])
|     >>> df.dot(s2)
|     0    -4
|     1     5
|     dtype: int64
|
| drop(self, labels=None, axis=0, index=None, columns=None, level=None,
| inplace=False, errors='raise')
|     Drop specified labels from rows or columns.
|
|     Remove rows or columns by specifying label names and corresponding
|     axis, or by specifying directly index or column names. When using a
|     multi-index, labels on different levels can be removed by specifying
|     the level.

```

```

|
| Parameters
| -----
| labels : single label or list-like
|         Index or column labels to drop.
| axis : {0 or 'index', 1 or 'columns'}, default 0
|         Whether to drop labels from the index (0 or 'index') or
|         columns (1 or 'columns').
| index : single label or list-like
|         Alternative to specifying axis (`labels, axis=0`)
|         is equivalent to `index=labels`).
|
| .. versionadded:: 0.21.0
| columns : single label or list-like
|         Alternative to specifying axis (`labels, axis=1`)
|         is equivalent to `columns=labels`).
|
| .. versionadded:: 0.21.0
| level : int or level name, optional
|         For MultiIndex, level from which the labels will be removed.
| inplace : bool, default False
|         If True, do operation inplace and return None.
| errors : {'ignore', 'raise'}, default 'raise'
|         If 'ignore', suppress error and only existing labels are
|         dropped.
|
| Returns
| -----
| DataFrame
|         DataFrame without the removed index or column labels.
|
| Raises
| -----
| KeyError
|         If any of the labels is not found in the selected axis.
|
| See Also
| -----
| DataFrame.loc : Label-location based indexer for selection by label.
| DataFrame.dropna : Return DataFrame with labels on given axis omitted
|                   where (all or any) data are missing.
| DataFrame.drop_duplicates : Return DataFrame with duplicate rows
|                             removed, optionally only considering certain columns.
| Series.drop : Return Series with specified index labels removed.

```

Examples

```
>>> df = pd.DataFrame(np.arange(12).reshape(3, 4),  
...                     columns=['A', 'B', 'C', 'D'])
```

```
>>> df  
   A  B  C  D  
0  0  1  2  3  
1  4  5  6  7  
2  8  9 10 11
```

Drop columns

```
>>> df.drop(['B', 'C'], axis=1)
```

```
   A  D  
0  0  3  
1  4  7  
2  8 11
```

```
>>> df.drop(columns=['B', 'C'])
```

```
   A  D  
0  0  3  
1  4  7  
2  8 11
```

Drop a row by index

```
>>> df.drop([0, 1])
```

```
   A  B  C  D  
2  8  9 10 11
```

Drop columns and/or rows of MultiIndex DataFrame

```
>>> midx = pd.MultiIndex(levels=[['lama', 'cow', 'falcon'],  
...                             ['speed', 'weight', 'length']],  
...                      codes=[[0, 0, 0, 1, 1, 1, 2, 2, 2],  
...                             [0, 1, 2, 0, 1, 2, 0, 1, 2]])  
>>> df = pd.DataFrame(index=midx, columns=['big', 'small'],  
...                   data=[[45, 30], [200, 100], [1.5, 1], [30, 20],  
...                         [250, 150], [1.5, 0.8], [320, 250],  
...                         [1, 0.8], [0.3, 0.2]])
```

```
>>> df  
           big  small  
lama  speed  45.0  30.0
```

```
|      weight  200.0  100.0
|      length  1.5    1.0
|      cow    speed  30.0  20.0
|      weight  250.0  150.0
|      length  1.5    0.8
|      falcon speed  320.0 250.0
|      weight  1.0    0.8
|      length  0.3    0.2
```

```
| >>> df.drop(index='cow', columns='small')
```

```
|      big
|      lama speed  45.0
|      weight 200.0
|      length 1.5
|      falcon speed 320.0
|      weight 1.0
|      length 0.3
```

```
| >>> df.drop(index='length', level=1)
```

```
|      big      small
|      lama speed  45.0  30.0
|      weight 200.0  100.0
|      cow    speed  30.0  20.0
|      weight 250.0  150.0
|      falcon speed  320.0 250.0
|      weight 1.0    0.8
```

```
| drop_duplicates(self, subset=None, keep='first', inplace=False)
```

```
| Return DataFrame with duplicate rows removed, optionally only
| considering certain columns. Indexes, including time indexes
| are ignored.
```

```
| Parameters
```

```
| -----
```

```
| subset : column label or sequence of labels, optional
```

```
| Only consider certain columns for identifying duplicates, by
| default use all of the columns
```

```
| keep : {'first', 'last', False}, default 'first'
```

```
| - ``first`` : Drop duplicates except for the first occurrence.
```

```
| - ``last`` : Drop duplicates except for the last occurrence.
```

```
| - False : Drop all duplicates.
```

```
| inplace : boolean, default False
```

```
| Whether to drop duplicates in place or to return a copy
```



```

|     Returns
|     -----
|     DataFrame
|
|     dropna(self, axis=0, how='any', thresh=None, subset=None, inplace=False)
|         Remove missing values.
|
|         See the :ref:`User Guide <missing_data>` for more on which values are
|         considered missing, and how to work with missing data.
|
|     Parameters
|     -----
|     axis : {0 or 'index', 1 or 'columns'}, default 0
|         Determine if rows or columns which contain missing values are
|         removed.
|
|         * 0, or 'index' : Drop rows which contain missing values.
|         * 1, or 'columns' : Drop columns which contain missing value.
|
|         .. deprecated:: 0.23.0
|
|         Pass tuple or list to drop on multiple axes.
|         Only a single axis is allowed.
|
|     how : {'any', 'all'}, default 'any'
|         Determine if row or column is removed from DataFrame, when we
have
|         at least one NA or all NA.
|
|         * 'any' : If any NA values are present, drop that row or column.
|         * 'all' : If all values are NA, drop that row or column.
|
|     thresh : int, optional
|         Require that many non-NA values.
|     subset : array-like, optional
|         Labels along other axis to consider, e.g. if you are dropping
rows
|         these would be a list of columns to include.
|     inplace : bool, default False
|         If True, do operation inplace and return None.
|
|     Returns
|     -----
|     DataFrame

```

```

|         DataFrame with NA entries dropped from it.
|
| See Also
| -----
| DataFrame.isna: Indicate missing values.
| DataFrame.notna : Indicate existing (non-missing) values.
| DataFrame.fillna : Replace missing values.
| Series.dropna : Drop missing values.
| Index.dropna : Drop missing indices.
|
| Examples
| -----
| >>> df = pd.DataFrame({"name": ['Alfred', 'Batman', 'Catwoman'],
| ...                    "toy": [np.nan, 'Batmobile', 'Bullwhip'],
| ...                    "born": [pd.NaT, pd.Timestamp("1940-04-25"),
| ...                             pd.NaT]})
| >>> df
|          name      toy      born
| 0    Alfred      NaN      NaT
| 1    Batman  Batmobile 1940-04-25
| 2  Catwoman  Bullwhip      NaT
|
| Drop the rows where at least one element is missing.
|
| >>> df.dropna()
|          name      toy      born
| 1    Batman  Batmobile 1940-04-25
|
| Drop the columns where at least one element is missing.
|
| >>> df.dropna(axis='columns')
|          name
| 0    Alfred
| 1    Batman
| 2  Catwoman
|
| Drop the rows where all elements are missing.
|
| >>> df.dropna(how='all')
|          name      toy      born
| 0    Alfred      NaN      NaT
| 1    Batman  Batmobile 1940-04-25
| 2  Catwoman  Bullwhip      NaT
|

```

```

|         Keep only the rows with at least 2 non-NA values.
|
|         >>> df.dropna(thresh=2)
|              name      toy      born
|         1   Batman  Batmobile 1940-04-25
|         2 Catwoman  Bullwhip      NaT
|
|         Define in which columns to look for missing values.
|
|         >>> df.dropna(subset=['name', 'born'])
|              name      toy      born
|         1   Batman  Batmobile 1940-04-25
|
|         Keep the DataFrame with valid entries in the same variable.
|
|         >>> df.dropna(inplace=True)
|         >>> df
|              name      toy      born
|         1   Batman  Batmobile 1940-04-25
|
| duplicated(self, subset=None, keep='first')
|         Return boolean Series denoting duplicate rows, optionally only
|         considering certain columns.
|
|         Parameters
|         -----
|         subset : column label or sequence of labels, optional
|             Only consider certain columns for identifying duplicates, by
|             default use all of the columns
|         keep : {'first', 'last', False}, default 'first'
|             - ``first`` : Mark duplicates as ``True`` except for the
|               first occurrence.
|             - ``last`` : Mark duplicates as ``True`` except for the
|               last occurrence.
|             - False : Mark all duplicates as ``True``.
|
|         Returns
|         -----
|         Series
|
|         eq(self, other, axis='columns', level=None)
|         Get Equal to of dataframe and other, element-wise (binary operator
|         `eq`).
|

```

```

|     Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to
comparison
|     operators.
|
|     Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose
axis
|     (rows or columns) and level for comparison.
|
|     Parameters
|     -----
|     other : scalar, sequence, Series, or DataFrame
|             Any single or multiple element data structure, or list-like
object.
|     axis : {0 or 'index', 1 or 'columns'}, default 'columns'
|             Whether to compare by the index (0 or 'index') or columns
|             (1 or 'columns').
|     level : int or label
|             Broadcast across a level, matching Index values on the passed
|             MultiIndex level.
|
|     Returns
|     -----
|     DataFrame of bool
|             Result of the comparison.
|
|     See Also
|     -----
|     DataFrame.eq : Compare DataFrames for equality elementwise.
|     DataFrame.ne : Compare DataFrames for inequality elementwise.
|     DataFrame.le : Compare DataFrames for less than inequality
|                   or equality elementwise.
|     DataFrame.lt : Compare DataFrames for strictly less than
|                   inequality elementwise.
|     DataFrame.ge : Compare DataFrames for greater than inequality
|                   or equality elementwise.
|     DataFrame.gt : Compare DataFrames for strictly greater than
|                   inequality elementwise.
|
|     Notes
|     -----
|     Mismatched indices will be unioned together.
|     `NaN` values are considered different (i.e. `NaN` != `NaN`).
|
|     Examples

```

```
|
| -----
| >>> df = pd.DataFrame({'cost': [250, 150, 100],
| ...                      'revenue': [100, 250, 300]},
| ...                      index=['A', 'B', 'C'])
| >>> df
|      cost  revenue
| A    250     100
| B    150     250
| C    100     300
|
| Comparison with a scalar, using either the operator or method:
|
| >>> df == 100
|      cost  revenue
| A  False     True
| B  False     False
| C   True     False
|
| >>> df.eq(100)
|      cost  revenue
| A  False     True
| B  False     False
| C   True     False
|
| When `other` is a :class:`Series`, the columns of a DataFrame are
aligned
| with the index of `other` and broadcast:
|
| >>> df != pd.Series([100, 250], index=["cost", "revenue"])
|      cost  revenue
| A   True     True
| B   True     False
| C  False     True
|
| Use the method to control the broadcast axis:
|
| >>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
|      cost  revenue
| A   True     False
| B   True     True
| C   True     True
| D   True     True
|
| When comparing to an arbitrary sequence, the number of columns must
```

```

|         match the number elements in `other`:
|
|         >>> df == [250, 100]
|
|         cost  revenue
|         A   True    True
|         B  False   False
|         C  False   False
|
|         Use the method to control the axis:
|
|         >>> df.eq([250, 250, 100], axis='index')
|
|         cost  revenue
|         A   True    False
|         B  False    True
|         C   True    False
|
|         Compare to a DataFrame of different shape.
|
|         >>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
|         ...                       index=['A', 'B', 'C', 'D'])
|         >>> other
|
|         revenue
|         A      300
|         B      250
|         C      100
|         D      150
|
|         >>> df.gt(other)
|
|         cost  revenue
|         A  False    False
|         B  False    False
|         C  False     True
|         D  False    False
|
|         Compare to a MultiIndex by level.
|
|         >>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300,
220],
|         ...                               'revenue': [100, 250, 300, 200, 175,
225]},
|         ...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2',
'Q2'],
|         ...                               ['A', 'B', 'C', 'A', 'B',
'C']])

```

```
|     >>> df_multindex
|           cost  revenue
|     Q1 A    250     100
|           B    150     250
|           C    100     300
|     Q2 A    150     200
|           B    300     175
|           C    220     225
```

```
|     >>> df.le(df_multindex, level=1)
|           cost  revenue
|     Q1 A   True    True
|           B   True    True
|           C   True    True
|     Q2 A  False    True
|           B   True   False
|           C   True   False
```

```
| eval(self, expr, inplace=False, **kwargs)
```

```
|     Evaluate a string describing operations on DataFrame columns.
```

```
|     Operates on columns only, not specific rows or elements. This allows
|     `eval` to run arbitrary code, which can make you vulnerable to code
|     injection if you pass user input to this function.
```

```
| Parameters
```

```
| -----
```

```
| expr : str
```

```
|     The expression string to evaluate.
```

```
| inplace : bool, default False
```

```
|     If the expression contains an assignment, whether to perform the
|     operation inplace and mutate the existing DataFrame. Otherwise,
|     a new DataFrame is returned.
```

```
| .. versionadded:: 0.18.0.
```

```
| kwargs : dict
```

```
|     See the documentation for :func:`eval` for complete details
|     on the keyword arguments accepted by
|     :meth:`~pandas.DataFrame.query`.
```

```
| Returns
```

```
| -----
```

```
| ndarray, scalar, or pandas object
```

```
|     The result of the evaluation.
```

```

|
| See Also
| -----
| DataFrame.query : Evaluates a boolean expression to query the columns
|                   of a frame.
| DataFrame.assign : Can evaluate an expression or function to create
new
|                   values for a column.
| eval : Evaluate a Python expression as a string using various
|                   backends.
|
| Notes
| -----
| For more details see the API documentation for :func:`~eval`.
| For detailed examples see :ref:`enhancing performance with eval`
| <enhancingperf.eval>`.
|
| Examples
| -----
| >>> df = pd.DataFrame({'A': range(1, 6), 'B': range(10, 0, -2)})
| >>> df
|      A  B
|  0  1  10
|  1  2   8
|  2  3   6
|  3  4   4
|  4  5   2
| >>> df.eval('A + B')
|      11
|      10
|       9
|       8
|       7
| dtype: int64
|
| Assignment is allowed though by default the original DataFrame is not
| modified.
|
| >>> df.eval('C = A + B')
|      A  B  C
|  0  1  10  11
|  1  2   8  10
|  2  3   6   9
|  3  4   4   8

```



```

|      4  5  2  7
|      >>> df
|          A  B
|      0  1  10
|      1  2   8
|      2  3   6
|      3  4   4
|      4  5   2
|
|      Use ``inplace=True`` to modify the original DataFrame.
|
|      >>> df.eval('C = A + B', inplace=True)
|      >>> df
|          A  B  C
|      0  1  10  11
|      1  2   8  10
|      2  3   6   9
|      3  4   4   8
|      4  5   2   7
|
|      ewm(self, com=None, span=None, halflife=None, alpha=None, min_periods=0,
adjust=True, ignore_na=False, axis=0)
|      Provide exponential weighted functions.
|
|      .. versionadded:: 0.18.0
|
|      Parameters
|      -----
|      com : float, optional
|          Specify decay in terms of center of mass,
|          :math:\alpha = 1 / (1 + com), \text{ for } com \geq 0.
|      span : float, optional
|          Specify decay in terms of span,
|          :math:\alpha = 2 / (span + 1), \text{ for } span \geq 1.
|      halflife : float, optional
|          Specify decay in terms of half-life,
|          :math:\alpha = 1 - \exp(\log(0.5) / halflife), \text{ for } halflife
> 0.
|      alpha : float, optional
|          Specify smoothing factor :math:\alpha` directly,
|          :math:0 < \alpha \leq 1.
|
|      .. versionadded:: 0.18.0
|

```

```

|     min_periods : int, default 0
|         Minimum number of observations in window required to have a value
|         (otherwise result is NA).
|     adjust : bool, default True
|         Divide by decaying adjustment factor in beginning periods to
account
|         for imbalance in relative weightings
|         (viewing EWMA as a moving average).
|     ignore_na : bool, default False
|         Ignore missing values when calculating weights;
|         specify True to reproduce pre-0.15.0 behavior.
|     axis : {0 or 'index', 1 or 'columns'}, default 0
|         The axis to use. The value 0 identifies the rows, and 1
|         identifies the columns.
|
| Returns
| -----
| DataFrame
|     A Window sub-classed for the particular operation.
|
| See Also
| -----
|     rolling : Provides rolling window calculations.
|     expanding : Provides expanding transformations.
|
| Notes
| -----
|     Exactly one of center of mass, span, half-life, and alpha must be
provided.
|     Allowed values and relationship between the parameters are specified
in the
|     parameter descriptions above; see the link at the end of this section
for
|     a detailed explanation.
|
|     When adjust is True (default), weighted averages are calculated using
|     weights  $(1-\alpha)^{(n-1)}$ ,  $(1-\alpha)^{(n-2)}$ , ...,  $1-\alpha$ , 1.
|
|     When adjust is False, weighted averages are calculated recursively
as:
|         weighted_average[0] = arg[0];
|         weighted_average[i] =  $(1-\alpha) \cdot \text{weighted\_average}[i-1] +$ 
alpha*arg[i].
|

```

```

|     When ignore_na is False (default), weights are based on absolute
positions.
|     For example, the weights of x and y used in calculating the final
weighted
|     average of [x, None, y] are (1-alpha)**2 and 1 (if adjust is True),
and
|     (1-alpha)**2 and alpha (if adjust is False).
|
|     When ignore_na is True (reproducing pre-0.15.0 behavior), weights are
based
|     on relative positions. For example, the weights of x and y used in
|     calculating the final weighted average of [x, None, y] are 1-alpha
and 1
|     (if adjust is True), and 1-alpha and alpha (if adjust is False).
|
|     More details can be found at
|     http://pandas.pydata.org/pandas-
docs/stable/user_guide/computation.html#exponentially-weighted-windows
|
|     Examples
|     -----
|
|     >>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]})
|     >>> df
|          B
|     0  0.0
|     1  1.0
|     2  2.0
|     3  NaN
|     4  4.0
|
|     >>> df.ewm(com=0.5).mean()
|          B
|     0  0.000000
|     1  0.750000
|     2  1.615385
|     3  1.615385
|     4  3.670213
|
|     expanding(self, min_periods=1, center=False, axis=0)
|         Provide expanding transformations.
|
|     .. versionadded:: 0.18.0
|

```

```

|     Parameters
|     -----
|     min_periods : int, default 1
|         Minimum number of observations in window required to have a value
|         (otherwise result is NA).
|     center : bool, default False
|         Set the labels at the center of the window.
|     axis : int or str, default 0
|
|     Returns
|     -----
|     a Window sub-classed for the particular operation
|
|     See Also
|     -----
|     rolling : Provides rolling window calculations.
|     ewm : Provides exponential weighted functions.
|
|     Notes
|     -----
|     By default, the result is set to the right edge of the window. This
can be
|     changed to the center of the window by setting ``center=True``.
|
|     Examples
|     -----
|
|     >>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]})
|
|         B
|     0  0.0
|     1  1.0
|     2  2.0
|     3  NaN
|     4  4.0
|
|     >>> df.expanding(2).sum()
|
|         B
|     0  NaN
|     1  1.0
|     2  3.0
|     3  3.0
|     4  7.0
|
|     explode(self, column: Union[str, Tuple]) -> 'DataFrame'

```

```

|     Transform each element of a list-like to a row, replicating the
|     index values.
|
|     .. versionadded:: 0.25.0
|
|     Parameters
|     -----
|     column : str or tuple
|
|     Returns
|     -----
|     DataFrame
|         Exploded lists to rows of the subset columns;
|         index will be duplicated for these rows.
|
|     Raises
|     -----
|     ValueError :
|         if columns of the frame are not unique.
|
|     See Also
|     -----
|     DataFrame.unstack : Pivot a level of the (necessarily hierarchical)
|         index labels
|     DataFrame.melt : Unpivot a DataFrame from wide format to long format
|     Series.explode : Explode a DataFrame from list-like columns to long
format.
|
|     Notes
|     -----
|     This routine will explode list-likes including lists, tuples,
|     Series, and np.ndarray. The result dtype of the subset rows will
|     be object. Scalars will be returned unchanged. Empty list-likes will
|     result in a np.nan for that row.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'A': [[1, 2, 3], 'foo', [], [3, 4]], 'B': 1})
|     >>> df
|
|           A  B
|     0  [1, 2, 3]  1
|     1         foo  1
|     2          []  1
|     3      [3, 4]  1

```

```

|
|     >>> df.explode('A')
|           A  B
|     0    1  1
|     0    2  1
|     0    3  1
|     1  foo  1
|     2  NaN  1
|     3    3  1
|     3    4  1
|
|     fillna(self, value=None, method=None, axis=None, inplace=False,
limit=None, downcast=None, **kwargs)
|         Fill NA/NaN values using the specified method.
|
|     Parameters
|     -----
|     value : scalar, dict, Series, or DataFrame
|         Value to use to fill holes (e.g. 0), alternately a
|         dict/Series/DataFrame of values specifying which value to use for
|         each index (for a Series) or column (for a DataFrame). Values
not
|         in the dict/Series/DataFrame will not be filled. This value
cannot
|         be a list.
|     method : {'backfill', 'bfill', 'pad', 'ffill', None}, default None
|         Method to use for filling holes in reindexed Series
|         pad / ffill: propagate last valid observation forward to next
valid
|         backfill / bfill: use next valid observation to fill gap.
|     axis : {0 or 'index', 1 or 'columns'}
|         Axis along which to fill missing values.
|     inplace : bool, default False
|         If True, fill in-place. Note: this will modify any
|         other views on this object (e.g., a no-copy slice for a column in
a
|         DataFrame).
|     limit : int, default None
|         If method is specified, this is the maximum number of consecutive
|         NaN values to forward/backward fill. In other words, if there is
|         a gap with more than this number of consecutive NaNs, it will
only
|         be partially filled. If method is not specified, this is the

```

```

|         maximum number of entries along the entire axis where NaNs will
be
|         filled. Must be greater than 0 if not None.
|         downcast : dict, default is None
|         A dict of item->dtype of what to downcast if possible,
|         or the string 'infer' which will try to downcast to an
appropriate
|         equal type (e.g. float64 to int64 if possible).
|
|     Returns
|     -----
|     DataFrame
|     Object with missing values filled.
|
|     See Also
|     -----
|     interpolate : Fill NaN values using interpolation.
|     reindex : Conform object to new index.
|     asfreq : Convert TimeSeries to specified frequency.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame([[np.nan, 2, np.nan, 0],
|     ...                    [3, 4, np.nan, 1],
|     ...                    [np.nan, np.nan, np.nan, 5],
|     ...                    [np.nan, 3, np.nan, 4]],
|     ...                    columns=list('ABCD'))
|     >>> df
|
|           A    B    C    D
|     0  NaN  2.0 NaN  0
|     1  3.0  4.0 NaN  1
|     2  NaN  NaN NaN  5
|     3  NaN  3.0 NaN  4
|
|     Replace all NaN elements with 0s.
|
|     >>> df.fillna(0)
|
|           A    B    C    D
|     0  0.0  2.0  0.0  0
|     1  3.0  4.0  0.0  1
|     2  0.0  0.0  0.0  5
|     3  0.0  3.0  0.0  4
|
|     We can also propagate non-null values forward or backward.

```

```

|
|     >>> df.fillna(method='ffill')
|           A    B    C    D
|     0    NaN  2.0  NaN  0
|     1    3.0  4.0  NaN  1
|     2    3.0  4.0  NaN  5
|     3    3.0  3.0  NaN  4
|
|     Replace all NaN elements in column 'A', 'B', 'C', and 'D', with 0, 1,
|     2, and 3 respectively.
|
|     >>> values = {'A': 0, 'B': 1, 'C': 2, 'D': 3}
|     >>> df.fillna(value=values)
|           A    B    C    D
|     0    0.0  2.0  2.0  0
|     1    3.0  4.0  2.0  1
|     2    0.0  1.0  2.0  5
|     3    0.0  3.0  2.0  4
|
|     Only replace the first NaN element.
|
|     >>> df.fillna(value=values, limit=1)
|           A    B    C    D
|     0    0.0  2.0  2.0  0
|     1    3.0  4.0  NaN  1
|     2    NaN  1.0  NaN  5
|     3    NaN  3.0  NaN  4
|
|     floordiv(self, other, axis='columns', level=None, fill_value=None)
|         Get Integer division of dataframe and other, element-wise (binary
operator `floordiv`).
|
|         Equivalent to ``dataframe // other``, but with support to substitute
a fill_value
|         for missing data in one of the inputs. With reverse version,
`rfloordiv`.
|
|         Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
|     Parameters
|     -----
|     other : scalar, sequence, Series, or DataFrame

```



```

|         Any single or multiple element data structure, or list-like
object.
|         axis : {0 or 'index', 1 or 'columns'}
|             Whether to compare by the index (0 or 'index') or columns
|             (1 or 'columns'). For Series input, axis to match Series index
on.
|         level : int or label
|             Broadcast across a level, matching Index values on the
|             passed MultiIndex level.
|         fill_value : float or None, default None
|             Fill existing missing (NaN) values, and any new element needed
for
|             successful DataFrame alignment, with this value before
computation.
|             If data in both corresponding DataFrame locations is missing
|             the result will be missing.
|
|         Returns
|         -----
|         DataFrame
|             Result of the arithmetic operation.
|
|         See Also
|         -----
|         DataFrame.add : Add DataFrames.
|         DataFrame.sub : Subtract DataFrames.
|         DataFrame.mul : Multiply DataFrames.
|         DataFrame.div : Divide DataFrames (float division).
|         DataFrame.truediv : Divide DataFrames (float division).
|         DataFrame.floordiv : Divide DataFrames (integer division).
|         DataFrame.mod : Calculate modulo (remainder after division).
|         DataFrame.pow : Calculate exponential power.
|
|         Notes
|         -----
|         Mismatched indices will be unioned together.
|
|         Examples
|         -----
|         >>> df = pd.DataFrame({'angles': [0, 3, 4],
|         ...                    'degrees': [360, 180, 360]},
|         ...                    index=['circle', 'triangle', 'rectangle'])
|         >>> df
|
|             angles  degrees

```

```

|   circle      0      360
|   triangle    3      180
|   rectangle   4      360
|
|   Add a scalar with operator version which return the same
|   results.
|
|   >>> df + 1
|           angles  degrees
|   circle      1      361
|   triangle    4      181
|   rectangle   5      361
|
|   >>> df.add(1)
|           angles  degrees
|   circle      1      361
|   triangle    4      181
|   rectangle   5      361
|
|   Divide by constant with reverse version.
|
|   >>> df.div(10)
|           angles  degrees
|   circle      0.0    36.0
|   triangle    0.3    18.0
|   rectangle   0.4    36.0
|
|   >>> df.rdiv(10)
|           angles  degrees
|   circle      inf    0.027778
|   triangle    3.333333 0.055556
|   rectangle   2.500000 0.027778
|
|   Subtract a list and Series by axis with operator version.
|
|   >>> df - [1, 2]
|           angles  degrees
|   circle      -1     358
|   triangle     2     178
|   rectangle     3     358
|
|   >>> df.sub([1, 2], axis='columns')
|           angles  degrees
|   circle      -1     358

```

```

|     triangle      2      178
|     rectangle     3      358
|
|     >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
'rectangle'])),
|     ...         axis='index')
|           angles  degrees
|     circle      -1      359
|     triangle     2      179
|     rectangle    3      359
|
|     Multiply a DataFrame of different shape with operator version.
|
|     >>> other = pd.DataFrame({'angles': [0, 3, 4]},
|     ...                       index=['circle', 'triangle', 'rectangle'])
|     >>> other
|           angles
|     circle      0
|     triangle     3
|     rectangle    4
|
|     >>> df * other
|           angles  degrees
|     circle      0      NaN
|     triangle     9      NaN
|     rectangle   16      NaN
|
|     >>> df.mul(other, fill_value=0)
|           angles  degrees
|     circle      0      0.0
|     triangle     9      0.0
|     rectangle   16      0.0
|
|     Divide by a MultiIndex by level.
|
|     >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
|     ...                               'degrees': [360, 180, 360, 360, 540,
720]}),
|     ...                               index=[['A', 'A', 'A', 'B', 'B',
'B'],
|     ...                               ['circle', 'triangle',
'rectangle',
|     ...                               'square', 'pentagon',
'hexagon']])

```

```

|     >>> df_multindex
|
|           angles  degrees
|   A circle      0      360
|   triangle      3      180
|   rectangle      4      360
|   B square      4      360
|   pentagon      5      540
|   hexagon       6      720
|
|
|     >>> df.div(df_multindex, level=1, fill_value=0)
|
|           angles  degrees
|   A circle      NaN      1.0
|   triangle      1.0      1.0
|   rectangle      1.0      1.0
|   B square      0.0      0.0
|   pentagon      0.0      0.0
|   hexagon       0.0      0.0
|
|
|   ge(self, other, axis='columns', level=None)
|       Get Greater than or equal to of dataframe and other, element-wise
(binary operator `ge`).
|
|       Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to
comparison
|       operators.
|
|       Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose
axis
|       (rows or columns) and level for comparison.
|
|       Parameters
|       -----
|       other : scalar, sequence, Series, or DataFrame
|           Any single or multiple element data structure, or list-like
object.
|       axis : {0 or 'index', 1 or 'columns'}, default 'columns'
|           Whether to compare by the index (0 or 'index') or columns
(1 or 'columns').
|       level : int or label
|           Broadcast across a level, matching Index values on the passed
MultiIndex level.
|
|       Returns
|       -----

```

```

|     DataFrame of bool
|         Result of the comparison.
|
|     See Also
|     -----
|     DataFrame.eq : Compare DataFrames for equality elementwise.
|     DataFrame.ne : Compare DataFrames for inequality elementwise.
|     DataFrame.le : Compare DataFrames for less than inequality
|         or equality elementwise.
|     DataFrame.lt : Compare DataFrames for strictly less than
|         inequality elementwise.
|     DataFrame.ge : Compare DataFrames for greater than inequality
|         or equality elementwise.
|     DataFrame.gt : Compare DataFrames for strictly greater than
|         inequality elementwise.
|
|     Notes
|     -----
|     Mismatched indices will be unioned together.
|     `NaN` values are considered different (i.e. `NaN` != `NaN`).
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'cost': [250, 150, 100],
|     ...                     'revenue': [100, 250, 300]},
|     ...                     index=['A', 'B', 'C'])
|     >>> df
|
|         cost  revenue
|     A    250     100
|     B    150     250
|     C    100     300
|
|     Comparison with a scalar, using either the operator or method:
|
|     >>> df == 100
|
|         cost  revenue
|     A  False     True
|     B  False     False
|     C   True     False
|
|     >>> df.eq(100)
|
|         cost  revenue
|     A  False     True
|     B  False     False

```

```

|      C    True    False
|
|      When `other` is a :class:`Series`, the columns of a DataFrame are
aligned
|      with the index of `other` and broadcast:
|
|      >>> df != pd.Series([100, 250], index=["cost", "revenue"])
|      cost  revenue
|      A    True    True
|      B    True    False
|      C    False   True
|
|      Use the method to control the broadcast axis:
|
|      >>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
|      cost  revenue
|      A    True    False
|      B    True     True
|      C    True     True
|      D    True     True
|
|      When comparing to an arbitrary sequence, the number of columns must
|      match the number elements in `other`:
|
|      >>> df == [250, 100]
|      cost  revenue
|      A    True     True
|      B   False    False
|      C   False    False
|
|      Use the method to control the axis:
|
|      >>> df.eq([250, 250, 100], axis='index')
|      cost  revenue
|      A    True    False
|      B   False     True
|      C    True    False
|
|      Compare to a DataFrame of different shape.
|
|      >>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
|      ...                      index=['A', 'B', 'C', 'D'])
|      >>> other
|      revenue

```

```

|      A      300
|      B      250
|      C      100
|      D      150
|
|      >>> df.gt(other)
|           cost  revenue
|      A  False   False
|      B  False   False
|      C  False    True
|      D  False   False
|
|      Compare to a MultiIndex by level.
|
|      >>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300,
220],
|      ...                               'revenue': [100, 250, 300, 200, 175,
225]},
|      ...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2',
'Q2'],
|      ...                               ['A', 'B', 'C', 'A', 'B',
'C']])
|
|      >>> df_multindex
|           cost  revenue
|      Q1 A    250     100
|          B    150     250
|          C    100     300
|      Q2 A    150     200
|          B    300     175
|          C    220     225
|
|      >>> df.le(df_multindex, level=1)
|           cost  revenue
|      Q1 A    True     True
|          B    True     True
|          C    True     True
|      Q2 A   False     True
|          B    True     False
|          C    True     False
|
|      get_value(self, index, col, takeable=False)
|      Quickly retrieve single value at passed column and index.
|
|      .. deprecated:: 0.21.0

```

```

|         Use .at[] or .iat[] accessors instead.
|
|     Parameters
|     -----
|     index : row label
|     col : column label
|     takeable : interpret the index/col as indexers, default False
|
|     Returns
|     -----
|     scalar
|
|     gt(self, other, axis='columns', level=None)
|         Get Greater than of dataframe and other, element-wise (binary
operator `gt`).
|
|     Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to
comparison
|     operators.
|
|     Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose
axis
|     (rows or columns) and level for comparison.
|
|     Parameters
|     -----
|     other : scalar, sequence, Series, or DataFrame
|         Any single or multiple element data structure, or list-like
object.
|     axis : {0 or 'index', 1 or 'columns'}, default 'columns'
|         Whether to compare by the index (0 or 'index') or columns
|         (1 or 'columns').
|     level : int or label
|         Broadcast across a level, matching Index values on the passed
|         MultiIndex level.
|
|     Returns
|     -----
|     DataFrame of bool
|         Result of the comparison.
|
|     See Also
|     -----
|     DataFrame.eq : Compare DataFrames for equality elementwise.

```



```

| DataFrame.ne : Compare DataFrames for inequality elementwise.
| DataFrame.le : Compare DataFrames for less than inequality
| or equality elementwise.
| DataFrame.lt : Compare DataFrames for strictly less than
| inequality elementwise.
| DataFrame.ge : Compare DataFrames for greater than inequality
| or equality elementwise.
| DataFrame.gt : Compare DataFrames for strictly greater than
| inequality elementwise.
|
| Notes
| -----
| Mismatched indices will be unioned together.
| `NaN` values are considered different (i.e. `NaN` != `NaN`).
|
| Examples
| -----
| >>> df = pd.DataFrame({'cost': [250, 150, 100],
| ...                    'revenue': [100, 250, 300]},
| ...                    index=['A', 'B', 'C'])
| >>> df
|      cost  revenue
| A    250     100
| B    150     250
| C    100     300
|
| Comparison with a scalar, using either the operator or method:
|
| >>> df == 100
|      cost  revenue
| A  False     True
| B  False    False
| C   True    False
|
| >>> df.eq(100)
|      cost  revenue
| A  False     True
| B  False    False
| C   True    False
|
| When `other` is a :class:`Series`, the columns of a DataFrame are
aligned
| with the index of `other` and broadcast:
|

```

```

|     >>> df != pd.Series([100, 250], index=["cost", "revenue"])
|
|     cost  revenue
|     A   True    True
|     B   True   False
|     C  False    True
|
|
|     Use the method to control the broadcast axis:
|
|     >>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
|
|     cost  revenue
|     A   True   False
|     B   True    True
|     C   True    True
|     D   True    True
|
|
|     When comparing to an arbitrary sequence, the number of columns must
|     match the number elements in `other`:
|
|     >>> df == [250, 100]
|
|     cost  revenue
|     A   True    True
|     B  False   False
|     C  False   False
|
|
|     Use the method to control the axis:
|
|     >>> df.eq([250, 250, 100], axis='index')
|
|     cost  revenue
|     A   True   False
|     B  False    True
|     C   True   False
|
|
|     Compare to a DataFrame of different shape.
|
|     >>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
|     ...                       index=['A', 'B', 'C', 'D'])
|     >>> other
|
|     revenue
|     A      300
|     B      250
|     C      100
|     D      150
|
|     >>> df.gt(other)

```

```

|         cost  revenue
|      A  False   False
|      B  False   False
|      C  False    True
|      D  False   False
|
|      Compare to a MultiIndex by level.
|
|      >>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300,
220],
|      ...                               'revenue': [100, 250, 300, 200, 175,
225]},
|      ...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2',
'Q2'],
|      ...                               ['A', 'B', 'C', 'A', 'B',
'C']])
|      >>> df_multindex
|          cost  revenue
|      Q1 A    250     100
|          B    150     250
|          C    100     300
|      Q2 A    150     200
|          B    300     175
|          C    220     225
|
|      >>> df.le(df_multindex, level=1)
|          cost  revenue
|      Q1 A    True     True
|          B    True     True
|          C    True     True
|      Q2 A   False     True
|          B    True     False
|          C    True     False
|
|      hist = hist_frame(data, column=None, by=None, grid=True, xlabelsize=None,
xrot=None, ylabelsize=None, yrot=None, ax=None, sharex=False, sharey=False,
figsize=None, layout=None, bins=10, **kws)
|      Make a histogram of the DataFrame's.
|
|      A `histogram`_ is a representation of the distribution of data.
|      This function calls :meth:`matplotlib.pyplot.hist`, on each series in
|      the DataFrame, resulting in one histogram per column.
|
|      .. _histogram: https://en.wikipedia.org/wiki/Histogram

```

```

|
| Parameters
| -----
| data : DataFrame
|     The pandas object holding the data.
| column : string or sequence
|     If passed, will be used to limit data to a subset of columns.
| by : object, optional
|     If passed, then used to form histograms for separate groups.
| grid : bool, default True
|     Whether to show axis grid lines.
| xlabelsize : int, default None
|     If specified changes the x-axis label size.
| xrot : float, default None
|     Rotation of x axis labels. For example, a value of 90 displays
the
|     x labels rotated 90 degrees clockwise.
| ylabelsize : int, default None
|     If specified changes the y-axis label size.
| yrot : float, default None
|     Rotation of y axis labels. For example, a value of 90 displays
the
|     y labels rotated 90 degrees clockwise.
| ax : Matplotlib axes object, default None
|     The axes to plot the histogram on.
| sharex : bool, default True if ax is None else False
|     In case subplots=True, share x axis and set some x axis labels to
|     invisible; defaults to True if ax is None otherwise False if an
ax
|     is passed in.
|     Note that passing in both an ax and sharex=True will alter all x
axis
|     labels for all subplots in a figure.
| sharey : bool, default False
|     In case subplots=True, share y axis and set some y axis labels to
|     invisible.
| figsize : tuple
|     The size in inches of the figure to create. Uses the value in
|     `matplotlib.rcParams` by default.
| layout : tuple, optional
|     Tuple of (rows, columns) for the layout of the histograms.
| bins : integer or sequence, default 10
|     Number of histogram bins to be used. If an integer is given, bins
+ 1

```

```

|         bin edges are calculated and returned. If bins is a sequence,
gives |
|         bin edges, including left edge of first bin and right edge of
last |
|         bin. In this case, bins is returned unmodified.
|     **kwds
|         All other plotting keyword arguments to be passed to
|         :meth:`matplotlib.pyplot.hist`.
|
|     Returns
|     -----
|     matplotlib.AxesSubplot or numpy.ndarray of them
|
|     See Also
|     -----
|     matplotlib.pyplot.hist : Plot a histogram using matplotlib.
|
|     Examples
|     -----
|
|     .. plot::
|
|         :context: close-figs
|
|         This example draws a histogram based on the length and width of
|         some animals, displayed in three bins
|
|         >>> df = pd.DataFrame({
|             'length': [1.5, 0.5, 1.2, 0.9, 3],
|             'width': [0.7, 0.2, 0.15, 0.2, 1.1]
|             }, index= ['pig', 'rabbit', 'duck', 'chicken', 'horse'])
|         >>> hist = df.hist(bins=3)
|
|     idxmax(self, axis=0, skipna=True)
|         Return index of first occurrence of maximum over requested axis.
|         NA/null values are excluded.
|
|     Parameters
|     -----
|     axis : {0 or 'index', 1 or 'columns'}, default 0
|         0 or 'index' for row-wise, 1 or 'columns' for column-wise
|     skipna : boolean, default True
|         Exclude NA/null values. If an entire row/column is NA, the result
|         will be NA.
|
|

```

```

|     Returns
|     -----
|     Series
|         Indexes of maxima along the specified axis.
|
|     Raises
|     -----
|     ValueError
|         * If the row/column is empty
|
|     See Also
|     -----
|     Series.idxmax
|
|     Notes
|     -----
|     This method is the DataFrame version of ``ndarray.argmax``.
|
|     idxmin(self, axis=0, skipna=True)
|         Return index of first occurrence of minimum over requested axis.
|         NA/null values are excluded.
|
|     Parameters
|     -----
|     axis : {0 or 'index', 1 or 'columns'}, default 0
|         0 or 'index' for row-wise, 1 or 'columns' for column-wise
|     skipna : boolean, default True
|         Exclude NA/null values. If an entire row/column is NA, the result
|         will be NA.
|
|     Returns
|     -----
|     Series
|         Indexes of minima along the specified axis.
|
|     Raises
|     -----
|     ValueError
|         * If the row/column is empty
|
|     See Also
|     -----
|     Series.idxmin
|

```

```

|     Notes
|     -----
|     This method is the DataFrame version of ``ndarray.argmax``.
|
| info(self, verbose=None, buf=None, max_cols=None, memory_usage=None,
null_counts=None)
|     Print a concise summary of a DataFrame.
|
|     This method prints information about a DataFrame including
|     the index dtype and column dtypes, non-null values and memory usage.
|
| Parameters
|     -----
|     verbose : bool, optional
|         Whether to print the full summary. By default, the setting in
|         ``pandas.options.display.max_info_columns`` is followed.
|     buf : writable buffer, defaults to sys.stdout
|         Where to send the output. By default, the output is printed to
|         sys.stdout. Pass a writable buffer if you need to further process
|         the output.
|     max_cols : int, optional
|         When to switch from the verbose to the truncated output. If the
|         DataFrame has more than `max_cols` columns, the truncated output
|         is used. By default, the setting in
|         ``pandas.options.display.max_info_columns`` is used.
|     memory_usage : bool, str, optional
|         Specifies whether total memory usage of the DataFrame
|         elements (including the index) should be displayed. By default,
|         this follows the ``pandas.options.display.memory_usage`` setting.
|
|         True always show memory usage. False never shows memory usage.
|         A value of 'deep' is equivalent to "True with deep
introspection".
|
|         Memory usage is shown in human-readable units (base-2
|         representation). Without deep introspection a memory estimation
is
|         made based in column dtype and number of rows assuming values
|         consume the same memory amount for corresponding dtypes. With
deep
|         memory introspection, a real memory usage calculation is
performed
|
|         at the cost of computational resources.
|     null_counts : bool, optional
|         Whether to show the non-null counts. By default, this is shown

```

```

|         only if the frame is smaller than
|         ``pandas.options.display.max_info_rows`` and
|         ``pandas.options.display.max_info_columns``. A value of True
always
|         shows the counts, and False never shows the counts.
|
| Returns
| -----
| None
|         This method prints a summary of a DataFrame and returns None.
|
| See Also
| -----
| DataFrame.describe: Generate descriptive statistics of DataFrame
|         columns.
| DataFrame.memory_usage: Memory usage of DataFrame columns.
|
| Examples
| -----
| >>> int_values = [1, 2, 3, 4, 5]
| >>> text_values = ['alpha', 'beta', 'gamma', 'delta', 'epsilon']
| >>> float_values = [0.0, 0.25, 0.5, 0.75, 1.0]
| >>> df = pd.DataFrame({"int_col": int_values, "text_col":
text_values,
|         ...                "float_col": float_values})
| >>> df
|
|    int_col text_col  float_col
|  0      1    alpha      0.00
|  1      2    beta      0.25
|  2      3   gamma      0.50
|  3      4   delta      0.75
|  4      5  epsilon      1.00
|
| Prints information of all columns:
|
| >>> df.info(verbose=True)
| <class 'pandas.core.frame.DataFrame'>
| RangeIndex: 5 entries, 0 to 4
| Data columns (total 3 columns):
| int_col      5 non-null int64
| text_col     5 non-null object
| float_col    5 non-null float64
| dtypes: float64(1), int64(1), object(1)
| memory usage: 248.0+ bytes

```



```

| Prints a summary of columns count and its dtypes but not per column
| information:
|
| >>> df.info(verbose=False)
| <class 'pandas.core.frame.DataFrame'>
| RangeIndex: 5 entries, 0 to 4
| Columns: 3 entries, int_col to float_col
| dtypes: float64(1), int64(1), object(1)
| memory usage: 248.0+ bytes
|
| Pipe output of DataFrame.info to buffer instead of sys.stdout, get
| buffer content and writes to a text file:
|
| >>> import io
| >>> buffer = io.StringIO()
| >>> df.info(buf=buffer)
| >>> s = buffer.getvalue()
| >>> with open("df_info.txt", "w",
| ...         encoding="utf-8") as f: # doctest: +SKIP
| ...     f.write(s)
| 260
|
| The `memory_usage` parameter allows deep introspection mode,
specially
| useful for big DataFrames and fine-tune memory optimization:
|
| >>> random_strings_array = np.random.choice(['a', 'b', 'c'], 10 ** 6)
| >>> df = pd.DataFrame({
| ...     'column_1': np.random.choice(['a', 'b', 'c'], 10 ** 6),
| ...     'column_2': np.random.choice(['a', 'b', 'c'], 10 ** 6),
| ...     'column_3': np.random.choice(['a', 'b', 'c'], 10 ** 6)
| ... })
| >>> df.info()
| <class 'pandas.core.frame.DataFrame'>
| RangeIndex: 1000000 entries, 0 to 999999
| Data columns (total 3 columns):
| column_1    1000000 non-null object
| column_2    1000000 non-null object
| column_3    1000000 non-null object
| dtypes: object(3)
| memory usage: 22.9+ MB
|
| >>> df.info(memory_usage='deep')

```

```

|     <class 'pandas.core.frame.DataFrame'>
|     RangeIndex: 1000000 entries, 0 to 999999
|     Data columns (total 3 columns):
|     column_1    1000000 non-null object
|     column_2    1000000 non-null object
|     column_3    1000000 non-null object
|     dtypes: object(3)
|     memory usage: 188.8 MB
|
|     insert(self, loc, column, value, allow_duplicates=False)
|         Insert column into DataFrame at specified location.
|
|         Raises a ValueError if `column` is already contained in the
DataFrame,
|         unless `allow_duplicates` is set to True.
|
|     Parameters
|     -----
|     loc : int
|         Insertion index. Must verify 0 <= loc <= len(columns)
|     column : string, number, or hashable object
|         label of the inserted column
|     value : int, Series, or array-like
|     allow_duplicates : bool, optional
|
|     isin(self, values)
|         Whether each element in the DataFrame is contained in values.
|
|     Parameters
|     -----
|     values : iterable, Series, DataFrame or dict
|         The result will only be true at a location if all the
|         labels match. If `values` is a Series, that's the index. If
|         `values` is a dict, the keys must be the column names,
|         which must match. If `values` is a DataFrame,
|         then both the index and column labels must match.
|
|     Returns
|     -----
|     DataFrame
|         DataFrame of booleans showing whether each element in the
DataFrame
|         is contained in values.
|

```

```
| See Also
| -----
| DataFrame.eq: Equality test for DataFrame.
| Series.isin: Equivalent method on Series.
| Series.str.contains: Test if pattern or regex is contained within a
| string of a Series or Index.
```

```
| Examples
| -----
```

```
| >>> df = pd.DataFrame({'num_legs': [2, 4], 'num_wings': [2, 0]},
| ...                    index=['falcon', 'dog'])
| >>> df
|          num_legs  num_wings
| falcon          2          2
| dog             4          0
```

```
| When ``values`` is a list check whether every value in the DataFrame
| is present in the list (which animals have 0 or 2 legs or wings)
```

```
| >>> df.isin([0, 2])
|          num_legs  num_wings
| falcon         True         True
| dog            False         True
```

```
| When ``values`` is a dict, we can pass values to check for each
| column separately:
```

```
| >>> df.isin({'num_wings': [0, 3]})
|          num_legs  num_wings
| falcon         False         False
| dog            False          True
```

```
| When ``values`` is a Series or DataFrame the index and column must
| match. Note that 'falcon' does not match based on the number of legs
| in df2.
```

```
| >>> other = pd.DataFrame({'num_legs': [8, 2], 'num_wings': [0, 2]},
| ...                    index=['spider', 'falcon'])
| >>> df.isin(other)
|          num_legs  num_wings
| falcon         True         True
| dog            False         False
```

```

|  isna(self)
|      Detect missing values.
|
|      Return a boolean same-sized object indicating if the values are NA.
|      NA values, such as None or :attr:`numpy.NaN`, gets mapped to True
|      values.
|      Everything else gets mapped to False values. Characters such as empty
|      strings ``''`` or :attr:`numpy.inf` are not considered NA values
|      (unless you set ``pandas.options.mode.use_inf_as_na = True``).
|
|  Returns
|  -----
|  DataFrame
|      Mask of bool values for each element in DataFrame that
|      indicates whether an element is not an NA value.
|
|  See Also
|  -----
|  DataFrame.isnull : Alias of isna.
|  DataFrame.notna : Boolean inverse of isna.
|  DataFrame.dropna : Omit axes labels with missing values.
|  isna : Top-level isna.
|
|  Examples
|  -----
|  Show which entries in a DataFrame are NA.
|
|  >>> df = pd.DataFrame({'age': [5, 6, np.NaN],
|  ...                     'born': [pd.NaT, pd.Timestamp('1939-05-27'),
|  ...                             pd.Timestamp('1940-04-25')],
|  ...                     'name': ['Alfred', 'Batman', ''],
|  ...                     'toy': [None, 'Batmobile', 'Joker']})
|  >>> df
|
|      age      born   name      toy
|  0  5.0         NaT  Alfred     None
|  1  6.0  1939-05-27  Batman  Batmobile
|  2  NaN  1940-04-25         Joker
|
|  >>> df.isna()
|
|      age  born  name  toy
|  0  False  True  False  True
|  1  False  False  False  False
|  2   True  False  False  False
|

```

```

|     Show which entries in a Series are NA.
|
|     >>> ser = pd.Series([5, 6, np.NaN])
|     >>> ser
|     0    5.0
|     1    6.0
|     2    NaN
|     dtype: float64
|
|     >>> ser.isna()
|     0    False
|     1    False
|     2     True
|     dtype: bool
|
|     isnull(self)
|         Detect missing values.
|
|         Return a boolean same-sized object indicating if the values are NA.
|         NA values, such as None or :attr:`numpy.NaN`, gets mapped to True
|         values.
|         Everything else gets mapped to False values. Characters such as empty
|         strings ``''`` or :attr:`numpy.inf` are not considered NA values
|         (unless you set ``pandas.options.mode.use_inf_as_na = True``).
|
|         Returns
|         -----
|         DataFrame
|             Mask of bool values for each element in DataFrame that
|             indicates whether an element is not an NA value.
|
|         See Also
|         -----
|         DataFrame.isnull : Alias of isna.
|         DataFrame.notna : Boolean inverse of isna.
|         DataFrame.dropna : Omit axes labels with missing values.
|         isna : Top-level isna.
|
|         Examples
|         -----
|         Show which entries in a DataFrame are NA.
|
|         >>> df = pd.DataFrame({'age': [5, 6, np.NaN],
|         ...                    'born': [pd.NaT, pd.Timestamp('1939-05-27'),

```

```

|         ...                                pd.Timestamp('1940-04-25')],
|         ...                                'name': ['Alfred', 'Batman', ''],
|         ...                                'toy': [None, 'Batmobile', 'Joker']})
|
| >>> df
|
|      age      born   name      toy
|    0  5.0      NaT  Alfred     None
|    1  6.0 1939-05-27  Batman  Batmobile
|    2  NaN 1940-04-25      Joker
|
|
| >>> df.isna()
|
|      age  born  name  toy
|    0  False  True  False  True
|    1  False False  False  False
|    2   True  False  False  False
|
|
|      Show which entries in a Series are NA.
|
| >>> ser = pd.Series([5, 6, np.NaN])
| >>> ser
|
|    0    5.0
|    1    6.0
|    2    NaN
|    dtype: float64
|
| >>> ser.isna()
|
|    0    False
|    1    False
|    2     True
|    dtype: bool
|
|
|      items(self)
|
|      Iterator over (column name, Series) pairs.
|
|
|      Iterates over the DataFrame columns, returning a tuple with
|      the column name and the content as a Series.
|
|
|      Yields
|      -----
|      label : object
|          The column names for the DataFrame being iterated over.
|      content : Series
|          The column entries belonging to each label, as a Series.
|
|
|      See Also

```

```

| -----
| DataFrame.iterrows : Iterate over DataFrame rows as
|   (index, Series) pairs.
| DataFrame.itertuples : Iterate over DataFrame rows as namedtuples
|   of the values.
|
| Examples
| -----
| >>> df = pd.DataFrame({'species': ['bear', 'bear', 'marsupial'],
| ...                    'population': [1864, 22000, 80000]},
| ...                    index=['panda', 'polar', 'koala'])
| >>> df
|           species  population
| panda    bear      1864
| polar    bear     22000
| koala    marsupial 80000
| >>> for label, content in df.items():
| ...     print('label:', label)
| ...     print('content:', content, sep='\n')
| ...
| label: species
| content:
| panda      bear
| polar      bear
| koala      marsupial
| Name: species, dtype: object
| label: population
| content:
| panda      1864
| polar      22000
| koala      80000
| Name: population, dtype: int64
|
| iteritems(self)
|   Iterator over (column name, Series) pairs.
|
|   Iterates over the DataFrame columns, returning a tuple with
|   the column name and the content as a Series.
|
| Returns
| -----
| label : object
|   The column names for the DataFrame being iterated over.
| content : Series

```

```

|         The column entries belonging to each label, as a Series.
|
|     See Also
|     -----
|     DataFrame.iterrows : Iterate over DataFrame rows as
|         (index, Series) pairs.
|     DataFrame.itertuples : Iterate over DataFrame rows as namedtuples
|         of the values.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'species': ['bear', 'bear', 'marsupial'],
|     ...                     'population': [1864, 22000, 80000]},
|     ...                     index=['panda', 'polar', 'koala'])
|     >>> df
|
|           species  population
|     panda    bear      1864
|     polar    bear     22000
|     koala  marsupial  80000
|     >>> for label, content in df.items():
|     ...     print('label:', label)
|     ...     print('content:', content, sep='\n')
|     ...
|     label: species
|     content:
|     panda      bear
|     polar      bear
|     koala  marsupial
|     Name: species, dtype: object
|     label: population
|     content:
|     panda      1864
|     polar      22000
|     koala      80000
|     Name: population, dtype: int64
|
|     iterrows(self)
|         Iterate over DataFrame rows as (index, Series) pairs.
|
|     Yields
|     -----
|     index : label or tuple of label
|         The index of the row. A tuple for a `MultiIndex`.
|     data : Series

```



```

|         The data of the row as a Series.
|
|         it : generator
|             A generator that iterates over the rows of the frame.
|
|         See Also
|         -----
|         itertuples : Iterate over DataFrame rows as namedtuples of the
values.
|         items : Iterate over (column name, Series) pairs.
|
|         Notes
|         -----
|
|         1. Because ``iterrows`` returns a Series for each row,
|            it does not preserve dtypes across the rows (dtypes are
|            preserved across columns for DataFrames). For example,
|
|            >>> df = pd.DataFrame([[1, 1.5]], columns=['int', 'float'])
|            >>> row = next(df.iterrows())[1]
|            >>> row
|            int      1.0
|            float    1.5
|            Name: 0, dtype: float64
|            >>> print(row['int'].dtype)
|            float64
|            >>> print(df['int'].dtype)
|            int64
|
|            To preserve dtypes while iterating over the rows, it is better
|            to use :meth:`itertuples` which returns namedtuples of the values
|            and which is generally faster than ``iterrows``.
|
|         2. You should never modify something you are iterating over.
|            This is not guaranteed to work in all cases. Depending on the
|            data types, the iterator returns a copy and not a view, and
writing
|            to it will have no effect.
|
|         itertuples(self, index=True, name='Pandas')
|             Iterate over DataFrame rows as namedtuples.
|
|         Parameters
|         -----

```

```

|     index : bool, default True
|         If True, return the index as the first element of the tuple.
|     name : str or None, default "Pandas"
|         The name of the returned namedtuples or None to return regular
|         tuples.
|
|
| Returns
| -----
| iterator
|     An object to iterate over namedtuples for each row in the
|     DataFrame with the first field possibly being the index and
|     following fields being the column values.
|
|
| See Also
| -----
| DataFrame.iterrows : Iterate over DataFrame rows as (index, Series)
|     pairs.
| DataFrame.items : Iterate over (column name, Series) pairs.
|
|
| Notes
| -----
|     The column names will be renamed to positional names if they are
|     invalid Python identifiers, repeated, or start with an underscore.
|     With a large number of columns (>255), regular tuples are returned.
|
|
| Examples
| -----
| >>> df = pd.DataFrame({'num_legs': [4, 2], 'num_wings': [0, 2]},
| ...                     index=['dog', 'hawk'])
| >>> df
|      num_legs  num_wings
|  dog         4         0
|  hawk        2         2
| >>> for row in df.itertuples():
| ...     print(row)
| ...
| Pandas(Index='dog', num_legs=4, num_wings=0)
| Pandas(Index='hawk', num_legs=2, num_wings=2)
|
| By setting the `index` parameter to False we can remove the index
| as the first element of the tuple:
|
| >>> for row in df.itertuples(index=False):
| ...     print(row)

```

```

|     ...
|     Pandas(num_legs=4, num_wings=0)
|     Pandas(num_legs=2, num_wings=2)
|
|     With the `name` parameter set we set a custom name for the yielded
|     namedtuples:
|
|     >>> for row in df.itertuples(name='Animal'):
|     ...     print(row)
|     ...
|     Animal(Index='dog', num_legs=4, num_wings=0)
|     Animal(Index='hawk', num_legs=2, num_wings=2)
|
|     join(self, other, on=None, how='left', lsuffix='', rsuffix='',
|sort=False)
|         Join columns of another DataFrame.
|
|         Join columns with `other` DataFrame either on index or on a key
|         column. Efficiently join multiple DataFrame objects by index at once
by
|         passing a list.
|
|         Parameters
|         -----
|         other : DataFrame, Series, or list of DataFrame
|             Index should be similar to one of the columns in this one. If a
|             Series is passed, its name attribute must be set, and that will
be
|             used as the column name in the resulting joined DataFrame.
|         on : str, list of str, or array-like, optional
|             Column or index level name(s) in the caller to join on the index
|             in `other`, otherwise joins index-on-index. If multiple
|             values given, the `other` DataFrame must have a MultiIndex. Can
|             pass an array as the join key if it is not already contained in
|             the calling DataFrame. Like an Excel VLOOKUP operation.
|         how : {'left', 'right', 'outer', 'inner'}, default 'left'
|             How to handle the operation of the two objects.
|
|         * left: use calling frame's index (or column if on is specified)
|         * right: use `other`'s index.
|         * outer: form union of calling frame's index (or column if on is
|             specified) with `other`'s index, and sort it.
|             lexicographically.
|         * inner: form intersection of calling frame's index (or column if

```

```

|         on is specified) with `other`'s index, preserving the order
|         of the calling's one.
|     lsuffix : str, default ''
|         Suffix to use from left frame's overlapping columns.
|     rsuffix : str, default ''
|         Suffix to use from right frame's overlapping columns.
|     sort : bool, default False
|         Order result DataFrame lexicographically by the join key. If
False,
|         the order of the join key depends on the join type (how keyword).
|
|     Returns
|     -----
|     DataFrame
|         A dataframe containing columns from both the caller and `other`.
|
|     See Also
|     -----
|     DataFrame.merge : For column(s)-on-columns(s) operations.
|
|     Notes
|     -----
|     Parameters `on`, `lsuffix`, and `rsuffix` are not supported when
|     passing a list of `DataFrame` objects.
|
|     Support for specifying index levels as the `on` parameter was added
|     in version 0.23.0.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'key': ['K0', 'K1', 'K2', 'K3', 'K4', 'K5'],
|     ...                    'A': ['A0', 'A1', 'A2', 'A3', 'A4', 'A5']})
|
|     >>> df
|         key  A
|     0  K0  A0
|     1  K1  A1
|     2  K2  A2
|     3  K3  A3
|     4  K4  A4
|     5  K5  A5
|
|     >>> other = pd.DataFrame({'key': ['K0', 'K1', 'K2'],
|     ...                      'B': ['B0', 'B1', 'B2']})

```

```

|
|     >>> other
|         key    B
|     0  K0    B0
|     1  K1    B1
|     2  K2    B2
|

```

Join DataFrames using their indexes.

```

|     >>> df.join(other, lsuffix='_caller', rsuffix='_other')
|         key_caller  A key_other  B
|     0          K0  A0          K0  B0
|     1          K1  A1          K1  B1
|     2          K2  A2          K2  B2
|     3          K3  A3          NaN NaN
|     4          K4  A4          NaN NaN
|     5          K5  A5          NaN NaN
|

```

If we want to join using the key columns, we need to set key to be the index in both `df` and `other`. The joined DataFrame will have key as its index.

```

|     >>> df.set_index('key').join(other.set_index('key'))
|         A    B
|     key
|     K0  A0  B0
|     K1  A1  B1
|     K2  A2  B2
|     K3  A3  NaN
|     K4  A4  NaN
|     K5  A5  NaN
|

```

Another option to join using the key columns is to use the `on` parameter. DataFrame.join always uses `other`'s index but we can use any column in `df`. This method preserves the original DataFrame's index in the result.

```

|     >>> df.join(other.set_index('key'), on='key')
|         key  A    B
|     0  K0  A0  B0
|     1  K1  A1  B1
|     2  K2  A2  B2
|     3  K3  A3  NaN
|     4  K4  A4  NaN
|

```

```

|      5  K5  A5  NaN
|
|      kurt(self, axis=None, skipna=None, level=None, numeric_only=None,
**kwargs)
|
|      Return unbiased kurtosis over requested axis using Fisher's
definition of
|
|      kurtosis (kurtosis of normal == 0.0). Normalized by N-1.
|
|
|      Parameters
|      -----
|
|      axis : {index (0), columns (1)}
|
|      Axis for the function to be applied on.
|
|      skipna : bool, default True
|
|      Exclude NA/null values when computing the result.
|
|      level : int or level name, default None
|
|      If the axis is a MultiIndex (hierarchical), count along a
|
|      particular level, collapsing into a Series.
|
|      numeric_only : bool, default None
|
|      Include only float, int, boolean columns. If None, will attempt
to use
|
|      everything, then use only numeric data. Not implemented for
Series.
|
|      **kwargs
|
|      Additional keyword arguments to be passed to the function.
|
|
|      Returns
|      -----
|
|      Series or DataFrame (if level specified)
|
|
|      kurtosis = kurt(self, axis=None, skipna=None, level=None,
numeric_only=None, **kwargs)
|
|
|      le(self, other, axis='columns', level=None)
|
|      Get Less than or equal to of dataframe and other, element-wise
(binary operator `le`).
|
|
|      Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to
comparison
|
|      operators.
|
|
|      Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose
axis
|
|      (rows or columns) and level for comparison.
|

```

```

|     Parameters
|     -----
|     other : scalar, sequence, Series, or DataFrame
|         Any single or multiple element data structure, or list-like
object.
|
|     axis : {0 or 'index', 1 or 'columns'}, default 'columns'
|         Whether to compare by the index (0 or 'index') or columns
|         (1 or 'columns').
|     level : int or label
|         Broadcast across a level, matching Index values on the passed
|         MultiIndex level.
|
|
|     Returns
|     -----
|     DataFrame of bool
|         Result of the comparison.
|
|
|     See Also
|     -----
|     DataFrame.eq : Compare DataFrames for equality elementwise.
|     DataFrame.ne : Compare DataFrames for inequality elementwise.
|     DataFrame.le : Compare DataFrames for less than inequality
|         or equality elementwise.
|     DataFrame.lt : Compare DataFrames for strictly less than
|         inequality elementwise.
|     DataFrame.ge : Compare DataFrames for greater than inequality
|         or equality elementwise.
|     DataFrame.gt : Compare DataFrames for strictly greater than
|         inequality elementwise.
|
|
|     Notes
|     -----
|     Mismatched indices will be unioned together.
|     `NaN` values are considered different (i.e. `NaN` != `NaN`).
|
|
|     Examples
|     -----
|
|     >>> df = pd.DataFrame({'cost': [250, 150, 100],
|     ...                   'revenue': [100, 250, 300]},
|     ...                   index=['A', 'B', 'C'])
|     >>> df
|
|         cost  revenue
|     A    250     100
|     B    150     250

```

```

|      C    100    300
|
|      Comparison with a scalar, using either the operator or method:
|
|      >>> df == 100
|           cost  revenue
|      A  False    True
|      B  False    False
|      C   True    False
|
|      >>> df.eq(100)
|           cost  revenue
|      A  False    True
|      B  False    False
|      C   True    False
|
|      When `other` is a :class:`Series`, the columns of a DataFrame are
aligned
|      with the index of `other` and broadcast:
|
|      >>> df != pd.Series([100, 250], index=["cost", "revenue"])
|           cost  revenue
|      A   True    True
|      B   True    False
|      C  False    True
|
|      Use the method to control the broadcast axis:
|
|      >>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
|           cost  revenue
|      A   True    False
|      B   True    True
|      C   True    True
|      D   True    True
|
|      When comparing to an arbitrary sequence, the number of columns must
|      match the number elements in `other`:
|
|      >>> df == [250, 100]
|           cost  revenue
|      A   True    True
|      B  False    False
|      C  False    False
|

```



```

|         Use the method to control the axis:
|
|         >>> df.eq([250, 250, 100], axis='index')
|             cost  revenue
|         A   True   False
|         B  False    True
|         C   True   False
|
|         Compare to a DataFrame of different shape.
|
|         >>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
|         ...                       index=['A', 'B', 'C', 'D'])
|         >>> other
|             revenue
|         A         300
|         B         250
|         C         100
|         D         150
|
|         >>> df.gt(other)
|             cost  revenue
|         A  False   False
|         B  False   False
|         C  False    True
|         D  False   False
|
|         Compare to a MultiIndex by level.
|
|         >>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300,
220],
|         ...                               'revenue': [100, 250, 300, 200, 175,
225]},
|         ...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2',
'Q2'],
|         ...                               ['A', 'B', 'C', 'A', 'B',
'C']])
|         >>> df_multindex
|             cost  revenue
|         Q1 A    250     100
|             B    150     250
|             C    100     300
|         Q2 A    150     200
|             B    300     175
|             C    220     225

```

```

|
|     >>> df.ile(df_multindex, level=1)
|
|         cost  revenue
|     Q1 A   True     True
|         B   True     True
|         C   True     True
|     Q2 A  False     True
|         B   True    False
|         C   True    False
|
|
|     lookup(self, row_labels, col_labels)
|         Label-based "fancy indexing" function for DataFrame.
|
|         Given equal-length arrays of row and column labels, return an
|         array of the values corresponding to each (row, col) pair.
|
|     Parameters
|     -----
|     row_labels : sequence
|         The row labels to use for lookup
|     col_labels : sequence
|         The column labels to use for lookup
|
|     Returns
|     -----
|     numpy.ndarray
|
|     Notes
|     -----
|     Akin to::
|
|         result = [df.get_value(row, col)
|                   for row, col in zip(row_labels, col_labels)]
|
|     Examples
|     -----
|     values : ndarray
|         The found values
|
|     lt(self, other, axis='columns', level=None)
|         Get Less than of dataframe and other, element-wise (binary operator
|         `lt`).
|
|

```

```

|     Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to
comparison
|     operators.
|
|     Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose
axis
|     (rows or columns) and level for comparison.
|
|     Parameters
|     -----
|     other : scalar, sequence, Series, or DataFrame
|             Any single or multiple element data structure, or list-like
object.
|
|     axis : {0 or 'index', 1 or 'columns'}, default 'columns'
|             Whether to compare by the index (0 or 'index') or columns
|             (1 or 'columns').
|     level : int or label
|             Broadcast across a level, matching Index values on the passed
|             MultiIndex level.
|
|     Returns
|     -----
|     DataFrame of bool
|             Result of the comparison.
|
|     See Also
|     -----
|     DataFrame.eq : Compare DataFrames for equality elementwise.
|     DataFrame.ne : Compare DataFrames for inequality elementwise.
|     DataFrame.le : Compare DataFrames for less than inequality
|                   or equality elementwise.
|     DataFrame.lt : Compare DataFrames for strictly less than
|                   inequality elementwise.
|     DataFrame.ge : Compare DataFrames for greater than inequality
|                   or equality elementwise.
|     DataFrame.gt : Compare DataFrames for strictly greater than
|                   inequality elementwise.
|
|     Notes
|     -----
|     Mismatched indices will be unioned together.
|     `NaN` values are considered different (i.e. `NaN` != `NaN`).
|
|     Examples

```

```
|
| -----
| >>> df = pd.DataFrame({'cost': [250, 150, 100],
| ...                     'revenue': [100, 250, 300]},
| ...                     index=['A', 'B', 'C'])
| >>> df
|      cost  revenue
| A    250     100
| B    150     250
| C    100     300
|
| Comparison with a scalar, using either the operator or method:
|
| >>> df == 100
|      cost  revenue
| A  False     True
| B  False     False
| C   True     False
|
| >>> df.eq(100)
|      cost  revenue
| A  False     True
| B  False     False
| C   True     False
|
| When `other` is a :class:`Series`, the columns of a DataFrame are
aligned
| with the index of `other` and broadcast:
|
| >>> df != pd.Series([100, 250], index=["cost", "revenue"])
|      cost  revenue
| A   True     True
| B   True     False
| C  False     True
|
| Use the method to control the broadcast axis:
|
| >>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
|      cost  revenue
| A   True     False
| B   True     True
| C   True     True
| D   True     True
|
| When comparing to an arbitrary sequence, the number of columns must
```

```

|         match the number elements in `other`:
|
|         >>> df == [250, 100]
|             cost  revenue
|         A   True    True
|         B  False   False
|         C  False   False
|
|         Use the method to control the axis:
|
|         >>> df.eq([250, 250, 100], axis='index')
|             cost  revenue
|         A   True    False
|         B  False    True
|         C   True    False
|
|         Compare to a DataFrame of different shape.
|
|         >>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
|         ...                       index=['A', 'B', 'C', 'D'])
|         >>> other
|             revenue
|         A         300
|         B         250
|         C         100
|         D         150
|
|         >>> df.gt(other)
|             cost  revenue
|         A  False   False
|         B  False   False
|         C  False    True
|         D  False   False
|
|         Compare to a MultiIndex by level.
|
|         >>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300,
220],
|         ...                               'revenue': [100, 250, 300, 200, 175,
225]},
|         ...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2',
'Q2'],
|         ...                               ['A', 'B', 'C', 'A', 'B',
'C']])

```

```

|     >>> df_multindex
|           cost  revenue
|   Q1 A    250    100
|       B    150    250
|       C    100    300
|   Q2 A    150    200
|       B    300    175
|       C    220    225
|
|     >>> df.le(df_multindex, level=1)
|           cost  revenue
|   Q1 A   True   True
|       B   True   True
|       C   True   True
|   Q2 A  False   True
|       B   True  False
|       C   True  False
|
|   mad(self, axis=None, skipna=None, level=None)
|       Return the mean absolute deviation of the values for the requested
axis.
|
|   Parameters
|   -----
|   axis : {index (0), columns (1)}
|       Axis for the function to be applied on.
|   skipna : bool, default True
|       Exclude NA/null values when computing the result.
|   level : int or level name, default None
|       If the axis is a MultiIndex (hierarchical), count along a
|       particular level, collapsing into a Series.
|   numeric_only : bool, default None
|       Include only float, int, boolean columns. If None, will attempt
to use
|       everything, then use only numeric data. Not implemented for
Series.
|   **kwargs
|       Additional keyword arguments to be passed to the function.
|
|   Returns
|   -----
|   Series or DataFrame (if level specified)
|

```

```

| max(self, axis=None, skipna=None, level=None, numeric_only=None,
**kwargs)
|     Return the maximum of the values for the requested axis.
|
|     If you want the *index* of the maximum, use ``idxmax``.
This is
|     the equivalent of the ``numpy.ndarray`` method
``argmax``.
|
| Parameters
| -----
| axis : {index (0), columns (1)}
|     Axis for the function to be applied on.
| skipna : bool, default True
|     Exclude NA/null values when computing the result.
| level : int or level name, default None
|     If the axis is a MultiIndex (hierarchical), count along a
|     particular level, collapsing into a Series.
| numeric_only : bool, default None
|     Include only float, int, boolean columns. If None, will attempt
to use
|     everything, then use only numeric data. Not implemented for
Series.
| **kwargs
|     Additional keyword arguments to be passed to the function.
|
| Returns
| -----
| Series or DataFrame (if level specified)
|
| See Also
| -----
| Series.sum : Return the sum.
| Series.min : Return the minimum.
| Series.max : Return the maximum.
| Series.idxmin : Return the index of the minimum.
| Series.idxmax : Return the index of the maximum.
| DataFrame.sum : Return the sum over the requested axis.
| DataFrame.min : Return the minimum over the requested axis.
| DataFrame.max : Return the maximum over the requested axis.
| DataFrame.idxmin : Return the index of the minimum over the requested
axis.
| DataFrame.idxmax : Return the index of the maximum over the requested
axis.

```

```

|
| Examples
| -----
|
| >>> idx = pd.MultiIndex.from_arrays([
| ...     ['warm', 'warm', 'cold', 'cold'],
| ...     ['dog', 'falcon', 'fish', 'spider']],
| ...     names=['blooded', 'animal'])
| >>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
| >>> s
| blooded  animal
| warm     dog      4
|           falcon   2
| cold     fish      0
|           spider   8
| Name: legs, dtype: int64
|
| >>> s.max()
| 8
|
| Max using level names, as well as indices.
|
| >>> s.max(level='blooded')
| blooded
| warm     4
| cold     8
| Name: legs, dtype: int64
|
| >>> s.max(level=0)
| blooded
| warm     4
| cold     8
| Name: legs, dtype: int64
|
| mean(self, axis=None, skipna=None, level=None, numeric_only=None,
**kwargs)
| Return the mean of the values for the requested axis.
|
| Parameters
| -----
| axis : {index (0), columns (1)}
|     Axis for the function to be applied on.
| skipna : bool, default True
|     Exclude NA/null values when computing the result.
| level : int or level name, default None

```



```

|         If the axis is a MultiIndex (hierarchical), count along a
|         particular level, collapsing into a Series.
|         numeric_only : bool, default None
|         Include only float, int, boolean columns. If None, will attempt
to use
|         everything, then use only numeric data. Not implemented for
Series.
|         **kwargs
|         Additional keyword arguments to be passed to the function.
|
|     Returns
|     -----
|     Series or DataFrame (if level specified)
|
|     median(self, axis=None, skipna=None, level=None, numeric_only=None,
**kwargs)
|     Return the median of the values for the requested axis.
|
|     Parameters
|     -----
|     axis : {index (0), columns (1)}
|     Axis for the function to be applied on.
|     skipna : bool, default True
|     Exclude NA/null values when computing the result.
|     level : int or level name, default None
|     If the axis is a MultiIndex (hierarchical), count along a
|     particular level, collapsing into a Series.
|     numeric_only : bool, default None
|     Include only float, int, boolean columns. If None, will attempt
to use
|     everything, then use only numeric data. Not implemented for
Series.
|     **kwargs
|     Additional keyword arguments to be passed to the function.
|
|     Returns
|     -----
|     Series or DataFrame (if level specified)
|
|     melt(self, id_vars=None, value_vars=None, var_name=None,
value_name='value', col_level=None)
|     Unpivot a DataFrame from wide format to long format, optionally
|     leaving identifier variables set.
|

```

```

|     This function is useful to massage a DataFrame into a format where
one |
|     or more columns are identifier variables (`id_vars`), while all other
|     columns, considered measured variables (`value_vars`), are
"unpivoted" to
|     the row axis, leaving just two non-identifier columns, 'variable' and
|     'value'.
|     .. versionadded:: 0.20.0
|
| Parameters
| -----
| frame : DataFrame
| id_vars : tuple, list, or ndarray, optional
|     Column(s) to use as identifier variables.
| value_vars : tuple, list, or ndarray, optional
|     Column(s) to unpivot. If not specified, uses all columns that
|     are not set as `id_vars`.
| var_name : scalar
|     Name to use for the 'variable' column. If None it uses
|     ``frame.columns.name`` or 'variable'.
| value_name : scalar, default 'value'
|     Name to use for the 'value' column.
| col_level : int or string, optional
|     If columns are a MultiIndex then use this level to melt.
|
| Returns
| -----
| DataFrame
|     Unpivoted DataFrame.
|
| See Also
| -----
| melt
| pivot_table
| DataFrame.pivot
| Series.explode
|
| Examples
| -----
| >>> df = pd.DataFrame({'A': {0: 'a', 1: 'b', 2: 'c'},
| ...                   'B': {0: 1, 1: 3, 2: 5},
| ...                   'C': {0: 2, 1: 4, 2: 6}})
| >>> df
|      A  B  C

```

```
|      0  a  1  2
|      1  b  3  4
|      2  c  5  6
```

```
|      >>> df.melt(id_vars=['A'], value_vars=['B'])
```

```
|      A variable  value
|      0  a      B      1
|      1  b      B      3
|      2  c      B      5
```

```
|      >>> df.melt(id_vars=['A'], value_vars=['B', 'C'])
```

```
|      A variable  value
|      0  a      B      1
|      1  b      B      3
|      2  c      B      5
|      3  a      C      2
|      4  b      C      4
|      5  c      C      6
```

```
|      The names of 'variable' and 'value' columns can be customized:
```

```
|      >>> df.melt(id_vars=['A'], value_vars=['B'],
|      ...      var_name='myVarname', value_name='myValname')
```

```
|      A myVarname  myValname
|      0  a      B      1
|      1  b      B      3
|      2  c      B      5
```

```
|      If you have multi-index columns:
```

```
|      >>> df.columns = [list('ABC'), list('DEF')]
|      >>> df
```

```
|      A  B  C
|      D  E  F
|      0  a  1  2
|      1  b  3  4
|      2  c  5  6
```

```
|      >>> df.melt(col_level=0, id_vars=['A'], value_vars=['B'])
```

```
|      A variable  value
|      0  a      B      1
|      1  b      B      3
|      2  c      B      5
```

```

|     >>> df.melt(id_vars=[('A', 'D')], value_vars=[('B', 'E')])
|           (A, D) variable_0 variable_1  value
|     0      a           B           E      1
|     1      b           B           E      3
|     2      c           B           E      5
|
|
|     memory_usage(self, index=True, deep=False)
|           Return the memory usage of each column in bytes.
|
|           The memory usage can optionally include the contribution of
|           the index and elements of `object` dtype.
|
|           This value is displayed in `DataFrame.info` by default. This can be
|           suppressed by setting ``pandas.options.display.memory_usage`` to
False.
|
|     Parameters
|     -----
|     index : bool, default True
|           Specifies whether to include the memory usage of the DataFrame's
|           index in returned Series. If ``index=True``, the memory usage of
|           the index is the first item in the output.
|     deep : bool, default False
|           If True, introspect the data deeply by interrogating
|           `object` dtypes for system-level memory consumption, and include
|           it in the returned values.
|
|     Returns
|     -----
|     Series
|           A Series whose index is the original column names and whose
values
|           is the memory usage of each column in bytes.
|
|     See Also
|     -----
|     numpy.ndarray.nbytes : Total bytes consumed by the elements of an
|     ndarray.
|     Series.memory_usage : Bytes consumed by a Series.
|     Categorical : Memory-efficient array for string values with
|     many repeated values.
|     DataFrame.info : Concise summary of a DataFrame.
|
|     Examples

```

```
| -----
| >>> dtypes = ['int64', 'float64', 'complex128', 'object', 'bool']
| >>> data = dict([(t, np.ones(shape=5000).astype(t))
| ...             for t in dtypes])
| >>> df = pd.DataFrame(data)
| >>> df.head()
|      int64  float64      complex128  object  bool
|  0      1      1.0  1.000000+0.000000j      1  True
|  1      1      1.0  1.000000+0.000000j      1  True
|  2      1      1.0  1.000000+0.000000j      1  True
|  3      1      1.0  1.000000+0.000000j      1  True
|  4      1      1.0  1.000000+0.000000j      1  True
|
| >>> df.memory_usage()
|      Index      128
|      int64     40000
|      float64    40000
|      complex128  80000
|      object     40000
|      bool       5000
|      dtype: int64
|
| >>> df.memory_usage(index=False)
|      int64     40000
|      float64    40000
|      complex128  80000
|      object     40000
|      bool       5000
|      dtype: int64
|
| The memory footprint of `object` dtype columns is ignored by default:
|
| >>> df.memory_usage(deep=True)
|      Index      128
|      int64     40000
|      float64    40000
|      complex128  80000
|      object    160000
|      bool       5000
|      dtype: int64
|
| Use a Categorical for efficient storage of an object-dtype column
with
| many repeated values.
```

```

|
|     >>> df['object'].astype('category').memory_usage(deep=True)
|     5216
|
|
|     merge(self, right, how='inner', on=None, left_on=None, right_on=None,
left_index=False, right_index=False, sort=False, suffixes=('_x', '_y'),
copy=True, indicator=False, validate=None)
|         Merge DataFrame or named Series objects with a database-style join.
|
|         The join is done on columns or indexes. If joining columns on
|         columns, the DataFrame indexes *will be ignored*. Otherwise if
joining indexes
|         on indexes or indexes on a column or columns, the index will be
passed on.
|
|         Parameters
|         -----
|         right : DataFrame or named Series
|             Object to merge with.
|         how : {'left', 'right', 'outer', 'inner'}, default 'inner'
|             Type of merge to be performed.
|
|         * left: use only keys from left frame, similar to a SQL left
outer join;
|             preserve key order.
|         * right: use only keys from right frame, similar to a SQL right
outer join;
|             preserve key order.
|         * outer: use union of keys from both frames, similar to a SQL
full outer
|             join; sort keys lexicographically.
|         * inner: use intersection of keys from both frames, similar to a
SQL inner
|             join; preserve the order of the left keys.
|         on : label or list
|             Column or index level names to join on. These must be found in
both
|             DataFrames. If `on` is None and not merging on indexes then this
defaults
|             to the intersection of the columns in both DataFrames.
|         left_on : label or list, or array-like
|             Column or index level names to join on in the left DataFrame. Can
also

```

```

|         be an array or list of arrays of the length of the left
DataFrame.
|         These arrays are treated as if they are columns.
|         right_on : label or list, or array-like
|         Column or index level names to join on in the right DataFrame.
Can also
|         be an array or list of arrays of the length of the right
DataFrame.
|         These arrays are treated as if they are columns.
|         left_index : bool, default False
|         Use the index from the left DataFrame as the join key(s). If it
is a
|         MultiIndex, the number of keys in the other DataFrame (either the
index
|         or a number of columns) must match the number of levels.
|         right_index : bool, default False
|         Use the index from the right DataFrame as the join key. Same
caveats as
|         left_index.
|         sort : bool, default False
|         Sort the join keys lexicographically in the result DataFrame. If
False,
|         the order of the join keys depends on the join type (how
keyword).
|         suffixes : tuple of (str, str), default ('_x', '_y')
|         Suffix to apply to overlapping column names in the left and right
|         side, respectively. To raise an exception on overlapping columns
use
|         (False, False).
|         copy : bool, default True
|         If False, avoid copy if possible.
|         indicator : bool or str, default False
|         If True, adds a column to output DataFrame called "_merge" with
|         information on the source of each row.
|         If string, column with information on source of each row will be
added to
|         output DataFrame, and column will be named value of string.
|         Information column is Categorical-type and takes on a value of
"left_only"
|         for observations whose merge key only appears in 'left'
DataFrame,
|         "right_only" for observations whose merge key only appears in
'right'

```

```

|         DataFrame, and "both" if the observation's merge key is found in
both.
|
|         validate : str, optional
|             If specified, checks if merge is of specified type.
|
|         * "one_to_one" or "1:1": check if merge keys are unique in both
|             left and right datasets.
|         * "one_to_many" or "1:m": check if merge keys are unique in left
|             dataset.
|         * "many_to_one" or "m:1": check if merge keys are unique in right
|             dataset.
|         * "many_to_many" or "m:m": allowed, but does not result in
checks.
|
|         .. versionadded:: 0.21.0
|
|     Returns
|     -----
|     DataFrame
|         A DataFrame of the two merged objects.
|
|     See Also
|     -----
|     merge_ordered : Merge with optional filling/interpolation.
|     merge_asof : Merge on nearest keys.
|     DataFrame.join : Similar method using indices.
|
|     Notes
|     -----
|     Support for specifying index levels as the `on`, `left_on`, and
|     `right_on` parameters was added in version 0.23.0
|     Support for merging named Series objects was added in version 0.24.0
|
|     Examples
|     -----
|
|     >>> df1 = pd.DataFrame({'lkey': ['foo', 'bar', 'baz', 'foo'],
|     ...                     'value': [1, 2, 3, 5]})
|     >>> df2 = pd.DataFrame({'rkey': ['foo', 'bar', 'baz', 'foo'],
|     ...                     'value': [5, 6, 7, 8]})
|     >>> df1
|         lkey value
|     0    foo     1

```



```

|      1  bar      2
|      2  baz      3
|      3  foo      5
|      >>> df2
|
|      rkey value
|      0  foo      5
|      1  bar      6
|      2  baz      7
|      3  foo      8
|
|
|      Merge df1 and df2 on the lkey and rkey columns. The value columns
have
|      the default suffixes, _x and _y, appended.
|
|      >>> df1.merge(df2, left_on='lkey', right_on='rkey')
|      lkey  value_x rkey  value_y
|      0  foo      1  foo      5
|      1  foo      1  foo      8
|      2  foo      5  foo      5
|      3  foo      5  foo      8
|      4  bar      2  bar      6
|      5  baz      3  baz      7
|
|
|      Merge DataFrames df1 and df2 with specified left and right suffixes
|      appended to any overlapping columns.
|
|      >>> df1.merge(df2, left_on='lkey', right_on='rkey',
|      ...           suffixes=('_left', '_right'))
|      lkey  value_left rkey  value_right
|      0  foo      1  foo      5
|      1  foo      1  foo      8
|      2  foo      5  foo      5
|      3  foo      5  foo      8
|      4  bar      2  bar      6
|      5  baz      3  baz      7
|
|
|      Merge DataFrames df1 and df2, but raise an exception if the
DataFrames have
|      any overlapping columns.
|
|      >>> df1.merge(df2, left_on='lkey', right_on='rkey', suffixes=(False,
False))
|      Traceback (most recent call last):
|      ...

```

```

|         ValueError: columns overlap but no suffix specified:
|         Index(['value'], dtype='object')
|
| min(self, axis=None, skipna=None, level=None, numeric_only=None,
**kwargs)
|         Return the minimum of the values for the requested axis.
|
|         If you want the *index* of the minimum, use ``idxmin``.
This is
|         the equivalent of the ``numpy.ndarray`` method
``argmin``.
|
| Parameters
| -----
| axis : {index (0), columns (1)}
|         Axis for the function to be applied on.
| skipna : bool, default True
|         Exclude NA/null values when computing the result.
| level : int or level name, default None
|         If the axis is a MultiIndex (hierarchical), count along a
|         particular level, collapsing into a Series.
| numeric_only : bool, default None
|         Include only float, int, boolean columns. If None, will attempt
to use
|         everything, then use only numeric data. Not implemented for
Series.
| **kwargs
|         Additional keyword arguments to be passed to the function.
|
| Returns
| -----
| Series or DataFrame (if level specified)
|
| See Also
| -----
| Series.sum : Return the sum.
| Series.min : Return the minimum.
| Series.max : Return the maximum.
| Series.idxmin : Return the index of the minimum.
| Series.idxmax : Return the index of the maximum.
| DataFrame.sum : Return the sum over the requested axis.
| DataFrame.min : Return the minimum over the requested axis.
| DataFrame.max : Return the maximum over the requested axis.

```

```

|     DataFrame.idxmin : Return the index of the minimum over the requested
axis.
|     DataFrame.idxmax : Return the index of the maximum over the requested
axis.
|
|     Examples
|     -----
|     >>> idx = pd.MultiIndex.from_arrays([
|         ['warm', 'warm', 'cold', 'cold'],
|         ['dog', 'falcon', 'fish', 'spider']],
|         names=['blooded', 'animal'])
|     >>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
|     >>> s
|
| blooded  animal
| warm     dog      4
|          falcon   2
| cold     fish     0
|          spider   8
| Name: legs, dtype: int64
|
|     >>> s.min()
|
| 0
|
|     Min using level names, as well as indices.
|
|     >>> s.min(level='blooded')
|
| blooded
| warm     2
| cold     0
| Name: legs, dtype: int64
|
|     >>> s.min(level=0)
|
| blooded
| warm     2
| cold     0
| Name: legs, dtype: int64
|
|     mod(self, other, axis='columns', level=None, fill_value=None)
|     Get Modulo of dataframe and other, element-wise (binary operator
`mod`).
|
|     Equivalent to ``dataframe % other``, but with support to substitute a
fill_value
|     for missing data in one of the inputs. With reverse version, `rmod`.

```

```

|
|   Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
|   arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
|   Parameters
|   -----
|
|   other : scalar, sequence, Series, or DataFrame
|           Any single or multiple element data structure, or list-like
object.
|   axis : {0 or 'index', 1 or 'columns'}
|           Whether to compare by the index (0 or 'index') or columns
|           (1 or 'columns'). For Series input, axis to match Series index
on.
|   level : int or label
|           Broadcast across a level, matching Index values on the
|           passed MultiIndex level.
|   fill_value : float or None, default None
|           Fill existing missing (NaN) values, and any new element needed
for
|           successful DataFrame alignment, with this value before
computation.
|           If data in both corresponding DataFrame locations is missing
|           the result will be missing.
|
|   Returns
|   -----
|
|   DataFrame
|           Result of the arithmetic operation.
|
|   See Also
|   -----
|
|   DataFrame.add : Add DataFrames.
|   DataFrame.sub : Subtract DataFrames.
|   DataFrame.mul : Multiply DataFrames.
|   DataFrame.div : Divide DataFrames (float division).
|   DataFrame.truediv : Divide DataFrames (float division).
|   DataFrame.floordiv : Divide DataFrames (integer division).
|   DataFrame.mod : Calculate modulo (remainder after division).
|   DataFrame.pow : Calculate exponential power.
|
|   Notes
|   -----
|
|   Mismatched indices will be unioned together.
|

```

```

|   Examples
|   -----
|   >>> df = pd.DataFrame({'angles': [0, 3, 4],
|   ...                     'degrees': [360, 180, 360]},
|   ...                     index=['circle', 'triangle', 'rectangle'])
|   >>> df
|
|           angles  degrees
|   circle         0      360
|   triangle        3      180
|   rectangle       4      360
|
|   Add a scalar with operator version which return the same
|   results.
|
|   >>> df + 1
|
|           angles  degrees
|   circle         1      361
|   triangle        4      181
|   rectangle       5      361
|
|   >>> df.add(1)
|
|           angles  degrees
|   circle         1      361
|   triangle        4      181
|   rectangle       5      361
|
|   Divide by constant with reverse version.
|
|   >>> df.div(10)
|
|           angles  degrees
|   circle        0.0      36.0
|   triangle       0.3      18.0
|   rectangle      0.4      36.0
|
|   >>> df.rdiv(10)
|
|           angles  degrees
|   circle         inf  0.027778
|   triangle  3.333333  0.055556
|   rectangle  2.500000  0.027778
|
|   Subtract a list and Series by axis with operator version.
|
|   >>> df - [1, 2]
|
|           angles  degrees

```

```

|      circle      -1      358
|      triangle      2      178
|      rectangle      3      358
|
|      >>> df.sub([1, 2], axis='columns')
|              angles  degrees
|      circle      -1      358
|      triangle      2      178
|      rectangle      3      358
|
|      >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
|      'rectangle'])),
|      ...      axis='index')
|              angles  degrees
|      circle      -1      359
|      triangle      2      179
|      rectangle      3      359
|
|      Multiply a DataFrame of different shape with operator version.
|
|      >>> other = pd.DataFrame({'angles': [0, 3, 4]},
|      ...      index=['circle', 'triangle', 'rectangle'])
|      >>> other
|              angles
|      circle      0
|      triangle      3
|      rectangle      4
|
|      >>> df * other
|              angles  degrees
|      circle      0      NaN
|      triangle      9      NaN
|      rectangle     16      NaN
|
|      >>> df.mul(other, fill_value=0)
|              angles  degrees
|      circle      0      0.0
|      triangle      9      0.0
|      rectangle     16      0.0
|
|      Divide by a MultiIndex by level.
|
|      >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],

```

```

|         ...                                'degrees': [360, 180, 360, 360, 540,
720]}},
|         ...                                index=[['A', 'A', 'A', 'B', 'B',
'B'],
|         ...                                ['circle', 'triangle',
'rectangle',
|         ...                                'square', 'pentagon',
'hexagon']])
|
|     >>> df_multindex
|
|           angles  degrees
|   A circle      0      360
|   triangle      3      180
|   rectangle      4      360
|   B square      4      360
|   pentagon      5      540
|   hexagon       6      720
|
|     >>> df.div(df_multindex, level=1, fill_value=0)
|
|           angles  degrees
|   A circle      NaN      1.0
|   triangle      1.0      1.0
|   rectangle      1.0      1.0
|   B square      0.0      0.0
|   pentagon      0.0      0.0
|   hexagon       0.0      0.0
|
|
|   mode(self, axis=0, numeric_only=False, dropna=True)
|       Get the mode(s) of each element along the selected axis.
|
|       The mode of a set of values is the value that appears most often.
|       It can be multiple values.
|
|       Parameters
|       -----
|       axis : {0 or 'index', 1 or 'columns'}, default 0
|           The axis to iterate over while searching for the mode:
|
|           * 0 or 'index' : get mode of each column
|           * 1 or 'columns' : get mode of each row
|       numeric_only : bool, default False
|           If True, only apply to numeric columns.
|       dropna : bool, default True
|           Don't consider counts of NaN/NaT.
|

```

```

|         .. versionadded:: 0.24.0
|
| Returns
| -----
| DataFrame
|     The modes of each column or row.
|
| See Also
| -----
| Series.mode : Return the highest frequency value in a Series.
| Series.value_counts : Return the counts of values in a Series.
|
| Examples
| -----
| >>> df = pd.DataFrame([('bird', 2, 2),
| ...                     ('mammal', 4, np.nan),
| ...                     ('arthropod', 8, 0),
| ...                     ('bird', 2, np.nan)],
| ...                     index=('falcon', 'horse', 'spider', 'ostrich'),
| ...                     columns=('species', 'legs', 'wings'))
| >>> df
|
|          species  legs  wings
| falcon      bird     2    2.0
| horse      mammal     4    NaN
| spider  arthropod     8    0.0
| ostrich     bird     2    NaN
|
| By default, missing values are not considered, and the mode of wings
| are both 0 and 2. The second row of species and legs contains
| ``NaN``,
| because they have only one mode, but the DataFrame has two rows.
|
| >>> df.mode()
|
|   species  legs  wings
| 0   bird    2.0    0.0
| 1   NaN    NaN    2.0
|
| Setting ``dropna=False`` ``NaN`` values are considered and they can
be
| the mode (like for wings).
|
| >>> df.mode(dropna=False)
|
|   species  legs  wings
| 0   bird     2    NaN

```



```

|
|     Setting ``numeric_only=True``, only the mode of numeric columns is
|     computed, and columns of other types are ignored.
|
|     >>> df.mode(numeric_only=True)
|          legs  wings
|     0     2.0    0.0
|     1    NaN    2.0
|
|     To compute the mode over columns and not rows, use the axis
parameter:
|
|     >>> df.mode(axis='columns', numeric_only=True)
|           0     1
|     falcon  2.0  NaN
|     horse   4.0  NaN
|     spider  0.0  8.0
|     ostrich 2.0  NaN
|
|     mul(self, other, axis='columns', level=None, fill_value=None)
|         Get Multiplication of dataframe and other, element-wise (binary
operator `mul`).
|
|         Equivalent to ``dataframe * other``, but with support to substitute a
fill_value
|         for missing data in one of the inputs. With reverse version, `rmul`.
|
|         Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
|         arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
|     Parameters
|     -----
|     other : scalar, sequence, Series, or DataFrame
|         Any single or multiple element data structure, or list-like
object.
|     axis : {0 or 'index', 1 or 'columns'}
|         Whether to compare by the index (0 or 'index') or columns
|         (1 or 'columns'). For Series input, axis to match Series index
on.
|     level : int or label
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level.
|     fill_value : float or None, default None

```

```

|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing.
|
|     Returns
|     -----
|     DataFrame
|         Result of the arithmetic operation.
|
|     See Also
|     -----
|     DataFrame.add : Add DataFrames.
|     DataFrame.sub : Subtract DataFrames.
|     DataFrame.mul : Multiply DataFrames.
|     DataFrame.div : Divide DataFrames (float division).
|     DataFrame.truediv : Divide DataFrames (float division).
|     DataFrame.floordiv : Divide DataFrames (integer division).
|     DataFrame.mod : Calculate modulo (remainder after division).
|     DataFrame.pow : Calculate exponential power.
|
|     Notes
|     -----
|     Mismatched indices will be unioned together.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'angles': [0, 3, 4],
|     ...                    'degrees': [360, 180, 360]},
|     ...                    index=['circle', 'triangle', 'rectangle'])
|     >>> df
|
|           angles  degrees
|     circle         0     360
|     triangle        3     180
|     rectangle       4     360
|
|     Add a scalar with operator version which return the same
|     results.
|
|     >>> df + 1
|
|           angles  degrees
|     circle         1     361

```

```

|     triangle      4      181
|     rectangle     5      361
|
|     >>> df.add(1)
|           angles  degrees
|     circle      1      361
|     triangle     4      181
|     rectangle     5      361
|
|     Divide by constant with reverse version.
|
|     >>> df.div(10)
|           angles  degrees
|     circle      0.0     36.0
|     triangle     0.3     18.0
|     rectangle     0.4     36.0
|
|     >>> df.rdiv(10)
|           angles  degrees
|     circle      inf  0.027778
|     triangle  3.333333  0.055556
|     rectangle  2.500000  0.027778
|
|     Subtract a list and Series by axis with operator version.
|
|     >>> df - [1, 2]
|           angles  degrees
|     circle      -1     358
|     triangle     2     178
|     rectangle     3     358
|
|     >>> df.sub([1, 2], axis='columns')
|           angles  degrees
|     circle      -1     358
|     triangle     2     178
|     rectangle     3     358
|
|     >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
|rectangle']),
|     ...         axis='index')
|           angles  degrees
|     circle      -1     359
|     triangle     2     179
|     rectangle     3     359

```

```

|
|     Multiply a DataFrame of different shape with operator version.
|
|     >>> other = pd.DataFrame({'angles': [0, 3, 4]},
|     ...                       index=['circle', 'triangle', 'rectangle'])
|     >>> other
|
|           angles
|     circle      0
|     triangle    3
|     rectangle   4
|
|     >>> df * other
|
|           angles  degrees
|     circle      0      NaN
|     triangle     9      NaN
|     rectangle   16      NaN
|
|     >>> df.mul(other, fill_value=0)
|
|           angles  degrees
|     circle      0      0.0
|     triangle     9      0.0
|     rectangle   16      0.0
|
|     Divide by a MultiIndex by level.
|
|     >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
|     ...                              'degrees': [360, 180, 360, 360, 540,
720]}),
|     ...                              index=[['A', 'A', 'A', 'B', 'B',
'B'],
|     ...                              ['circle', 'triangle',
'rectangle',
|     ...                              'square', 'pentagon',
'hexagon']])
|     >>> df_multindex
|
|           angles  degrees
|     A circle      0      360
|       triangle    3      180
|       rectangle   4      360
|     B square      4      360
|       pentagon    5      540
|       hexagon     6      720
|
|     >>> df.div(df_multindex, level=1, fill_value=0)

```

	angles	degrees
A circle	NaN	1.0
triangle	1.0	1.0
rectangle	1.0	1.0
B square	0.0	0.0
pentagon	0.0	0.0
hexagon	0.0	0.0

multiply = mul(self, other, axis='columns', level=None, fill_value=None)

ne(self, other, axis='columns', level=None)

Get Not equal to of dataframe and other, element-wise (binary operator `ne`).

Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to comparison operators.

Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose axis (rows or columns) and level for comparison.

Parameters

other : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

axis : {0 or 'index', 1 or 'columns'}, default 'columns'

Whether to compare by the index (0 or 'index') or columns (1 or 'columns').

level : int or label

Broadcast across a level, matching Index values on the passed MultiIndex level.

Returns

DataFrame of bool

Result of the comparison.

See Also

DataFrame.eq : Compare DataFrames for equality elementwise.

DataFrame.ne : Compare DataFrames for inequality elementwise.

DataFrame.le : Compare DataFrames for less than inequality

```

|         or equality elementwise.
|     DataFrame.lt : Compare DataFrames for strictly less than
|         inequality elementwise.
|     DataFrame.ge : Compare DataFrames for greater than inequality
|         or equality elementwise.
|     DataFrame.gt : Compare DataFrames for strictly greater than
|         inequality elementwise.
|
|     Notes
|     -----
|     Mismatched indices will be unioned together.
|     `NaN` values are considered different (i.e. `NaN` != `NaN`).
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'cost': [250, 150, 100],
|     ...                     'revenue': [100, 250, 300]},
|     ...                     index=['A', 'B', 'C'])
|     >>> df
|          cost  revenue
|     A    250     100
|     B    150     250
|     C    100     300
|
|     Comparison with a scalar, using either the operator or method:
|
|     >>> df == 100
|          cost  revenue
|     A  False     True
|     B  False     False
|     C   True     False
|
|     >>> df.eq(100)
|          cost  revenue
|     A  False     True
|     B  False     False
|     C   True     False
|
|     When `other` is a :class:`Series`, the columns of a DataFrame are
aligned
|     with the index of `other` and broadcast:
|
|     >>> df != pd.Series([100, 250], index=["cost", "revenue"])
|          cost  revenue

```

A	True	True
B	True	False
C	False	True

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
```

	cost	revenue
A	True	False
B	True	True
C	True	True
D	True	True

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
```

	cost	revenue
A	True	True
B	False	False
C	False	False

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
```

	cost	revenue
A	True	False
B	False	True
C	True	False

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                        index=['A', 'B', 'C', 'D'])
>>> other
```

	revenue
A	300
B	250
C	100
D	150

```
>>> df.gt(other)
```

	cost	revenue
A	False	False

```

|      B  False  False
|      C  False   True
|      D  False  False
|
|      Compare to a MultiIndex by level.
|
|      >>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300,
220],
|      ...                               'revenue': [100, 250, 300, 200, 175,
225]},
|      ...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2',
'Q2'],
|      ...                               ['A', 'B', 'C', 'A', 'B',
'C']])
|
|      >>> df_multindex
|
|          cost  revenue
|      Q1 A    250     100
|          B    150     250
|          C    100     300
|      Q2 A    150     200
|          B    300     175
|          C    220     225
|
|      >>> df.le(df_multindex, level=1)
|
|          cost  revenue
|      Q1 A    True     True
|          B    True     True
|          C    True     True
|      Q2 A   False     True
|          B    True     False
|          C    True     False
|
|      nlargest(self, n, columns, keep='first')
|      Return the first `n` rows ordered by `columns` in descending order.
|
|      Return the first `n` rows with the largest values in `columns`, in
|      descending order. The columns that are not specified are returned as
|      well, but not used for ordering.
|
|      This method is equivalent to
|      ``df.sort_values(columns, ascending=False).head(n)`` , but more
|      performant.
|
|      Parameters

```



```

|     -----
|     n : int
|         Number of rows to return.
|     columns : label or list of labels
|         Column label(s) to order by.
|     keep : {'first', 'last', 'all'}, default 'first'
|         Where there are duplicate values:
|
|         - `first` : prioritize the first occurrence(s)
|         - `last` : prioritize the last occurrence(s)
|         - ``all`` : do not drop any duplicates, even it means
|                     selecting more than `n` items.
|
|     .. versionadded:: 0.24.0
|
|     Returns
|     -----
|     DataFrame
|         The first `n` rows ordered by the given columns in descending
|         order.
|
|     See Also
|     -----
|     DataFrame.nsmallest : Return the first `n` rows ordered by `columns`
in
|         ascending order.
|     DataFrame.sort_values : Sort DataFrame by the values.
|     DataFrame.head : Return the first `n` rows without re-ordering.
|
|     Notes
|     -----
|     This function cannot be used with all column types. For example, when
|     specifying columns with `object` or `category` dtypes, ``TypeError``
is
|     raised.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'population': [59000000, 65000000, 434000,
|     ...                               434000, 434000, 337000, 11300,
|     ...                               11300, 11300],
|     ...                   'GDP': [1937894, 2583560 , 12011, 4520, 12128,
|     ...                           17036, 182, 38, 311],
|     ...                   'alpha-2': ["IT", "FR", "MT", "MV", "BN",

```

```
|     ...     "IS", "NR", "TV", "AI"]},
|     ...     index=["Italy", "France", "Malta",
|     ...     "Maldives", "Brunei", "Iceland",
|     ...     "Nauru", "Tuvalu", "Anguilla"]])
|
| >>> df
|
|      population      GDP alpha-2
|  Italy      59000000  1937894    IT
|  France     65000000  2583560    FR
|  Malta       434000   12011     MT
|  Maldives    434000    4520     MV
|  Brunei      434000   12128     BN
|  Iceland     337000   17036     IS
|  Nauru        11300    182      NR
|  Tuvalu       11300    38       TV
|  Anguilla     11300    311      AI
|
|
| In the following example, we will use ``nlargest`` to select the
three
|
| rows having the largest values in column "population".
|
| >>> df.nlargest(3, 'population')
|
|      population      GDP alpha-2
|  France     65000000  2583560    FR
|  Italy      59000000  1937894    IT
|  Malta       434000   12011     MT
|
|
| When using ``keep='last'``, ties are resolved in reverse order:
|
| >>> df.nlargest(3, 'population', keep='last')
|
|      population      GDP alpha-2
|  France     65000000  2583560    FR
|  Italy      59000000  1937894    IT
|  Brunei      434000   12128     BN
|
|
| When using ``keep='all'``, all duplicate items are maintained:
|
| >>> df.nlargest(3, 'population', keep='all')
|
|      population      GDP alpha-2
|  France     65000000  2583560    FR
|  Italy      59000000  1937894    IT
|  Malta       434000   12011     MT
|  Maldives    434000    4520     MV
|  Brunei      434000   12128     BN
|
```

```

|         To order by the largest values in column "population" and then "GDP",
|         we can specify multiple columns like in the next example.
|
|         >>> df.nlargest(3, ['population', 'GDP'])
|
|              population      GDP alpha-2
|    France    65000000    2583560      FR
|    Italy     59000000    1937894      IT
|    Brunei     434000     12128      BN
|
|         notna(self)
|             Detect existing (non-missing) values.
|
|             Return a boolean same-sized object indicating if the values are not
NA.
|
|             Non-missing values get mapped to True. Characters such as empty
|             strings ``''`` or :attr:`numpy.inf` are not considered NA values
|             (unless you set ``pandas.options.mode.use_inf_as_na = True``).
|             NA values, such as None or :attr:`numpy.NaN`, get mapped to False
|             values.
|
|             Returns
|             -----
|             DataFrame
|                 Mask of bool values for each element in DataFrame that
|                 indicates whether an element is not an NA value.
|
|             See Also
|             -----
|             DataFrame.notnull : Alias of notna.
|             DataFrame.isna : Boolean inverse of notna.
|             DataFrame.dropna : Omit axes labels with missing values.
|             notna : Top-level notna.
|
|             Examples
|             -----
|             Show which entries in a DataFrame are not NA.
|
|             >>> df = pd.DataFrame({'age': [5, 6, np.NaN],
| ...                               'born': [pd.NaT, pd.Timestamp('1939-05-27'),
| ...                                       pd.Timestamp('1940-04-25')],
| ...                               'name': ['Alfred', 'Batman', ''],
| ...                               'toy': [None, 'Batmobile', 'Joker']})
|             >>> df
|
|              age      born      name      toy

```

```

|      0  5.0      NaT  Alfred      None
|      1  6.0 1939-05-27  Batman  Batmobile
|      2  NaN 1940-04-25      Joker
|
|      >>> df.notna()
|
|          age  born  name  toy
|      0   True  False  True  False
|      1   True   True  True   True
|      2  False   True  True   True
|
|      Show which entries in a Series are not NA.
|
|      >>> ser = pd.Series([5, 6, np.NaN])
|      >>> ser
|      0    5.0
|      1    6.0
|      2    NaN
|      dtype: float64
|
|      >>> ser.notna()
|      0     True
|      1     True
|      2    False
|      dtype: bool
|
|      notnull(self)
|      Detect existing (non-missing) values.
|
|      Return a boolean same-sized object indicating if the values are not
NA.
|
|      Non-missing values get mapped to True. Characters such as empty
|      strings ``''`` or :attr:`numpy.inf` are not considered NA values
|      (unless you set ``pandas.options.mode.use_inf_as_na = True``).
|      NA values, such as None or :attr:`numpy.NaN`, get mapped to False
|      values.
|
|      Returns
|      -----
|      DataFrame
|          Mask of bool values for each element in DataFrame that
|          indicates whether an element is not an NA value.
|
|      See Also
|      -----

```

```

| DataFrame.notnull : Alias of notna.
| DataFrame.isna : Boolean inverse of notna.
| DataFrame.dropna : Omit axes labels with missing values.
| notna : Top-level notna.
|
| Examples
| -----
| Show which entries in a DataFrame are not NA.
|
| >>> df = pd.DataFrame({'age': [5, 6, np.NaN],
| ...                     'born': [pd.NaT, pd.Timestamp('1939-05-27'),
| ...                               pd.Timestamp('1940-04-25')],
| ...                     'name': ['Alfred', 'Batman', ''],
| ...                     'toy': [None, 'Batmobile', 'Joker']})
| >>> df
|      age      born  name      toy
|  0  5.0      NaT  Alfred    None
|  1  6.0  1939-05-27  Batman  Batmobile
|  2  NaN  1940-04-25      Joker
|
| >>> df.notna()
|      age  born  name  toy
|  0   True  False  True  False
|  1   True   True  True   True
|  2  False   True  True   True
|
| Show which entries in a Series are not NA.
|
| >>> ser = pd.Series([5, 6, np.NaN])
| >>> ser
|  0    5.0
|  1    6.0
|  2    NaN
| dtype: float64
|
| >>> ser.notna()
|  0    True
|  1    True
|  2   False
| dtype: bool
|
| nsmallest(self, n, columns, keep='first')
| Return the first `n` rows ordered by `columns` in ascending order.
|

```

```

|     Return the first `n` rows with the smallest values in `columns`, in
|     ascending order. The columns that are not specified are returned as
|     well, but not used for ordering.
|
|     This method is equivalent to
|     ``df.sort_values(columns, ascending=True).head(n)`` , but more
|     performant.
|
|     Parameters
|     -----
|     n : int
|         Number of items to retrieve.
|     columns : list or str
|         Column name or names to order by.
|     keep : {'first', 'last', 'all'}, default 'first'
|         Where there are duplicate values:
|
|         - ``first`` : take the first occurrence.
|         - ``last`` : take the last occurrence.
|         - ``all`` : do not drop any duplicates, even it means
|           selecting more than `n` items.
|
|     .. versionadded:: 0.24.0
|
|     Returns
|     -----
|     DataFrame
|
|     See Also
|     -----
|     DataFrame.nlargest : Return the first `n` rows ordered by `columns`
in
|     descending order.
|     DataFrame.sort_values : Sort DataFrame by the values.
|     DataFrame.head : Return the first `n` rows without re-ordering.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'population': [59000000, 65000000, 434000,
|     ...                                     434000, 434000, 337000, 11300,
|     ...                                     11300, 11300],
|     ...                     'GDP': [1937894, 2583560 , 12011, 4520, 12128,
|     ...                               17036, 182, 38, 311],
|     ...                     'alpha-2': ["IT", "FR", "MT", "MV", "BN",

```

```
|     ...     "IS", "NR", "TV", "AI"]},
|     ...     index=["Italy", "France", "Malta",
|     ...     "Maldives", "Brunei", "Iceland",
|     ...     "Nauru", "Tuvalu", "Anguilla"])
```

```
| >>> df
|           population      GDP alpha-2
| Italy      59000000  1937894      IT
| France     65000000  2583560      FR
| Malta       434000   12011      MT
| Maldives   434000    4520      MV
| Brunei     434000   12128      BN
| Iceland    337000   17036      IS
| Nauru       11300    182      NR
| Tuvalu     11300     38      TV
| Anguilla   11300    311      AI
```

| In the following example, we will use ``nsmallest`` to select the three rows having the smallest values in column "a".

```
| >>> df.nsmallest(3, 'population')
|           population      GDP alpha-2
| Nauru       11300    182      NR
| Tuvalu     11300     38      TV
| Anguilla   11300    311      AI
```

| When using ``keep='last'``, ties are resolved in reverse order:

```
| >>> df.nsmallest(3, 'population', keep='last')
|           population      GDP alpha-2
| Anguilla   11300    311      AI
| Tuvalu     11300     38      TV
| Nauru       11300    182      NR
```

| When using ``keep='all'``, all duplicate items are maintained:

```
| >>> df.nsmallest(3, 'population', keep='all')
|           population      GDP alpha-2
| Nauru       11300    182      NR
| Tuvalu     11300     38      TV
| Anguilla   11300    311      AI
```

| To order by the largest values in column "a" and then "c", we can specify multiple columns like in the next example.

```

|     >>> df.nsmallest(3, ['population', 'GDP'])
|
|           population  GDP alpha-2
|     Tuvalu          11300    38      TV
|     Nauru           11300   182      NR
|     Anguilla        11300   311      AI
|
|
|     nunique(self, axis=0, dropna=True)
|         Count distinct observations over requested axis.
|
|     Return Series with number of distinct observations. Can ignore NaN
|     values.
|
|     .. versionadded:: 0.20.0
|
|     Parameters
|     -----
|     axis : {0 or 'index', 1 or 'columns'}, default 0
|         The axis to use. 0 or 'index' for row-wise, 1 or 'columns' for
|         column-wise.
|     dropna : bool, default True
|         Don't include NaN in the counts.
|
|     Returns
|     -----
|     Series
|
|     See Also
|     -----
|     Series.nunique: Method nunique for Series.
|     DataFrame.count: Count non-NA cells for each column or row.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'A': [1, 2, 3], 'B': [1, 1, 1]})
|     >>> df.nunique()
|     A      3
|     B      1
|     dtype: int64
|
|     >>> df.nunique(axis=1)
|     0      1
|     1      2
|     2      2
|     dtype: int64

```



```

| pivot(self, index=None, columns=None, values=None)
|     Return reshaped DataFrame organized by given index / column values.
|
| Reshape data (produce a "pivot" table) based on column values. Uses
| unique values from specified `index` / `columns` to form axes of the
| resulting DataFrame. This function does not support data
| aggregation, multiple values will result in a MultiIndex in the
| columns. See the :ref:`User Guide <reshaping>` for more on reshaping.
|
| Parameters
| -----
| index : string or object, optional
|     Column to use to make new frame's index. If None, uses
|     existing index.
| columns : string or object
|     Column to use to make new frame's columns.
| values : string, object or a list of the previous, optional
|     Column(s) to use for populating new frame's values. If not
|     specified, all remaining columns will be used and the result will
|     have hierarchically indexed columns.
|
| .. versionchanged :: 0.23.0
|     Also accept list of column names.
|
| Returns
| -----
| DataFrame
|     Returns reshaped DataFrame.
|
| Raises
| -----
| ValueError:
|     When there are any `index`, `columns` combinations with multiple
|     values. `DataFrame.pivot_table` when you need to aggregate.
|
| See Also
| -----
| DataFrame.pivot_table : Generalization of pivot that can handle
|     duplicate values for one index/column pair.
| DataFrame.unstack : Pivot based on the index values instead of a
|     column.
|
| Notes

```

```

|     -----
|     For finer-tuned control, see hierarchical indexing documentation
along
|     with the related stack/unstack methods.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'foo': ['one', 'one', 'one', 'two', 'two',
|     ...                          'two'],
|     ...                    'bar': ['A', 'B', 'C', 'A', 'B', 'C'],
|     ...                    'baz': [1, 2, 3, 4, 5, 6],
|     ...                    'zoo': ['x', 'y', 'z', 'q', 'w', 't']})
|     >>> df
|          foo  bar  baz  zoo
|     0  one   A    1    x
|     1  one   B    2    y
|     2  one   C    3    z
|     3  two   A    4    q
|     4  two   B    5    w
|     5  two   C    6    t
|
|     >>> df.pivot(index='foo', columns='bar', values='baz')
|     bar  A  B  C
|     foo
|     one  1  2  3
|     two  4  5  6
|
|     >>> df.pivot(index='foo', columns='bar')['baz']
|     bar  A  B  C
|     foo
|     one  1  2  3
|     two  4  5  6
|
|     >>> df.pivot(index='foo', columns='bar', values=['baz', 'zoo'])
|          baz      zoo
|     bar  A  B  C  A  B  C
|     foo
|     one  1  2  3  x  y  z
|     two  4  5  6  q  w  t
|
|     A ValueError is raised if there are any duplicates.
|
|     >>> df = pd.DataFrame({"foo": ['one', 'one', 'two', 'two'],
|     ...                    "bar": ['A', 'A', 'B', 'C'],

```

```
|     ...                                     "baz": [1, 2, 3, 4]})
|
|     >>> df
|           foo bar  baz
|     0  one   A    1
|     1  one   A    2
|     2  two   B    3
|     3  two   C    4
|
|     Notice that the first two rows are the same for our `index`
|     and `columns` arguments.
|
|     >>> df.pivot(index='foo', columns='bar', values='baz')
|     Traceback (most recent call last):
|         ...
|     ValueError: Index contains duplicate entries, cannot reshape
|
|     pivot_table(self, values=None, index=None, columns=None, aggfunc='mean',
| fill_value=None, margins=False, dropna=True, margins_name='All',
| observed=False)
|
|     Create a spreadsheet-style pivot table as a DataFrame. The levels in
|     the pivot table will be stored in MultiIndex objects (hierarchical
|     indexes) on the index and columns of the result DataFrame.
|
|     Parameters
|     -----
|     values : column to aggregate, optional
|     index : column, Grouper, array, or list of the previous
|         If an array is passed, it must be the same length as the data.
|
| The
|     list can contain any of the other types (except list).
|     Keys to group by on the pivot table index. If an array is
| passed,
|     it is being used as the same manner as column values.
|     columns : column, Grouper, array, or list of the previous
|         If an array is passed, it must be the same length as the data.
|
| The
|     list can contain any of the other types (except list).
|     Keys to group by on the pivot table column. If an array is
| passed,
|     it is being used as the same manner as column values.
|     aggfunc : function, list of functions, dict, default numpy.mean
|         If list of functions passed, the resulting pivot table will have
|         hierarchical columns whose top level are the function names
|         (inferred from the function objects themselves)
```

```

|         If dict is passed, the key is column to aggregate and value
|         is function or list of functions
|     fill_value : scalar, default None
|         Value to replace missing values with
|     margins : boolean, default False
|         Add all row / columns (e.g. for subtotal / grand totals)
|     dropna : boolean, default True
|         Do not include columns whose entries are all NaN
|     margins_name : string, default 'All'
|         Name of the row / column that will contain the totals
|         when margins is True.
|     observed : boolean, default False
|         This only applies if any of the groupers are Categoricals.
|         If True: only show observed values for categorical groupers.
|         If False: show all values for categorical groupers.

```

```

| .. versionchanged :: 0.25.0

```

Returns

```

| -----

```

```

| DataFrame

```

See Also

```

| -----

```

```

| DataFrame.pivot : Pivot without aggregation that can handle
| non-numeric data.

```

Examples

```

| -----

```

```

| >>> df = pd.DataFrame({"A": ["foo", "foo", "foo", "foo", "foo",
| ...                          "bar", "bar", "bar", "bar"],
| ...                    "B": ["one", "one", "one", "two", "two",
| ...                          "one", "one", "two", "two"],
| ...                    "C": ["small", "large", "large", "small",
| ...                          "small", "large", "small", "small",
| ...                          "large"],
| ...                    "D": [1, 2, 2, 3, 3, 4, 5, 6, 7],
| ...                    "E": [2, 4, 5, 5, 6, 6, 8, 9, 9]})

```

```

| >>> df

```

	A	B	C	D	E
0	foo	one	small	1	2
1	foo	one	large	2	4
2	foo	one	large	2	5
3	foo	two	small	3	5

```
|      4  foo  two  small  3  6
|      5  bar  one  large  4  6
|      6  bar  one  small  5  8
|      7  bar  two  small  6  9
|      8  bar  two  large  7  9
```

| This first example aggregates values by taking the sum.

```
| >>> table = pd.pivot_table(df, values='D', index=['A', 'B'],
| ...                          columns=['C'], aggfunc=np.sum)
| >>> table
|      C      large  small
|  A  B
|  bar one    4.0    5.0
|      two    7.0    6.0
|  foo one    4.0    1.0
|      two    NaN    6.0
```

| We can also fill missing values using the `fill_value` parameter.

```
| >>> table = pd.pivot_table(df, values='D', index=['A', 'B'],
| ...                          columns=['C'], aggfunc=np.sum, fill_value=0)
| >>> table
|      C      large  small
|  A  B
|  bar one     4     5
|      two     7     6
|  foo one     4     1
|      two     0     6
```

| The next example aggregates by taking the mean across multiple columns.

```
| >>> table = pd.pivot_table(df, values=['D', 'E'], index=['A', 'C'],
| ...                          aggfunc={'D': np.mean,
| ...                          'E': np.mean})
| >>> table
|      D      E
|  A  C
|  bar large  5.500000  7.500000
|      small  5.500000  8.500000
|  foo large  2.000000  4.500000
|      small  2.333333  4.333333
```

```

|     We can also calculate multiple types of aggregations for any given
|     value column.
|
|     >>> table = pd.pivot_table(df, values=['D', 'E'], index=['A', 'C'],
|     ...                        aggfunc={'D': np.mean,
|     ...                        'E': [min, max, np.mean]})
|     >>> table
|
|           D      E
|           mean  max      mean  min
|  A  C
|  bar large  5.500000  9.0  7.500000  6.0
|        small  5.500000  9.0  8.500000  8.0
|  foo large  2.000000  5.0  4.500000  4.0
|        small  2.333333  6.0  4.333333  2.0
|
|     pow(self, other, axis='columns', level=None, fill_value=None)
|     Get Exponential power of dataframe and other, element-wise (binary
operator `pow`).
|
|     Equivalent to ``dataframe ** other``, but with support to substitute
a fill_value
|     for missing data in one of the inputs. With reverse version, `rpow`.
|
|     Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
|     arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
|     Parameters
|     -----
|     other : scalar, sequence, Series, or DataFrame
|           Any single or multiple element data structure, or list-like
object.
|     axis : {0 or 'index', 1 or 'columns'}
|           Whether to compare by the index (0 or 'index') or columns
|           (1 or 'columns'). For Series input, axis to match Series index
on.
|     level : int or label
|           Broadcast across a level, matching Index values on the
|           passed MultiIndex level.
|     fill_value : float or None, default None
|           Fill existing missing (NaN) values, and any new element needed
for
|           successful DataFrame alignment, with this value before
computation.
|           If data in both corresponding DataFrame locations is missing

```

```

|         the result will be missing.
|
| Returns
| -----
| DataFrame
|         Result of the arithmetic operation.
|
| See Also
| -----
| DataFrame.add : Add DataFrames.
| DataFrame.sub : Subtract DataFrames.
| DataFrame.mul : Multiply DataFrames.
| DataFrame.div : Divide DataFrames (float division).
| DataFrame.truediv : Divide DataFrames (float division).
| DataFrame.floordiv : Divide DataFrames (integer division).
| DataFrame.mod : Calculate modulo (remainder after division).
| DataFrame.pow : Calculate exponential power.
|
| Notes
| -----
| Mismatched indices will be unioned together.
|
| Examples
| -----
| >>> df = pd.DataFrame({'angles': [0, 3, 4],
| ...                    'degrees': [360, 180, 360]},
| ...                    index=['circle', 'triangle', 'rectangle'])
| >>> df
|           angles  degrees
| circle         0      360
| triangle       3      180
| rectangle      4      360
|
| Add a scalar with operator version which return the same
| results.
|
| >>> df + 1
|           angles  degrees
| circle         1      361
| triangle       4      181
| rectangle      5      361
|
| >>> df.add(1)
|           angles  degrees

```

```

|     circle      1      361
|     triangle    4      181
|     rectangle   5      361
|
|     Divide by constant with reverse version.
|
|     >>> df.div(10)
|           angles  degrees
|     circle      0.0     36.0
|     triangle    0.3     18.0
|     rectangle   0.4     36.0
|
|     >>> df.rdiv(10)
|           angles  degrees
|     circle      inf    0.027778
|     triangle    3.333333 0.055556
|     rectangle   2.500000 0.027778
|
|     Subtract a list and Series by axis with operator version.
|
|     >>> df - [1, 2]
|           angles  degrees
|     circle      -1     358
|     triangle     2     178
|     rectangle     3     358
|
|     >>> df.sub([1, 2], axis='columns')
|           angles  degrees
|     circle      -1     358
|     triangle     2     178
|     rectangle     3     358
|
|     >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
'rectangle']),
|     ...         axis='index')
|           angles  degrees
|     circle      -1     359
|     triangle     2     179
|     rectangle     3     359
|
|     Multiply a DataFrame of different shape with operator version.
|
|     >>> other = pd.DataFrame({'angles': [0, 3, 4]},
|     ...                       index=['circle', 'triangle', 'rectangle'])

```



```

|     >>> other
|
|           angles
| circle         0
| triangle        3
| rectangle       4
|
|
|     >>> df * other
|           angles  degrees
| circle          0      NaN
| triangle         9      NaN
| rectangle       16      NaN
|
|
|     >>> df.mul(other, fill_value=0)
|           angles  degrees
| circle          0      0.0
| triangle         9      0.0
| rectangle       16      0.0
|
|
|     Divide by a MultiIndex by level.
|
|     >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ...                                'degrees': [360, 180, 360, 360, 540,
720]}),
| ...                                index=[['A', 'A', 'A', 'B', 'B',
'B'],
| ...                                ['circle', 'triangle',
'rectangle',
| ...                                'square', 'pentagon',
'hexagon']])
|
|     >>> df_multindex
|           angles  degrees
| A circle         0      360
|   triangle        3      180
|   rectangle       4      360
| B square         4      360
|   pentagon        5      540
|   hexagon         6      720
|
|
|     >>> df.div(df_multindex, level=1, fill_value=0)
|           angles  degrees
| A circle        NaN      1.0
|   triangle       1.0      1.0
|   rectangle       1.0      1.0
| B square         0.0      0.0

```

```

|         pentagon      0.0      0.0
|         hexagon       0.0      0.0
|
| prod(self, axis=None, skipna=None, level=None, numeric_only=None,
min_count=0, **kwargs)
|     Return the product of the values for the requested axis.
|
| Parameters
| -----
|     axis : {index (0), columns (1)}
|         Axis for the function to be applied on.
|     skipna : bool, default True
|         Exclude NA/null values when computing the result.
|     level : int or level name, default None
|         If the axis is a MultiIndex (hierarchical), count along a
|         particular level, collapsing into a Series.
|     numeric_only : bool, default None
|         Include only float, int, boolean columns. If None, will attempt
to use
|         everything, then use only numeric data. Not implemented for
Series.
|     min_count : int, default 0
|         The required number of valid values to perform the operation. If
fewer than
|         ``min_count`` non-NA values are present the result will be NA.
|
| .. versionadded :: 0.22.0
|
|     Added with the default being 0. This means the sum of an all-
NA
|     or empty Series is 0, and the product of an all-NA or empty
|     Series is 1.
|     **kwargs
|         Additional keyword arguments to be passed to the function.
|
| Returns
| -----
|     Series or DataFrame (if level specified)
|
| Examples
| -----
|     By default, the product of an empty or all-NA Series is ``1``
|
|     >>> pd.Series([]).prod()

```

```

|     1.0
|
|     This can be controlled with the ``min_count`` parameter
|
|     >>> pd.Series([]).prod(min_count=1)
|     nan
|
|     Thanks to the ``skipna`` parameter, ``min_count`` handles all-NA and
|     empty series identically.
|
|     >>> pd.Series([np.nan]).prod()
|     1.0
|
|     >>> pd.Series([np.nan]).prod(min_count=1)
|     nan
|
|     product = prod(self, axis=None, skipna=None, level=None,
numeric_only=None, min_count=0, **kwargs)
|
|     quantile(self, q=0.5, axis=0, numeric_only=True, interpolation='linear')
|     Return values at the given quantile over requested axis.
|
|     Parameters
|     -----
|     q : float or array-like, default 0.5 (50% quantile)
|         Value between 0 <= q <= 1, the quantile(s) to compute.
|     axis : {0, 1, 'index', 'columns'} (default 0)
|         Equals 0 or 'index' for row-wise, 1 or 'columns' for column-wise.
|     numeric_only : bool, default True
|         If False, the quantile of datetime and timedelta data will be
|         computed as well.
|     interpolation : {'linear', 'lower', 'higher', 'midpoint', 'nearest'}
|         This optional parameter specifies the interpolation method to
use,
|         when the desired quantile lies between two data points `i` and
`j`:
|
|         * linear: `i + (j - i) * fraction`, where `fraction` is the
|           fractional part of the index surrounded by `i` and `j`.
|         * lower: `i`.
|         * higher: `j`.
|         * nearest: `i` or `j` whichever is nearest.
|         * midpoint: `(i + j) / 2.
|

```

```

|         .. versionadded:: 0.18.0
|
|     Returns
|     -----
|     Series or DataFrame
|
|     If ``q`` is an array, a DataFrame will be returned where the
|     index is ``q``, the columns are the columns of self, and the
|     values are the quantiles.
|     If ``q`` is a float, a Series will be returned where the
|     index is the columns of self and the values are the quantiles.
|
|     See Also
|     -----
|     core.window.Rolling.quantile: Rolling quantile.
|     numpy.percentile: Numpy function to compute the percentile.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame(np.array([[1, 1], [2, 10], [3, 100], [4,
100]]),
|         ...                 columns=['a', 'b'])
|     >>> df.quantile(.1)
|     a    1.3
|     b    3.7
|     Name: 0.1, dtype: float64
|     >>> df.quantile([.1, .5])
|           a      b
|     0.1  1.3   3.7
|     0.5  2.5  55.0
|
|     Specifying `numeric_only=False` will also compute the quantile of
|     datetime and timedelta data.
|
|     >>> df = pd.DataFrame({'A': [1, 2],
|     ...                    'B': [pd.Timestamp('2010'),
|     ...                    pd.Timestamp('2011')],
|     ...                    'C': [pd.Timedelta('1 days'),
|     ...                    pd.Timedelta('2 days')]])
|     >>> df.quantile(0.5, numeric_only=False)
|     A          1.5
|     B    2010-07-02 12:00:00
|     C          1 days 12:00:00
|     Name: 0.5, dtype: object

```

```

|
| query(self, expr, inplace=False, **kwargs)
|     Query the columns of a DataFrame with a boolean expression.
|
| Parameters
| -----
| expr : str
|     The query string to evaluate. You can refer to variables
|     in the environment by prefixing them with an '@' character like
|     ``@a + b``.
|
| .. versionadded:: 0.25.0
|
|     You can refer to column names that contain spaces by surrounding
|     them in backticks.
|
|     For example, if one of your columns is called ``a a`` and you
want
|     to sum it with ``b``, your query should be ``a a + b``.
|
| inplace : bool
|     Whether the query should modify the data in place or return
|     a modified copy.
| **kwargs
|     See the documentation for :func:`eval` for complete details
|     on the keyword arguments accepted by :meth:`DataFrame.query`.
|
| .. versionadded:: 0.18.0
|
| Returns
| -----
| DataFrame
|     DataFrame resulting from the provided query expression.
|
| See Also
| -----
| eval : Evaluate a string describing operations on
|     DataFrame columns.
| DataFrame.eval : Evaluate a string describing operations on
|     DataFrame columns.
|
| Notes
| -----
|     The result of the evaluation of this expression is first passed to

```

```

|     :attr:`DataFrame.loc` and if that fails because of a
|     multidimensional key (e.g., a DataFrame) then the result will be
passed
|     to :meth:`DataFrame.__getitem__`.
|
|     This method uses the top-level :func:`eval` function to
|     evaluate the passed query.
|
|     The :meth:`~pandas.DataFrame.query` method uses a slightly
|     modified Python syntax by default. For example, the ``&`` and ``|``
|     (bitwise) operators have the precedence of their boolean cousins,
|     :keyword:`and` and :keyword:`or`. This is syntactically valid
Python,
|     however the semantics are different.
|
|     You can change the semantics of the expression by passing the keyword
|     argument ``parser='python'``. This enforces the same semantics as
|     evaluation in Python space. Likewise, you can pass
``engine='python'``
|     to evaluate an expression using Python itself as a backend. This is
not
|     recommended as it is inefficient compared to using ``numexpr`` as the
|     engine.
|
|     The :attr:`DataFrame.index` and
|     :attr:`DataFrame.columns` attributes of the
|     :class:`~pandas.DataFrame` instance are placed in the query namespace
|     by default, which allows you to treat both the index and columns of
the
|     frame as a column in the frame.
|     The identifier ``index`` is used for the frame index; you can also
|     use the name of the index to identify it in a query. Please note that
|     Python keywords may not be used as identifiers.
|
|     For further details and examples see the ``query`` documentation in
|     :ref:`indexing <indexing.query>`.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'A': range(1, 6),
|     ...                     'B': range(10, 0, -2),
|     ...                     'C C': range(10, 5, -1)})
|     >>> df
|          A   B  C C

```

```
|      0  1  10  10
|      1  2   8   9
|      2  3   6   8
|      3  4   4   7
|      4  5   2   6
|      >>> df.query('A > B')
|          A  B  C  C
|          4  5  2   6
|
|      The previous expression is equivalent to
|
|      >>> df[df.A > df.B]
|          A  B  C  C
|          4  5  2   6
|
|      For columns with spaces in their name, you can use backtick quoting.
|
|      >>> df.query('B == `C C`')
|          A  B  C  C
|          0  1  10  10
|
|      The previous expression is equivalent to
|
|      >>> df[df.B == df['C C']]
|          A  B  C  C
|          0  1  10  10
|
|      radd(self, other, axis='columns', level=None, fill_value=None)
|      Get Addition of dataframe and other, element-wise (binary operator
`radd`).
|
|      Equivalent to ``other + dataframe``, but with support to substitute a
fill_value
|      for missing data in one of the inputs. With reverse version, `add`.
|
|      Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
|      Parameters
|      -----
|      other : scalar, sequence, Series, or DataFrame
|          Any single or multiple element data structure, or list-like
object.
|      axis : {0 or 'index', 1 or 'columns'}
```

```

|         Whether to compare by the index (0 or 'index') or columns
|         (1 or 'columns'). For Series input, axis to match Series index
on.
|         level : int or label
|         Broadcast across a level, matching Index values on the
|         passed MultiIndex level.
|         fill_value : float or None, default None
|         Fill existing missing (NaN) values, and any new element needed
for
|         successful DataFrame alignment, with this value before
computation.
|         If data in both corresponding DataFrame locations is missing
|         the result will be missing.
|
|     Returns
|     -----
|     DataFrame
|     Result of the arithmetic operation.
|
|     See Also
|     -----
|     DataFrame.add : Add DataFrames.
|     DataFrame.sub : Subtract DataFrames.
|     DataFrame.mul : Multiply DataFrames.
|     DataFrame.div : Divide DataFrames (float division).
|     DataFrame.truediv : Divide DataFrames (float division).
|     DataFrame.floordiv : Divide DataFrames (integer division).
|     DataFrame.mod : Calculate modulo (remainder after division).
|     DataFrame.pow : Calculate exponential power.
|
|     Notes
|     -----
|     Mismatched indices will be unioned together.
|
|     Examples
|     -----
|     >>> df = pd.DataFrame({'angles': [0, 3, 4],
|     ...                    'degrees': [360, 180, 360]},
|     ...                    index=['circle', 'triangle', 'rectangle'])
|     >>> df
|
|           angles  degrees
|     circle      0      360
|     triangle    3      180
|     rectangle   4      360

```



```
|
| Add a scalar with operator version which return the same
| results.
```

```
| >>> df + 1
|          angles  degrees
| circle         1     361
| triangle        4     181
| rectangle       5     361
```

```
| >>> df.add(1)
|          angles  degrees
| circle         1     361
| triangle        4     181
| rectangle       5     361
```

```
| Divide by constant with reverse version.
```

```
| >>> df.div(10)
|          angles  degrees
| circle        0.0     36.0
| triangle       0.3     18.0
| rectangle      0.4     36.0
```

```
| >>> df.rdiv(10)
|          angles  degrees
| circle         inf  0.027778
| triangle  3.333333  0.055556
| rectangle  2.500000  0.027778
```

```
| Subtract a list and Series by axis with operator version.
```

```
| >>> df - [1, 2]
|          angles  degrees
| circle        -1     358
| triangle        2     178
| rectangle        3     358
```

```
| >>> df.sub([1, 2], axis='columns')
|          angles  degrees
| circle        -1     358
| triangle        2     178
| rectangle        3     358
|
```

```
|
|     >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
| 'rectangle'])),
|
|     ...         axis='index')
|
|         angles  degrees
|
|     circle      -1      359
|     triangle      2      179
|     rectangle      3      359
|
|
|     Multiply a DataFrame of different shape with operator version.
|
|
|     >>> other = pd.DataFrame({'angles': [0, 3, 4]},
| ...                           index=['circle', 'triangle', 'rectangle'])
|
|     >>> other
|
|         angles
|
|     circle      0
|     triangle      3
|     rectangle      4
|
|
|     >>> df * other
|
|         angles  degrees
|
|     circle      0      NaN
|     triangle      9      NaN
|     rectangle     16      NaN
|
|
|     >>> df.mul(other, fill_value=0)
|
|         angles  degrees
|
|     circle      0      0.0
|     triangle      9      0.0
|     rectangle     16      0.0
|
|
|     Divide by a MultiIndex by level.
|
|
|     >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ...                                   'degrees': [360, 180, 360, 360, 540,
720]},
|
|     ...                                   index=[['A', 'A', 'A', 'B', 'B',
'B'],
|
|     ...                                   ['circle', 'triangle',
'rectangle',
|
|     ...                                   'square', 'pentagon',
'hexagon']])
|
|     >>> df_multindex
|
|         angles  degrees
|
|     A circle      0      360
```

```

|         triangle      3      180
|         rectangle     4      360
|         B square      4      360
|         pentagon      5      540
|         hexagon       6      720
|
|
|         >>> df.div(df_multindex, level=1, fill_value=0)
|
|             angles  degrees
|         A circle   NaN     1.0
|         triangle   1.0     1.0
|         rectangle   1.0     1.0
|         B square    0.0     0.0
|         pentagon    0.0     0.0
|         hexagon     0.0     0.0
|
|
|         rdiv = rtruediv(self, other, axis='columns', level=None, fill_value=None)
|
|
|         reindex(self, labels=None, index=None, columns=None, axis=None,
method=None, copy=True, level=None, fill_value=nan, limit=None,
tolerance=None)
|
|         Conform DataFrame to new index with optional filling logic, placing
|         NA/NaN in locations having no value in the previous index. A new
object
|
|         is produced unless the new index is equivalent to the current one and
|         ``copy=False``.
|
|
|         Parameters
|         -----
|         labels : array-like, optional
|                 New labels / index to conform the axis specified by
'axis' to.
|         index, columns : array-like, optional
|                 New labels / index to conform to, should be specified using
|                 keywords. Preferably an Index object to avoid duplicating data
|         axis : int or str, optional
|                 Axis to target. Can be either the axis name ('index',
'columns')
|
|                 or number (0, 1).
|         method : {None, 'backfill'/'bfill', 'pad'/'ffill', 'nearest'}
|                 Method to use for filling holes in reindexed DataFrame.
|                 Please note: this is only applicable to DataFrames/Series with a
|                 monotonically increasing/decreasing index.
|
|
|         * None (default): don't fill gaps

```

```

|         * pad / ffill: propagate last valid observation forward to next
|         valid
|         * backfill / bfill: use next valid observation to fill gap
|         * nearest: use nearest valid observations to fill gap
|
|         copy : bool, default True
|             Return a new object, even if the passed indexes are the same.
|         level : int or name
|             Broadcast across a level, matching Index values on the
|             passed MultiIndex level.
|         fill_value : scalar, default np.NaN
|             Value to use for missing values. Defaults to NaN, but can be any
|             "compatible" value.
|         limit : int, default None
|             Maximum number of consecutive elements to forward or backward
fill.
|         tolerance : optional
|             Maximum distance between original and new labels for inexact
|             matches. The values of the index at the matching locations must
|             satisfy the equation ``abs(index[indexer] - target) <=
tolerance``.
|
|             Tolerance may be a scalar value, which applies the same tolerance
|             to all values, or list-like, which applies variable tolerance per
|             element. List-like includes list, tuple, array, Series, and must
be
|             the same size as the index and its dtype must exactly match the
|             index's type.
|
|             .. versionadded:: 0.21.0 (list-like tolerance)
|
|         Returns
|         -----
|         DataFrame with changed index.
|
|         See Also
|         -----
|         DataFrame.set_index : Set row labels.
|         DataFrame.reset_index : Remove row labels or move them to new
columns.
|         DataFrame.reindex_like : Change to same indices as other DataFrame.
|
|         Examples
|         -----

```

```DataFrame.reindex``` supports two calling conventions

\* ```(index=index_labels, columns=column_labels, ...)```

\* ```(labels, axis={'index', 'columns'}, ...)```

We *highly* recommend using keyword arguments to clarify your intent.

Create a dataframe with some fictional data.

```
>>> index = ['Firefox', 'Chrome', 'Safari', 'IE10', 'Konqueror']
```

```
>>> df = pd.DataFrame({
... 'http_status': [200,200,404,404,301],
... 'response_time': [0.04, 0.02, 0.07, 0.08, 1.0]},
... index=index)
```

```
>>> df
```

	http_status	response_time
Firefox	200	0.04
Chrome	200	0.02
Safari	404	0.07
IE10	404	0.08
Konqueror	301	1.00

Create a new index and reindex the dataframe. By default values in the new index that do not have corresponding records in the dataframe are assigned ```NaN```.

```
>>> new_index= ['Safari', 'Iceweasel', 'Comodo Dragon', 'IE10',
... 'Chrome']
>>> df.reindex(new_index)
```

	http_status	response_time
Safari	404.0	0.07
Iceweasel	NaN	NaN
Comodo Dragon	NaN	NaN
IE10	404.0	0.08
Chrome	200.0	0.02

We can fill in the missing values by passing a value to the keyword ```fill_value```. Because the index is not monotonically increasing or decreasing, we cannot use arguments to the keyword ```method``` to fill the ```NaN``` values.

```
>>> df.reindex(new_index, fill_value=0)
```

	http_status	response_time
Safari	404	0.07
Iceweasel	0	0.00
Comodo Dragon	0	0.00
IE10	404	0.08
Chrome	200	0.02

```
>>> df.reindex(new_index, fill_value='missing')
```

	http_status	response_time
Safari	404	0.07
Iceweasel	missing	missing
Comodo Dragon	missing	missing
IE10	404	0.08
Chrome	200	0.02

We can also reindex the columns.

```
>>> df.reindex(columns=['http_status', 'user_agent'])
```

	http_status	user_agent
Firefox	200	NaN
Chrome	200	NaN
Safari	404	NaN
IE10	404	NaN
Konqueror	301	NaN

Or we can use "axis-style" keyword arguments

```
>>> df.reindex(['http_status', 'user_agent'], axis="columns")
```

	http_status	user_agent
Firefox	200	NaN
Chrome	200	NaN
Safari	404	NaN
IE10	404	NaN
Konqueror	301	NaN

To further illustrate the filling functionality in ``reindex``, we will create a dataframe with a monotonically increasing index (for example, a sequence of dates).

```
>>> date_index = pd.date_range('1/1/2010', periods=6, freq='D')
>>> df2 = pd.DataFrame({"prices": [100, 101, np.nan, 100, 89, 88]},
... index=date_index)
>>> df2
```

```

| prices
| 2010-01-01 100.0
| 2010-01-02 101.0
| 2010-01-03 NaN
| 2010-01-04 100.0
| 2010-01-05 89.0
| 2010-01-06 88.0
|
| Suppose we decide to expand the dataframe to cover a wider
| date range.
|
| >>> date_index2 = pd.date_range('12/29/2009', periods=10, freq='D')
| >>> df2.reindex(date_index2)
| prices
| 2009-12-29 NaN
| 2009-12-30 NaN
| 2009-12-31 NaN
| 2010-01-01 100.0
| 2010-01-02 101.0
| 2010-01-03 NaN
| 2010-01-04 100.0
| 2010-01-05 89.0
| 2010-01-06 88.0
| 2010-01-07 NaN
|
| The index entries that did not have a value in the original data
frame
| (for example, '2009-12-29') are by default filled with ``NaN``.
| If desired, we can fill in the missing values using one of several
| options.
|
| For example, to back-propagate the last valid value to fill the
``NaN``
| values, pass ``bfill`` as an argument to the ``method`` keyword.
|
| >>> df2.reindex(date_index2, method='bfill')
| prices
| 2009-12-29 100.0
| 2009-12-30 100.0
| 2009-12-31 100.0
| 2010-01-01 100.0
| 2010-01-02 101.0
| 2010-01-03 NaN
| 2010-01-04 100.0

```

```

| 2010-01-05 89.0
| 2010-01-06 88.0
| 2010-01-07 NaN
|
| Please note that the ``NaN`` value present in the original dataframe
| (at index value 2010-01-03) will not be filled by any of the
| value propagation schemes. This is because filling while reindexing
| does not look at dataframe values, but only compares the original and
| desired indexes. If you do want to fill in the ``NaN`` values present
| in the original dataframe, use the ``fillna()`` method.
|
| See the :ref:`user guide <basics.reindexing>` for more.
|
| rename(self, mapper=None, index=None, columns=None, axis=None, copy=True,
inplace=False, level=None, errors='ignore')
| Alter axes labels.
|
| Function / dict values must be unique (1-to-1). Labels not contained
in
| a dict / Series will be left as-is. Extra labels listed don't throw
an
| error.
|
| See the :ref:`user guide <basics.rename>` for more.
|
| Parameters
| -----
| mapper : dict-like or function
| Dict-like or functions transformations to apply to
| that axis' values. Use either ``mapper`` and ``axis`` to
| specify the axis to target with ``mapper``, or ``index`` and
| ``columns``.
| index : dict-like or function
| Alternative to specifying axis (``mapper, axis=0``
| is equivalent to ``index=mapper``).
| columns : dict-like or function
| Alternative to specifying axis (``mapper, axis=1``
| is equivalent to ``columns=mapper``).
| axis : int or str
| Axis to target with ``mapper``. Can be either the axis name
| ('index', 'columns') or number (0, 1). The default is 'index'.
| copy : bool, default True
| Also copy underlying data.
| inplace : bool, default False

```



```

| Whether to return a new DataFrame. If True then value of copy is
| ignored.
| level : int or level name, default None
| In case of a MultiIndex, only rename labels in the specified
| level.
| errors : {'ignore', 'raise'}, default 'ignore'
| If 'raise', raise a `KeyError` when a dict-like `mapper`,
| `index`,
| or `columns` contains labels that are not present in the Index
| being transformed.
| If 'ignore', existing keys will be renamed and extra keys will be
| ignored.
|
| Returns
| -----
| DataFrame
| DataFrame with the renamed axis labels.
|
| Raises
| -----
| KeyError
| If any of the labels is not found in the selected axis and
| "errors='raise'".
|
| See Also
| -----
| DataFrame.rename_axis : Set the name of the axis.
|
| Examples
| -----
|
| ``DataFrame.rename`` supports two calling conventions
|
| * ``(index=index_mapper, columns=columns_mapper, ...)``
| * ``(mapper, axis={'index', 'columns'}, ...)``
|
| We highly recommend using keyword arguments to clarify your
| intent.
|
| Rename columns using a mapping:
|
| >>> df = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})
| >>> df.rename(columns={"A": "a", "B": "c"})
|
| a c

```

```

| 0 1 4
| 1 2 5
| 2 3 6
|
| Rename index using a mapping:
|
| >>> df.rename(index={0: "x", 1: "y", 2: "z"})
|
| A B
| x 1 4
| y 2 5
| z 3 6
|
| Cast index labels to a different type:
|
| >>> df.index
| RangeIndex(start=0, stop=3, step=1)
| >>> df.rename(index=str).index
| Index(['0', '1', '2'], dtype='object')
|
| >>> df.rename(columns={"A": "a", "B": "b", "C": "c"}, errors="raise")
| Traceback (most recent call last):
| KeyError: ['C'] not found in axis
|
| Using axis-style parameters
|
| >>> df.rename(str.lower, axis='columns')
|
| a b
| 0 1 4
| 1 2 5
| 2 3 6
|
| >>> df.rename({1: 2, 2: 4}, axis='index')
|
| A B
| 0 1 4
| 2 2 5
| 4 3 6
|
| reorder_levels(self, order, axis=0)
| Rearrange index levels using input order. May not drop or
| duplicate levels.
|
| Parameters
| -----
| order : list of int or list of str

```

```

| List representing new level order. Reference level by number
| (position) or by key (label).
| axis : int
| Where to reorder levels.
|
| Returns
| -----
| type of caller (new object)
|
| replace(self, to_replace=None, value=None, inplace=False, limit=None,
| regex=False, method='pad')
| Replace values given in `to_replace` with `value`.
|
| Values of the DataFrame are replaced with other values dynamically.
| This differs from updating with ``.loc`` or ``.iloc``, which require
| you to specify a location to update with some value.
|
| Parameters
| -----
| to_replace : str, regex, list, dict, Series, int, float, or None
| How to find the values that will be replaced.
|
| * numeric, str or regex:
|
| - numeric: numeric values equal to `to_replace` will be
| replaced with `value`
| - str: string exactly matching `to_replace` will be replaced
| with `value`
| - regex: regexs matching `to_replace` will be replaced with
| `value`
|
| * list of str, regex, or numeric:
|
| - First, if `to_replace` and `value` are both lists, they
| **must** be the same length.
| - Second, if ``regex=True`` then all of the strings in
| **both**
| lists will be interpreted as regexs otherwise they will
| match
| directly. This doesn't matter much for `value` since there
| are only a few possible substitution regexes you can use.
| - str, regex and numeric rules apply as above.
|
| * dict:

```

```

|
| - Dicts can be used to specify different replacement values
| for different existing values. For example,
| ``{'a': 'b', 'y': 'z'}`` replaces the value 'a' with 'b'
and
| 'y' with 'z'. To use a dict in this way the `value`
| parameter should be `None`.
| - For a DataFrame a dict can specify that different values
| should be replaced in different columns. For example,
| ``{'a': 1, 'b': 'z'}`` looks for the value 1 in column 'a'
| and the value 'z' in column 'b' and replaces these values
| with whatever is specified in `value`. The `value`
parameter
| should not be ``None`` in this case. You can treat this as
a
| special case of passing two lists except that you are
| specifying the column to search in.
| - For a DataFrame nested dictionaries, e.g.,
| ``{'a': {'b': np.nan}}``, are read as follows: look in
column
| 'a' for the value 'b' and replace it with NaN. The `value`
| parameter should be ``None`` to use a nested dict in this
| way. You can nest regular expressions as well. Note that
| column names (the top-level dictionary keys in a nested
| dictionary) cannot be regular expressions.
|
| * None:
|
| - This means that the `regex` argument must be a string,
| compiled regular expression, or list, dict, ndarray or
| Series of such elements. If `value` is also ``None`` then
| this must be a nested dictionary or Series.
|
| See the examples section for examples of each of these.
| value : scalar, dict, list, str, regex, default None
| Value to replace any values matching `to_replace` with.
| For a DataFrame a dict of values can be used to specify which
| value to use for each column (columns not in the dict will not be
| filled). Regular expressions, strings and lists or dicts of such
| objects are also allowed.
| inplace : bool, default False
| If True, in place. Note: this will modify any
| other views on this object (e.g. a column from a DataFrame).
| Returns the caller if this is True.

```

```

| limit : int, default None
| Maximum size gap to forward or backward fill.
| regex : bool or same types as `to_replace`, default False
| Whether to interpret `to_replace` and/or `value` as regular
| expressions. If this is ``True`` then `to_replace` *must* be a
| string. Alternatively, this could be a regular expression or a
| list, dict, or array of regular expressions in which case
| `to_replace` must be ``None``.
| method : {'pad', 'ffill', 'bfill', 'None'}
| The method to use when for replacement, when `to_replace` is a
| scalar, list or tuple and `value` is ``None``.
|
| .. versionchanged:: 0.23.0
| Added to DataFrame.
|
| Returns
| -----
| DataFrame
| Object after replacement.
|
| Raises
| -----
| AssertionError
| * If `regex` is not a ``bool`` and `to_replace` is not
| ``None``.
| TypeError
| * If `to_replace` is a ``dict`` and `value` is not a ``list``,
| ``dict``, ``ndarray``, or ``Series``
| * If `to_replace` is ``None`` and `regex` is not compilable
| into a regular expression or is a list, dict, ndarray, or
| Series.
| * When replacing multiple ``bool`` or ``datetime64`` objects and
| the arguments to `to_replace` does not match the type of the
| value being replaced
| ValueError
| * If a ``list`` or an ``ndarray`` is passed to `to_replace` and
| `value` but they are not the same length.
|
| See Also
| -----
| DataFrame.fillna : Fill NA values.
| DataFrame.where : Replace values based on boolean condition.
| Series.str.replace : Simple string replacement.
|

```

```

| Notes
| -----
| * Regex substitution is performed under the hood with ``re.sub``. The
| rules for substitution for ``re.sub`` are the same.
| * Regular expressions will only substitute on strings, meaning you
| cannot provide, for example, a regular expression matching floating
| point numbers and expect the columns in your frame that have a
| numeric dtype to be matched. However, if those floating point
| numbers are strings, then you can do this.
| * This method has a lot of options. You are encouraged to
experiment
| and play with this method to gain intuition about how it works.
| * When dict is used as the `to_replace` value, it is like
| key(s) in the dict are the to_replace part and
| value(s) in the dict are the value parameter.
|
| Examples
| -----
|
| **Scalar `to_replace` and `value`**
|
| >>> s = pd.Series([0, 1, 2, 3, 4])
| >>> s.replace(0, 5)
| 0 5
| 1 1
| 2 2
| 3 3
| 4 4
| dtype: int64
|
| >>> df = pd.DataFrame({'A': [0, 1, 2, 3, 4],
| ... 'B': [5, 6, 7, 8, 9],
| ... 'C': ['a', 'b', 'c', 'd', 'e']})
| >>> df.replace(0, 5)
| A B C
| 0 5 5 a
| 1 1 6 b
| 2 2 7 c
| 3 3 8 d
| 4 4 9 e
|
| **List-like `to_replace`**
|
| >>> df.replace([0, 1, 2, 3], 4)

```

```

| A B C
| 0 4 5 a
| 1 4 6 b
| 2 4 7 c
| 3 4 8 d
| 4 4 9 e
|
| >>> df.replace([0, 1, 2, 3], [4, 3, 2, 1])
| A B C
| 0 4 5 a
| 1 3 6 b
| 2 2 7 c
| 3 1 8 d
| 4 4 9 e
|
| >>> s.replace([1, 2], method='bfill')
| 0 0
| 1 3
| 2 3
| 3 3
| 4 4
| dtype: int64
|
| **dict-like `to_replace`**
|
| >>> df.replace({0: 10, 1: 100})
| A B C
| 0 10 5 a
| 1 100 6 b
| 2 2 7 c
| 3 3 8 d
| 4 4 9 e
|
| >>> df.replace({'A': 0, 'B': 5}, 100)
| A B C
| 0 100 100 a
| 1 1 6 b
| 2 2 7 c
| 3 3 8 d
| 4 4 9 e
|
| >>> df.replace({'A': {0: 100, 4: 400}})
| A B C
| 0 100 5 a

```

```
| 1 1 6 b
| 2 2 7 c
| 3 3 8 d
| 4 400 9 e
```

```
| **Regular expression `to_replace`**
```

```
| >>> df = pd.DataFrame({'A': ['bat', 'foo', 'bait'],
| ... 'B': ['abc', 'bar', 'xyz']})
| >>> df.replace(to_replace=r'^ba.$', value='new', regex=True)
```

```
| A B
| 0 new abc
| 1 foo new
| 2 bait xyz
```

```
| >>> df.replace({'A': r'^ba.$'}, {'A': 'new'}, regex=True)
```

```
| A B
| 0 new abc
| 1 foo bar
| 2 bait xyz
```

```
| >>> df.replace(regex=r'^ba.$', value='new')
```

```
| A B
| 0 new abc
| 1 foo new
| 2 bait xyz
```

```
| >>> df.replace(regex={r'^ba.$': 'new', 'foo': 'xyz'})
```

```
| A B
| 0 new abc
| 1 xyz new
| 2 bait xyz
```

```
| >>> df.replace(regex=[r'^ba.$', 'foo'], value='new')
```

```
| A B
| 0 new abc
| 1 new new
| 2 bait xyz
```

```
| Note that when replacing multiple ``bool`` or ``datetime64`` objects,
| the data types in the `to_replace` parameter must match the data
| type of the value being replaced:
```

```
| >>> df = pd.DataFrame({'A': [True, False, True],
```



```

| ... 'B': [False, True, False]])
| >>> df.replace({'a string': 'new value', True: False}) # raises
| Traceback (most recent call last):
| ...
| TypeError: Cannot compare types 'ndarray(dtype=bool)' and 'str'
|
| This raises a ``TypeError`` because one of the ``dict`` keys is not
of
| the correct type for replacement.
|
| Compare the behavior of ``s.replace({'a': None})`` and
| ``s.replace('a', None)`` to understand the peculiarities
| of the `to_replace` parameter:
|
| >>> s = pd.Series([10, 'a', 'a', 'b', 'a'])
|
| When one uses a dict as the `to_replace` value, it is like the
| value(s) in the dict are equal to the `value` parameter.
| ``s.replace({'a': None})`` is equivalent to
| ``s.replace(to_replace={'a': None}, value=None, method=None)``:
|
| >>> s.replace({'a': None})
| 0 10
| 1 None
| 2 None
| 3 b
| 4 None
| dtype: object
|
| When ``value=None`` and `to_replace` is a scalar, list or
| tuple, `replace` uses the method parameter (default 'pad') to do the
| replacement. So this is why the 'a' values are being replaced by 10
| in rows 1 and 2 and 'b' in row 4 in this case.
| The command ``s.replace('a', None)`` is actually equivalent to
| ``s.replace(to_replace='a', value=None, method='pad')``:
|
| >>> s.replace('a', None)
| 0 10
| 1 10
| 2 10
| 3 b
| 4 b
| dtype: object
|

```

```

| reset_index(self, level=None, drop=False, inplace=False, col_level=0,
col_fill='')
| Reset the index, or a level of it.
|
| Reset the index of the DataFrame, and use the default one instead.
| If the DataFrame has a MultiIndex, this method can remove one or more
| levels.
|
| Parameters
| -----
| level : int, str, tuple, or list, default None
| Only remove the given levels from the index. Removes all levels
by
| default.
| drop : bool, default False
| Do not try to insert index into dataframe columns. This resets
| the index to the default integer index.
| inplace : bool, default False
| Modify the DataFrame in place (do not create a new object).
| col_level : int or str, default 0
| If the columns have multiple levels, determines which level the
| labels are inserted into. By default it is inserted into the
first
| level.
| col_fill : object, default ''
| If the columns have multiple levels, determines how the other
| levels are named. If None then the index name is repeated.
|
| Returns
| -----
| DataFrame
| DataFrame with the new index.
|
| See Also
| -----
| DataFrame.set_index : Opposite of reset_index.
| DataFrame.reindex : Change to new indices or expand indices.
| DataFrame.reindex_like : Change to same indices as other DataFrame.
|
| Examples
| -----
| >>> df = pd.DataFrame([('bird', 389.0),
| ... ('bird', 24.0),
| ... ('mammal', 80.5),

```

```
| ('mammal', np.nan)],
| index=['falcon', 'parrot', 'lion', 'monkey'],
| columns=('class', 'max_speed'))
| >>> df
| class max_speed
| falcon bird 389.0
| parrot bird 24.0
| lion mammal 80.5
| monkey mammal NaN
```

| When we reset the index, the old index is added as a column, and a new sequential index is used:

```
| >>> df.reset_index()
| index class max_speed
| 0 falcon bird 389.0
| 1 parrot bird 24.0
| 2 lion mammal 80.5
| 3 monkey mammal NaN
```

| We can use the `drop` parameter to avoid the old index being added as a column:

```
| >>> df.reset_index(drop=True)
| class max_speed
| 0 bird 389.0
| 1 bird 24.0
| 2 mammal 80.5
| 3 mammal NaN
```

| You can also use `reset\_index` with `MultiIndex`.

```
| >>> index = pd.MultiIndex.from_tuples([('bird', 'falcon'),
| ... ('bird', 'parrot'),
| ... ('mammal', 'lion'),
| ... ('mammal', 'monkey')],
| ... names=['class', 'name'])
| >>> columns = pd.MultiIndex.from_tuples([('speed', 'max'),
| ... ('species', 'type')])
| >>> df = pd.DataFrame([(389.0, 'fly'),
| ... (24.0, 'fly'),
| ... (80.5, 'run'),
| ... (np.nan, 'jump')],
| ... index=index,
```

```
| ... columns=columns)
```

```
| >>> df
| speed species
| max type
| class name
| bird falcon 389.0 fly
| parrot 24.0 fly
| mammal lion 80.5 run
| monkey NaN jump
```

| If the index has multiple levels, we can reset a subset of them:

```
| >>> df.reset_index(level='class')
| class speed species
| max type
| name
| falcon bird 389.0 fly
| parrot bird 24.0 fly
| lion mammal 80.5 run
| monkey mammal NaN jump
```

| If we are not dropping the index, by default, it is placed in the top level. We can place it in another level:

```
| >>> df.reset_index(level='class', col_level=1)
| speed species
| class max type
| name
| falcon bird 389.0 fly
| parrot bird 24.0 fly
| lion mammal 80.5 run
| monkey mammal NaN jump
```

| When the index is inserted under another level, we can specify under which one with the parameter `col\_fill`:

```
| >>> df.reset_index(level='class', col_level=1, col_fill='species')
| species speed species
| class max type
| name
| falcon bird 389.0 fly
| parrot bird 24.0 fly
| lion mammal 80.5 run
| monkey mammal NaN jump
```

```

|
| If we specify a nonexistent level for `col_fill`, it is created:
|
| >>> df.reset_index(level='class', col_level=1, col_fill='genus')
| genus speed species
| class max type
|
| name
| falcon bird 389.0 fly
| parrot bird 24.0 fly
| lion mammal 80.5 run
| monkey mammal NaN jump
|
| rfloordiv(self, other, axis='columns', level=None, fill_value=None)
| Get Integer division of dataframe and other, element-wise (binary
operator `rfloordiv`).
|
| Equivalent to ``other // dataframe``, but with support to substitute
a fill_value
| for missing data in one of the inputs. With reverse version,
`floordiv`.
|
| Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
| arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
| Parameters
| -----
| other : scalar, sequence, Series, or DataFrame
| Any single or multiple element data structure, or list-like
object.
| axis : {0 or 'index', 1 or 'columns'}
| Whether to compare by the index (0 or 'index') or columns
| (1 or 'columns'). For Series input, axis to match Series index
on.
| level : int or label
| Broadcast across a level, matching Index values on the
| passed MultiIndex level.
| fill_value : float or None, default None
| Fill existing missing (NaN) values, and any new element needed
for
| successful DataFrame alignment, with this value before
computation.
| If data in both corresponding DataFrame locations is missing
| the result will be missing.
|

```

```

| Returns
| -----
| DataFrame
| Result of the arithmetic operation.
|
| See Also
| -----
| DataFrame.add : Add DataFrames.
| DataFrame.sub : Subtract DataFrames.
| DataFrame.mul : Multiply DataFrames.
| DataFrame.div : Divide DataFrames (float division).
| DataFrame.truediv : Divide DataFrames (float division).
| DataFrame.floordiv : Divide DataFrames (integer division).
| DataFrame.mod : Calculate modulo (remainder after division).
| DataFrame.pow : Calculate exponential power.
|
| Notes
| -----
| Mismatched indices will be unioned together.
|
| Examples
| -----
| >>> df = pd.DataFrame({'angles': [0, 3, 4],
| ... 'degrees': [360, 180, 360]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> df
| angles degrees
| circle 0 360
| triangle 3 180
| rectangle 4 360
|
| Add a scalar with operator version which return the same
| results.
|
| >>> df + 1
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| >>> df.add(1)
| angles degrees
| circle 1 361
| triangle 4 181

```

```

| rectangle 5 361
|
| Divide by constant with reverse version.
|
| >>> df.div(10)
| angles degrees
| circle 0.0 36.0
| triangle 0.3 18.0
| rectangle 0.4 36.0
|
| >>> df.rdiv(10)
| angles degrees
| circle inf 0.027778
| triangle 3.333333 0.055556
| rectangle 2.500000 0.027778
|
| Subtract a list and Series by axis with operator version.
|
| >>> df - [1, 2]
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub([1, 2], axis='columns')
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
|rectangle']),
| ... axis='index')
| angles degrees
| circle -1 359
| triangle 2 179
| rectangle 3 359
|
| Multiply a DataFrame of different shape with operator version.
|
| >>> other = pd.DataFrame({'angles': [0, 3, 4]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> other
| angles

```

```

| circle 0
| triangle 3
| rectangle 4
|
| >>> df * other
|
| angles degrees
| circle 0 NaN
| triangle 9 NaN
| rectangle 16 NaN
|
| >>> df.mul(other, fill_value=0)
|
| angles degrees
| circle 0 0.0
| triangle 9 0.0
| rectangle 16 0.0
|
| Divide by a MultiIndex by level.
|
| >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ... 'degrees': [360, 180, 360, 360, 540,
720]}),
| ... index=[['A', 'A', 'A', 'B', 'B',
'B'],
| ... ['circle', 'triangle',
'rectangle',
| ... 'square', 'pentagon',
'hexagon']])
|
| >>> df_multindex
|
| angles degrees
| A circle 0 360
| triangle 3 180
| rectangle 4 360
| B square 4 360
| pentagon 5 540
| hexagon 6 720
|
| >>> df.div(df_multindex, level=1, fill_value=0)
|
| angles degrees
| A circle NaN 1.0
| triangle 1.0 1.0
| rectangle 1.0 1.0
| B square 0.0 0.0
| pentagon 0.0 0.0
| hexagon 0.0 0.0

```



```

|
| rmod(self, other, axis='columns', level=None, fill_value=None)
| Get Modulo of dataframe and other, element-wise (binary operator
`rmod`).
|
| Equivalent to ``other % dataframe``, but with support to substitute a
fill_value
| for missing data in one of the inputs. With reverse version, `mod`.
|
| Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
| arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
| Parameters
| -----
| other : scalar, sequence, Series, or DataFrame
| Any single or multiple element data structure, or list-like
object.
| axis : {0 or 'index', 1 or 'columns'}
| Whether to compare by the index (0 or 'index') or columns
| (1 or 'columns'). For Series input, axis to match Series index
on.
| level : int or label
| Broadcast across a level, matching Index values on the
| passed MultiIndex level.
| fill_value : float or None, default None
| Fill existing missing (NaN) values, and any new element needed
for
| successful DataFrame alignment, with this value before
computation.
| If data in both corresponding DataFrame locations is missing
| the result will be missing.
|
| Returns
| -----
| DataFrame
| Result of the arithmetic operation.
|
| See Also
| -----
| DataFrame.add : Add DataFrames.
| DataFrame.sub : Subtract DataFrames.
| DataFrame.mul : Multiply DataFrames.
| DataFrame.div : Divide DataFrames (float division).
| DataFrame.truediv : Divide DataFrames (float division).

```

```

| DataFrame.floordiv : Divide DataFrames (integer division).
| DataFrame.mod : Calculate modulo (remainder after division).
| DataFrame.pow : Calculate exponential power.
|
| Notes
| -----
| Mismatched indices will be unioned together.
|
| Examples
| -----
| >>> df = pd.DataFrame({'angles': [0, 3, 4],
| ... 'degrees': [360, 180, 360]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> df
| angles degrees
| circle 0 360
| triangle 3 180
| rectangle 4 360
|
| Add a scalar with operator version which return the same
| results.
|
| >>> df + 1
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| >>> df.add(1)
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| Divide by constant with reverse version.
|
| >>> df.div(10)
| angles degrees
| circle 0.0 36.0
| triangle 0.3 18.0
| rectangle 0.4 36.0
|
| >>> df.rdiv(10)
| angles degrees

```

```

| circle inf 0.027778
| triangle 3.333333 0.055556
| rectangle 2.500000 0.027778
|
| Subtract a list and Series by axis with operator version.
|
| >>> df - [1, 2]
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub([1, 2], axis='columns')
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
|'rectangle'])),
| ... axis='index')
| angles degrees
| circle -1 359
| triangle 2 179
| rectangle 3 359
|
| Multiply a DataFrame of different shape with operator version.
|
| >>> other = pd.DataFrame({'angles': [0, 3, 4]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> other
| angles
| circle 0
| triangle 3
| rectangle 4
|
| >>> df * other
| angles degrees
| circle 0 NaN
| triangle 9 NaN
| rectangle 16 NaN
|
| >>> df.mul(other, fill_value=0)
| angles degrees

```

```

| circle 0 0.0
| triangle 9 0.0
| rectangle 16 0.0
|
| Divide by a MultiIndex by level.
|
| >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ... 'degrees': [360, 180, 360, 360, 540,
720]}),
| ... index=[['A', 'A', 'A', 'B', 'B',
'B'],
| ... ['circle', 'triangle',
'rectangle',
| ... 'square', 'pentagon',
'hexagon']])
|
| >>> df_multindex
|
| angles degrees
| A circle 0 360
| triangle 3 180
| rectangle 4 360
| B square 4 360
| pentagon 5 540
| hexagon 6 720
|
| >>> df.div(df_multindex, level=1, fill_value=0)
|
| angles degrees
| A circle NaN 1.0
| triangle 1.0 1.0
| rectangle 1.0 1.0
| B square 0.0 0.0
| pentagon 0.0 0.0
| hexagon 0.0 0.0
|
| rmul(self, other, axis='columns', level=None, fill_value=None)
| Get Multiplication of dataframe and other, element-wise (binary
operator `rmul`).
|
| Equivalent to ``other * dataframe``, but with support to substitute a
fill_value
| for missing data in one of the inputs. With reverse version, `mul`.
|
| Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
| arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|

```

```

| Parameters
| -----
| other : scalar, sequence, Series, or DataFrame
| Any single or multiple element data structure, or list-like
object.
| axis : {0 or 'index', 1 or 'columns'}
| Whether to compare by the index (0 or 'index') or columns
| (1 or 'columns'). For Series input, axis to match Series index
on.
| level : int or label
| Broadcast across a level, matching Index values on the
| passed MultiIndex level.
| fill_value : float or None, default None
| Fill existing missing (NaN) values, and any new element needed
for
| successful DataFrame alignment, with this value before
computation.
| If data in both corresponding DataFrame locations is missing
| the result will be missing.
|
| Returns
| -----
| DataFrame
| Result of the arithmetic operation.
|
| See Also
| -----
| DataFrame.add : Add DataFrames.
| DataFrame.sub : Subtract DataFrames.
| DataFrame.mul : Multiply DataFrames.
| DataFrame.div : Divide DataFrames (float division).
| DataFrame.truediv : Divide DataFrames (float division).
| DataFrame.floordiv : Divide DataFrames (integer division).
| DataFrame.mod : Calculate modulo (remainder after division).
| DataFrame.pow : Calculate exponential power.
|
| Notes
| -----
| Mismatched indices will be unioned together.
|
| Examples
| -----
| >>> df = pd.DataFrame({'angles': [0, 3, 4],
| ... 'degrees': [360, 180, 360]}),

```

```
| ... index=['circle', 'triangle', 'rectangle'])
```

```
| >>> df
```

```
| angles degrees
| circle 0 360
| triangle 3 180
| rectangle 4 360
```

```
| Add a scalar with operator version which return the same
| results.
```

```
| >>> df + 1
```

```
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
```

```
| >>> df.add(1)
```

```
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
```

```
| Divide by constant with reverse version.
```

```
| >>> df.div(10)
```

```
| angles degrees
| circle 0.0 36.0
| triangle 0.3 18.0
| rectangle 0.4 36.0
```

```
| >>> df.rdiv(10)
```

```
| angles degrees
| circle inf 0.027778
| triangle 3.333333 0.055556
| rectangle 2.500000 0.027778
```

```
| Subtract a list and Series by axis with operator version.
```

```
| >>> df - [1, 2]
```

```
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
```

```

| >>> df.sub([1, 2], axis='columns')
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
|rectangle']),
| ... axis='index')
| angles degrees
| circle -1 359
| triangle 2 179
| rectangle 3 359
|
| Multiply a DataFrame of different shape with operator version.
|
| >>> other = pd.DataFrame({'angles': [0, 3, 4]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> other
| angles
| circle 0
| triangle 3
| rectangle 4
|
| >>> df * other
| angles degrees
| circle 0 NaN
| triangle 9 NaN
| rectangle 16 NaN
|
| >>> df.mul(other, fill_value=0)
| angles degrees
| circle 0 0.0
| triangle 9 0.0
| rectangle 16 0.0
|
| Divide by a MultiIndex by level.
|
| >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ... 'degrees': [360, 180, 360, 360, 540,
720]},
| ... index=[['A', 'A', 'A', 'B', 'B',
'B'],

```

```
| ... ['circle', 'triangle',
'rectangle',
| ... 'square', 'pentagon',
'hexagon']]
|
| >>> df_multindex
|
| angles degrees
| A circle 0 360
| triangle 3 180
| rectangle 4 360
| B square 4 360
| pentagon 5 540
| hexagon 6 720
|
| >>> df.div(df_multindex, level=1, fill_value=0)
|
| angles degrees
| A circle NaN 1.0
| triangle 1.0 1.0
| rectangle 1.0 1.0
| B square 0.0 0.0
| pentagon 0.0 0.0
| hexagon 0.0 0.0
|
| rolling(self, window, min_periods=None, center=False, win_type=None,
on=None, axis=0, closed=None)
|
| Provide rolling window calculations.
|
| .. versionadded:: 0.18.0
|
| Parameters
| -----
|
| window : int, or offset
| Size of the moving window. This is the number of observations
used for
| calculating the statistic. Each window will be a fixed size.
|
| If its an offset then this will be the time period of each
window. Each
| window will be a variable sized based on the observations
included in
| the time-period. This is only valid for datetimelike indexes.
This is
| new in 0.19.0
|
| min_periods : int, default None
| Minimum number of observations in window required to have a value
```



```

| (otherwise result is NA). For a window that is specified by an
offset,
| `min_periods` will default to 1. Otherwise, `min_periods` will
default
| to the size of the window.
| center : bool, default False
| Set the labels at the center of the window.
| win_type : str, default None
| Provide a window type. If ``None``, all points are evenly
weighted.
| See the notes below for further information.
| on : str, optional
| For a DataFrame, a datetime-like column on which to calculate the
rolling
| window, rather than the DataFrame's index. Provided integer
column is
| ignored and excluded from result since an integer index is not
used to
| calculate the rolling window.
| axis : int or str, default 0
| closed : str, default None
| Make the interval closed on the 'right', 'left', 'both' or
| 'neither' endpoints.
| For offset-based windows, it defaults to 'right'.
| For fixed windows, defaults to 'both'. Remaining cases not
implemented
| for fixed windows.
|
| .. versionadded:: 0.20.0
|
| Returns
| -----
| a Window or Rolling sub-classed for the particular operation
|
| See Also
| -----
| expanding : Provides expanding transformations.
| ewm : Provides exponential weighted functions.
|
| Notes
| -----
| By default, the result is set to the right edge of the window. This
can be
| changed to the center of the window by setting ``center=True``.

```

```

|
| To learn more about the offsets & frequency strings, please see `this
link
| <http://pandas.pydata.org/pandas-
docs/stable/user_guide/timeseries.html#offset-aliases>`___.
|
| The recognized win_types are:
|
| * ``boxcar``
| * ``triang``
| * ``blackman``
| * ``hamming``
| * ``bartlett``
| * ``parzen``
| * ``bohman``
| * ``blackmanharris``
| * ``nuttall``
| * ``barthann``
| * ``kaiser`` (needs beta)
| * ``gaussian`` (needs std)
| * ``general_gaussian`` (needs power, width)
| * ``slepian`` (needs width)
| * ``exponential`` (needs tau), center is set to None.
|
| If ``win_type=None`` all points are evenly weighted. To learn more
about
| different window types see `scipy.signal window functions
| <https://docs.scipy.org/doc/scipy/reference/signal.html#window-
functions>`___.
|
| Examples
| -----
|
| >>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]})
| >>> df
| B
| 0 0.0
| 1 1.0
| 2 2.0
| 3 NaN
| 4 4.0
|
| Rolling sum with a window length of 2, using the 'triang'
| window type.

```

```

|
| >>> df.rolling(2, win_type='triang').sum()
| B
| 0 NaN
| 1 0.5
| 2 1.5
| 3 NaN
| 4 NaN
|
| Rolling sum with a window length of 2, min_periods defaults
| to the window length.
|
| >>> df.rolling(2).sum()
| B
| 0 NaN
| 1 1.0
| 2 3.0
| 3 NaN
| 4 NaN
|
| Same as above, but explicitly set the min_periods
|
| >>> df.rolling(2, min_periods=1).sum()
| B
| 0 0.0
| 1 1.0
| 2 3.0
| 3 2.0
| 4 4.0
|
| A ragged (meaning not-a-regular frequency), time-indexed DataFrame
|
| >>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]},
| ... index = [pd.Timestamp('20130101 09:00:00'),
| ... pd.Timestamp('20130101 09:00:02'),
| ... pd.Timestamp('20130101 09:00:03'),
| ... pd.Timestamp('20130101 09:00:05'),
| ... pd.Timestamp('20130101 09:00:06')])
|
| >>> df
| B
| 2013-01-01 09:00:00 0.0
| 2013-01-01 09:00:02 1.0
| 2013-01-01 09:00:03 2.0

```

```

| 2013-01-01 09:00:05 NaN
| 2013-01-01 09:00:06 4.0
|
| Contrasting to an integer rolling window, this will roll a variable
| length window corresponding to the time period.
| The default for min_periods is 1.
|
| >>> df.rolling('2s').sum()
|
| B
| 2013-01-01 09:00:00 0.0
| 2013-01-01 09:00:02 1.0
| 2013-01-01 09:00:03 3.0
| 2013-01-01 09:00:05 NaN
| 2013-01-01 09:00:06 4.0
|
| round(self, decimals=0, *args, **kwargs)
| Round a DataFrame to a variable number of decimal places.
|
| Parameters
| -----
| decimals : int, dict, Series
| Number of decimal places to round each column to. If an int is
| given, round each column to the same number of places.
| Otherwise dict and Series round to variable numbers of places.
| Column names should be in the keys if `decimals` is a
| dict-like, or in the index if `decimals` is a Series. Any
| columns not included in `decimals` will be left as is. Elements
| of `decimals` which are not columns of the input will be
| ignored.
| *args
| Additional keywords have no effect but might be accepted for
| compatibility with numpy.
| **kwargs
| Additional keywords have no effect but might be accepted for
| compatibility with numpy.
|
| Returns
| -----
| DataFrame
| A DataFrame with the affected columns rounded to the specified
| number of decimal places.
|
| See Also
| -----

```

```

| numpy.around : Round a numpy array to the given number of decimals.
| Series.round : Round a Series to the given number of decimals.
|
| Examples
| -----
| >>> df = pd.DataFrame([(0.21, 0.32), (0.01, 0.67), (0.66, 0.03), (0.21,
.18)],
| ... columns=['dogs', 'cats'])
| >>> df
| dogs cats
| 0 0.21 0.32
| 1 0.01 0.67
| 2 0.66 0.03
| 3 0.21 0.18
|
| By providing an integer each column is rounded to the same number
| of decimal places
|
| >>> df.round(1)
| dogs cats
| 0 0.2 0.3
| 1 0.0 0.7
| 2 0.7 0.0
| 3 0.2 0.2
|
| With a dict, the number of places for specific columns can be
| specified with the column names as key and the number of decimal
| places as value
|
| >>> df.round({'dogs': 1, 'cats': 0})
| dogs cats
| 0 0.2 0.0
| 1 0.0 1.0
| 2 0.7 0.0
| 3 0.2 0.0
|
| Using a Series, the number of places for specific columns can be
| specified with the column names as index and the number of
| decimal places as value
|
| >>> decimals = pd.Series([0, 1], index=['cats', 'dogs'])
| >>> df.round(decimals)
| dogs cats
| 0 0.2 0.0

```

```

| 1 0.0 1.0
| 2 0.7 0.0
| 3 0.2 0.0
|
| rpow(self, other, axis='columns', level=None, fill_value=None)
| Get Exponential power of dataframe and other, element-wise (binary
operator `rpow`).
|
| Equivalent to ``other ** dataframe``, but with support to substitute
a fill_value
| for missing data in one of the inputs. With reverse version, `pow`.
|
| Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
| Parameters
| -----
| other : scalar, sequence, Series, or DataFrame
| Any single or multiple element data structure, or list-like
object.
| axis : {0 or 'index', 1 or 'columns'}
| Whether to compare by the index (0 or 'index') or columns
(1 or 'columns'). For Series input, axis to match Series index
on.
| level : int or label
| Broadcast across a level, matching Index values on the
passed MultiIndex level.
| fill_value : float or None, default None
| Fill existing missing (NaN) values, and any new element needed
for
successful DataFrame alignment, with this value before
computation.
| If data in both corresponding DataFrame locations is missing
the result will be missing.
|
| Returns
| -----
| DataFrame
| Result of the arithmetic operation.
|
| See Also
| -----
| DataFrame.add : Add DataFrames.
| DataFrame.sub : Subtract DataFrames.

```

```

| DataFrame.mul : Multiply DataFrames.
| DataFrame.div : Divide DataFrames (float division).
| DataFrame.truediv : Divide DataFrames (float division).
| DataFrame.floordiv : Divide DataFrames (integer division).
| DataFrame.mod : Calculate modulo (remainder after division).
| DataFrame.pow : Calculate exponential power.
|
| Notes
| -----
| Mismatched indices will be unioned together.
|
| Examples
| -----
| >>> df = pd.DataFrame({'angles': [0, 3, 4],
| ... 'degrees': [360, 180, 360]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> df
| angles degrees
| circle 0 360
| triangle 3 180
| rectangle 4 360
|
| Add a scalar with operator version which return the same
| results.
|
| >>> df + 1
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| >>> df.add(1)
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| Divide by constant with reverse version.
|
| >>> df.div(10)
| angles degrees
| circle 0.0 36.0
| triangle 0.3 18.0
| rectangle 0.4 36.0

```

```
|
| >>> df.rdiv(10)
| angles degrees
| circle inf 0.027778
| triangle 3.333333 0.055556
| rectangle 2.500000 0.027778
|
| Subtract a list and Series by axis with operator version.
|
| >>> df - [1, 2]
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub([1, 2], axis='columns')
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
|'rectangle'])),
| ... axis='index')
| angles degrees
| circle -1 359
| triangle 2 179
| rectangle 3 359
|
| Multiply a DataFrame of different shape with operator version.
|
| >>> other = pd.DataFrame({'angles': [0, 3, 4]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> other
| angles
| circle 0
| triangle 3
| rectangle 4
|
| >>> df * other
| angles degrees
| circle 0 NaN
| triangle 9 NaN
| rectangle 16 NaN
```



```
|
| >>> df.mul(other, fill_value=0)
| angles degrees
| circle 0 0.0
| triangle 9 0.0
| rectangle 16 0.0
|
| Divide by a MultiIndex by level.
|
| >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ... 'degrees': [360, 180, 360, 360, 540,
720]}),
| ... index=[['A', 'A', 'A', 'B', 'B',
'B'],
| ... ['circle', 'triangle',
'rectangle',
| ... 'square', 'pentagon',
'hexagon']])
| >>> df_multindex
| angles degrees
| A circle 0 360
| triangle 3 180
| rectangle 4 360
| B square 4 360
| pentagon 5 540
| hexagon 6 720
|
| >>> df.div(df_multindex, level=1, fill_value=0)
| angles degrees
| A circle NaN 1.0
| triangle 1.0 1.0
| rectangle 1.0 1.0
| B square 0.0 0.0
| pentagon 0.0 0.0
| hexagon 0.0 0.0
|
| rsub(self, other, axis='columns', level=None, fill_value=None)
| Get Subtraction of dataframe and other, element-wise (binary operator
`rsub`).
|
| Equivalent to ``other - dataframe``, but with support to substitute a
fill_value
| for missing data in one of the inputs. With reverse version, `sub`.
|
```

```

| Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
| arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
| Parameters
| -----
| other : scalar, sequence, Series, or DataFrame
| Any single or multiple element data structure, or list-like
object.
| axis : {0 or 'index', 1 or 'columns'}
| Whether to compare by the index (0 or 'index') or columns
| (1 or 'columns'). For Series input, axis to match Series index
on.
| level : int or label
| Broadcast across a level, matching Index values on the
| passed MultiIndex level.
| fill_value : float or None, default None
| Fill existing missing (NaN) values, and any new element needed
for
| successful DataFrame alignment, with this value before
computation.
| If data in both corresponding DataFrame locations is missing
| the result will be missing.
|
| Returns
| -----
| DataFrame
| Result of the arithmetic operation.
|
| See Also
| -----
| DataFrame.add : Add DataFrames.
| DataFrame.sub : Subtract DataFrames.
| DataFrame.mul : Multiply DataFrames.
| DataFrame.div : Divide DataFrames (float division).
| DataFrame.truediv : Divide DataFrames (float division).
| DataFrame.floordiv : Divide DataFrames (integer division).
| DataFrame.mod : Calculate modulo (remainder after division).
| DataFrame.pow : Calculate exponential power.
|
| Notes
| -----
| Mismatched indices will be unioned together.
|
| Examples

```

```

| -----
| >>> df = pd.DataFrame({'angles': [0, 3, 4],
| ... 'degrees': [360, 180, 360]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> df
| angles degrees
| circle 0 360
| triangle 3 180
| rectangle 4 360
|
| Add a scalar with operator version which return the same
| results.
|
| >>> df + 1
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| >>> df.add(1)
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| Divide by constant with reverse version.
|
| >>> df.div(10)
| angles degrees
| circle 0.0 36.0
| triangle 0.3 18.0
| rectangle 0.4 36.0
|
| >>> df.rdiv(10)
| angles degrees
| circle inf 0.027778
| triangle 3.333333 0.055556
| rectangle 2.500000 0.027778
|
| Subtract a list and Series by axis with operator version.
|
| >>> df - [1, 2]
| angles degrees
| circle -1 358

```

```

| triangle 2 178
| rectangle 3 358
|
| >>> df.sub([1, 2], axis='columns')
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
|rectangle']),
| ... axis='index')
| angles degrees
| circle -1 359
| triangle 2 179
| rectangle 3 359
|
| Multiply a DataFrame of different shape with operator version.
|
| >>> other = pd.DataFrame({'angles': [0, 3, 4]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> other
| angles
| circle 0
| triangle 3
| rectangle 4
|
| >>> df * other
| angles degrees
| circle 0 NaN
| triangle 9 NaN
| rectangle 16 NaN
|
| >>> df.mul(other, fill_value=0)
| angles degrees
| circle 0 0.0
| triangle 9 0.0
| rectangle 16 0.0
|
| Divide by a MultiIndex by level.
|
| >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ... 'degrees': [360, 180, 360, 360, 540,
720]}),

```

```

| ... index=[['A', 'A', 'A', 'B', 'B',
'B'],
| ... ['circle', 'triangle',
'rectangle',
| ... 'square', 'pentagon',
'hexagon']])
|
| >>> df_multindex
|
| angles degrees
| A circle 0 360
| triangle 3 180
| rectangle 4 360
| B square 4 360
| pentagon 5 540
| hexagon 6 720
|
|
| >>> df.div(df_multindex, level=1, fill_value=0)
|
| angles degrees
| A circle NaN 1.0
| triangle 1.0 1.0
| rectangle 1.0 1.0
| B square 0.0 0.0
| pentagon 0.0 0.0
| hexagon 0.0 0.0
|
|
| rtruediv(self, other, axis='columns', level=None, fill_value=None)
| Get Floating division of dataframe and other, element-wise (binary
operator `rtruediv`).
|
| Equivalent to ``other / dataframe``, but with support to substitute a
fill_value
| for missing data in one of the inputs. With reverse version,
`truediv`.
|
| Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
| Parameters
| -----
| other : scalar, sequence, Series, or DataFrame
| Any single or multiple element data structure, or list-like
object.
| axis : {0 or 'index', 1 or 'columns'}
| Whether to compare by the index (0 or 'index') or columns

```

```

| (1 or 'columns'). For Series input, axis to match Series index
on.
| level : int or label
| Broadcast across a level, matching Index values on the
| passed MultiIndex level.
| fill_value : float or None, default None
| Fill existing missing (NaN) values, and any new element needed
for
| successful DataFrame alignment, with this value before
computation.
| If data in both corresponding DataFrame locations is missing
| the result will be missing.
|
| Returns
| -----
| DataFrame
| Result of the arithmetic operation.
|
| See Also
| -----
| DataFrame.add : Add DataFrames.
| DataFrame.sub : Subtract DataFrames.
| DataFrame.mul : Multiply DataFrames.
| DataFrame.div : Divide DataFrames (float division).
| DataFrame.truediv : Divide DataFrames (float division).
| DataFrame.floordiv : Divide DataFrames (integer division).
| DataFrame.mod : Calculate modulo (remainder after division).
| DataFrame.pow : Calculate exponential power.
|
| Notes
| -----
| Mismatched indices will be unioned together.
|
| Examples
| -----
| >>> df = pd.DataFrame({'angles': [0, 3, 4],
| ... 'degrees': [360, 180, 360]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> df
|
| angles degrees
| circle 0 360
| triangle 3 180
| rectangle 4 360
|

```

```

| Add a scalar with operator version which return the same
| results.
|
| >>> df + 1
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| >>> df.add(1)
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| Divide by constant with reverse version.
|
| >>> df.div(10)
| angles degrees
| circle 0.0 36.0
| triangle 0.3 18.0
| rectangle 0.4 36.0
|
| >>> df.rdiv(10)
| angles degrees
| circle inf 0.027778
| triangle 3.333333 0.055556
| rectangle 2.500000 0.027778
|
| Subtract a list and Series by axis with operator version.
|
| >>> df - [1, 2]
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub([1, 2], axis='columns')
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|

```

```
|
| >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
| 'rectangle'])),
| ... axis='index')
| angles degrees
| circle -1 359
| triangle 2 179
| rectangle 3 359
|
|
| Multiply a DataFrame of different shape with operator version.
|
| >>> other = pd.DataFrame({'angles': [0, 3, 4]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> other
| angles
| circle 0
| triangle 3
| rectangle 4
|
| >>> df * other
| angles degrees
| circle 0 NaN
| triangle 9 NaN
| rectangle 16 NaN
|
| >>> df.mul(other, fill_value=0)
| angles degrees
| circle 0 0.0
| triangle 9 0.0
| rectangle 16 0.0
|
| Divide by a MultiIndex by level.
|
| >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ... 'degrees': [360, 180, 360, 360, 540,
720]},
| ... index=[['A', 'A', 'A', 'B', 'B',
'B'],
| ... ['circle', 'triangle',
'rectangle',
| ... 'square', 'pentagon',
'hexagon']])
| >>> df_multindex
| angles degrees
| A circle 0 360
```



```

| triangle 3 180
| rectangle 4 360
| B square 4 360
| pentagon 5 540
| hexagon 6 720
|
|
| >>> df.div(df_multindex, level=1, fill_value=0)
|
| angles degrees
| A circle NaN 1.0
| triangle 1.0 1.0
| rectangle 1.0 1.0
| B square 0.0 0.0
| pentagon 0.0 0.0
| hexagon 0.0 0.0
|
|
| select_dtypes(self, include=None, exclude=None)
|
| Return a subset of the DataFrame's columns based on the column
dtypes.
|
|
| Parameters
| -----
| include, exclude : scalar or list-like
|
| A selection of dtypes or strings to be included/excluded. At
least
|
| one of these parameters must be supplied.
|
|
| Returns
| -----
| DataFrame
|
| The subset of the frame including the dtypes in ``include`` and
| excluding the dtypes in ``exclude``.
|
|
| Raises
| -----
| ValueError
|
| * If both of ``include`` and ``exclude`` are empty
| * If ``include`` and ``exclude`` have overlapping elements
| * If any kind of string dtype is passed in.
|
|
| Notes
| -----
|
| * To select all *numeric* types, use ``np.number`` or ``'number'``
| * To select strings you must use the ``object`` dtype, but note that
| this will return *all* object dtype columns

```

```
| * See the `numpy dtype hierarchy`
| <http://docs.scipy.org/doc/numpy/reference/arrays.scalars.html>`
| * To select datetimes, use ``np.datetime64``, ``'datetime'`` or
| ``'datetime64'``
| * To select timedeltas, use ``np.timedelta64``, ``'timedelta'`` or
| ``'timedelta64'``
| * To select Pandas categorical dtypes, use ``'category'``
| * To select Pandas datetimetz dtypes, use ``'datetimetz'`` (new in
| 0.20.0) or ``'datetime64[ns, tz]'``
```

#### | Examples

```
| -----
```

```
| >>> df = pd.DataFrame({'a': [1, 2] * 3,
| ... 'b': [True, False] * 3,
| ... 'c': [1.0, 2.0] * 3})
```

```
| >>> df
|
```

	a	b	c
0	1	True	1.0
1	2	False	2.0
2	1	True	1.0
3	2	False	2.0
4	1	True	1.0
5	2	False	2.0

```
| >>> df.select_dtypes(include='bool')
|
```

	b
0	True
1	False
2	True
3	False
4	True
5	False

```
| >>> df.select_dtypes(include=['float64'])
|
```

	c
0	1.0
1	2.0
2	1.0
3	2.0
4	1.0
5	2.0

```
| >>> df.select_dtypes(exclude=['int'])
|
```

	b	c
--	---	---

```

| 0 True 1.0
| 1 False 2.0
| 2 True 1.0
| 3 False 2.0
| 4 True 1.0
| 5 False 2.0
|
|
| sem(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None,
**kwargs)
| Return unbiased standard error of the mean over requested axis.
|
| Normalized by N-1 by default. This can be changed using the ddof
argument
|
| Parameters
| -----
| axis : {index (0), columns (1)}
| skipna : bool, default True
| Exclude NA/null values. If an entire row/column is NA, the result
| will be NA
| level : int or level name, default None
| If the axis is a MultiIndex (hierarchical), count along a
| particular level, collapsing into a Series
| ddof : int, default 1
| Delta Degrees of Freedom. The divisor used in calculations is N -
ddof,
| where N represents the number of elements.
| numeric_only : bool, default None
| Include only float, int, boolean columns. If None, will attempt
to use
| everything, then use only numeric data. Not implemented for
Series.
|
| Returns
| -----
| Series or DataFrame (if level specified)
|
| set_index(self, keys, drop=True, append=False, inplace=False,
verify_integrity=False)
| Set the DataFrame index using existing columns.
|
| Set the DataFrame index (row labels) using one or more existing
| columns or arrays (of the correct length). The index can replace the
| existing index or expand on it.

```

```

|
| Parameters
| -----
| keys : label or array-like or list of labels/arrays
| This parameter can be either a single column key, a single array
of
| the same length as the calling DataFrame, or a list containing an
| arbitrary combination of column keys and arrays. Here, "array"
| encompasses :class:`Series`, :class:`Index`, ``np.ndarray``, and
| instances of :class:`~collections.abc.Iterator`.
| drop : bool, default True
| Delete columns to be used as the new index.
| append : bool, default False
| Whether to append columns to existing index.
| inplace : bool, default False
| Modify the DataFrame in place (do not create a new object).
| verify_integrity : bool, default False
| Check the new index for duplicates. Otherwise defer the check
until
| necessary. Setting to False will improve the performance of this
| method.
|
| Returns
| -----
| DataFrame
| Changed row labels.
|
| See Also
| -----
| DataFrame.reset_index : Opposite of set_index.
| DataFrame.reindex : Change to new indices or expand indices.
| DataFrame.reindex_like : Change to same indices as other DataFrame.
|
| Examples
| -----
| >>> df = pd.DataFrame({'month': [1, 4, 7, 10],
| ... 'year': [2012, 2014, 2013, 2014],
| ... 'sale': [55, 40, 84, 31]})
| >>> df
| month year sale
| 0 1 2012 55
| 1 4 2014 40
| 2 7 2013 84
| 3 10 2014 31

```

Set the index to become the 'month' column:

```
>>> df.set_index('month')
```

	year	sale
month		
1	2012	55
4	2014	40
7	2013	84
10	2014	31

Create a MultiIndex using columns 'year' and 'month':

```
>>> df.set_index(['year', 'month'])
```

		sale
year	month	
2012	1	55
2014	4	40
2013	7	84
2014	10	31

Create a MultiIndex using an Index and a column:

```
>>> df.set_index([pd.Index([1, 2, 3, 4]), 'year'])
```

		month	sale
	year		
1	2012	1	55
2	2014	4	40
3	2013	7	84
4	2014	10	31

Create a MultiIndex using two Series:

```
>>> s = pd.Series([1, 2, 3, 4])
```

```
>>> df.set_index([s, s**2])
```

		month	year	sale
1	1	1	2012	55
2	4	4	2014	40
3	9	7	2013	84
4	16	10	2014	31

```
set_value(self, index, col, value, takeable=False)
```

Put single value at passed column and index.

```

| .. deprecated:: 0.21.0
| Use .at[] or .iat[] accessors instead.
|
| Parameters
| -----
| index : row label
| col : column label
| value : scalar
| takeable : interpret the index/col as indexers, default False
|
| Returns
| -----
| DataFrame
| If label pair is contained, will be reference to calling
DataFrame,
| otherwise a new object.
|
| shift(self, periods=1, freq=None, axis=0, fill_value=None)
| Shift index by desired number of periods with an optional time
`freq`.
|
| When `freq` is not passed, shift the index without realigning the
data.
| If `freq` is passed (in this case, the index must be date or
datetime,
| or it will raise a `NotImplementedError`), the index will be
| increased using the periods and the `freq`.
|
| Parameters
| -----
| periods : int
| Number of periods to shift. Can be positive or negative.
| freq : DateOffset, tseries.offsets, timedelta, or str, optional
| Offset to use from the tseries module or time rule (e.g. 'EOM').
| If `freq` is specified then the index values are shifted but the
| data is not realigned. That is, use `freq` if you would like to
| extend the index when shifting and preserve the original data.
| axis : {0 or 'index', 1 or 'columns', None}, default None
| Shift direction.
| fill_value : object, optional
| The scalar value to use for newly introduced missing values.
| the default depends on the dtype of `self`.
| For numeric data, ``np.nan`` is used.

```

```

| For datetime, timedelta, or period data, etc. :attr:`NaT` is
used.
| For extension dtypes, ``self.dtype.na_value`` is used.
|
| .. versionchanged:: 0.24.0
|
| Returns
| -----
| DataFrame
| Copy of input object, shifted.
|
| See Also
| -----
| Index.shift : Shift values of Index.
| DatetimeIndex.shift : Shift values of DatetimeIndex.
| PeriodIndex.shift : Shift values of PeriodIndex.
| tshift : Shift the time index, using the index's frequency if
| available.
|
| Examples
| -----
| >>> df = pd.DataFrame({'Col1': [10, 20, 15, 30, 45],
| ... 'Col2': [13, 23, 18, 33, 48],
| ... 'Col3': [17, 27, 22, 37, 52]})
|
| >>> df.shift(periods=3)
|
| Col1 Col2 Col3
| 0 NaN NaN NaN
| 1 NaN NaN NaN
| 2 NaN NaN NaN
| 3 10.0 13.0 17.0
| 4 20.0 23.0 27.0
|
| >>> df.shift(periods=1, axis='columns')
|
| Col1 Col2 Col3
| 0 NaN 10.0 13.0
| 1 NaN 20.0 23.0
| 2 NaN 15.0 18.0
| 3 NaN 30.0 33.0
| 4 NaN 45.0 48.0
|
| >>> df.shift(periods=3, fill_value=0)
|
| Col1 Col2 Col3
| 0 0 0 0

```

```

| 1 0 0 0
| 2 0 0 0
| 3 10 13 17
| 4 20 23 27
|
| skew(self, axis=None, skipna=None, level=None, numeric_only=None,
**kwargs)
| Return unbiased skew over requested axis
| Normalized by N-1.
|
| Parameters
| -----
| axis : {index (0), columns (1)}
| Axis for the function to be applied on.
| skipna : bool, default True
| Exclude NA/null values when computing the result.
| level : int or level name, default None
| If the axis is a MultiIndex (hierarchical), count along a
| particular level, collapsing into a Series.
| numeric_only : bool, default None
| Include only float, int, boolean columns. If None, will attempt
to use
| everything, then use only numeric data. Not implemented for
Series.
| **kwargs
| Additional keyword arguments to be passed to the function.
|
| Returns
| -----
| Series or DataFrame (if level specified)
|
| sort_index(self, axis=0, level=None, ascending=True, inplace=False,
kind='quicksort', na_position='last', sort_remaining=True, by=None)
| Sort object by labels (along an axis).
|
| Parameters
| -----
| axis : {0 or 'index', 1 or 'columns'}, default 0
| The axis along which to sort. The value 0 identifies the rows,
| and 1 identifies the columns.
| level : int or level name or list of ints or list of level names
| If not None, sort on values in specified index level(s).
| ascending : bool, default True
| Sort ascending vs. descending.

```



```

| inplace : bool, default False
| If True, perform operation in-place.
| kind : {'quicksort', 'mergesort', 'heapsort'}, default 'quicksort'
| Choice of sorting algorithm. See also ndarray.np.sort for more
| information. `mergesort` is the only stable algorithm. For
| DataFrames, this option is only applied when sorting on a single
| column or label.
| na_position : {'first', 'last'}, default 'last'
| Puts NaNs at the beginning if `first`; `last` puts NaNs at the
end.
|
| Not implemented for MultiIndex.
| sort_remaining : bool, default True
| If True and sorting by level and index is multilevel, sort by
other
| levels too (in order) after sorting by specified level.
|
| Returns
| -----
| sorted_obj : DataFrame or None
| DataFrame with sorted index if inplace=False, None otherwise.
|
| sort_values(self, by, axis=0, ascending=True, inplace=False,
kind='quicksort', na_position='last')
| Sort by the values along either axis.
|
| Parameters
| -----
| by : str or list of str
| Name or list of names to sort by.
|
| - if `axis` is 0 or `'index'` then `by` may contain index
| levels and/or column labels
| - if `axis` is 1 or `'columns'` then `by` may contain
column
| levels and/or index labels
|
| .. versionchanged:: 0.23.0
| Allow specifying index or column level names.
| axis : {0 or 'index', 1 or 'columns'}, default 0
| Axis to be sorted.
| ascending : bool or list of bool, default True
| Sort ascending vs. descending. Specify list for multiple sort
| orders. If this is a list of bools, must match the length of
| the by.

```

```

| inplace : bool, default False
| If True, perform operation in-place.
| kind : {'quicksort', 'mergesort', 'heapsort'}, default 'quicksort'
| Choice of sorting algorithm. See also ndarray.np.sort for more
| information. `mergesort` is the only stable algorithm. For
| DataFrames, this option is only applied when sorting on a single
| column or label.
| na_position : {'first', 'last'}, default 'last'
| Puts NaNs at the beginning if `first`; `last` puts NaNs at the
| end.

```

```

| Returns
| -----
| sorted_obj : DataFrame or None
| DataFrame with sorted values if inplace=False, None otherwise.

```

#### Examples

```

| -----
| >>> df = pd.DataFrame({
| ... 'col1': ['A', 'A', 'B', np.nan, 'D', 'C'],
| ... 'col2': [2, 1, 9, 8, 7, 4],
| ... 'col3': [0, 1, 9, 4, 2, 3],
| ... })
| >>> df
| col1 col2 col3
| 0 A 2 0
| 1 A 1 1
| 2 B 9 9
| 3 NaN 8 4
| 4 D 7 2
| 5 C 4 3

```

#### Sort by col1

```

| >>> df.sort_values(by=['col1'])
| col1 col2 col3
| 0 A 2 0
| 1 A 1 1
| 2 B 9 9
| 5 C 4 3
| 4 D 7 2
| 3 NaN 8 4

```

#### Sort by multiple columns

```

|
| >>> df.sort_values(by=['col1', 'col2'])
| col1 col2 col3
| 1 A 1 1
| 0 A 2 0
| 2 B 9 9
| 5 C 4 3
| 4 D 7 2
| 3 NaN 8 4
|
|
| Sort Descending
|
| >>> df.sort_values(by='col1', ascending=False)
| col1 col2 col3
| 4 D 7 2
| 5 C 4 3
| 2 B 9 9
| 0 A 2 0
| 1 A 1 1
| 3 NaN 8 4
|
|
| Putting NAs first
|
| >>> df.sort_values(by='col1', ascending=False, na_position='first')
| col1 col2 col3
| 3 NaN 8 4
| 4 D 7 2
| 5 C 4 3
| 2 B 9 9
| 0 A 2 0
| 1 A 1 1
|
|
| stack(self, level=-1, dropna=True)
| Stack the prescribed level(s) from columns to index.
|
|
| Return a reshaped DataFrame or Series having a multi-level
| index with one or more new inner-most levels compared to the current
| DataFrame. The new inner-most levels are created by pivoting the
| columns of the current dataframe:
|
|
| - if the columns have a single level, the output is a Series;
| - if the columns have multiple levels, the new index
| level(s) is (are) taken from the prescribed level(s) and
| the output is a DataFrame.

```

```

|
| The new index levels are sorted.
|
| Parameters
| -----
| level : int, str, list, default -1
| Level(s) to stack from the column axis onto the index
| axis, defined as one index or label, or a list of indices
| or labels.
| dropna : bool, default True
| Whether to drop rows in the resulting Frame/Series with
| missing values. Stacking a column level onto the index
| axis can create combinations of index and column values
| that are missing from the original dataframe. See Examples
| section.
|
| Returns
| -----
| DataFrame or Series
| Stacked dataframe or series.
|
| See Also
| -----
| DataFrame.unstack : Unstack prescribed level(s) from index axis
| onto column axis.
| DataFrame.pivot : Reshape dataframe from long format to wide
| format.
| DataFrame.pivot_table : Create a spreadsheet-style pivot table
| as a DataFrame.
|
| Notes
| -----
| The function is named by analogy with a collection of books
| being reorganized from being side by side on a horizontal
| position (the columns of the dataframe) to being stacked
| vertically on top of each other (in the index of the
| dataframe).
|
| Examples
| -----
| **Single level columns**
|
| >>> df_single_level_cols = pd.DataFrame([[0, 1], [2, 3]],
| ... index=['cat', 'dog'],

```

```
| ... columns=['weight', 'height'])
|
| Stacking a dataframe with a single level column axis returns a
Series:
|
| >>> df_single_level_cols
| weight height
| cat 0 1
| dog 2 3
| >>> df_single_level_cols.stack()
| cat weight 0
| height 1
| dog weight 2
| height 3
| dtype: int64
|
| **Multi level columns: simple case**
|
| >>> multicoll = pd.MultiIndex.from_tuples([('weight', 'kg'),
| ... ('weight', 'pounds')])
| >>> df_multi_level_cols1 = pd.DataFrame([[1, 2], [2, 4]],
| ... index=['cat', 'dog'],
| ... columns=multicoll)
|
| Stacking a dataframe with a multi-level column axis:
|
| >>> df_multi_level_cols1
| weight
| kg pounds
| cat 1 2
| dog 2 4
| >>> df_multi_level_cols1.stack()
| weight
| cat kg 1
| pounds 2
| dog kg 2
| pounds 4
|
| **Missing values**
|
| >>> multicol2 = pd.MultiIndex.from_tuples([('weight', 'kg'),
| ... ('height', 'm')])
| >>> df_multi_level_cols2 = pd.DataFrame([[1.0, 2.0], [3.0, 4.0]],
| ... index=['cat', 'dog'],
```

```
| ... columns=multicol2)
```

| It is common to have missing values when stacking a dataframe  
| with multi-level columns, as the stacked dataframe typically  
| has more values than the original dataframe. Missing values  
| are filled with NaNs:

```
|
| >>> df_multi_level_cols2
| weight height
| kg m
| cat 1.0 2.0
| dog 3.0 4.0
| >>> df_multi_level_cols2.stack()
| height weight
| cat kg NaN 1.0
| m 2.0 NaN
| dog kg NaN 3.0
| m 4.0 NaN
|
| **Prescribing the level(s) to be stacked**
|
| The first parameter controls which level or levels are stacked:
|
| >>> df_multi_level_cols2.stack(0)
| kg m
| cat height NaN 2.0
| weight 1.0 NaN
| dog height NaN 4.0
| weight 3.0 NaN
| >>> df_multi_level_cols2.stack([0, 1])
| cat height m 2.0
| weight kg 1.0
| dog height m 4.0
| weight kg 3.0
| dtype: float64
|
| **Dropping missing values**
|
| >>> df_multi_level_cols3 = pd.DataFrame([[None, 1.0], [2.0, 3.0]],
| ... index=['cat', 'dog'],
| ... columns=multicol2)
```

| Note that rows where all values are missing are dropped by  
| default but this behaviour can be controlled via the dropna

```

| keyword parameter:
|
| >>> df_multi_level_cols3
| weight height
| kg m
| cat NaN 1.0
| dog 2.0 3.0
| >>> df_multi_level_cols3.stack(dropna=False)
| height weight
| cat kg NaN NaN
| m 1.0 NaN
| dog kg NaN 2.0
| m 3.0 NaN
| >>> df_multi_level_cols3.stack(dropna=True)
| height weight
| cat m 1.0 NaN
| dog kg NaN 2.0
| m 3.0 NaN
|
| std(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None,
**kwargs)
| Return sample standard deviation over requested axis.
|
| Normalized by N-1 by default. This can be changed using the ddof
argument
|
| Parameters
| -----
| axis : {index (0), columns (1)}
| skipna : bool, default True
| Exclude NA/null values. If an entire row/column is NA, the result
| will be NA
| level : int or level name, default None
| If the axis is a MultiIndex (hierarchical), count along a
| particular level, collapsing into a Series
| ddof : int, default 1
| Delta Degrees of Freedom. The divisor used in calculations is N -
ddof,
| where N represents the number of elements.
| numeric_only : bool, default None
| Include only float, int, boolean columns. If None, will attempt
to use
| everything, then use only numeric data. Not implemented for
Series.

```

```

|
| Returns
| -----
| Series or DataFrame (if level specified)
|
| sub(self, other, axis='columns', level=None, fill_value=None)
| Get Subtraction of dataframe and other, element-wise (binary operator
`sub`).
|
| Equivalent to ``dataframe - other``, but with support to substitute a
fill_value
| for missing data in one of the inputs. With reverse version, `rsub`.
|
| Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
| arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
| Parameters
| -----
| other : scalar, sequence, Series, or DataFrame
| Any single or multiple element data structure, or list-like
object.
|
| axis : {0 or 'index', 1 or 'columns'}
| Whether to compare by the index (0 or 'index') or columns
| (1 or 'columns'). For Series input, axis to match Series index
on.
|
| level : int or label
| Broadcast across a level, matching Index values on the
| passed MultiIndex level.
|
| fill_value : float or None, default None
| Fill existing missing (NaN) values, and any new element needed
for
| successful DataFrame alignment, with this value before
computation.
|
| If data in both corresponding DataFrame locations is missing
| the result will be missing.
|
| Returns
| -----
| DataFrame
| Result of the arithmetic operation.
|
| See Also
| -----
| DataFrame.add : Add DataFrames.

```



```

| DataFrame.sub : Subtract DataFrames.
| DataFrame.mul : Multiply DataFrames.
| DataFrame.div : Divide DataFrames (float division).
| DataFrame.truediv : Divide DataFrames (float division).
| DataFrame.floordiv : Divide DataFrames (integer division).
| DataFrame.mod : Calculate modulo (remainder after division).
| DataFrame.pow : Calculate exponential power.
|
| Notes
| -----
| Mismatched indices will be unioned together.
|
| Examples
| -----
| >>> df = pd.DataFrame({'angles': [0, 3, 4],
| ... 'degrees': [360, 180, 360]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> df
| angles degrees
| circle 0 360
| triangle 3 180
| rectangle 4 360
|
| Add a scalar with operator version which return the same
| results.
|
| >>> df + 1
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| >>> df.add(1)
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| Divide by constant with reverse version.
|
| >>> df.div(10)
| angles degrees
| circle 0.0 36.0
| triangle 0.3 18.0

```

```
| rectangle 0.4 36.0
|
| >>> df.rdiv(10)
| angles degrees
| circle inf 0.027778
| triangle 3.333333 0.055556
| rectangle 2.500000 0.027778
|
| Subtract a list and Series by axis with operator version.
|
| >>> df - [1, 2]
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub([1, 2], axis='columns')
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
'rectangle'])),
| ... axis='index')
| angles degrees
| circle -1 359
| triangle 2 179
| rectangle 3 359
|
| Multiply a DataFrame of different shape with operator version.
|
| >>> other = pd.DataFrame({'angles': [0, 3, 4]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> other
| angles
| circle 0
| triangle 3
| rectangle 4
|
| >>> df * other
| angles degrees
| circle 0 NaN
| triangle 9 NaN
```

```

| rectangle 16 NaN
|
| >>> df.mul(other, fill_value=0)
|
| angles degrees
| circle 0 0.0
| triangle 9 0.0
| rectangle 16 0.0
|
| Divide by a MultiIndex by level.
|
| >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ... 'degrees': [360, 180, 360, 360, 540,
720]}),
| ... index=[['A', 'A', 'A', 'B', 'B',
'B'],
| ... ['circle', 'triangle',
'rectangle',
| ... 'square', 'pentagon',
'hexagon']])
|
| >>> df_multindex
|
| angles degrees
| A circle 0 360
| triangle 3 180
| rectangle 4 360
| B square 4 360
| pentagon 5 540
| hexagon 6 720
|
| >>> df.div(df_multindex, level=1, fill_value=0)
|
| angles degrees
| A circle NaN 1.0
| triangle 1.0 1.0
| rectangle 1.0 1.0
| B square 0.0 0.0
| pentagon 0.0 0.0
| hexagon 0.0 0.0
|
| subtract = sub(self, other, axis='columns', level=None, fill_value=None)
|
| sum(self, axis=None, skipna=None, level=None, numeric_only=None,
min_count=0, **kwargs)
|
| Return the sum of the values for the requested axis.
|
| This is equivalent to the method ``numpy.sum``.

```

```

|
| Parameters
| -----
| axis : {index (0), columns (1)}
| Axis for the function to be applied on.
| skipna : bool, default True
| Exclude NA/null values when computing the result.
| level : int or level name, default None
| If the axis is a MultiIndex (hierarchical), count along a
| particular level, collapsing into a Series.
| numeric_only : bool, default None
| Include only float, int, boolean columns. If None, will attempt
to use
| everything, then use only numeric data. Not implemented for
Series.
| min_count : int, default 0
| The required number of valid values to perform the operation. If
fewer than
| ``min_count`` non-NA values are present the result will be NA.
|
| .. versionadded :: 0.22.0
|
| Added with the default being 0. This means the sum of an all-
NA
| or empty Series is 0, and the product of an all-NA or empty
| Series is 1.
| **kwargs
| Additional keyword arguments to be passed to the function.
|
| Returns
| -----
| Series or DataFrame (if level specified)
|
| See Also
| -----
| Series.sum : Return the sum.
| Series.min : Return the minimum.
| Series.max : Return the maximum.
| Series.idxmin : Return the index of the minimum.
| Series.idxmax : Return the index of the maximum.
| DataFrame.sum : Return the sum over the requested axis.
| DataFrame.min : Return the minimum over the requested axis.
| DataFrame.max : Return the maximum over the requested axis.

```

```

| DataFrame.idxmin : Return the index of the minimum over the requested
axis.
| DataFrame.idxmax : Return the index of the maximum over the requested
axis.
|
| Examples
| -----
| >>> idx = pd.MultiIndex.from_arrays([
| ... ['warm', 'warm', 'cold', 'cold'],
| ... ['dog', 'falcon', 'fish', 'spider']],
| ... names=['blooded', 'animal'])
| >>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
| >>> s
| blooded animal
| warm dog 4
| falcon 2
| cold fish 0
| spider 8
| Name: legs, dtype: int64
|
| >>> s.sum()
| 14
|
| Sum using level names, as well as indices.
|
| >>> s.sum(level='blooded')
| blooded
| warm 6
| cold 8
| Name: legs, dtype: int64
|
| >>> s.sum(level=0)
| blooded
| warm 6
| cold 8
| Name: legs, dtype: int64
|
| By default, the sum of an empty or all-NA Series is ``0``.
|
| >>> pd.Series([]).sum() # min_count=0 is the default
| 0.0
|
| This can be controlled with the ``min_count`` parameter. For example,
if

```

```

| you'd like the sum of an empty series to be NaN, pass
| ``min_count=1``.
|
| >>> pd.Series([]).sum(min_count=1)
| nan
|
| Thanks to the ``skipna`` parameter, ``min_count`` handles all-NA and
| empty series identically.
|
| >>> pd.Series([np.nan]).sum()
| 0.0
|
| >>> pd.Series([np.nan]).sum(min_count=1)
| nan
|
| swaplevel(self, i=-2, j=-1, axis=0)
| Swap levels i and j in a MultiIndex on a particular axis.
|
| Parameters
| -----
| i, j : int, string (can be mixed)
| Level of index to be swapped. Can pass level name as string.
|
| Returns
| -----
| DataFrame
|
| .. versionchanged:: 0.18.1
|
| The indexes ``i`` and ``j`` are now optional, and default to
| the two innermost levels of the index.
|
| to_dict(self, orient='dict', into=<class 'dict'>)
| Convert the DataFrame to a dictionary.
|
| The type of the key-value pairs can be customized with the parameters
| (see below).
|
| Parameters
| -----
| orient : str {'dict', 'list', 'series', 'split', 'records', 'index'}
| Determines the type of the values of the dictionary.
|
| - 'dict' (default) : dict like {column -> {index -> value}}

```

```

| - 'list' : dict like {column -> [values]}
| - 'series' : dict like {column -> Series(values)}
| - 'split' : dict like
| {'index' -> [index], 'columns' -> [columns], 'data' ->
[values]}
| - 'records' : list like
| [{column -> value}, ... , {column -> value}]
| - 'index' : dict like {index -> {column -> value}}
|
| Abbreviations are allowed. `s` indicates `series` and `sp`
| indicates `split`.
|
| into : class, default dict
| The collections.abc.Mapping subclass used for all Mappings
| in the return value. Can be the actual class or an empty
| instance of the mapping type you want. If you want a
| collections.defaultdict, you must pass it initialized.
|
| .. versionadded:: 0.21.0
|
| Returns
| -----
| dict, list or collections.abc.Mapping
| Return a collections.abc.Mapping object representing the
DataFrame.
| The resulting transformation depends on the `orient` parameter.
|
| See Also
| -----
| DataFrame.from_dict: Create a DataFrame from a dictionary.
| DataFrame.to_json: Convert a DataFrame to JSON format.
|
| Examples
| -----
| >>> df = pd.DataFrame({'col1': [1, 2],
| ... 'col2': [0.5, 0.75]},
| ... index=['row1', 'row2'])
| >>> df
| col1 col2
| row1 1 0.50
| row2 2 0.75
| >>> df.to_dict()
| {'col1': {'row1': 1, 'row2': 2}, 'col2': {'row1': 0.5, 'row2': 0.75}}
|

```

```

| You can specify the return orientation.
|
| >>> df.to_dict('series')
| {'col1': row1 1
| row2 2
| Name: col1, dtype: int64,
| 'col2': row1 0.50
| row2 0.75
| Name: col2, dtype: float64}
|
| >>> df.to_dict('split')
| {'index': ['row1', 'row2'], 'columns': ['col1', 'col2'],
| 'data': [[1, 0.5], [2, 0.75]]}
|
| >>> df.to_dict('records')
| [{'col1': 1, 'col2': 0.5}, {'col1': 2, 'col2': 0.75}]
|
| >>> df.to_dict('index')
| {'row1': {'col1': 1, 'col2': 0.5}, 'row2': {'col1': 2, 'col2': 0.75}}
|
| You can also specify the mapping type.
|
| >>> from collections import OrderedDict, defaultdict
| >>> df.to_dict(into=OrderedDict)
| OrderedDict([('col1', OrderedDict([('row1', 1), ('row2', 2)])),
| ('col2', OrderedDict([('row1', 0.5), ('row2', 0.75)]))])
|
| If you want a `defaultdict`, you need to initialize it:
|
| >>> dd = defaultdict(list)
| >>> df.to_dict('records', into=dd)
| [defaultdict(<class 'list'>, {'col1': 1, 'col2': 0.5}),
| defaultdict(<class 'list'>, {'col1': 2, 'col2': 0.75})]
|
| to_feather(self, fname)
| Write out the binary feather-format for DataFrames.
|
| .. versionadded:: 0.20.0
|
| Parameters
| -----
| fname : str
| string file path
|

```



```

| to_gbq(self, destination_table, project_id=None, chunksize=None,
reauth=False, if_exists='fail', auth_local_webserver=False,
table_schema=None, location=None, progress_bar=True, credentials=None,
verbose=None, private_key=None)
|
| Write a DataFrame to a Google BigQuery table.
|
| This function requires the `pandas-gbq` package
| <https://pandas-gbq.readthedocs.io>`__`.
|
| See the `How to authenticate with Google BigQuery
| <https://pandas-gbq.readthedocs.io/en/latest/howto/authentication.html>`
| guide for authentication instructions.
|
| Parameters
| -----
| destination_table : str
| Name of table to be written, in the form ``dataset.tablename``.
| project_id : str, optional
| Google BigQuery Account project ID. Optional when available from
| the environment.
| chunksize : int, optional
| Number of rows to be inserted in each chunk from the dataframe.
| Set to ``None`` to load the whole dataframe at once.
| reauth : bool, default False
| Force Google BigQuery to re-authenticate the user. This is useful
| if multiple accounts are used.
| if_exists : str, default 'fail'
| Behavior when the destination table exists. Value can be one of:
|
| ``'fail'``
| If table exists, do nothing.
| ``'replace'``
| If table exists, drop it, recreate it, and insert data.
| ``'append'``
| If table exists, insert data. Create if does not exist.
| auth_local_webserver : bool, default False
| Use the `local webserver flow` instead of the `console flow`
| when getting user credentials.
|
| .. _local webserver flow:
| http://google-auth-oauthlib.readthedocs.io/en/latest/reference/google_auth_oauthlib.flow.InstalledAppFlow.run_local_server

```

```

| .. _console flow:
| http://google-auth-
oauthlib.readthedocs.io/en/latest/reference/google_auth_oauthlib.flow.html#go
ogle_auth_oauthlib.flow.InstalledAppFlow.run_console
|
| *New in version 0.2.0 of pandas-gbq*.
| table_schema : list of dicts, optional
| List of BigQuery table fields to which according DataFrame
| columns conform to, e.g. ``[{'name': 'col1', 'type':
| 'STRING'},...]``. If schema is not provided, it will be
| generated according to dtypes of DataFrame columns. See
| BigQuery API documentation on available names of a field.
|
| *New in version 0.3.1 of pandas-gbq*.
| location : str, optional
| Location where the load job should run. See the `BigQuery
locations
| documentation
| <https://cloud.google.com/bigquery/docs/dataset-locations>`__ for
a
| list of available locations. The location must match that of the
| target dataset.
|
| *New in version 0.5.0 of pandas-gbq*.
| progress_bar : bool, default True
| Use the library `tqdm` to show the progress bar for the upload,
| chunk by chunk.
|
| *New in version 0.5.0 of pandas-gbq*.
| credentials : google.auth.credentials.Credentials, optional
| Credentials for accessing Google APIs. Use this parameter to
| override default credentials, such as to use Compute Engine
| :class:`google.auth.compute_engine.Credentials` or Service
| Account :class:`google.oauth2.service_account.Credentials`
| directly.
|
| *New in version 0.8.0 of pandas-gbq*.
|
| .. versionadded:: 0.24.0
| verbose : bool, deprecated
| Deprecated in pandas-gbq version 0.4.0. Use the `logging module
| to adjust verbosity instead
| <https://pandas-
gbq.readthedocs.io/en/latest/intro.html#logging>`__.

```

```

| private_key : str, deprecated
| Deprecated in pandas-gbq version 0.8.0. Use the ``credentials``
| parameter and
|
:func:`google.oauth2.service_account.Credentials.from_service_account_info`
| or
|
:func:`google.oauth2.service_account.Credentials.from_service_account_file`
| instead.
|
| Service account private key in JSON format. Can be file path
| or string contents. This is useful for remote server
| authentication (eg. Jupyter/IPython notebook on remote host).
|
| See Also
| -----
| pandas_gbq.to_gbq : This function in the pandas-gbq library.
| read_gbq : Read a DataFrame from Google BigQuery.
|
| to_html(self, buf=None, columns=None, col_space=None, header=True,
index=True, na_rep='NaN', formatters=None, float_format=None, sparsify=None,
index_names=True, justify=None, max_rows=None, max_cols=None,
show_dimensions=False, decimal='.', bold_rows=True, classes=None,
escape=True, notebook=False, border=None, table_id=None, render_links=False)
| Render a DataFrame as an HTML table.
|
| Parameters
| -----
| buf : StringIO-like, optional
| Buffer to write to.
| columns : sequence, optional, default None
| The subset of columns to write. Writes all columns by default.
| col_space : str or int, optional
| The minimum width of each column in CSS length units. An int is
assumed to be px units.
|
| .. versionadded:: 0.25.0
| Ability to use str.
| header : bool, optional
| Whether to print column labels, default True.
| index : bool, optional, default True
| Whether to print index (row) labels.
| na_rep : str, optional, default 'NaN'
| String representation of NAN to use.

```

```

| formatters : list or dict of one-param. functions, optional
| Formatter functions to apply to columns' elements by position or
| name.
| The result of each function must be a unicode string.
| List must be of length equal to the number of columns.
| float_format : one-parameter function, optional, default None
| Formatter function to apply to columns' elements if they are
| floats. The result of this function must be a unicode string.
| sparsify : bool, optional, default True
| Set to False for a DataFrame with a hierarchical index to print
| every multiindex key at each row.
| index_names : bool, optional, default True
| Prints the names of the indexes.
| justify : str, default None
| How to justify the column labels. If None uses the option from
| the print configuration (controlled by set_option), 'right' out
| of the box. Valid values are
|
| * left
| * right
| * center
| * justify
| * justify-all
| * start
| * end
| * inherit
| * match-parent
| * initial
| * unset.
| max_rows : int, optional
| Maximum number of rows to display in the console.
| min_rows : int, optional
| The number of rows to display in the console in a truncated repr
| (when number of rows is above `max_rows`).
| max_cols : int, optional
| Maximum number of columns to display in the console.
| show_dimensions : bool, default False
| Display DataFrame dimensions (number of rows by number of
columns).
| decimal : str, default '.'
| Character recognized as decimal separator, e.g. ',' in Europe.
|
| .. versionadded:: 0.18.0
|

```

```

| bold_rows : bool, default True
| Make the row labels bold in the output.
| classes : str or list or tuple, default None
| CSS class(es) to apply to the resulting html table.
| escape : bool, default True
| Convert the characters <, >, and & to HTML-safe sequences.
| notebook : {True, False}, default False
| Whether the generated HTML is for IPython Notebook.
| border : int
| A ``border=border`` attribute is included in the opening
| ``<table>`` tag. Default ``pd.options.display.html.border``.
|
| .. versionadded:: 0.19.0
|
| table_id : str, optional
| A css id is included in the opening ``<table>`` tag if specified.
|
| .. versionadded:: 0.23.0
|
| render_links : bool, default False
| Convert URLs to HTML links.
|
| .. versionadded:: 0.24.0
|
| Returns
| -----
| str (or unicode, depending on data and options)
| String representation of the dataframe.
|
| See Also
| -----
| to_string : Convert DataFrame to a string.
|
| to_numpy(self, dtype=None, copy=False)
| Convert the DataFrame to a NumPy array.
|
| .. versionadded:: 0.24.0
|
| By default, the dtype of the returned array will be the common NumPy
| dtype of all types in the DataFrame. For example, if the dtypes are
| ``float16`` and ``float32``, the results dtype will be ``float32``.
| This may require copying data and coercing values, which may be
| expensive.
|
|

```

```

| Parameters
| -----
| dtype : str or numpy.dtype, optional
| The dtype to pass to :meth:`numpy.asarray`
| copy : bool, default False
| Whether to ensure that the returned value is a not a view on
| another array. Note that ``copy=False`` does not *ensure* that
| ``to_numpy()`` is no-copy. Rather, ``copy=True`` ensure that
| a copy is made, even if not strictly necessary.
|
| Returns
| -----
| numpy.ndarray
|
| See Also
| -----
| Series.to_numpy : Similar method for Series.
|
| Examples
| -----
| >>> pd.DataFrame({"A": [1, 2], "B": [3, 4]}).to_numpy()
| array([[1, 3],
| [2, 4]])
|
| With heterogenous data, the lowest common type will have to
| be used.
|
| >>> df = pd.DataFrame({"A": [1, 2], "B": [3.0, 4.5]})
| >>> df.to_numpy()
| array([[1. , 3.],
| [2. , 4.5]])
|
| For a mix of numeric and non-numeric types, the output array will
| have object dtype.
|
| >>> df['C'] = pd.date_range('2000', periods=2)
| >>> df.to_numpy()
| array([[1, 3.0, Timestamp('2000-01-01 00:00:00')],
| [2, 4.5, Timestamp('2000-01-02 00:00:00')]], dtype=object)
|
| to_parquet(self, fname, engine='auto', compression='snappy', index=None,
partition_cols=None, **kwargs)
| Write a DataFrame to the binary parquet format.
|

```

```

| .. versionadded:: 0.21.0
|
| This function writes the dataframe as a `parquet file
| <https://parquet.apache.org/>`_. You can choose different parquet
| backends, and have the option of compression. See
| :ref:`the user guide <io.parquet>` for more details.
|
| Parameters
| -----
| fname : str
| File path or Root Directory path. Will be used as Root Directory
| path while writing a partitioned dataset.
|
| .. versionchanged:: 0.24.0
|
| engine : {'auto', 'pyarrow', 'fastparquet'}, default 'auto'
| Parquet library to use. If 'auto', then the option
| ``io.parquet.engine`` is used. The default ``io.parquet.engine``
| behavior is to try 'pyarrow', falling back to 'fastparquet' if
| 'pyarrow' is unavailable.
| compression : {'snappy', 'gzip', 'brotli', None}, default 'snappy'
| Name of the compression to use. Use ``None`` for no compression.
| index : bool, default None
| If ``True``, include the dataframe's index(es) in the file
output.
| If ``False``, they will not be written to the file. If ``None``,
| the behavior depends on the chosen engine.
|
| .. versionadded:: 0.24.0
|
| partition_cols : list, optional, default None
| Column names by which to partition the dataset
| Columns are partitioned in the order they are given
|
| .. versionadded:: 0.24.0
|
| **kwargs
| Additional arguments passed to the parquet library. See
| :ref:`pandas io <io.parquet>` for more details.
|
| See Also
| -----
| read_parquet : Read a parquet file.
| DataFrame.to_csv : Write a csv file.

```

```

| DataFrame.to_sql : Write to a sql table.
| DataFrame.to_hdf : Write to hdf.
|
| Notes
| -----
| This function requires either the `fastparquet`
| <https://pypi.org/project/fastparquet> or `pyarrow`
| <https://arrow.apache.org/docs/python/> library.
|
| Examples
| -----
| >>> df = pd.DataFrame(data={'col1': [1, 2], 'col2': [3, 4]})
| >>> df.to_parquet('df.parquet.gzip',
| ... compression='gzip') # doctest: +SKIP
| >>> pd.read_parquet('df.parquet.gzip') # doctest: +SKIP
|
| col1 col2
| 0 1 3
| 1 2 4
|
| to_period(self, freq=None, axis=0, copy=True)
| Convert DataFrame from DatetimeIndex to PeriodIndex with desired
| frequency (inferred from index if not passed).
|
| Parameters
| -----
| freq : str, default
| Frequency of the PeriodIndex.
| axis : {0 or 'index', 1 or 'columns'}, default 0
| The axis to convert (the index by default).
| copy : bool, default True
| If False then underlying input data is not copied.
|
| Returns
| -----
| TimeSeries with PeriodIndex
|
| to_records(self, index=True, convert_datetime64=None, column_dtypes=None,
| index_dtypes=None)
| Convert DataFrame to a NumPy record array.
|
| Index will be included as the first field of the record array if
| requested.
|
| Parameters

```



```

| -----
| index : bool, default True
| Include index in resulting record array, stored in 'index'
| field or using the index label, if set.
| convert_datetime64 : bool, default None
| .. deprecated:: 0.23.0
|
| Whether to convert the index to datetime.datetime if it is a
| DatetimeIndex.
| column_dtypes : str, type, dict, default None
| .. versionadded:: 0.24.0
|
| If a string or type, the data type to store all columns. If
| a dictionary, a mapping of column names and indices (zero-
indexed)
| to specific data types.
| index_dtypes : str, type, dict, default None
| .. versionadded:: 0.24.0
|
| If a string or type, the data type to store all index levels. If
| a dictionary, a mapping of index level names and indices
| (zero-indexed) to specific data types.
|
| This mapping is applied only if `index=True`.
|
| Returns
| -----
| numpy.recarray
| NumPy ndarray with the DataFrame labels as fields and each row
| of the DataFrame as entries.
|
| See Also
| -----
| DataFrame.from_records: Convert structured or record ndarray
| to DataFrame.
| numpy.recarray: An ndarray that allows field access using
| attributes, analogous to typed columns in a
| spreadsheet.
|
| Examples
| -----
| >>> df = pd.DataFrame({'A': [1, 2], 'B': [0.5, 0.75]},
| ... index=['a', 'b'])
| >>> df

```

```

| A B
| a 1 0.50
| b 2 0.75
| >>> df.to_records()
| rec.array([('a', 1, 0.5), ('b', 2, 0.75)],
| dtype=[('index', 'O'), ('A', '<i8'), ('B', '<f8')])
|
| If the DataFrame index has no label then the recarray field name
| is set to 'index'. If the index has a label then this is used as the
| field name:
|
| >>> df.index = df.index.rename("I")
| >>> df.to_records()
| rec.array([('a', 1, 0.5), ('b', 2, 0.75)],
| dtype=[('I', 'O'), ('A', '<i8'), ('B', '<f8')])
|
| The index can be excluded from the record array:
|
| >>> df.to_records(index=False)
| rec.array([(1, 0.5), (2, 0.75)],
| dtype=[('A', '<i8'), ('B', '<f8')])
|
| Data types can be specified for the columns:
|
| >>> df.to_records(column_dtypes={"A": "int32"})
| rec.array([('a', 1, 0.5), ('b', 2, 0.75)],
| dtype=[('I', 'O'), ('A', '<i4'), ('B', '<f8')])
|
| As well as for the index:
|
| >>> df.to_records(index_dtypes="<S2")
| rec.array([(b'a', 1, 0.5), (b'b', 2, 0.75)],
| dtype=[('I', 'S2'), ('A', '<i8'), ('B', '<f8')])
|
| >>> index_dtypes = "<S{}".format(df.index.str.len().max())
| >>> df.to_records(index_dtypes=index_dtypes)
| rec.array([(b'a', 1, 0.5), (b'b', 2, 0.75)],
| dtype=[('I', 'S1'), ('A', '<i8'), ('B', '<f8')])
|
| to_sparse(self, fill_value=None, kind='block')
| Convert to SparseDataFrame.
|
| .. deprecated:: 0.25.0
|

```

```

| Implement the sparse version of the DataFrame meaning that any data
| matching a specific value it's omitted in the representation.
| The sparse DataFrame allows for a more efficient storage.
|
| Parameters
| -----
| fill_value : float, default None
| The specific value that should be omitted in the representation.
| kind : {'block', 'integer'}, default 'block'
| The kind of the SparseIndex tracking where data is not equal to
| the fill value:
|
| - 'block' tracks only the locations and sizes of blocks of data.
| - 'integer' keeps an array with all the locations of the data.
|
| In most cases 'block' is recommended, since it's more memory
| efficient.
|
| Returns
| -----
| SparseDataFrame
| The sparse representation of the DataFrame.
|
| See Also
| -----
| DataFrame.to_dense :
| Converts the DataFrame back to the its dense form.
|
| Examples
| -----
| >>> df = pd.DataFrame([(np.nan, np.nan),
| ... (1., np.nan),
| ... (np.nan, 1.)])
| >>> df
| 0 1
| 0 NaN NaN
| 1 1.0 NaN
| 2 NaN 1.0
| >>> type(df)
| <class 'pandas.core.frame.DataFrame'>
|
| >>> sdf = df.to_sparse() # doctest: +SKIP
| >>> sdf # doctest: +SKIP
| 0 1

```

```

| 0 NaN NaN
| 1 1.0 NaN
| 2 NaN 1.0
| >>> type(sdf) # doctest: +SKIP
| <class 'pandas.core.sparse.frame.SparseDataFrame'>
|
| to_stata(self, fname, convert_dates=None, write_index=True,
encoding='latin-1', byteorder=None, time_stamp=None, data_label=None,
variable_labels=None, version=114, convert_str1=None)
| Export DataFrame object to Stata dta format.
|
| Writes the DataFrame to a Stata dataset file.
| "dta" files contain a Stata dataset.
|
| Parameters
| -----
| fname : str, buffer or path object
| String, path object (pathlib.Path or py._path.local.LocalPath) or
| object implementing a binary write() function. If using a buffer
| then the buffer will not be automatically closed after the file
| data has been written.
| convert_dates : dict
| Dictionary mapping columns containing datetime types to stata
| internal format to use when writing the dates. Options are 'tc',
| 'td', 'tm', 'tw', 'th', 'tq', 'ty'. Column can be either an
integer
| or a name. Datetime columns that do not have a conversion type
| specified will be converted to 'tc'. Raises NotImplementedError
if
| a datetime column has timezone information.
| write_index : bool
| Write the index to Stata dataset.
| encoding : str
| Default is latin-1. Unicode is not supported.
| byteorder : str
| Can be ">", "<", "little", or "big". default is `sys.byteorder`.
| time_stamp : datetime
| A datetime to use as file creation date. Default is the current
| time.
| data_label : str, optional
| A label for the data set. Must be 80 characters or smaller.
| variable_labels : dict
| Dictionary containing columns as keys and variable labels as
| values. Each label must be 80 characters or smaller.

```

```

|
| .. versionadded:: 0.19.0
|
| version : {114, 117}, default 114
| Version to use in the output dta file. Version 114 can be used
| read by Stata 10 and later. Version 117 can be read by Stata 13
| or later. Version 114 limits string variables to 244 characters
or
| fewer while 117 allows strings with lengths up to 2,000,000
| characters.
|
| .. versionadded:: 0.23.0
|
| convert_strl : list, optional
| List of column names to convert to string columns to Stata StrL
| format. Only available if version is 117. Storing strings in the
| StrL format can produce smaller dta files if strings have more
than
| 8 characters and values are repeated.
|
| .. versionadded:: 0.23.0
|
| Raises
| -----
| NotImplementedError
| * If datetimes contain timezone information
| * Column dtype is not representable in Stata
| ValueError
| * Columns listed in convert_dates are neither datetime64[ns]
| or datetime.datetime
| * Column listed in convert_dates is not in DataFrame
| * Categorical label contains more than 32,000 characters
|
| .. versionadded:: 0.19.0
|
| See Also
| -----
| read_stata : Import Stata data files.
| io.stata.StataWriter : Low-level writer for Stata data files.
| io.stata.StataWriter117 : Low-level writer for version 117 files.
|
| Examples
| -----
| >>> df = pd.DataFrame({'animal': ['falcon', 'parrot', 'falcon',

```

```

| ... 'parrot'],
| ... 'speed': [350, 18, 361, 15]})
| >>> df.to_stata('animals.dta') # doctest: +SKIP
|
| to_string(self, buf=None, columns=None, col_space=None, header=True,
index=True, na_rep='NaN', formatters=None, float_format=None, sparsify=None,
index_names=True, justify=None, max_rows=None, min_rows=None, max_cols=None,
show_dimensions=False, decimal='.', line_width=None)
|
| Render a DataFrame to a console-friendly tabular output.
|
| Parameters
| -----
| buf : StringIO-like, optional
| Buffer to write to.
| columns : sequence, optional, default None
| The subset of columns to write. Writes all columns by default.
| col_space : int, optional
| The minimum width of each column.
| header : bool, optional
| Write out the column names. If a list of strings is given, it is
assumed to be aliases for the column names.
| index : bool, optional, default True
| Whether to print index (row) labels.
| na_rep : str, optional, default 'NaN'
| String representation of NAN to use.
| formatters : list or dict of one-param. functions, optional
| Formatter functions to apply to columns' elements by position or
name.
| The result of each function must be a unicode string.
| List must be of length equal to the number of columns.
| float_format : one-parameter function, optional, default None
| Formatter function to apply to columns' elements if they are
floats. The result of this function must be a unicode string.
| sparsify : bool, optional, default True
| Set to False for a DataFrame with a hierarchical index to print
every multiindex key at each row.
| index_names : bool, optional, default True
| Prints the names of the indexes.
| justify : str, default None
| How to justify the column labels. If None uses the option from
the print configuration (controlled by set_option), 'right' out
of the box. Valid values are
|
| * left

```

```

| * right
| * center
| * justify
| * justify-all
| * start
| * end
| * inherit
| * match-parent
| * initial
| * unset.
|
| max_rows : int, optional
| Maximum number of rows to display in the console.
|
| min_rows : int, optional
| The number of rows to display in the console in a truncated repr
| (when number of rows is above `max_rows`).
|
| max_cols : int, optional
| Maximum number of columns to display in the console.
|
| show_dimensions : bool, default False
| Display DataFrame dimensions (number of rows by number of
columns).
|
| decimal : str, default '.'
| Character recognized as decimal separator, e.g. ',' in Europe.
|
| .. versionadded:: 0.18.0
|
| line_width : int, optional
| Width to wrap a line in characters.
|
| Returns
| -----
| str (or unicode, depending on data and options)
| String representation of the dataframe.
|
| See Also
| -----
| to_html : Convert DataFrame to HTML.
|
| Examples
| -----
| >>> d = {'col1': [1, 2, 3], 'col2': [4, 5, 6]}
| >>> df = pd.DataFrame(d)
| >>> print(df.to_string())
|
| col1 col2
| 0 1 4

```

```

| 1 2 5
| 2 3 6
|
| to_timestamp(self, freq=None, how='start', axis=0, copy=True)
| Cast to DatetimeIndex of timestamps, at *beginning* of period.
|
| Parameters
| -----
| freq : str, default frequency of PeriodIndex
| Desired frequency.
| how : {'s', 'e', 'start', 'end'}
| Convention for converting period to timestamp; start of period
| vs. end.
| axis : {0 or 'index', 1 or 'columns'}, default 0
| The axis to convert (the index by default).
| copy : bool, default True
| If False then underlying input data is not copied.
|
| Returns
| -----
| DataFrame with DatetimeIndex
|
| transform(self, func, axis=0, *args, **kwargs)
| Call ``func`` on self producing a DataFrame with transformed values
| and that has the same axis length as self.
|
| .. versionadded:: 0.20.0
|
| Parameters
| -----
| func : function, str, list or dict
| Function to use for transforming the data. If a function, must
either
| work when passed a DataFrame or when passed to DataFrame.apply.
|
| Accepted combinations are:
|
| - function
| - string function name
| - list of functions and/or function names, e.g. ``[np.exp.
'sqrt']``
|
| - dict of axis labels -> functions, function names or list of
such.
|
| axis : {0 or 'index', 1 or 'columns'}, default 0

```



```

| If 0 or 'index': apply function to each column.
| If 1 or 'columns': apply function to each row.
| *args
| Positional arguments to pass to `func`.
| **kwargs
| Keyword arguments to pass to `func`.
|
| Returns
| -----
| DataFrame
| A DataFrame that must have the same length as self.
|
| Raises
| -----
| ValueError : If the returned DataFrame has a different length than
self.
|
| See Also
| -----
| DataFrame.agg : Only perform aggregating type operations.
| DataFrame.apply : Invoke function on a DataFrame.
|
| Examples
| -----
| >>> df = pd.DataFrame({'A': range(3), 'B': range(1, 4)})
| >>> df
| A B
| 0 0 1
| 1 1 2
| 2 2 3
| >>> df.transform(lambda x: x + 1)
| A B
| 0 1 2
| 1 2 3
| 2 3 4
|
| Even though the resulting DataFrame must have the same length as the
| input DataFrame, it is possible to provide several input functions:
|
| >>> s = pd.Series(range(3))
| >>> s
| 0 0
| 1 1
| 2 2

```

```

| dtype: int64
| >>> s.transform([np.sqrt, np.exp])
| sqrt exp
| 0 0.000000 1.000000
| 1 1.000000 2.718282
| 2 1.414214 7.389056
|
| transpose(self, *args, **kwargs)
| Transpose index and columns.
|
| Reflect the DataFrame over its main diagonal by writing rows as
columns
| and vice-versa. The property :attr:`.T` is an accessor to the method
| :meth:`transpose`.
|
| Parameters
| -----
| copy : bool, default False
| If True, the underlying data is copied. Otherwise (default), no
| copy is made if possible.
| *args, **kwargs
| Additional keywords have no effect but might be accepted for
| compatibility with numpy.
|
| Returns
| -----
| DataFrame
| The transposed DataFrame.
|
| See Also
| -----
| numpy.transpose : Permute the dimensions of a given array.
|
| Notes
| -----
| Transposing a DataFrame with mixed dtypes will result in a
homogeneous
| DataFrame with the `object` dtype. In such a case, a copy of the data
| is always made.
|
| Examples
| -----
| **Square DataFrame with homogeneous dtype**
|

```

```
| >>> d1 = {'col1': [1, 2], 'col2': [3, 4]}
| >>> df1 = pd.DataFrame(data=d1)
| >>> df1
| col1 col2
| 0 1 3
| 1 2 4
|
| >>> df1_transposed = df1.T # or df1.transpose()
| >>> df1_transposed
| 0 1
| col1 1 2
| col2 3 4
```

| When the dtype is homogeneous in the original DataFrame, we get a  
| transposed DataFrame with the same dtype:

```
| >>> df1.dtypes
| col1 int64
| col2 int64
| dtype: object
| >>> df1_transposed.dtypes
| 0 int64
| 1 int64
| dtype: object
```

| **\*\*Non-square DataFrame with mixed dtypes\*\***

```
| >>> d2 = {'name': ['Alice', 'Bob'],
| ... 'score': [9.5, 8],
| ... 'employed': [False, True],
| ... 'kids': [0, 0]}
| >>> df2 = pd.DataFrame(data=d2)
| >>> df2
| name score employed kids
| 0 Alice 9.5 False 0
| 1 Bob 8.0 True 0
|
| >>> df2_transposed = df2.T # or df2.transpose()
| >>> df2_transposed
| 0 1
| name Alice Bob
| score 9.5 8
| employed False True
| kids 0 0
```

```

|
| When the DataFrame has mixed dtypes, we get a transposed DataFrame
with
| the `object` dtype:
|
| >>> df2.dtypes
| name object
| score float64
| employed bool
| kids int64
| dtype: object
| >>> df2_transposed.dtypes
| 0 object
| 1 object
| dtype: object
|
| truediv(self, other, axis='columns', level=None, fill_value=None)
| Get Floating division of dataframe and other, element-wise (binary
operator `truediv`).
|
| Equivalent to ``dataframe / other``, but with support to substitute a
fill_value
| for missing data in one of the inputs. With reverse version,
`rtruediv`.
|
| Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.
|
| Parameters
| -----
| other : scalar, sequence, Series, or DataFrame
| Any single or multiple element data structure, or list-like
object.
| axis : {0 or 'index', 1 or 'columns'}
| Whether to compare by the index (0 or 'index') or columns
| (1 or 'columns'). For Series input, axis to match Series index
on.
| level : int or label
| Broadcast across a level, matching Index values on the
| passed MultiIndex level.
| fill_value : float or None, default None
| Fill existing missing (NaN) values, and any new element needed
for

```

```

| successful DataFrame alignment, with this value before
computation.
| If data in both corresponding DataFrame locations is missing
| the result will be missing.
|
| Returns
| -----
| DataFrame
| Result of the arithmetic operation.
|
| See Also
| -----
| DataFrame.add : Add DataFrames.
| DataFrame.sub : Subtract DataFrames.
| DataFrame.mul : Multiply DataFrames.
| DataFrame.div : Divide DataFrames (float division).
| DataFrame.truediv : Divide DataFrames (float division).
| DataFrame.floordiv : Divide DataFrames (integer division).
| DataFrame.mod : Calculate modulo (remainder after division).
| DataFrame.pow : Calculate exponential power.
|
| Notes
| -----
| Mismatched indices will be unioned together.
|
| Examples
| -----
| >>> df = pd.DataFrame({'angles': [0, 3, 4],
| ... 'degrees': [360, 180, 360]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> df
|
| angles degrees
| circle 0 360
| triangle 3 180
| rectangle 4 360
|
| Add a scalar with operator version which return the same
| results.
|
| >>> df + 1
|
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361

```

```
|
| >>> df.add(1)
| angles degrees
| circle 1 361
| triangle 4 181
| rectangle 5 361
|
| Divide by constant with reverse version.
|
| >>> df.div(10)
| angles degrees
| circle 0.0 36.0
| triangle 0.3 18.0
| rectangle 0.4 36.0
|
| >>> df.rdiv(10)
| angles degrees
| circle inf 0.027778
| triangle 3.333333 0.055556
| rectangle 2.500000 0.027778
|
| Subtract a list and Series by axis with operator version.
|
| >>> df - [1, 2]
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub([1, 2], axis='columns')
| angles degrees
| circle -1 358
| triangle 2 178
| rectangle 3 358
|
| >>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle',
|'rectangle'])),
| ... axis='index')
| angles degrees
| circle -1 359
| triangle 2 179
| rectangle 3 359
|
| Multiply a DataFrame of different shape with operator version.
```

```
|
| >>> other = pd.DataFrame({'angles': [0, 3, 4]},
| ... index=['circle', 'triangle', 'rectangle'])
| >>> other
| angles
| circle 0
| triangle 3
| rectangle 4
|
| >>> df * other
| angles degrees
| circle 0 NaN
| triangle 9 NaN
| rectangle 16 NaN
|
| >>> df.mul(other, fill_value=0)
| angles degrees
| circle 0 0.0
| triangle 9 0.0
| rectangle 16 0.0
|
| Divide by a MultiIndex by level.
|
| >>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
| ... 'degrees': [360, 180, 360, 360, 540,
720]},
| ... index=[['A', 'A', 'A', 'B', 'B',
'B'],
| ... ['circle', 'triangle',
'rectangle',
| ... 'square', 'pentagon',
'hexagon']])
| >>> df_multindex
| angles degrees
| A circle 0 360
| triangle 3 180
| rectangle 4 360
| B square 4 360
| pentagon 5 540
| hexagon 6 720
|
| >>> df.div(df_multindex, level=1, fill_value=0)
| angles degrees
| A circle NaN 1.0
```

```

| triangle 1.0 1.0
| rectangle 1.0 1.0
| B square 0.0 0.0
| pentagon 0.0 0.0
| hexagon 0.0 0.0
|
| unstack(self, level=-1, fill_value=None)
| Pivot a level of the (necessarily hierarchical) index labels,
returning
| a DataFrame having a new level of column labels whose inner-most
level
| consists of the pivoted index labels.
|
| If the index is not a MultiIndex, the output will be a Series
| (the analogue of stack when the columns are not a MultiIndex).
|
| The level involved will automatically get sorted.
|
| Parameters
| -----
| level : int, string, or list of these, default -1 (last level)
| Level(s) of index to unstack, can pass level name
| fill_value : replace NaN with this value if the unstack produces
| missing values
|
| .. versionadded:: 0.18.0
|
| Returns
| -----
| Series or DataFrame
|
| See Also
| -----
| DataFrame.pivot : Pivot a table based on column values.
| DataFrame.stack : Pivot a level of the column labels (inverse
operation
| from `unstack`).
|
| Examples
| -----
| >>> index = pd.MultiIndex.from_tuples([('one', 'a'), ('one', 'b'),
| ... ('two', 'a'), ('two', 'b')])
| >>> s = pd.Series(np.arange(1.0, 5.0), index=index)
| >>> s

```



```

| one a 1.0
| b 2.0
| two a 3.0
| b 4.0
| dtype: float64
|
| >>> s.unstack(level=-1)
| a b
| one 1.0 2.0
| two 3.0 4.0
|
| >>> s.unstack(level=0)
| one two
| a 1.0 3.0
| b 2.0 4.0
|
| >>> df = s.unstack(level=0)
| >>> df.unstack()
| one a 1.0
| b 2.0
| two a 3.0
| b 4.0
| dtype: float64
|
| update(self, other, join='left', overwrite=True, filter_func=None,
errors='ignore')
| Modify in place using non-NA values from another DataFrame.
|
| Aligns on indices. There is no return value.
|
| Parameters
| -----
| other : DataFrame, or object coercible into a DataFrame
| Should have at least one matching index/column label
| with the original DataFrame. If a Series is passed,
| its name attribute must be set, and that will be
| used as the column name to align with the original DataFrame.
| join : {'left'}, default 'left'
| Only left join is implemented, keeping the index and columns of
the
| original object.
| overwrite : bool, default True
| How to handle non-NA values for overlapping keys:
|

```

```

| * True: overwrite original DataFrame's values
| with values from `other`.
| * False: only update values that are NA in
| the original DataFrame.
|
| filter_func : callable(1d-array) -> bool 1d-array, optional
| Can choose to replace values other than NA. Return True for
values
| that should be updated.
| errors : {'raise', 'ignore'}, default 'ignore'
| If 'raise', will raise a ValueError if the DataFrame and `other`
| both contain non-NA data in the same place.
|
| .. versionchanged :: 0.24.0
| Changed from `raise_conflict=False|True`
| to `errors='ignore'|'raise'`.
|
| Returns
| -----
| None : method directly changes calling object
|
| Raises
| -----
| ValueError
| * When `errors='raise'` and there's overlapping non-NA data.
| * When `errors` is not either `ignore` or `raise`
| NotImplementedError
| * If `join != 'left'`
|
| See Also
| -----
| dict.update : Similar method for dictionaries.
| DataFrame.merge : For column(s)-on-columns(s) operations.
|
| Examples
| -----
| >>> df = pd.DataFrame({'A': [1, 2, 3],
| ... 'B': [400, 500, 600]})
| >>> new_df = pd.DataFrame({'B': [4, 5, 6],
| ... 'C': [7, 8, 9]})
| >>> df.update(new_df)
| >>> df
|
| A B
| 0 1 4

```

```
| 1 2 5
| 2 3 6
|
```

| The DataFrame's length does not increase as a result of the update,  
| only values at matching index/column labels are updated.

```
| >>> df = pd.DataFrame({'A': ['a', 'b', 'c'],
| ... 'B': ['x', 'y', 'z']})
| >>> new_df = pd.DataFrame({'B': ['d', 'e', 'f', 'g', 'h', 'i']})
| >>> df.update(new_df)
| >>> df
| A B
| 0 a d
| 1 b e
| 2 c f
|
```

| For Series, it's name attribute must be set.

```
| >>> df = pd.DataFrame({'A': ['a', 'b', 'c'],
| ... 'B': ['x', 'y', 'z']})
| >>> new_column = pd.Series(['d', 'e'], name='B', index=[0, 2])
| >>> df.update(new_column)
| >>> df
| A B
| 0 a d
| 1 b y
| 2 c e
| >>> df = pd.DataFrame({'A': ['a', 'b', 'c'],
| ... 'B': ['x', 'y', 'z']})
| >>> new_df = pd.DataFrame({'B': ['d', 'e']}, index=[1, 2])
| >>> df.update(new_df)
| >>> df
| A B
| 0 a x
| 1 b d
| 2 c e
|
```

| If `other` contains NaNs the corresponding values are not updated  
| in the original dataframe.

```
| >>> df = pd.DataFrame({'A': [1, 2, 3],
| ... 'B': [400, 500, 600]})
| >>> new_df = pd.DataFrame({'B': [4, np.nan, 6]})
| >>> df.update(new_df)
```

```

| >>> df
| A B
| 0 1 4.0
| 1 2 500.0
| 2 3 6.0
|
| var(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None,
**kwargs)
| Return unbiased variance over requested axis.
|
| Normalized by N-1 by default. This can be changed using the ddof
argument
|
| Parameters
| -----
| axis : {index (0), columns (1)}
| skipna : bool, default True
| Exclude NA/null values. If an entire row/column is NA, the result
| will be NA
| level : int or level name, default None
| If the axis is a MultiIndex (hierarchical), count along a
| particular level, collapsing into a Series
| ddof : int, default 1
| Delta Degrees of Freedom. The divisor used in calculations is N -
ddof,
| where N represents the number of elements.
| numeric_only : bool, default None
| Include only float, int, boolean columns. If None, will attempt
to use
| everything, then use only numeric data. Not implemented for
Series.
|
| Returns
| -----
| Series or DataFrame (if level specified)
|
| -----
| Class methods defined here:
|
| from_dict(data, orient='columns', dtype=None, columns=None) from
builtins.type
| Construct DataFrame from dict of array-like or dicts.
|
| Creates DataFrame object from dictionary by columns or by index

```

```

| allowing dtype specification.
|
| Parameters
| -----
| data : dict
| Of the form {field : array-like} or {field : dict}.
| orient : {'columns', 'index'}, default 'columns'
| The "orientation" of the data. If the keys of the passed dict
| should be the columns of the resulting DataFrame, pass 'columns'
| (default). Otherwise if the keys should be rows, pass 'index'.
| dtype : dtype, default None
| Data type to force, otherwise infer.
| columns : list, default None
| Column labels to use when ``orient='index'``. Raises a ValueError
| if used with ``orient='columns'``.
|
| .. versionadded:: 0.23.0
|
| Returns
| -----
| DataFrame
|
| See Also
| -----
| DataFrame.from_records : DataFrame from ndarray (structured
| dtype), list of tuples, dict, or DataFrame.
| DataFrame : DataFrame object creation using constructor.
|
| Examples
| -----
| By default the keys of the dict become the DataFrame columns:
|
| >>> data = {'col_1': [3, 2, 1, 0], 'col_2': ['a', 'b', 'c', 'd']}
| >>> pd.DataFrame.from_dict(data)
| col_1 col_2
| 0 3 a
| 1 2 b
| 2 1 c
| 3 0 d
|
| Specify ``orient='index'`` to create the DataFrame using dictionary
| keys as rows:
|
| >>> data = {'row_1': [3, 2, 1, 0], 'row_2': ['a', 'b', 'c', 'd']}

```

```

| >>> pd.DataFrame.from_dict(data, orient='index')
|
| 0 1 2 3
| row_1 3 2 1 0
| row_2 a b c d
|
|
| When using the 'index' orientation, the column names can be
| specified manually:
|
|
| >>> pd.DataFrame.from_dict(data, orient='index',
| ... columns=['A', 'B', 'C', 'D'])
|
| A B C D
| row_1 3 2 1 0
| row_2 a b c d
|
|
| from_items(items, columns=None, orient='columns') from builtins.type
| Construct a DataFrame from a list of tuples.
|
|
| .. deprecated:: 0.23.0
| `from_items` is deprecated and will be removed in a future version.
| Use :meth:`DataFrame.from_dict(dict(items))` <DataFrame.from_dict>`
| instead.
| :meth:`DataFrame.from_dict(OrderedDict(items))`
| <DataFrame.from_dict>`
| may be used to preserve the key order.
|
|
| Convert (key, value) pairs to DataFrame. The keys will be the axis
| index (usually the columns, but depends on the specified
| orientation). The values should be arrays or Series.
|
|
| Parameters
| -----
| items : sequence of (key, value) pairs
| Values should be arrays or Series.
| columns : sequence of column labels, optional
| Must be passed if orient='index'.
| orient : {'columns', 'index'}, default 'columns'
| The "orientation" of the data. If the keys of the
| input correspond to column labels, pass 'columns'
| (default). Otherwise if the keys correspond to the index,
| pass 'index'.
|
|
| Returns
| -----
| DataFrame

```

```

|
| from_records(data, index=None, exclude=None, columns=None,
coerce_float=False, nrows=None) from builtins.type
| Convert structured or record ndarray to DataFrame.
|
| Parameters
| -----
| data : ndarray (structured dtype), list of tuples, dict, or DataFrame
| index : string, list of fields, array-like
| Field of array to use as the index, alternately a specific set of
| input labels to use
| exclude : sequence, default None
| Columns or fields to exclude
| columns : sequence, default None
| Column names to use. If the passed data do not have names
| associated with them, this argument provides names for the
| columns. Otherwise this argument indicates the order of the
columns
| in the result (any names not found in the data will become all-NA
| columns)
| coerce_float : boolean, default False
| Attempt to convert values of non-string, non-numeric objects
| (like
| decimal.Decimal) to floating point, useful for SQL result sets
| nrows : int, default None
| Number of rows to read if data is an iterator
|
| Returns
| -----
| DataFrame
|
| -----
| Data descriptors defined here:
|
| T
| Transpose index and columns.
|
| Reflect the DataFrame over its main diagonal by writing rows as
columns
| and vice-versa. The property :attr:`.T` is an accessor to the method
| :meth:`transpose`.
|
| Parameters
| -----

```

```

| copy : bool, default False
| If True, the underlying data is copied. Otherwise (default), no
| copy is made if possible.
| *args, **kwargs
| Additional keywords have no effect but might be accepted for
| compatibility with numpy.
|
| Returns
| -----
| DataFrame
| The transposed DataFrame.
|
| See Also
| -----
| numpy.transpose : Permute the dimensions of a given array.
|
| Notes
| -----
| Transposing a DataFrame with mixed dtypes will result in a
homogeneous
| DataFrame with the `object` dtype. In such a case, a copy of the data
| is always made.
|
| Examples
| -----
| **Square DataFrame with homogeneous dtype**
|
| >>> d1 = {'col1': [1, 2], 'col2': [3, 4]}
| >>> df1 = pd.DataFrame(data=d1)
| >>> df1
| col1 col2
| 0 1 3
| 1 2 4
|
| >>> df1_transposed = df1.T # or df1.transpose()
| >>> df1_transposed
| 0 1
| col1 1 2
| col2 3 4
|
| When the dtype is homogeneous in the original DataFrame, we get a
| transposed DataFrame with the same dtype:
|
| >>> df1.dtypes

```



```

| col1 int64
| col2 int64
| dtype: object
| >>> df1_transposed.dtypes
| 0 int64
| 1 int64
| dtype: object
|
| **Non-square DataFrame with mixed dtypes**
|
| >>> d2 = {'name': ['Alice', 'Bob'],
| ... 'score': [9.5, 8],
| ... 'employed': [False, True],
| ... 'kids': [0, 0]}
| >>> df2 = pd.DataFrame(data=d2)
| >>> df2
| name score employed kids
| 0 Alice 9.5 False 0
| 1 Bob 8.0 True 0
|
| >>> df2_transposed = df2.T # or df2.transpose()
| >>> df2_transposed
| 0 1
| name Alice Bob
| score 9.5 8
| employed False True
| kids 0 0
|
| When the DataFrame has mixed dtypes, we get a transposed DataFrame
with
| the `object` dtype:
|
| >>> df2.dtypes
| name object
| score float64
| employed bool
| kids int64
| dtype: object
| >>> df2_transposed.dtypes
| 0 object
| 1 object
| dtype: object
|
| axes

```

```

| Return a list representing the axes of the DataFrame.
|
| It has the row axis labels and column axis labels as the only
members.
| They are returned in that order.
|
| Examples
| -----
| >>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
| >>> df.axes
| [RangeIndex(start=0, stop=2, step=1), Index(['col1', 'col2'],
| dtype='object')]
|
| columns
| The column labels of the DataFrame.
|
| index
| The index (row labels) of the DataFrame.
|
| shape
| Return a tuple representing the dimensionality of the DataFrame.
|
| See Also
| -----
| ndarray.shape
|
| Examples
| -----
| >>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
| >>> df.shape
| (2, 2)
|
| >>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4],
| ... 'col3': [5, 6]})
| >>> df.shape
| (2, 3)
|
| style
| Property returning a Styler object containing methods for
| building a styled HTML representation fo the DataFrame.
|
| See Also
| -----
| io.formats.style.Styler

```

```

|
| -----
| Data and other attributes defined here:
|
| plot = <class 'pandas.plotting._core.PlotAccessor'>
| Make plots of Series or DataFrame using the backend specified by the
| option ``plotting.backend``. By default, matplotlib is used.
|
| Parameters
| -----
| data : Series or DataFrame
| The object for which the method is called
| x : label or position, default None
| Only used if data is a DataFrame.
| y : label, position or list of label, positions, default None
| Allows plotting of one column versus another. Only used if data
is a
| DataFrame.
| kind : str
| - 'line' : line plot (default)
| - 'bar' : vertical bar plot
| - 'barh' : horizontal bar plot
| - 'hist' : histogram
| - 'box' : boxplot
| - 'kde' : Kernel Density Estimation plot
| - 'density' : same as 'kde'
| - 'area' : area plot
| - 'pie' : pie plot
| - 'scatter' : scatter plot
| - 'hexbin' : hexbin plot
| figsize : a tuple (width, height) in inches
| use_index : bool, default True
| Use index as ticks for x axis
| title : string or list
| Title to use for the plot. If a string is passed, print the
string
| at the top of the figure. If a list is passed and `subplots` is
| True, print each item in the list above the corresponding
subplot.
| grid : bool, default None (matlab style default)
| Axis grid lines
| legend : False/True/'reverse'
| Place legend on axis subplots
| style : list or dict

```

```

| matplotlib line style per column
| logx : bool or 'sym', default False
| Use log scaling or symlog scaling on x axis
| .. versionchanged:: 0.25.0
|
| logy : bool or 'sym' default False
| Use log scaling or symlog scaling on y axis
| .. versionchanged:: 0.25.0
|
| loglog : bool or 'sym', default False
| Use log scaling or symlog scaling on both x and y axes
| .. versionchanged:: 0.25.0
|
| xticks : sequence
| Values to use for the xticks
| yticks : sequence
| Values to use for the yticks
| xlim : 2-tuple/list
| ylim : 2-tuple/list
| rot : int, default None
| Rotation for ticks (xticks for vertical, yticks for horizontal
| plots)
| fontsize : int, default None
| Font size for xticks and yticks
| colormap : str or matplotlib colormap object, default None
| Colormap to select colors from. If string, load colormap with
that
| name from matplotlib.
| colorbar : bool, optional
| If True, plot colorbar (only relevant for 'scatter' and 'hexbin'
| plots)
| position : float
| Specify relative alignments for bar plot layout.
| From 0 (left/bottom-end) to 1 (right/top-end). Default is 0.5
| (center)
| table : bool, Series or DataFrame, default False
| If True, draw a table using the data in the DataFrame and the
data
| will be transposed to meet matplotlib's default layout.
| If a Series or DataFrame is passed, use passed data to draw a
| table.
| yerr : DataFrame, Series, array-like, dict and str
| See :ref:`Plotting with Error Bars <visualization.errorbars>` for
| detail.

```

```

| xerr : DataFrame, Series, array-like, dict and str
| Equivalent to yerr.
| mark_right : bool, default True
| When using a secondary_y axis, automatically mark the column
| labels with "(right)" in the legend
| `**kwds` : keywords
| Options to pass to matplotlib plotting method
|
| Returns
| -----
| :class:`matplotlib.axes.Axes` or numpy.ndarray of them
| If the backend is not the default matplotlib one, the return
value
| will be the object returned by the backend.
|
| Notes
| -----
| - See matplotlib documentation online for more on this subject
| - If `kind` = 'bar' or 'barh', you can specify relative alignments
| for bar plot layout by `position` keyword.
| From 0 (left/bottom-end) to 1 (right/top-end). Default is 0.5
| (center)
|
| sparse = <class 'pandas.core.arrays.sparse.SparseFrameAccessor'>
| DataFrame accessor for sparse data.
|
| .. versionadded :: 0.25.0
|
| -----
| Methods inherited from pandas.core.generic.NDFrame:
|
| __abs__ (self)
|
| __array__ (self, dtype=None)
|
| __array_wrap__ (self, result, context=None)
|
| __bool__ = __nonzero__ (self)
|
| __contains__ (self, key)
| True if the key is in the info axis
|
| __copy__ (self, deep=True)
|
|

```

```

| __deepcopy__(self, memo=None)
| Parameters
| -----
| memo, default None
| Standard signature. Unused
|
| __delitem__(self, key)
| Delete item
|
| __finalize__(self, other, method=None, **kwargs)
| Propagate metadata from other to self.
|
| Parameters
| -----
| other : the object from which to get the attributes that we are going
| to propagate
| method : optional, a passed method name ; possibly to take different
| types of propagation actions based on this
|
| __getattr__(self, name)
| After regular attribute access, try looking up the name
| This allows simpler access to columns for interactive use.
|
| __getstate__(self)
|
| __hash__(self)
| Return hash(self).
|
| __invert__(self)
|
| __iter__(self)
| Iterate over info axis.
|
| Returns
| -----
| iterator
| Info axis as iterator.
|
| __neg__(self)
|
| __nonzero__(self)
|
| __pos__(self)
|
|

```

```

| __round__(self, decimals=0)
|
| __setattr__(self, name, value)
| After regular attribute access, try setting the name
| This allows simpler access to columns for interactive use.
|
| __setstate__(self, state)
|
| abs(self)
| Return a Series/DataFrame with absolute numeric value of each
element.
|
| This function only applies to elements that are all numeric.
|
| Returns
| -----
| abs
| Series/DataFrame containing the absolute value of each element.
|
| See Also
| -----
| numpy.absolute : Calculate the absolute value element-wise.
|
| Notes
| -----
| For ``complex`` inputs, ``1.2 + 1j``, the absolute value is
| :math:\sqrt{a^2 + b^2}``.
|
| Examples
| -----
| Absolute numeric values in a Series.
|
| >>> s = pd.Series([-1.10, 2, -3.33, 4])
| >>> s.abs()
| 0 1.10
| 1 2.00
| 2 3.33
| 3 4.00
| dtype: float64
|
| Absolute numeric values in a Series with complex numbers.
|
| >>> s = pd.Series([1.2 + 1j])
| >>> s.abs()

```

```

| 0 1.56205
| dtype: float64
|
| Absolute numeric values in a Series with a Timedelta element.
|
| >>> s = pd.Series([pd.Timedelta('1 days')])
| >>> s.abs()
| 0 1 days
| dtype: timedelta64[ns]
|
| Select rows with data closest to certain value using argsort (from
| `StackOverflow <https://stackoverflow.com/a/17758115>`__).
|
| >>> df = pd.DataFrame({
| ... 'a': [4, 5, 6, 7],
| ... 'b': [10, 20, 30, 40],
| ... 'c': [100, 50, -30, -50]
| ... })
| >>> df
|
| a b c
| 0 4 10 100
| 1 5 20 50
| 2 6 30 -30
| 3 7 40 -50
| >>> df.loc[(df.c - 43).abs().argsort()]
|
| a b c
| 1 5 20 50
| 0 4 10 100
| 2 6 30 -30
| 3 7 40 -50
|
| add_prefix(self, prefix)
| Prefix labels with string `prefix`.
|
| For Series, the row labels are prefixed.
| For DataFrame, the column labels are prefixed.
|
| Parameters
| -----
| prefix : str
| The string to add before each label.
|
| Returns
| -----

```



```

| Series or DataFrame
| New Series or DataFrame with updated labels.
|
| See Also
| -----
| Series.add_suffix: Suffix row labels with string `suffix`.
| DataFrame.add_suffix: Suffix column labels with string `suffix`.
|
| Examples
| -----
| >>> s = pd.Series([1, 2, 3, 4])
| >>> s
| 0 1
| 1 2
| 2 3
| 3 4
| dtype: int64
|
| >>> s.add_prefix('item_')
| item_0 1
| item_1 2
| item_2 3
| item_3 4
| dtype: int64
|
| >>> df = pd.DataFrame({'A': [1, 2, 3, 4], 'B': [3, 4, 5, 6]})
| >>> df
| A B
| 0 1 3
| 1 2 4
| 2 3 5
| 3 4 6
|
| >>> df.add_prefix('col_')
| col_A col_B
| 0 1 3
| 1 2 4
| 2 3 5
| 3 4 6
|
| add_suffix(self, suffix)
| Suffix labels with string `suffix`.
|
| For Series, the row labels are suffixed.

```

```

| For DataFrame, the column labels are suffixed.
|
| Parameters
| -----
| suffix : str
| The string to add after each label.
|
| Returns
| -----
| Series or DataFrame
| New Series or DataFrame with updated labels.
|
| See Also
| -----
| Series.add_prefix: Prefix row labels with string `prefix`.
| DataFrame.add_prefix: Prefix column labels with string `prefix`.
|
| Examples
| -----
| >>> s = pd.Series([1, 2, 3, 4])
| >>> s
| 0 1
| 1 2
| 2 3
| 3 4
| dtype: int64
|
| >>> s.add_suffix('_item')
| 0_item 1
| 1_item 2
| 2_item 3
| 3_item 4
| dtype: int64
|
| >>> df = pd.DataFrame({'A': [1, 2, 3, 4], 'B': [3, 4, 5, 6]})
| >>> df
| A B
| 0 1 3
| 1 2 4
| 2 3 5
| 3 4 6
|
| >>> df.add_suffix('_col')
| A_col B_col

```

```

| 0 1 3
| 1 2 4
| 2 3 5
| 3 4 6
|
| as_blocks(self, copy=True)
| Convert the frame to a dict of dtype -> Constructor Types that each
has
| a homogeneous dtype.
|
| .. deprecated:: 0.21.0
|
| NOTE: the dtypes of the blocks WILL BE PRESERVED HERE (unlike in
| as_matrix)
|
| Parameters
| -----
| copy : boolean, default True
|
| Returns
| -----
| values : a dict of dtype -> Constructor Types
|
| as_matrix(self, columns=None)
| Convert the frame to its Numpy-array representation.
|
| .. deprecated:: 0.23.0
| Use :meth:`DataFrame.values` instead.
|
| Parameters
| -----
| columns : list, optional, default:None
| If None, return all columns, otherwise, returns specified
columns.
|
| Returns
| -----
| values : ndarray
| If the caller is heterogeneous and contains booleans or objects,
| the result will be of dtype=object. See Notes.
|
| See Also
| -----
| DataFrame.values

```

```

|
| Notes
| -----
| Return is NOT a Numpy-matrix, rather, a Numpy-array.
|
| The dtype will be a lower-common-denominator dtype (implicit
| upcasting); that is to say if the dtypes (even of numeric types)
| are mixed, the one that accommodates all will be chosen. Use this
| with care if you are not dealing with the blocks.
|
| e.g. If the dtypes are float16 and float32, dtype will be upcast to
| float32. If dtypes are int32 and uint8, dtype will be upcase to
| int32. By numpy.find_common_type convention, mixing int64 and uint64
| will result in a float64 dtype.
|
| This method is provided for backwards compatibility. Generally,
| it is recommended to use '.values'.
|
| asfreq(self, freq, method=None, how=None, normalize=False,
fill_value=None)
| Convert TimeSeries to specified frequency.
|
| Optionally provide filling method to pad/backfill missing values.
|
| Returns the original data conformed to a new index with the specified
| frequency. ``resample`` is more appropriate if an operation, such as
| summarization, is necessary to represent the data at the new
frequency.
|
| Parameters
| -----
| freq : DateOffset object, or string
| method : {'backfill'/'bfill', 'pad'/'ffill'}, default None
| Method to use for filling holes in reindexed Series (note this
| does not fill NaNs that already were present):
|
| * 'pad' / 'ffill': propagate last valid observation forward to
next
| valid
| * 'backfill' / 'bfill': use NEXT valid observation to fill
| how : {'start', 'end'}, default end
| For PeriodIndex only, see PeriodIndex.asfreq
| normalize : bool, default False
| Whether to reset output index to midnight

```

```

| fill_value : scalar, optional
| Value to use for missing values, applied during upsampling (note
| this does not fill NaNs that already were present).
|
| .. versionadded:: 0.20.0
|
| Returns
| -----
| converted : same type as caller
|
| See Also
| -----
| reindex
|
| Notes
| -----
| To learn more about the frequency strings, please see `this link
| <http://pandas.pydata.org/pandas-
docs/stable/user_guide/timeseries.html#offset-aliases>`.
|
| Examples
| -----
|
| Start by creating a series with 4 one minute timestamps.
|
| >>> index = pd.date_range('1/1/2000', periods=4, freq='T')
| >>> series = pd.Series([0.0, None, 2.0, 3.0], index=index)
| >>> df = pd.DataFrame({'s':series})
| >>> df
|
| s
| 2000-01-01 00:00:00 0.0
| 2000-01-01 00:01:00 NaN
| 2000-01-01 00:02:00 2.0
| 2000-01-01 00:03:00 3.0
|
| Upsample the series into 30 second bins.
|
| >>> df.asfreq(freq='30S')
|
| s
| 2000-01-01 00:00:00 0.0
| 2000-01-01 00:00:30 NaN
| 2000-01-01 00:01:00 NaN
| 2000-01-01 00:01:30 NaN
| 2000-01-01 00:02:00 2.0

```

```
| 2000-01-01 00:02:30 NaN
| 2000-01-01 00:03:00 3.0
```

```
| Upsample again, providing a ``fill value``.
```

```
| >>> df.asfreq(freq='30S', fill_value=9.0)
```

```
| s
| 2000-01-01 00:00:00 0.0
| 2000-01-01 00:00:30 9.0
| 2000-01-01 00:01:00 NaN
| 2000-01-01 00:01:30 9.0
| 2000-01-01 00:02:00 2.0
| 2000-01-01 00:02:30 9.0
| 2000-01-01 00:03:00 3.0
```

```
| Upsample again, providing a ``method``.
```

```
| >>> df.asfreq(freq='30S', method='bfill')
```

```
| s
| 2000-01-01 00:00:00 0.0
| 2000-01-01 00:00:30 NaN
| 2000-01-01 00:01:00 NaN
| 2000-01-01 00:01:30 2.0
| 2000-01-01 00:02:00 2.0
| 2000-01-01 00:02:30 3.0
| 2000-01-01 00:03:00 3.0
```

```
| asof(self, where, subset=None)
```

```
| Return the last row(s) without any NaNs before `where`.
```

```
| The last row (for each element in `where`, if list) without any
| NaN is taken.
```

```
| In case of a :class:`~pandas.DataFrame`, the last row without NaN
| considering only the subset of columns (if not `None`)
```

```
| .. versionadded:: 0.19.0 For DataFrame
```

```
| If there is no good value, NaN is returned for a Series or
| a Series of NaN values for a DataFrame
```

```
| Parameters
```

```
| -----
```

```
| where : date or array-like of dates
```

```
| Date(s) before which the last row(s) are returned.
```

```

| subset : str or array-like of str, default `None`
| For DataFrame, if not `None`, only use these columns to
| check for NaNs.
|
| Returns
| -----
| scalar, Series, or DataFrame
|
| The return can be:
|
| * scalar : when `self` is a Series and `where` is a scalar
| * Series: when `self` is a Series and `where` is an array-like,
| or when `self` is a DataFrame and `where` is a scalar
| * DataFrame : when `self` is a DataFrame and `where` is an
| array-like
|
| Return scalar, Series, or DataFrame.
|
| See Also
| -----
| merge_asof : Perform an asof merge. Similar to left join.
|
| Notes
| -----
| Dates are assumed to be sorted. Raises if this is not the case.
|
| Examples
| -----
| A Series and a scalar `where`.
|
| >>> s = pd.Series([1, 2, np.nan, 4], index=[10, 20, 30, 40])
| >>> s
| 10 1.0
| 20 2.0
| 30 NaN
| 40 4.0
| dtype: float64
|
| >>> s.asof(20)
| 2.0
|
| For a sequence `where`, a Series is returned. The first value is
| NaN, because the first element of `where` is before the first
| index value.

```

```

|
| >>> s.asof([5, 20])
| 5 NaN
| 20 2.0
| dtype: float64
|
| Missing values are not considered. The following is ``2.0``, not
| NaN, even though NaN is at the index location for ``30``.
|
| >>> s.asof(30)
| 2.0
|
| Take all columns into consideration
|
| >>> df = pd.DataFrame({'a': [10, 20, 30, 40, 50],
| ... 'b': [None, None, None, None, 500]},
| ... index=pd.DatetimeIndex(['2018-02-27 09:01:00',
| ... '2018-02-27 09:02:00',
| ... '2018-02-27 09:03:00',
| ... '2018-02-27 09:04:00',
| ... '2018-02-27
09:05:00'])))
| >>> df.asof(pd.DatetimeIndex(['2018-02-27 09:03:30',
| ... '2018-02-27 09:04:30']))
|
| a b
| 2018-02-27 09:03:30 NaN NaN
| 2018-02-27 09:04:30 NaN NaN
|
| Take a single column into consideration
|
| >>> df.asof(pd.DatetimeIndex(['2018-02-27 09:03:30',
| ... '2018-02-27 09:04:30']),
| ... subset=['a'])
|
| a b
| 2018-02-27 09:03:30 30.0 NaN
| 2018-02-27 09:04:30 40.0 NaN
|
| astype(self, dtype, copy=True, errors='raise', **kwargs)
| Cast a pandas object to a specified dtype ``dtype``.
|
| Parameters
| -----
|
| dtype : data type, or dict of column name -> data type
| Use a numpy.dtype or Python type to cast entire pandas object to

```



```

| the same type. Alternatively, use {col: dtype, ...}, where col is
a
| column label and dtype is a numpy.dtype or Python type to cast
one
| or more of the DataFrame's columns to column-specific types.
| copy : bool, default True
| Return a copy when ``copy=True`` (be very careful setting
| ``copy=False`` as changes to values then may propagate to other
| pandas objects).
| errors : {'raise', 'ignore'}, default 'raise'
| Control raising of exceptions on invalid data for provided dtype.
|
| - ``raise`` : allow exceptions to be raised
| - ``ignore`` : suppress exceptions. On error return original
object
|
| .. versionadded:: 0.20.0
|
| kwargs : keyword arguments to pass on to the constructor
|
| Returns
| -----
| casted : same type as caller
|
| See Also
| -----
| to_datetime : Convert argument to datetime.
| to_timedelta : Convert argument to timedelta.
| to_numeric : Convert argument to a numeric type.
| numpy.ndarray.astype : Cast a numpy array to a specified type.
|
| Examples
| -----
| Create a DataFrame:
|
| >>> d = {'col1': [1, 2], 'col2': [3, 4]}
| >>> df = pd.DataFrame(data=d)
| >>> df.dtypes
| col1 int64
| col2 int64
| dtype: object
|
| Cast all columns to int32:
|

```

```

| >>> df.astype('int32').dtypes
| col1 int32
| col2 int32
| dtype: object
|
| Cast col1 to int32 using a dictionary:
|
| >>> df.astype({'col1': 'int32'}).dtypes
| col1 int32
| col2 int64
| dtype: object
|
| Create a series:
|
| >>> ser = pd.Series([1, 2], dtype='int32')
| >>> ser
| 0 1
| 1 2
| dtype: int32
| >>> ser.astype('int64')
| 0 1
| 1 2
| dtype: int64
|
| Convert to categorical type:
|
| >>> ser.astype('category')
| 0 1
| 1 2
| dtype: category
| Categories (2, int64): [1, 2]
|
| Convert to ordered categorical type with custom ordering:
|
| >>> cat_dtype = pd.api.types.CategoricalDtype(
| ... categories=[2, 1], ordered=True)
| >>> ser.astype(cat_dtype)
| 0 1
| 1 2
| dtype: category
| Categories (2, int64): [2 < 1]
|
| Note that using ``copy=False`` and changing data on a new
| pandas object may propagate changes:

```

```

|
| >>> s1 = pd.Series([1,2])
| >>> s2 = s1.astype('int64', copy=False)
| >>> s2[0] = 10
| >>> s1 # note that s1[0] has changed too
| 0 10
| 1 2
| dtype: int64
|
| at_time(self, time, asof=False, axis=None)
| Select values at particular time of day (e.g. 9:30AM).
|
| Parameters
| -----
| time : datetime.time or str
| axis : {0 or 'index', 1 or 'columns'}, default 0
|
| .. versionadded:: 0.24.0
|
| Returns
| -----
| Series or DataFrame
|
| Raises
| -----
| TypeError
| If the index is not a :class:`DatetimeIndex`
|
| See Also
| -----
| between_time : Select values between particular times of the day.
| first : Select initial periods of time series based on a date offset.
| last : Select final periods of time series based on a date offset.
| DatetimeIndex.indexer_at_time : Get just the index locations for
| values at particular time of the day.
|
| Examples
| -----
| >>> i = pd.date_range('2018-04-09', periods=4, freq='12H')
| >>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
| >>> ts
|
| A
| 2018-04-09 00:00:00 1
| 2018-04-09 12:00:00 2

```

```

| 2018-04-10 00:00:00 3
| 2018-04-10 12:00:00 4
|
| >>> ts.at_time('12:00')
|
| A
| 2018-04-09 12:00:00 2
| 2018-04-10 12:00:00 4
|
| between_time(self, start_time, end_time, include_start=True,
include_end=True, axis=None)
|
| Select values between particular times of the day (e.g., 9:00-9:30
AM).
|
| By setting ``start_time`` to be later than ``end_time``,
| you can get the times that are *not* between the two times.
|
| Parameters
| -----
| start_time : datetime.time or str
| end_time : datetime.time or str
| include_start : bool, default True
| include_end : bool, default True
| axis : {0 or 'index', 1 or 'columns'}, default 0
|
| .. versionadded:: 0.24.0
|
| Returns
| -----
| Series or DataFrame
|
| Raises
| -----
| TypeError
| If the index is not a :class:`DatetimeIndex`
|
| See Also
| -----
| at_time : Select values at a particular time of the day.
| first : Select initial periods of time series based on a date offset.
| last : Select final periods of time series based on a date offset.
| DatetimeIndex.indexer_between_time : Get just the index locations for
| values between particular times of the day.
|
| Examples

```

```

| -----
| >>> i = pd.date_range('2018-04-09', periods=4, freq='1D20min')
| >>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
| >>> ts
|
| A
| 2018-04-09 00:00:00 1
| 2018-04-10 00:20:00 2
| 2018-04-11 00:40:00 3
| 2018-04-12 01:00:00 4
|
| >>> ts.between_time('0:15', '0:45')
|
| A
| 2018-04-10 00:20:00 2
| 2018-04-11 00:40:00 3
|
| You get the times that are *not* between two times by setting
| ``start_time`` later than ``end_time``:
|
| >>> ts.between_time('0:45', '0:15')
|
| A
| 2018-04-09 00:00:00 1
| 2018-04-12 01:00:00 4
|
| bfill(self, axis=None, inplace=False, limit=None, downcast=None)
| Synonym for :meth:`DataFrame.fillna` with ``method='bfill'``.
|
| Returns
| -----
| %(klass)s
| Object with missing values filled.
|
| bool(self)
| Return the bool of a single element PandasObject.
|
| This must be a boolean scalar value, either True or False. Raise a
| ValueError if the PandasObject does not have exactly 1 element, or
that
| element is not boolean
|
| Returns
| -----
| bool
| Same single boolean value converted to bool type.
|

```

```

| clip(self, lower=None, upper=None, axis=None, inplace=False, *args,
**kwargs)
| Trim values at input threshold(s).
|
| Assigns values outside boundary to boundary values. Thresholds
| can be singular values or array like, and in the latter case
| the clipping is performed element-wise in the specified axis.
|
| Parameters
| -----
| lower : float or array_like, default None
| Minimum threshold value. All values below this
| threshold will be set to it.
| upper : float or array_like, default None
| Maximum threshold value. All values above this
| threshold will be set to it.
| axis : int or str axis name, optional
| Align object with lower and upper along the given axis.
| inplace : bool, default False
| Whether to perform the operation in place on the data.
|
| .. versionadded:: 0.21.0
| *args, **kwargs
| Additional keywords have no effect but might be accepted
| for compatibility with numpy.
|
| Returns
| -----
| Series or DataFrame
| Same type as calling object with the values outside the
| clip boundaries replaced.
|
| Examples
| -----
| >>> data = {'col_0': [9, -3, 0, -1, 5], 'col_1': [-2, -7, 6, 8, -5]}
| >>> df = pd.DataFrame(data)
| >>> df
|
| col_0 col_1
| 0 9 -2
| 1 -3 -7
| 2 0 6
| 3 -1 8
| 4 5 -5
|

```

```

| Clips per column using lower and upper thresholds:
|
| >>> df.clip(-4, 6)
| col_0 col_1
| 0 6 -2
| 1 -3 -4
| 2 0 6
| 3 -1 6
| 4 5 -4
|
| Clips using specific lower and upper thresholds per column element:
|
| >>> t = pd.Series([2, -4, -1, 6, 3])
| >>> t
| 0 2
| 1 -4
| 2 -1
| 3 6
| 4 3
| dtype: int64
|
| >>> df.clip(t, t + 4, axis=0)
| col_0 col_1
| 0 6 2
| 1 -3 -4
| 2 0 3
| 3 6 8
| 4 5 3
|
| clip_lower(self, threshold, axis=None, inplace=False)
| Trim values below a given threshold.
|
| .. deprecated:: 0.24.0
| Use clip(lower=threshold) instead.
|
| Elements below the `threshold` will be changed to match the
| `threshold` value(s). Threshold can be a single value or an array,
| in the latter case it performs the truncation element-wise.
|
| Parameters
| -----
| threshold : numeric or array-like
| Minimum value allowed. All values below threshold will be set to
| this value.

```

```

|
| * float : every value is compared to `threshold`.
| * array-like : The shape of `threshold` should match the object
| it's compared to. When `self` is a Series, `threshold` should
be
| the length. When `self` is a DataFrame, `threshold` should 2-D
| and the same shape as `self` for ``axis=None``, or 1-D and the
| same length as the axis being compared.
|
| axis : {0 or 'index', 1 or 'columns'}, default 0
| Align `self` with `threshold` along the given axis.
|
| inplace : bool, default False
| Whether to perform the operation in place on the data.
|
| .. versionadded:: 0.21.0
|
| Returns
| -----
| Series or DataFrame
| Original data with values trimmed.
|
| See Also
| -----
| Series.clip : General purpose method to trim Series values to given
| threshold(s).
| DataFrame.clip : General purpose method to trim DataFrame values to
| given threshold(s).
|
| Examples
| -----
|
| Series single threshold clipping:
|
| >>> s = pd.Series([5, 6, 7, 8, 9])
| >>> s.clip(lower=8)
| 0 8
| 1 8
| 2 8
| 3 8
| 4 9
| dtype: int64
|

```



```

| Series clipping element-wise using an array of thresholds.
| `threshold`
| should be the same length as the Series.
|
| >>> elemwise_thresholds = [4, 8, 7, 2, 5]
| >>> s.clip(lower=elemwise_thresholds)
| 0 5
| 1 8
| 2 7
| 3 8
| 4 9
| dtype: int64
|
| DataFrames can be compared to a scalar.
|
| >>> df = pd.DataFrame({"A": [1, 3, 5], "B": [2, 4, 6]})
| >>> df
| A B
| 0 1 2
| 1 3 4
| 2 5 6
|
| >>> df.clip(lower=3)
| A B
| 0 3 3
| 1 3 4
| 2 5 6
|
| Or to an array of values. By default, `threshold` should be the same
| shape as the DataFrame.
|
| >>> df.clip(lower=np.array([[3, 4], [2, 2], [6, 2]]))
| A B
| 0 3 4
| 1 3 4
| 2 6 6
|
| Control how `threshold` is broadcast with `axis`. In this case
| `threshold` should be the same length as the axis specified by
| `axis`.
|
| >>> df.clip(lower=[3, 3, 5], axis='index')
| A B
| 0 3 3

```

```

| 1 3 4
| 2 5 6
|
| >>> df.clip(lower=[4, 5], axis='columns')
| A B
| 0 4 5
| 1 4 5
| 2 5 6
|
| clip_upper(self, threshold, axis=None, inplace=False)
| Trim values above a given threshold.
|
| .. deprecated:: 0.24.0
| Use clip(upper=threshold) instead.
|
| Elements above the `threshold` will be changed to match the
| `threshold` value(s). Threshold can be a single value or an array,
| in the latter case it performs the truncation element-wise.
|
| Parameters
| -----
| threshold : numeric or array-like
| Maximum value allowed. All values above threshold will be set to
| this value.
|
| * float : every value is compared to `threshold`.
| * array-like : The shape of `threshold` should match the object
| it's compared to. When `self` is a Series, `threshold` should
be
| the length. When `self` is a DataFrame, `threshold` should 2-D
| and the same shape as `self` for ``axis=None``, or 1-D and the
| same length as the axis being compared.
|
| axis : {0 or 'index', 1 or 'columns'}, default 0
| Align object with `threshold` along the given axis.
| inplace : bool, default False
| Whether to perform the operation in place on the data.
|
| .. versionadded:: 0.21.0
|
| Returns
| -----
| Series or DataFrame
| Original data with values trimmed.

```

| See Also

| -----

| Series.clip : General purpose method to trim Series values to given  
| threshold(s).

| DataFrame.clip : General purpose method to trim DataFrame values to  
| given threshold(s).

|

| Examples

| -----

| >>> s = pd.Series([1, 2, 3, 4, 5])

| >>> s

| 0 1

| 1 2

| 2 3

| 3 4

| 4 5

| dtype: int64

|

| >>> s.clip(upper=3)

| 0 1

| 1 2

| 2 3

| 3 3

| 4 3

| dtype: int64

|

| >>> elemwise\_thresholds = [5, 4, 3, 2, 1]

| >>> elemwise\_thresholds

| [5, 4, 3, 2, 1]

|

| >>> s.clip(upper=elemwise\_thresholds)

| 0 1

| 1 2

| 2 3

| 3 2

| 4 1

| dtype: int64

|

| copy(self, deep=True)

| Make a copy of this object's indices and data.

|

| When ``deep=True`` (default), a new object will be created with a

| copy of the calling object's data and indices. Modifications to

```

| the data or indices of the copy will not be reflected in the
| original object (see notes below).
|
| When ``deep=False``, a new object will be created without copying
| the calling object's data or index (only references to the data
| and index are copied). Any changes to the data of the original
| will be reflected in the shallow copy (and vice versa).
|
| Parameters
| -----
|
| deep : bool, default True
| Make a deep copy, including a copy of the data and the indices.
| With ``deep=False`` neither the indices nor the data are copied.
|
| Returns
| -----
|
| copy : Series or DataFrame
| Object type matches caller.
|
| Notes
| -----
|
| When ``deep=True``, data is copied but actual Python objects
| will not be copied recursively, only the reference to the object.
| This is in contrast to `copy.deepcopy` in the Standard Library,
| which recursively copies object data (see examples below).
|
| While ``Index`` objects are copied when ``deep=True``, the underlying
| numpy array is not copied for performance reasons. Since ``Index`` is
| immutable, the underlying data can be safely shared and a copy
| is not needed.
|
| Examples
| -----
|
| >>> s = pd.Series([1, 2], index=["a", "b"])
| >>> s
| a 1
| b 2
| dtype: int64
|
| >>> s_copy = s.copy()
| >>> s_copy
| a 1
| b 2
| dtype: int64

```

```

|
| **Shallow copy versus default (deep) copy:**
|
| >>> s = pd.Series([1, 2], index=["a", "b"])
| >>> deep = s.copy()
| >>> shallow = s.copy(deep=False)
|
| Shallow copy shares data and index with original.
|
| >>> s is shallow
| False
| >>> s.values is shallow.values and s.index is shallow.index
| True
|
| Deep copy has own copy of data and index.
|
| >>> s is deep
| False
| >>> s.values is deep.values or s.index is deep.index
| False
|
| Updates to the data shared by shallow copy and original is reflected
| in both; deep copy remains unchanged.
|
| >>> s[0] = 3
| >>> shallow[1] = 4
| >>> s
| a 3
| b 4
| dtype: int64
| >>> shallow
| a 3
| b 4
| dtype: int64
| >>> deep
| a 1
| b 2
| dtype: int64
|
| Note that when copying an object containing Python objects, a deep
copy
| will copy the data, but will not do so recursively. Updating a nested
| data object will be reflected in the deep copy.
|

```

```

| >>> s = pd.Series([[1, 2], [3, 4]])
| >>> deep = s.copy()
| >>> s[0][0] = 10
| >>> s
| 0 [10, 2]
| 1 [3, 4]
| dtype: object
| >>> deep
| 0 [10, 2]
| 1 [3, 4]
| dtype: object
|
| describe(self, percentiles=None, include=None, exclude=None)
| Generate descriptive statistics that summarize the central tendency,
| dispersion and shape of a dataset's distribution, excluding
| ``NaN`` values.
|
| Analyzes both numeric and object series, as well
| as ``DataFrame`` column sets of mixed data types. The output
| will vary depending on what is provided. Refer to the notes
| below for more detail.
|
| Parameters
| -----
| percentiles : list-like of numbers, optional
| The percentiles to include in the output. All should
| fall between 0 and 1. The default is
| ``[.25, .5, .75]``, which returns the 25th, 50th, and
| 75th percentiles.
| include : 'all', list-like of dtypes or None (default), optional
| A white list of data types to include in the result. Ignored
| for ``Series``. Here are the options:
|
| - 'all' : All columns of the input will be included in the
output.
|
| - A list-like of dtypes : Limits the results to the
| provided data types.
| To limit the result to numeric types submit
| ``numpy.number``. To limit it instead to object columns submit
| the ``numpy.object`` data type. Strings
| can also be used in the style of
| ``select_dtypes`` (e.g. ``df.describe(include=['O'])``). To
| select pandas categorical columns, use ``'category'``
|
| - None (default) : The result will include all numeric columns.

```

```

| exclude : list-like of dtypes or None (default), optional,
| A black list of data types to omit from the result. Ignored
| for ``Series``. Here are the options:
|
| - A list-like of dtypes : Excludes the provided data types
| from the result. To exclude numeric types submit
| ``numpy.number``. To exclude object columns submit the data
| type ``numpy.object``. Strings can also be used in the style of
| ``select_dtypes`` (e.g. ``df.describe(include=['O'])``). To
| exclude pandas categorical columns, use ``'category'``
| - None (default) : The result will exclude nothing.
|
| Returns
| -----
| Series or DataFrame
| Summary statistics of the Series or Dataframe provided.
|
| See Also
| -----
| DataFrame.count: Count number of non-NA/null observations.
| DataFrame.max: Maximum of the values in the object.
| DataFrame.min: Minimum of the values in the object.
| DataFrame.mean: Mean of the values.
| DataFrame.std: Standard deviation of the observations.
| DataFrame.select_dtypes: Subset of a DataFrame including/excluding
| columns based on their dtype.
|
| Notes
| -----
| For numeric data, the result's index will include ``count``,
| ``mean``, ``std``, ``min``, ``max`` as well as lower, ``50`` and
| upper percentiles. By default the lower percentile is ``25`` and the
| upper percentile is ``75``. The ``50`` percentile is the
| same as the median.
|
| For object data (e.g. strings or timestamps), the result's index
| will include ``count``, ``unique``, ``top``, and ``freq``. The
| ``top``
| is the most common value. The ``freq`` is the most common value's
| frequency. Timestamps also include the ``first`` and ``last`` items.
|
| If multiple object values have the highest count, then the
| ``count`` and ``top`` results will be arbitrarily chosen from
| among those with the highest count.

```

For mixed data types provided via a ``DataFrame``, the default is to return only an analysis of numeric columns. If the dataframe consists only of object and categorical data without any numeric columns, the default is to return an analysis of both the object and categorical columns. If ``include='all'`` is provided as an option, the result will include a union of attributes of each type.

The ``include`` and ``exclude`` parameters can be used to limit which columns in a ``DataFrame`` are analyzed for the output. The parameters are ignored when analyzing a ``Series``.

#### Examples

-----

Describing a numeric ``Series``.

```
>>> s = pd.Series([1, 2, 3])
>>> s.describe()
count 3.0
mean 2.0
std 1.0
min 1.0
25% 1.5
50% 2.0
75% 2.5
max 3.0
dtype: float64
```

Describing a categorical ``Series``.

```
>>> s = pd.Series(['a', 'a', 'b', 'c'])
>>> s.describe()
count 4
unique 3
top a
freq 2
dtype: object
```

Describing a timestamp ``Series``.

```
>>> s = pd.Series([
... np.datetime64("2000-01-01"),
... np.datetime64("2010-01-01"),
... np.datetime64("2010-01-01")])
```



```
| ...])
```

```
| >>> s.describe()
```

```
| count 3
```

```
| unique 2
```

```
| top 2010-01-01 00:00:00
```

```
| freq 2
```

```
| first 2000-01-01 00:00:00
```

```
| last 2010-01-01 00:00:00
```

```
| dtype: object
```

```
|
```

```
| Describing a ``DataFrame``. By default only numeric fields
```

```
| are returned.
```

```
|
```

```
| >>> df = pd.DataFrame({'categorical': pd.Categorical(['d','e','f']),
```

```
| ... 'numeric': [1, 2, 3],
```

```
| ... 'object': ['a', 'b', 'c']
```

```
| ... })
```

```
| >>> df.describe()
```

```
|
```

	numeric
count	3.0
mean	2.0
std	1.0
min	1.0
25%	1.5
50%	2.0
75%	2.5
max	3.0

```
|
```

```
| Describing all columns of a ``DataFrame`` regardless of data type.
```

```
|
```

```
| >>> df.describe(include='all')
```

	categorical	numeric	object
count	3	3.0	3
unique	3	NaN	3
top	f	NaN	c
freq	1	NaN	1
mean	NaN	2.0	NaN
std	NaN	1.0	NaN
min	NaN	1.0	NaN
25%	NaN	1.5	NaN
50%	NaN	2.0	NaN
75%	NaN	2.5	NaN
max	NaN	3.0	NaN

```
|
```

```
| Describing a column from a ``DataFrame`` by accessing it as
| an attribute.
```

```
| >>> df.numeric.describe()
| count 3.0
| mean 2.0
| std 1.0
| min 1.0
| 25% 1.5
| 50% 2.0
| 75% 2.5
| max 3.0
| Name: numeric, dtype: float64
```

```
| Including only numeric columns in a ``DataFrame`` description.
```

```
| >>> df.describe(include=[np.number])
| numeric
| count 3.0
| mean 2.0
| std 1.0
| min 1.0
| 25% 1.5
| 50% 2.0
| 75% 2.5
| max 3.0
```

```
| Including only string columns in a ``DataFrame`` description.
```

```
| >>> df.describe(include=[np.object])
| object
| count 3
| unique 3
| top c
| freq 1
```

```
| Including only categorical columns from a ``DataFrame`` description.
```

```
| >>> df.describe(include=['category'])
| categorical
| count 3
| unique 3
| top f
| freq 1
```

```
| Excluding numeric columns from a ``DataFrame`` description.
```

```
| >>> df.describe(exclude=[np.number])
```

```
| categorical object
| count 3 3
| unique 3 3
| top f c
| freq 1 1
```

```
| Excluding object columns from a ``DataFrame`` description.
```

```
| >>> df.describe(exclude=[np.object])
```

```
| categorical numeric
| count 3 3.0
| unique 3 NaN
| top f NaN
| freq 1 NaN
| mean NaN 2.0
| std NaN 1.0
| min NaN 1.0
| 25% NaN 1.5
| 50% NaN 2.0
| 75% NaN 2.5
| max NaN 3.0
```

```
| droplevel(self, level, axis=0)
```

```
| Return DataFrame with requested index / column level(s) removed.
```

```
| .. versionadded:: 0.24.0
```

```
| Parameters
```

```
| -----
```

```
| level : int, str, or list-like
```

```
| If a string is given, must be the name of a level
```

```
| If list-like, elements must be names or positional indexes
| of levels.
```

```
| axis : {0 or 'index', 1 or 'columns'}, default 0
```

```
| Returns
```

```
| -----
```

```
| DataFrame.droplevel()
```

|       Examples

|       -----

```
| >>> df = pd.DataFrame([
| ... [1, 2, 3, 4],
| ... [5, 6, 7, 8],
| ... [9, 10, 11, 12]
| ...]).set_index([0, 1]).rename_axis(['a', 'b'])
```

```
| >>> df.columns = pd.MultiIndex.from_tuples([
| ... ('c', 'e'), ('d', 'f')
| ...], names=['level_1', 'level_2'])
```

```
| >>> df
| level_1 c d
| level_2 e f
| a b
| 1 2 3 4
| 5 6 7 8
| 9 10 11 12
```

```
| >>> df.droplevel('a')
| level_1 c d
| level_2 e f
| b
| 2 3 4
| 6 7 8
| 10 11 12
```

```
| >>> df.droplevel('level2', axis=1)
| level_1 c d
| a b
| 1 2 3 4
| 5 6 7 8
| 9 10 11 12
```

```
| equals(self, other)
```

|       Test whether two objects contain the same elements.

|       This function allows two Series or DataFrames to be compared against each other to see if they have the same shape and elements. NaNs in the same location are considered equal. The column headers do not need to have the same type, but the elements within the columns must be the same dtype.

```

| Parameters
| -----
| other : Series or DataFrame
| The other Series or DataFrame to be compared with the first.
|
| Returns
| -----
| bool
| True if all elements are the same in both objects, False
| otherwise.
|
| See Also
| -----
| Series.eq : Compare two Series objects of the same length
| and return a Series where each element is True if the element
| in each Series is equal, False otherwise.
| DataFrame.eq : Compare two DataFrame objects of the same shape and
| return a DataFrame where each element is True if the respective
| element in each DataFrame is equal, False otherwise.
| assert_series_equal : Return True if left and right Series are equal,
| False otherwise.
| assert_frame_equal : Return True if left and right DataFrames are
| equal, False otherwise.
| numpy.array_equal : Return True if two arrays have the same shape
| and elements, False otherwise.
|
| Notes
| -----
| This function requires that the elements have the same dtype as their
| respective elements in the other Series or DataFrame. However, the
| column labels do not need to have the same type, as long as they are
| still considered equal.
|
| Examples
| -----
| >>> df = pd.DataFrame({'1': [10], '2': [20]})
| >>> df
| 1 2
| 0 10 20
|
| DataFrames df and exactly_equal have the same types and values for
| their elements and column labels, which will return True.
|
| >>> exactly_equal = pd.DataFrame({'1': [10], '2': [20]})

```

```

| >>> exactly_equal
| 1 2
| 0 10 20
| >>> df.equals(exactly_equal)
| True
|
| DataFrames df and different_column_type have the same element
| types and values, but have different types for the column labels,
| which will still return True.
|
| >>> different_column_type = pd.DataFrame({1.0: [10], 2.0: [20]})
| >>> different_column_type
| 1.0 2.0
| 0 10 20
| >>> df.equals(different_column_type)
| True
|
| DataFrames df and different_data_type have different types for the
| same values for their elements, and will return False even though
| their column labels are the same values and types.
|
| >>> different_data_type = pd.DataFrame({1: [10.0], 2: [20.0]})
| >>> different_data_type
| 1 2
| 0 10.0 20.0
| >>> df.equals(different_data_type)
| False
|
| ffill(self, axis=None, inplace=False, limit=None, downcast=None)
| Synonym for :meth:`DataFrame.fillna` with ``method='ffill'``.
|
| Returns
| -----
| %(klass)s
| Object with missing values filled.
|
| filter(self, items=None, like=None, regex=None, axis=None)
| Subset rows or columns of dataframe according to labels in
| the specified index.
|
| Note that this routine does not filter a dataframe on its
| contents. The filter is applied to the labels of the index.
|
| Parameters

```

```

| -----
| items : list-like
| Keep labels from axis which are in items.
| like : string
| Keep labels from axis for which "like in label == True".
| regex : string (regular expression)
| Keep labels from axis for which re.search(regex, label) == True.
| axis : int or string axis name
| The axis to filter on. By default this is the info axis,
| 'index' for Series, 'columns' for DataFrame.
|
| Returns
| -----
| same type as input object
|
| See Also
| -----
| DataFrame.loc
|
| Notes
| -----
| The ``items``, ``like``, and ``regex`` parameters are
| enforced to be mutually exclusive.
|
| ``axis`` defaults to the info axis that is used when indexing
| with ``[]``.
|
| Examples
| -----
| >>> df = pd.DataFrame(np.array([[1, 2, 3], [4, 5, 6]]),
| ... index=['mouse', 'rabbit'],
| ... columns=['one', 'two', 'three'])
|
| >>> # select columns by name
| >>> df.filter(items=['one', 'three'])
|
| one three
| mouse 1 3
| rabbit 4 6
|
| >>> # select columns by regular expression
| >>> df.filter(regex='e$', axis=1)
|
| one three
| mouse 1 3
| rabbit 4 6

```

```

|
| >>> # select rows containing 'bbi'
| >>> df.filter(like='bbi', axis=0)
| one two three
| rabbit 4 5 6
|
| first(self, offset)
| Convenience method for subsetting initial periods of time series data
| based on a date offset.
|
| Parameters
| -----
| offset : string, DateOffset, dateutil.relativedelta
|
| Returns
| -----
| subset : same type as caller
|
| Raises
| -----
| TypeError
| If the index is not a :class:`DatetimeIndex`
|
| See Also
| -----
| last : Select final periods of time series based on a date offset.
| at_time : Select values at a particular time of the day.
| between_time : Select values between particular times of the day.
|
| Examples
| -----
| >>> i = pd.date_range('2018-04-09', periods=4, freq='2D')
| >>> ts = pd.DataFrame({'A': [1,2,3,4]}, index=i)
| >>> ts
|
| A
| 2018-04-09 1
| 2018-04-11 2
| 2018-04-13 3
| 2018-04-15 4
|
| Get the rows for the first 3 days:
|
| >>> ts.first('3D')
|
| A

```



```

| 2018-04-09 1
| 2018-04-11 2
|
| Notice the data for 3 first calendar days were returned, not the
first
| 3 days observed in the dataset, and therefore data for 2018-04-13 was
| not returned.
|
| first_valid_index(self)
| Return index for first non-NA/null value.
|
| Returns
| -----
| scalar : type of index
|
| Notes
| -----
| If all elements are non-NA/null, returns None.
| Also returns None for empty Series/DataFrame.
|
| get(self, key, default=None)
| Get item from object for given key (ex: DataFrame column).
|
| Returns default value if not found.
|
| Parameters
| -----
| key : object
|
| Returns
| -----
| value : same type as items contained in object
|
| get_dtype_counts(self)
| Return counts of unique dtypes in this object.
|
| .. deprecated:: 0.25.0
|
| Use ``.dtypes.value_counts()`` instead.
|
| Returns
| -----
| dtype : Series
| Series with the count of columns with each dtype.

```

```

|
| See Also
| -----
| dtypes : Return the dtypes in this object.
|
| Examples
| -----
| >>> a = [['a', 1, 1.0], ['b', 2, 2.0], ['c', 3, 3.0]]
| >>> df = pd.DataFrame(a, columns=['str', 'int', 'float'])
| >>> df
| str int float
| 0 a 1 1.0
| 1 b 2 2.0
| 2 c 3 3.0
|
| >>> df.get_dtype_counts()
| float64 1
| int64 1
| object 1
| dtype: int64
|
| get_ftype_counts(self)
| Return counts of unique ftypes in this object.
|
| .. deprecated:: 0.23.0
|
| This is useful for SparseDataFrame or for DataFrames containing
| sparse arrays.
|
| Returns
| -----
| dtype : Series
| Series with the count of columns with each type and
| sparsity (dense/sparse).
|
| See Also
| -----
| ftypes : Return ftypes (indication of sparse/dense and dtype) in
| this object.
|
| Examples
| -----
| >>> a = [['a', 1, 1.0], ['b', 2, 2.0], ['c', 3, 3.0]]
| >>> df = pd.DataFrame(a, columns=['str', 'int', 'float'])

```

```

| >>> df
| str int float
| 0 a 1 1.0
| 1 b 2 2.0
| 2 c 3 3.0
|
| >>> df.get_ftype_counts() # doctest: +SKIP
| float64:dense 1
| int64:dense 1
| object:dense 1
| dtype: int64
|
| get_values(self)
| Return an ndarray after converting sparse values to dense.
|
| .. deprecated:: 0.25.0
| Use ``np.asarray(..)`` or :meth:`DataFrame.values` instead.
|
| This is the same as ``.values`` for non-sparse data. For sparse
| data contained in a `SparseArray`, the data are first
| converted to a dense representation.
|
| Returns
| -----
| numpy.ndarray
| Numpy representation of DataFrame.
|
| See Also
| -----
| values : Numpy representation of DataFrame.
| SparseArray : Container for sparse data.
|
| Examples
| -----
| >>> df = pd.DataFrame({'a': [1, 2], 'b': [True, False],
| ... 'c': [1.0, 2.0]})
| >>> df
| a b c
| 0 1 True 1.0
| 1 2 False 2.0
|
| >>> df.get_values()
| array([[1, True, 1.0], [2, False, 2.0]], dtype=object)
|

```

```

| >>> df = pd.DataFrame({"a": pd.SparseArray([1, None, None]),
| ... "c": [1.0, 2.0, 3.0]})
| >>> df
| a c
| 0 1.0 1.0
| 1 NaN 2.0
| 2 NaN 3.0
|
| >>> df.get_values()
| array([[1., 1.],
| [nan, 2.],
| [nan, 3.]])
|
| groupby(self, by=None, axis=0, level=None, as_index=True, sort=True,
group_keys=True, squeeze=False, observed=False, **kwargs)
| Group DataFrame or Series using a mapper or by a Series of columns.
|
| A groupby operation involves some combination of splitting the
| object, applying a function, and combining the results. This can be
| used to group large amounts of data and compute operations on these
| groups.
|
| Parameters
| -----
| by : mapping, function, label, or list of labels
| Used to determine the groups for the groupby.
| If ``by`` is a function, it's called on each value of the
object's
| index. If a dict or Series is passed, the Series or dict VALUES
| will be used to determine the groups (the Series' values are
first
| aligned; see ``.align()`` method). If an ndarray is passed, the
| values are used as-is determine the groups. A label or list of
| labels may be passed to group by the columns in ``self``. Notice
| that a tuple is interpreted as a (single) key.
| axis : {0 or 'index', 1 or 'columns'}, default 0
| Split along rows (0) or columns (1).
| level : int, level name, or sequence of such, default None
| If the axis is a MultiIndex (hierarchical), group by a particular
| level or levels.
| as_index : bool, default True
| For aggregated output, return object with group labels as the
| index. Only relevant for DataFrame input. as_index=False is
| effectively "SQL-style" grouped output.

```

```

| sort : bool, default True
| Sort group keys. Get better performance by turning this off.
| Note this does not influence the order of observations within
each
| group. Groupby preserves the order of rows within each group.
| group_keys : bool, default True
| When calling apply, add group keys to index to identify pieces.
| squeeze : bool, default False
| Reduce the dimensionality of the return type if possible,
| otherwise return a consistent type.
| observed : bool, default False
| This only applies if any of the groupers are Categoricals.
| If True: only show observed values for categorical groupers.
| If False: show all values for categorical groupers.
|
| .. versionadded:: 0.23.0
|
| **kwargs
| Optional, only accepts keyword argument 'mutated' and is passed
| to groupby.
|
| Returns
| -----
| DataFrameGroupBy or SeriesGroupBy
| Depends on the calling object and returns groupby object that
| contains information about the groups.
|
| See Also
| -----
| resample : Convenience method for frequency conversion and resampling
| of time series.
|
| Notes
| -----
| See the `user guide
| <http://pandas.pydata.org/pandas-docs/stable/groupby.html>`_ for
more.
|
| Examples
| -----
| >>> df = pd.DataFrame({'Animal': ['Falcon', 'Falcon',
| ... 'Parrot', 'Parrot'],
| ... 'Max Speed': [380., 370., 24., 26.]})
| >>> df

```

```

| Animal Max Speed
| 0 Falcon 380.0
| 1 Falcon 370.0
| 2 Parrot 24.0
| 3 Parrot 26.0
| >>> df.groupby(['Animal']).mean()
| Max Speed
| Animal
| Falcon 375.0
| Parrot 25.0
|
| **Hierarchical Indexes**
|
| We can groupby different levels of a hierarchical index
| using the `level` parameter:
|
| >>> arrays = [['Falcon', 'Falcon', 'Parrot', 'Parrot'],
| ... ['Captive', 'Wild', 'Captive', 'Wild']]
| >>> index = pd.MultiIndex.from_arrays(arrays, names=('Animal',
| 'Type'))
| >>> df = pd.DataFrame({'Max Speed': [390., 350., 30., 20.]},
| ... index=index)
| >>> df
| Max Speed
| Animal Type
| Falcon Captive 390.0
| Wild 350.0
| Parrot Captive 30.0
| Wild 20.0
| >>> df.groupby(level=0).mean()
| Max Speed
| Animal
| Falcon 370.0
| Parrot 25.0
| >>> df.groupby(level=1).mean()
| Max Speed
| Type
| Captive 210.0
| Wild 185.0
|
| head(self, n=5)
| Return the first `n` rows.
|
| This function returns the first `n` rows for the object based

```

```

| on position. It is useful for quickly testing if your object
| has the right type of data in it.
|
| Parameters
| -----
| n : int, default 5
| Number of rows to select.
|
| Returns
| -----
| obj_head : same type as caller
| The first `n` rows of the caller object.
|
| See Also
| -----
| DataFrame.tail: Returns the last `n` rows.
|
| Examples
| -----
| >>> df = pd.DataFrame({'animal':['alligator', 'bee', 'falcon',
'lion',
| ... 'monkey', 'parrot', 'shark', 'whale',
'lion', 'zebra']})
| >>> df
| animal
| 0 alligator
| 1 bee
| 2 falcon
| 3 lion
| 4 monkey
| 5 parrot
| 6 shark
| 7 whale
| 8 zebra
|
| Viewing the first 5 lines
|
| >>> df.head()
| animal
| 0 alligator
| 1 bee
| 2 falcon
| 3 lion
| 4 monkey

```

```

|
| Viewing the first `n` lines (three in this case)
|
| >>> df.head(3)
| animal
| 0 alligator
| 1 bee
| 2 falcon
|
| infer_objects(self)
| Attempt to infer better dtypes for object columns.
|
| Attempts soft conversion of object-dtyped
| columns, leaving non-object and unconvertible
| columns unchanged. The inference rules are the
| same as during normal Series/DataFrame construction.
|
| .. versionadded:: 0.21.0
|
| Returns
| -----
| converted : same type as input object
|
| See Also
| -----
| to_datetime : Convert argument to datetime.
| to_timedelta : Convert argument to timedelta.
| to_numeric : Convert argument to numeric type.
|
| Examples
| -----
| >>> df = pd.DataFrame({"A": ["a", 1, 2, 3]})
| >>> df = df.iloc[1:]
| >>> df
| A
| 1 1
| 2 2
| 3 3
|
| >>> df.dtypes
| A object
| dtype: object
|
| >>> df.infer_objects().dtypes

```



```

| A int64
| dtype: object
|
| interpolate(self, method='linear', axis=0, limit=None, inplace=False,
limit_direction='forward', limit_area=None, downcast=None, **kwargs)
| Interpolate values according to different methods.
|
| Please note that only ``method='linear'`` is supported for
| DataFrame/Series with a MultiIndex.
|
| Parameters
| -----
| method : str, default 'linear'
| Interpolation technique to use. One of:
|
| * 'linear': Ignore the index and treat the values as equally
| spaced. This is the only method supported on MultiIndexes.
| * 'time': Works on daily and higher resolution data to
interpolate
| given length of interval.
| * 'index', 'values': use the actual numerical values of the
index.
| * 'pad': Fill in NaNs using existing values.
| * 'nearest', 'zero', 'slinear', 'quadratic', 'cubic', 'spline',
| 'barycentric', 'polynomial': Passed to
| `scipy.interpolate.interpld`. These methods use the numerical
| values of the index. Both 'polynomial' and 'spline' require
that
| you also specify an `order` (int), e.g.
| ``df.interpolate(method='polynomial', order=5)``.
| * 'krogh', 'piecewise_polynomial', 'spline', 'pchip', 'akima':
| Wrappers around the SciPy interpolation methods of similar
| names. See `Notes`.
| * 'from_derivatives': Refers to
| `scipy.interpolate.BPoly.from_derivatives` which
| replaces 'piecewise_polynomial' interpolation method in
| scipy 0.18.
|
| .. versionadded:: 0.18.1
|
| Added support for the 'akima' method.
| Added interpolate method 'from_derivatives' which replaces
| 'piecewise_polynomial' in SciPy 0.18; backwards-compatible
with

```

```

| SciPy < 0.18
|
| axis : {0 or 'index', 1 or 'columns', None}, default None
| Axis to interpolate along.
| limit : int, optional
| Maximum number of consecutive NaNs to fill. Must be greater than
| 0.
| inplace : bool, default False
| Update the data in place if possible.
| limit_direction : {'forward', 'backward', 'both'}, default 'forward'
| If limit is specified, consecutive NaNs will be filled in this
| direction.
| limit_area : {'None', 'inside', 'outside'}, default None
| If limit is specified, consecutive NaNs will be filled with this
| restriction.
|
| * ``None``: No fill restriction.
| * 'inside': Only fill NaNs surrounded by valid values
| (interpolate).
| * 'outside': Only fill NaNs outside valid values (extrapolate).
|
| .. versionadded:: 0.23.0
|
| downcast : optional, 'infer' or None, defaults to None
| Downcast dtypes if possible.
| **kwargs
| Keyword arguments to pass on to the interpolating function.
|
| Returns
| -----
| Series or DataFrame
| Returns the same object type as the caller, interpolated at
| some or all ``NaN`` values.
|
| See Also
| -----
| fillna : Fill missing values using different methods.
| scipy.interpolate.Akima1DInterpolator : Piecewise cubic polynomials
| (Akima interpolator).
| scipy.interpolate.BPoly.from_derivatives : Piecewise polynomial in
the
| Bernstein basis.
| scipy.interpolate.interpld : Interpolate a 1-D function.
| scipy.interpolate.KroghInterpolator : Interpolate polynomial (Krogh

```

```

| interpolator).
| scipy.interpolate.PchipInterpolator : PCHIP 1-d monotonic cubic
| interpolation.
| scipy.interpolate.CubicSpline : Cubic spline data interpolator.
|
| Notes
| -----
| The 'krogh', 'piecewise_polynomial', 'spline', 'pchip' and 'akima'
| methods are wrappers around the respective SciPy implementations of
| similar names. These use the actual numerical values of the index.
| For more information on their behavior, see the
| `SciPy documentation
|
| <http://docs.scipy.org/doc/scipy/reference/interpolate.html#univariate-
interpolation>`__
| and `SciPy tutorial
|
| <http://docs.scipy.org/doc/scipy/reference/tutorial/interpolate.html>`___.
|
| Examples
| -----
| Filling in ``NaN`` in a :class:`~pandas.Series` via linear
| interpolation.
|
| >>> s = pd.Series([0, 1, np.nan, 3])
| >>> s
| 0 0.0
| 1 1.0
| 2 NaN
| 3 3.0
| dtype: float64
| >>> s.interpolate()
| 0 0.0
| 1 1.0
| 2 2.0
| 3 3.0
| dtype: float64
|
| Filling in ``NaN`` in a Series by padding, but filling at most two
| consecutive ``NaN`` at a time.
|
| >>> s = pd.Series([np.nan, "single_one", np.nan,
| ... "fill_two_more", np.nan, np.nan, np.nan,
| ... 4.71, np.nan])

```

```

| >>> s
| 0 NaN
| 1 single_one
| 2 NaN
| 3 fill_two_more
| 4 NaN
| 5 NaN
| 6 NaN
| 7 4.71
| 8 NaN
| dtype: object
| >>> s.interpolate(method='pad', limit=2)
| 0 NaN
| 1 single_one
| 2 single_one
| 3 fill_two_more
| 4 fill_two_more
| 5 fill_two_more
| 6 NaN
| 7 4.71
| 8 4.71
| dtype: object
|
| Filling in ``NaN`` in a Series via polynomial interpolation or
splines:
| Both 'polynomial' and 'spline' methods require that you also specify
| an ``order`` (int).
|
| >>> s = pd.Series([0, 2, np.nan, 8])
| >>> s.interpolate(method='polynomial', order=2)
| 0 0.000000
| 1 2.000000
| 2 4.666667
| 3 8.000000
| dtype: float64
|
| Fill the DataFrame forward (that is, going down) along each column
| using linear interpolation.
|
| Note how the last entry in column 'a' is interpolated differently,
| because there is no entry after it to use for interpolation.
| Note how the first entry in column 'b' remains ``NaN``, because there
| is no entry before it to use for interpolation.
|

```

```

| >>> df = pd.DataFrame([(0.0, np.nan, -1.0, 1.0),
| ... (np.nan, 2.0, np.nan, np.nan),
| ... (2.0, 3.0, np.nan, 9.0),
| ... (np.nan, 4.0, -4.0, 16.0)],
| ... columns=list('abcd'))
|
| >>> df
|
| a b c d
| 0 0.0 NaN -1.0 1.0
| 1 NaN 2.0 NaN NaN
| 2 2.0 3.0 NaN 9.0
| 3 NaN 4.0 -4.0 16.0
|
| >>> df.interpolate(method='linear', limit_direction='forward',
axis=0)
|
| a b c d
| 0 0.0 NaN -1.0 1.0
| 1 1.0 2.0 -2.0 5.0
| 2 2.0 3.0 -3.0 9.0
| 3 2.0 4.0 -4.0 16.0
|
|
| Using polynomial interpolation.
|
| >>> df['d'].interpolate(method='polynomial', order=2)
|
| 0 1.0
| 1 4.0
| 2 9.0
| 3 16.0
|
| Name: d, dtype: float64
|
|
| keys(self)
|
| Get the 'info axis' (see Indexing for more)
|
|
| This is index for Series, columns for DataFrame.
|
|
| Returns
| -----
|
| Index
|
| Info axis.
|
|
| last(self, offset)
|
| Convenience method for subsetting final periods of time series data
| based on a date offset.
|
|
| Parameters
| -----

```

```

| offset : string, DateOffset, dateutil.relativedelta
|
| Returns
| -----
| subset : same type as caller
|
| Raises
| -----
| TypeError
| If the index is not a :class:`DatetimeIndex`
|
| See Also
| -----
| first : Select initial periods of time series based on a date offset.
| at_time : Select values at a particular time of the day.
| between_time : Select values between particular times of the day.
|
| Examples
| -----
| >>> i = pd.date_range('2018-04-09', periods=4, freq='2D')
| >>> ts = pd.DataFrame({'A': [1,2,3,4]}, index=i)
| >>> ts
|
| A
| 2018-04-09 1
| 2018-04-11 2
| 2018-04-13 3
| 2018-04-15 4
|
| Get the rows for the last 3 days:
|
| >>> ts.last('3D')
|
| A
| 2018-04-13 3
| 2018-04-15 4
|
| Notice the data for 3 last calendar days were returned, not the last
| 3 observed days in the dataset, and therefore data for 2018-04-11 was
| not returned.
|
| last_valid_index(self)
| Return index for last non-NA/null value.
|
| Returns
| -----

```

```

| scalar : type of index
|
| Notes
| -----
| If all elements are non-NA/null, returns None.
| Also returns None for empty Series/DataFrame.
|
| mask(self, cond, other=nan, inplace=False, axis=None, level=None,
errors='raise', try_cast=False)
| Replace values where the condition is True.
|
| Parameters
| -----
| cond : boolean Series/DataFrame, array-like, or callable
| Where `cond` is False, keep the original value. Where
| True, replace with corresponding value from `other`.
| If `cond` is callable, it is computed on the Series/DataFrame and
| should return boolean Series/DataFrame or array. The callable
must
| not change input Series/DataFrame (though pandas doesn't check
it).
|
| .. versionadded:: 0.18.1
| A callable can be used as cond.
|
| other : scalar, Series/DataFrame, or callable
| Entries where `cond` is True are replaced with
| corresponding value from `other`.
| If other is callable, it is computed on the Series/DataFrame and
| should return scalar or Series/DataFrame. The callable must not
| change input Series/DataFrame (though pandas doesn't check it).
|
| .. versionadded:: 0.18.1
| A callable can be used as other.
|
| inplace : bool, default False
| Whether to perform the operation in place on the data.
| axis : int, default None
| Alignment axis if needed.
| level : int, default None
| Alignment level if needed.
| errors : str, {'raise', 'ignore'}, default 'raise'
| Note that currently this parameter won't affect
| the results and will always coerce to a suitable dtype.

```

```

|
| - 'raise' : allow exceptions to be raised.
| - 'ignore' : suppress exceptions. On error return original
object.
|
| try_cast : bool, default False
| Try to cast the result back to the input type (if possible).
|
| Returns
| -----
| Same type as caller
|
| See Also
| -----
| :func:`DataFrame.where` : Return an object of same shape as
| self.
|
| Notes
| -----
| The mask method is an application of the if-then idiom. For each
| element in the calling DataFrame, if ``cond`` is ``False`` the
| element is used; otherwise the corresponding element from the
DataFrame
| ``other`` is used.
|
| The signature for :func:`DataFrame.where` differs from
| :func:`numpy.where`. Roughly ``df1.where(m, df2)`` is equivalent to
| ``np.where(m, df1, df2)``.
|
| For further details and examples see the ``mask`` documentation in
| :ref:`indexing <indexing.where_mask>`.
|
| Examples
| -----
| >>> s = pd.Series(range(5))
| >>> s.where(s > 0)
| 0 NaN
| 1 1.0
| 2 2.0
| 3 3.0
| 4 4.0
| dtype: float64
|
| >>> s.mask(s > 0)

```



```

| 0 0.0
| 1 NaN
| 2 NaN
| 3 NaN
| 4 NaN
| dtype: float64
|
| >>> s.where(s > 1, 10)
| 0 10
| 1 10
| 2 2
| 3 3
| 4 4
| dtype: int64
|
| >>> df = pd.DataFrame(np.arange(10).reshape(-1, 2), columns=['A',
'B'])
| >>> df
| A B
| 0 0 1
| 1 2 3
| 2 4 5
| 3 6 7
| 4 8 9
| >>> m = df % 3 == 0
| >>> df.where(m, -df)
| A B
| 0 0 -1
| 1 -2 3
| 2 -4 -5
| 3 6 -7
| 4 -8 9
| >>> df.where(m, -df) == np.where(m, df, -df)
| A B
| 0 True True
| 1 True True
| 2 True True
| 3 True True
| 4 True True
| >>> df.where(m, -df) == df.mask(~m, -df)
| A B
| 0 True True
| 1 True True
| 2 True True

```

```

| 3 True True
| 4 True True
|
| pct_change(self, periods=1, fill_method='pad', limit=None, freq=None,
**kwargs)
| Percentage change between the current and a prior element.
|
| Computes the percentage change from the immediately previous row by
| default. This is useful in comparing the percentage of change in a
time
| series of elements.
|
| Parameters
| -----
| periods : int, default 1
| Periods to shift for forming percent change.
| fill_method : str, default 'pad'
| How to handle NAs before computing percent changes.
| limit : int, default None
| The number of consecutive NAs to fill before stopping.
| freq : DateOffset, timedelta, or offset alias string, optional
| Increment to use from time series API (e.g. 'M' or BDay()).
| **kwargs
| Additional keyword arguments are passed into
| `DataFrame.shift` or `Series.shift`.
|
| Returns
| -----
| chg : Series or DataFrame
| The same type as the calling object.
|
| See Also
| -----
| Series.diff : Compute the difference of two elements in a Series.
| DataFrame.diff : Compute the difference of two elements in a
DataFrame.
| Series.shift : Shift the index by some number of periods.
| DataFrame.shift : Shift the index by some number of periods.
|
| Examples
| -----
| **Series**
|
| >>> s = pd.Series([90, 91, 85])

```

```

| >>> s
| 0 90
| 1 91
| 2 85
| dtype: int64
|
| >>> s.pct_change()
| 0 NaN
| 1 0.011111
| 2 -0.065934
| dtype: float64
|
| >>> s.pct_change(periods=2)
| 0 NaN
| 1 NaN
| 2 -0.055556
| dtype: float64
|
| See the percentage change in a Series where filling NAs with last
| valid observation forward to next valid.
|
| >>> s = pd.Series([90, 91, None, 85])
| >>> s
| 0 90.0
| 1 91.0
| 2 NaN
| 3 85.0
| dtype: float64
|
| >>> s.pct_change(fill_method='ffill')
| 0 NaN
| 1 0.011111
| 2 0.000000
| 3 -0.065934
| dtype: float64
|
| **DataFrame**
|
| Percentage change in French franc, Deutsche Mark, and Italian lira
from
| 1980-01-01 to 1980-03-01.
|
| >>> df = pd.DataFrame({
| ... 'FR': [4.0405, 4.0963, 4.3149],

```

```

| ... 'GR': [1.7246, 1.7482, 1.8519],
| ... 'IT': [804.74, 810.01, 860.13]},
| ... index=['1980-01-01', '1980-02-01', '1980-03-01'])
|
| >>> df
|
| FR GR IT
| 1980-01-01 4.0405 1.7246 804.74
| 1980-02-01 4.0963 1.7482 810.01
| 1980-03-01 4.3149 1.8519 860.13
|
|
| >>> df.pct_change()
|
| FR GR IT
| 1980-01-01 NaN NaN NaN
| 1980-02-01 0.013810 0.013684 0.006549
| 1980-03-01 0.053365 0.059318 0.061876
|
|
| Percentage of change in GOOG and APPL stock volume. Shows computing
| the percentage change between columns.
|
| >>> df = pd.DataFrame({
| ... '2016': [1769950, 30586265],
| ... '2015': [1500923, 40912316],
| ... '2014': [1371819, 41403351]},
| ... index=['GOOG', 'APPL'])
| >>> df
|
| 2016 2015 2014
| GOOG 1769950 1500923 1371819
| APPL 30586265 40912316 41403351
|
| >>> df.pct_change(axis='columns')
|
| 2016 2015 2014
| GOOG NaN -0.151997 -0.086016
| APPL NaN 0.337604 0.012002
|
|
| pipe(self, func, *args, **kwargs)
| Apply func(self, *args, **kwargs).
|
| Parameters
| -----
| func : function
| function to apply to the Series/DataFrame.
| ``args``, and ``kwargs`` are passed into ``func``.
| Alternatively a ``(callable, data_keyword)`` tuple where
| ``data_keyword`` is a string indicating the keyword of
| ``callable`` that expects the Series/DataFrame.

```

```

| args : iterable, optional
| positional arguments passed into ``func``.
| kwargs : mapping, optional
| a dictionary of keyword arguments passed into ``func``.
|
| Returns
| -----
| object : the return type of ``func``.
|
| See Also
| -----
| DataFrame.apply
| DataFrame.applymap
| Series.map
|
| Notes
| -----
|
| Use ``.pipe`` when chaining together functions that expect
| Series, DataFrames or GroupBy objects. Instead of writing
|
| >>> f(g(h(df), arg1=a), arg2=b, arg3=c)
|
| You can write
|
| >>> (df.pipe(h)
|pipe(g, arg1=a)
|pipe(f, arg2=b, arg3=c)
| ...)
|
| If you have a function that takes the data as (say) the second
| argument, pass a tuple indicating which keyword expects the
| data. For example, suppose ``f`` takes its data as ``arg2``:
|
| >>> (df.pipe(h)
|pipe(g, arg1=a)
|pipe((f, 'arg2'), arg1=a, arg3=c)
| ...)
|
| pop(self, item)
| Return item and drop from frame. Raise KeyError if not found.
|
| Parameters
| -----

```

```

| item : str
| Label of column to be popped.
|
| Returns
| -----
| Series
|
| Examples
| -----
| >>> df = pd.DataFrame([('falcon', 'bird', 389.0),
| ... ('parrot', 'bird', 24.0),
| ... ('lion', 'mammal', 80.5),
| ... ('monkey', 'mammal', np.nan)],
| ... columns=('name', 'class', 'max_speed'))
| >>> df
| name class max_speed
| 0 falcon bird 389.0
| 1 parrot bird 24.0
| 2 lion mammal 80.5
| 3 monkey mammal NaN
|
| >>> df.pop('class')
| 0 bird
| 1 bird
| 2 mammal
| 3 mammal
| Name: class, dtype: object
|
| >>> df
| name max_speed
| 0 falcon 389.0
| 1 parrot 24.0
| 2 lion 80.5
| 3 monkey NaN
|
| rank(self, axis=0, method='average', numeric_only=None, na_option='keep',
ascending=True, pct=False)
| Compute numerical data ranks (1 through n) along axis.
|
| By default, equal values are assigned a rank that is the average of
the
| ranks of those values.
|
| Parameters

```

```

| -----
| axis : {0 or 'index', 1 or 'columns'}, default 0
| Index to direct ranking.
| method : {'average', 'min', 'max', 'first', 'dense'}, default
'average'
| How to rank the group of records that have the same value
| (i.e. ties):
|
| * average: average rank of the group
| * min: lowest rank in the group
| * max: highest rank in the group
| * first: ranks assigned in order they appear in the array
| * dense: like 'min', but rank always increases by 1 between
groups
| numeric_only : bool, optional
| For DataFrame objects, rank only numeric columns if set to True.
| na_option : {'keep', 'top', 'bottom'}, default 'keep'
| How to rank NaN values:
|
| * keep: assign NaN rank to NaN values
| * top: assign smallest rank to NaN values if ascending
| * bottom: assign highest rank to NaN values if ascending
| ascending : bool, default True
| Whether or not the elements should be ranked in ascending order.
| pct : bool, default False
| Whether or not to display the returned rankings in percentile
| form.
|
| Returns
| -----
| same type as caller
| Return a Series or DataFrame with data ranks as values.
|
| See Also
| -----
| core.groupby.GroupBy.rank : Rank of values within each group.
|
| Examples
| -----
|
| >>> df = pd.DataFrame(data={'Animal': ['cat', 'penguin', 'dog',
| ... 'spider', 'snake'],
| ... 'Number_legs': [4, 2, 4, 8, np.nan]})
| >>> df

```

```

| Animal Number_legs
| 0 cat 4.0
| 1 penguin 2.0
| 2 dog 4.0
| 3 spider 8.0
| 4 snake NaN
|
|
| The following example shows how the method behaves with the above
| parameters:
|
| * default_rank: this is the default behaviour obtained without using
| any parameter.
| * max_rank: setting ``method = 'max'`` the records that have the
| same values are ranked using the highest rank (e.g.: since 'cat'
| and 'dog' are both in the 2nd and 3rd position, rank 3 is
assigned.)
| * NA_bottom: choosing ``na_option = 'bottom'``, if there are records
| with NaN values they are placed at the bottom of the ranking.
| * pct_rank: when setting ``pct = True``, the ranking is expressed as
| percentile rank.
|
|
| >>> df['default_rank'] = df['Number_legs'].rank()
| >>> df['max_rank'] = df['Number_legs'].rank(method='max')
| >>> df['NA_bottom'] = df['Number_legs'].rank(na_option='bottom')
| >>> df['pct_rank'] = df['Number_legs'].rank(pct=True)
| >>> df
|
| Animal Number_legs default_rank max_rank NA_bottom pct_rank
| 0 cat 4.0 2.5 3.0 2.5 0.625
| 1 penguin 2.0 1.0 1.0 1.0 0.250
| 2 dog 4.0 2.5 3.0 2.5 0.625
| 3 spider 8.0 4.0 4.0 4.0 1.000
| 4 snake NaN NaN NaN 5.0 NaN
|
|
| reindex_like(self, other, method=None, copy=True, limit=None,
tolerance=None)
|
| Return an object with matching indices as other object.
|
|
| Conform the object to the same index on all axes. Optional
| filling logic, placing NaN in locations having no value
| in the previous index. A new object is produced unless the
| new index is equivalent to the current one and copy=False.
|
|
| Parameters
| -----

```



```

| other : Object of the same data type
| Its row and column indices are used to define the new indices
| of this object.
| method : {None, 'backfill'/'bfill', 'pad'/'ffill', 'nearest'}
| Method to use for filling holes in reindexed DataFrame.
| Please note: this is only applicable to DataFrames/Series with a
| monotonically increasing/decreasing index.
|
| * None (default): don't fill gaps
| * pad / ffill: propagate last valid observation forward to next
| valid
| * backfill / bfill: use next valid observation to fill gap
| * nearest: use nearest valid observations to fill gap
|
| copy : bool, default True
| Return a new object, even if the passed indexes are the same.
| limit : int, default None
| Maximum number of consecutive labels to fill for inexact matches.
| tolerance : optional
| Maximum distance between original and new labels for inexact
| matches. The values of the index at the matching locations must
| satisfy the equation ``abs(index[indexer] - target) <=
tolerance``.
|
| Tolerance may be a scalar value, which applies the same tolerance
| to all values, or list-like, which applies variable tolerance per
| element. List-like includes list, tuple, array, Series, and must
be
| the same size as the index and its dtype must exactly match the
| index's type.
|
| .. versionadded:: 0.21.0 (list-like tolerance)
|
| Returns
| -----
| Series or DataFrame
| Same type as caller, but with changed indices on each axis.
|
| See Also
| -----
| DataFrame.set_index : Set row labels.
| DataFrame.reset_index : Remove row labels or move them to new
columns.
| DataFrame.reindex : Change to new indices or expand indices.

```

| Notes

| -----

| Same as calling

| ```.reindex(index=other.index, columns=other.columns,...)``.`

| Examples

| -----

```
>>> df1 = pd.DataFrame([[24.3, 75.7, 'high'],
... [31, 87.8, 'high'],
... [22, 71.6, 'medium'],
... [35, 95, 'medium']],
... columns=['temp_celsius', 'temp_fahrenheit', 'windspeed'],
... index=pd.date_range(start='2014-02-12',
... end='2014-02-15', freq='D'))
```

```
>>> df1
```

	temp_celsius	temp_fahrenheit	windspeed
2014-02-12	24.3	75.7	high
2014-02-13	31.0	87.8	high
2014-02-14	22.0	71.6	medium
2014-02-15	35.0	95.0	medium

```
>>> df2 = pd.DataFrame([[28, 'low'],
... [30, 'low'],
... [35.1, 'medium']],
... columns=['temp_celsius', 'windspeed'],
... index=pd.DatetimeIndex(['2014-02-12', '2014-02-13',
... '2014-02-15']))
```

```
>>> df2
```

	temp_celsius	windspeed
2014-02-12	28.0	low
2014-02-13	30.0	low
2014-02-15	35.1	medium

```
>>> df2.reindex_like(df1)
```

	temp_celsius	temp_fahrenheit	windspeed
2014-02-12	28.0	NaN	low
2014-02-13	30.0	NaN	low
2014-02-14	NaN	NaN	NaN
2014-02-15	35.1	NaN	medium

```

| rename_axis(self, mapper=None, index=None, columns=None, axis=None,
copy=True, inplace=False)
| Set the name of the axis for the index or columns.
|
| Parameters
| -----
| mapper : scalar, list-like, optional
| Value to set the axis name attribute.
| index, columns : scalar, list-like, dict-like or function, optional
| A scalar, list-like, dict-like or functions transformations to
| apply to that axis' values.
|
| Use either ``mapper`` and ``axis`` to
| specify the axis to target with ``mapper``, or ``index``
| and/or ``columns``.
|
| .. versionchanged:: 0.24.0
|
| axis : {0 or 'index', 1 or 'columns'}, default 0
| The axis to rename.
| copy : bool, default True
| Also copy underlying data.
| inplace : bool, default False
| Modifies the object directly, instead of creating a new Series
| or DataFrame.
|
| Returns
| -----
| Series, DataFrame, or None
| The same type as the caller or None if `inplace` is True.
|
| See Also
| -----
| Series.rename : Alter Series index labels or name.
| DataFrame.rename : Alter DataFrame index labels or name.
| Index.rename : Set new names on index.
|
| Notes
| -----
| ``DataFrame.rename_axis`` supports two calling conventions
|
| * ``(index=index_mapper, columns=columns_mapper, ...)``
| * ``(mapper, axis={'index', 'columns'}, ...)``
|

```

| The first calling convention will only modify the names of  
| the index and/or the names of the Index object that is the columns.  
| In this case, the parameter ``copy`` is ignored.

| The second calling convention will modify the names of the  
| the corresponding index if mapper is a list or a scalar.  
| However, if mapper is dict-like or a function, it will use the  
| deprecated behavior of modifying the axis \*labels\*.

| We *highly* recommend using keyword arguments to clarify your  
| intent.

#### | Examples

| -----

#### | **\*\*Series\*\***

| >>> s = pd.Series(["dog", "cat", "monkey"])

| >>> s

| 0 dog

| 1 cat

| 2 monkey

| dtype: object

| >>> s.rename\_axis("animal")

| animal

| 0 dog

| 1 cat

| 2 monkey

| dtype: object

#### | **\*\*DataFrame\*\***

| >>> df = pd.DataFrame({"num\_legs": [4, 4, 2],

| ... "num\_arms": [0, 0, 2]},

| ... ["dog", "cat", "monkey"])

| >>> df

| num\_legs num\_arms

| dog 4 0

| cat 4 0

| monkey 2 2

| >>> df = df.rename\_axis("animal")

| >>> df

| num\_legs num\_arms

| animal

| dog 4 0

```

| cat 4 0
| monkey 2 2
| >>> df = df.rename_axis("limbs", axis="columns")
| >>> df
| limbs num_legs num_arms
| animal
| dog 4 0
| cat 4 0
| monkey 2 2
|
| **MultiIndex**
|
| >>> df.index = pd.MultiIndex.from_product(['mammal'],
| ... ['dog', 'cat', 'monkey']],
| ... names=['type', 'name'])
| >>> df
| limbs num_legs num_arms
| type name
| mammal dog 4 0
| cat 4 0
| monkey 2 2
|
| >>> df.rename_axis(index={'type': 'class'})
| limbs num_legs num_arms
| class name
| mammal dog 4 0
| cat 4 0
| monkey 2 2
|
| >>> df.rename_axis(columns=str.upper)
| LIMBS num_legs num_arms
| type name
| mammal dog 4 0
| cat 4 0
| monkey 2 2
|
| resample(self, rule, how=None, axis=0, fill_method=None, closed=None,
label=None, convention='start', kind=None, loffset=None, limit=None, base=0,
on=None, level=None)
| Resample time-series data.
|
| Convenience method for frequency conversion and resampling of time
| series. Object must have a datetime-like index (`DatetimeIndex`,
| `PeriodIndex`, or `TimedeltaIndex`), or pass datetime-like values

```

```

| to the `on` or `level` keyword.
|
| Parameters
| -----
| rule : DateOffset, Timedelta or str
| The offset string or object representing target conversion.
| how : str
| Method for down/re-sampling, default to 'mean' for downsampling.
|
| .. deprecated:: 0.18.0
| The new syntax is ``.resample(...).mean()`` , or
| ``.resample(...).apply(<func>)``
| axis : {0 or 'index', 1 or 'columns'}, default 0
| Which axis to use for up- or down-sampling. For `Series` this
| will default to 0, i.e. along the rows. Must be
| `DatetimeIndex`, `TimedeltaIndex` or `PeriodIndex`.
| fill_method : str, default None
| Filling method for upsampling.
|
| .. deprecated:: 0.18.0
| The new syntax is ``.resample(...).<func>()`` ,
| e.g. ``.resample(...).pad()``
| closed : {'right', 'left'}, default None
| Which side of bin interval is closed. The default is 'left'
| for all frequency offsets except for 'M', 'A', 'Q', 'BM',
| 'BA', 'BQ', and 'W' which all have a default of 'right'.
| label : {'right', 'left'}, default None
| Which bin edge label to label bucket with. The default is 'left'
| for all frequency offsets except for 'M', 'A', 'Q', 'BM',
| 'BA', 'BQ', and 'W' which all have a default of 'right'.
| convention : {'start', 'end', 's', 'e'}, default 'start'
| For `PeriodIndex` only, controls whether to use the start or
| end of `rule`.
| kind : {'timestamp', 'period'}, optional, default None
| Pass 'timestamp' to convert the resulting index to a
| `DateTimeIndex` or 'period' to convert it to a `PeriodIndex`.
| By default the input representation is retained.
| loffset : timedelta, default None
| Adjust the resampled time labels.
| limit : int, default None
| Maximum size gap when reindexing with `fill_method`.
|
| .. deprecated:: 0.18.0
| base : int, default 0

```

```

| For frequencies that evenly subdivide 1 day, the "origin" of the
| aggregated intervals. For example, for '5min' frequency, base
could
| range from 0 through 4. Defaults to 0.
| on : str, optional
| For a DataFrame, column to use instead of index for resampling.
| Column must be datetime-like.
|
| .. versionadded:: 0.19.0
|
| level : str or int, optional
| For a MultiIndex, level (name or number) to use for
| resampling. `level` must be datetime-like.
|
| .. versionadded:: 0.19.0
|
| Returns
| -----
| Resampler object
|
| See Also
| -----
| groupby : Group by mapping, function, label, or list of labels.
| Series.resample : Resample a Series.
| DataFrame.resample: Resample a DataFrame.
|
| Notes
| -----
| See the `user guide
| <https://pandas.pydata.org/pandas-
docs/stable/user_guide/timeseries.html#resampling>`
| for more.
|
| To learn more about the offset strings, please see `this link
| <http://pandas.pydata.org/pandas-
docs/stable/user_guide/timeseries.html#dateoffset-objects>`.
|
| Examples
| -----
|
| Start by creating a series with 9 one minute timestamps.
|
| >>> index = pd.date_range('1/1/2000', periods=9, freq='T')
| >>> series = pd.Series(range(9), index=index)

```

```

| >>> series
| 2000-01-01 00:00:00 0
| 2000-01-01 00:01:00 1
| 2000-01-01 00:02:00 2
| 2000-01-01 00:03:00 3
| 2000-01-01 00:04:00 4
| 2000-01-01 00:05:00 5
| 2000-01-01 00:06:00 6
| 2000-01-01 00:07:00 7
| 2000-01-01 00:08:00 8
| Freq: T, dtype: int64
|
| Downsample the series into 3 minute bins and sum the values
| of the timestamps falling into a bin.
|
| >>> series.resample('3T').sum()
| 2000-01-01 00:00:00 3
| 2000-01-01 00:03:00 12
| 2000-01-01 00:06:00 21
| Freq: 3T, dtype: int64
|
| Downsample the series into 3 minute bins as above, but label each
| bin using the right edge instead of the left. Please note that the
| value in the bucket used as the label is not included in the bucket,
| which it labels. For example, in the original series the
| bucket ``2000-01-01 00:03:00`` contains the value 3, but the summed
| value in the resampled bucket with the label ``2000-01-01 00:03:00``
| does not include 3 (if it did, the summed value would be 6, not 3).
| To include this value close the right side of the bin interval as
| illustrated in the example below this one.
|
| >>> series.resample('3T', label='right').sum()
| 2000-01-01 00:03:00 3
| 2000-01-01 00:06:00 12
| 2000-01-01 00:09:00 21
| Freq: 3T, dtype: int64
|
| Downsample the series into 3 minute bins as above, but close the
right
| side of the bin interval.
|
| >>> series.resample('3T', label='right', closed='right').sum()
| 2000-01-01 00:00:00 0
| 2000-01-01 00:03:00 6

```



```

| 2000-01-01 00:06:00 15
| 2000-01-01 00:09:00 15
| Freq: 3T, dtype: int64
|
| Upsample the series into 30 second bins.
|
| >>> series.resample('30S').asfreq()[0:5] # Select first 5 rows
| 2000-01-01 00:00:00 0.0
| 2000-01-01 00:00:30 NaN
| 2000-01-01 00:01:00 1.0
| 2000-01-01 00:01:30 NaN
| 2000-01-01 00:02:00 2.0
| Freq: 30S, dtype: float64
|
| Upsample the series into 30 second bins and fill the ``NaN``
| values using the ``pad`` method.
|
| >>> series.resample('30S').pad()[0:5]
| 2000-01-01 00:00:00 0
| 2000-01-01 00:00:30 0
| 2000-01-01 00:01:00 1
| 2000-01-01 00:01:30 1
| 2000-01-01 00:02:00 2
| Freq: 30S, dtype: int64
|
| Upsample the series into 30 second bins and fill the
| ``NaN`` values using the ``bfill`` method.
|
| >>> series.resample('30S').bfill()[0:5]
| 2000-01-01 00:00:00 0
| 2000-01-01 00:00:30 1
| 2000-01-01 00:01:00 1
| 2000-01-01 00:01:30 2
| 2000-01-01 00:02:00 2
| Freq: 30S, dtype: int64
|
| Pass a custom function via ``apply``
|
| >>> def custom_resampler(array_like):
| ... return np.sum(array_like) + 5
| ...
| >>> series.resample('3T').apply(custom_resampler)
| 2000-01-01 00:00:00 8
| 2000-01-01 00:03:00 17

```

```

| 2000-01-01 00:06:00 26
| Freq: 3T, dtype: int64
|
| For a Series with a PeriodIndex, the keyword `convention` can be
| used to control whether to use the start or end of `rule`.
|
| Resample a year by quarter using 'start' `convention`. Values are
| assigned to the first quarter of the period.
|
| >>> s = pd.Series([1, 2], index=pd.period_range('2012-01-01',
| ... freq='A',
| ... periods=2))
| >>> s
| 2012 1
| 2013 2
| Freq: A-DEC, dtype: int64
| >>> s.resample('Q', convention='start').asfreq()
| 2012Q1 1.0
| 2012Q2 NaN
| 2012Q3 NaN
| 2012Q4 NaN
| 2013Q1 2.0
| 2013Q2 NaN
| 2013Q3 NaN
| 2013Q4 NaN
| Freq: Q-DEC, dtype: float64
|
| Resample quarters by month using 'end' `convention`. Values are
| assigned to the last month of the period.
|
| >>> q = pd.Series([1, 2, 3, 4], index=pd.period_range('2018-01-01',
| ... freq='Q',
| ... periods=4))
| >>> q
| 2018Q1 1
| 2018Q2 2
| 2018Q3 3
| 2018Q4 4
| Freq: Q-DEC, dtype: int64
| >>> q.resample('M', convention='end').asfreq()
| 2018-03 1.0
| 2018-04 NaN
| 2018-05 NaN
| 2018-06 2.0

```

```

| 2018-07 NaN
| 2018-08 NaN
| 2018-09 3.0
| 2018-10 NaN
| 2018-11 NaN
| 2018-12 4.0
| Freq: M, dtype: float64
|
| For DataFrame objects, the keyword `on` can be used to specify the
| column instead of the index for resampling.
|
| >>> d = dict({'price': [10, 11, 9, 13, 14, 18, 17, 19],
| ... 'volume': [50, 60, 40, 100, 50, 100, 40, 50]})
| >>> df = pd.DataFrame(d)
| >>> df['week_starting'] = pd.date_range('01/01/2018',
| ... periods=8,
| ... freq='W')
| >>> df
|
| price volume week_starting
| 0 10 50 2018-01-07
| 1 11 60 2018-01-14
| 2 9 40 2018-01-21
| 3 13 100 2018-01-28
| 4 14 50 2018-02-04
| 5 18 100 2018-02-11
| 6 17 40 2018-02-18
| 7 19 50 2018-02-25
| >>> df.resample('M', on='week_starting').mean()
|
| price volume
| week_starting
| 2018-01-31 10.75 62.5
| 2018-02-28 17.00 60.0
|
| For a DataFrame with MultiIndex, the keyword `level` can be used to
| specify on which level the resampling needs to take place.
|
| >>> days = pd.date_range('1/1/2000', periods=4, freq='D')
| >>> d2 = dict({'price': [10, 11, 9, 13, 14, 18, 17, 19],
| ... 'volume': [50, 60, 40, 100, 50, 100, 40, 50]})
| >>> df2 = pd.DataFrame(d2,
| ... index=pd.MultiIndex.from_product([days,
| ... ['morning',
| ... 'afternoon']])

```

```

| ...
| >>> df2
|
| price volume
| 2000-01-01 morning 10 50
| afternoon 11 60
| 2000-01-02 morning 9 40
| afternoon 13 100
| 2000-01-03 morning 14 50
| afternoon 18 100
| 2000-01-04 morning 17 40
| afternoon 19 50
| >>> df2.resample('D', level=0).sum()
|
| price volume
| 2000-01-01 21 110
| 2000-01-02 22 140
| 2000-01-03 32 150
| 2000-01-04 36 90
|
| sample(self, n=None, frac=None, replace=False, weights=None,
random_state=None, axis=None)
|
| Return a random sample of items from an axis of object.
|
| You can use `random_state` for reproducibility.
|
| Parameters
| -----
|
| n : int, optional
| Number of items from axis to return. Cannot be used with `frac`.
| Default = 1 if `frac` = None.
| frac : float, optional
| Fraction of axis items to return. Cannot be used with `n`.
| replace : bool, default False
| Sample with or without replacement.
| weights : str or ndarray-like, optional
| Default 'None' results in equal probability weighting.
| If passed a Series, will align with target object on index. Index
| values in weights not found in sampled object will be ignored and
| index values in sampled object not in weights will be assigned
| weights of zero.
| If called on a DataFrame, will accept the name of a column
| when axis = 0.
| Unless weights are a Series, weights must be same length as axis
| being sampled.
| If weights do not sum to 1, they will be normalized to sum to 1.

```

```

| Missing values in the weights column will be treated as zero.
| Infinite values not allowed.
| random_state : int or numpy.random.RandomState, optional
| Seed for the random number generator (if int), or numpy
RandomState
| object.
| axis : int or string, optional
| Axis to sample. Accepts axis number or name. Default is stat axis
| for given data type (0 for Series and DataFrames).
|
| Returns
| -----
| Series or DataFrame
| A new object of same type as caller containing `n` items randomly
| sampled from the caller object.
|
| See Also
| -----
| numpy.random.choice: Generates a random sample from a given 1-D numpy
| array.
|
| Examples
| -----
| >>> df = pd.DataFrame({'num_legs': [2, 4, 8, 0],
| ... 'num_wings': [2, 0, 0, 0],
| ... 'num_specimen_seen': [10, 2, 1, 8]},
| ... index=['falcon', 'dog', 'spider', 'fish'])
| >>> df
| num_legs num_wings num_specimen_seen
| falcon 2 2 10
| dog 4 0 2
| spider 8 0 1
| fish 0 0 8
|
| Extract 3 random elements from the ``Series`` ``df['num_legs']``:
| Note that we use `random_state` to ensure the reproducibility of
| the examples.
|
| >>> df['num_legs'].sample(n=3, random_state=1)
| fish 0
| spider 8
| falcon 2
| Name: num_legs, dtype: int64
|

```

```
| A random 50% sample of the ``DataFrame`` with replacement:
|
| >>> df.sample(frac=0.5, replace=True, random_state=1)
| num_legs num_wings num_specimen_seen
| dog 4 0 2
| fish 0 0 8
|
| Using a DataFrame column as weights. Rows with larger value in the
| `num_specimen_seen` column are more likely to be sampled.
|
| >>> df.sample(n=2, weights='num_specimen_seen', random_state=1)
| num_legs num_wings num_specimen_seen
| falcon 2 2 10
| fish 0 0 8
|
| set_axis(self, labels, axis=0, inplace=None)
| Assign desired index to given axis.
|
| Indexes for column or row labels can be changed by assigning
| a list-like or Index.
|
| .. versionchanged:: 0.21.0
|
| The signature is now `labels` and `axis`, consistent with
| the rest of pandas API. Previously, the `axis` and `labels`
| arguments were respectively the first and second positional
| arguments.
|
| Parameters
| -----
| labels : list-like, Index
| The values for the new index.
|
| axis : {0 or 'index', 1 or 'columns'}, default 0
| The axis to update. The value 0 identifies the rows, and 1
| identifies the columns.
|
| inplace : bool, default None
| Whether to return a new %(klass)s instance.
|
| .. warning::
|
| ``inplace=None`` currently falls back to to True, but in a
| future version, will default to False. Use inplace=True
```

```

| explicitly rather than relying on the default.
|
| Returns
| -----
| renamed : %(klass)s or None
|
| An object of same type as caller if inplace=False, None
otherwise.
|
| See Also
| -----
| DataFrame.rename_axis : Alter the name of the index or columns.
|
| Examples
| -----
| **Series**
|
| >>> s = pd.Series([1, 2, 3])
| >>> s
| 0 1
| 1 2
| 2 3
| dtype: int64
|
| >>> s.set_axis(['a', 'b', 'c'], axis=0, inplace=False)
| a 1
| b 2
| c 3
| dtype: int64
|
| The original object is not modified.
|
| >>> s
| 0 1
| 1 2
| 2 3
| dtype: int64
|
| **DataFrame**
|
| >>> df = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})
|
| Change the row labels.
|
| >>> df.set_axis(['a', 'b', 'c'], axis='index', inplace=False)

```

```
| A B
| a 1 4
| b 2 5
| c 3 6
```

```
| Change the column labels.
```

```
| >>> df.set_axis(['I', 'II'], axis='columns', inplace=False)
```

```
| I II
| 0 1 4
| 1 2 5
| 2 3 6
```

```
| Now, update the labels inplace.
```

```
| >>> df.set_axis(['i', 'ii'], axis='columns', inplace=True)
```

```
| >>> df
| i ii
| 0 1 4
| 1 2 5
| 2 3 6
```

```
| slice_shift(self, periods=1, axis=0)
```

```
| Equivalent to `shift` without copying data. The shifted data will
| not include the dropped periods and the shifted axis will be smaller
| than the original.
```

```
| Parameters
```

```
| -----
```

```
| periods : int
```

```
| Number of periods to move, can be positive or negative
```

```
| Returns
```

```
| -----
```

```
| shifted : same type as caller
```

```
| Notes
```

```
| -----
```

```
| While the `slice_shift` is faster than `shift`, you may pay for it
| later during alignment.
```

```
| squeeze(self, axis=None)
```

```
| Squeeze 1 dimensional axis objects into scalars.
```



```

| Series or DataFrames with a single element are squeezed to a scalar.
| DataFrames with a single column or a single row are squeezed to a
| Series. Otherwise the object is unchanged.
|
| This method is most useful when you don't know if your
| object is a Series or DataFrame, but you do know it has just a single
| column. In that case you can safely call `squeeze` to ensure you have
a
| Series.
|
| Parameters
| -----
| axis : {0 or 'index', 1 or 'columns', None}, default None
| A specific axis to squeeze. By default, all length-1 axes are
| squeezed.
|
| .. versionadded:: 0.20.0
|
| Returns
| -----
| DataFrame, Series, or scalar
| The projection after squeezing `axis` or all the axes.
|
| See Also
| -----
| Series.iloc : Integer-location based indexing for selecting scalars.
| DataFrame.iloc : Integer-location based indexing for selecting
Series.
| Series.to_frame : Inverse of DataFrame.squeeze for a
| single-column DataFrame.
|
| Examples
| -----
| >>> primes = pd.Series([2, 3, 5, 7])
|
| Slicing might produce a Series with a single value:
|
| >>> even_primes = primes[primes % 2 == 0]
| >>> even_primes
| 0 2
| dtype: int64
|
| >>> even_primes.squeeze()
| 2

```

|  
| Squeezing objects with more than one value in every axis does  
nothing:

```
|
| >>> odd_primes = primes[primes % 2 == 1]
| >>> odd_primes
| 1 3
| 2 5
| 3 7
| dtype: int64
```

```
|
| >>> odd_primes.squeeze()
| 1 3
| 2 5
| 3 7
| dtype: int64
```

| Squeezing is even more effective when used with DataFrames.

```
|
| >>> df = pd.DataFrame([[1, 2], [3, 4]], columns=['a', 'b'])
| >>> df
| a b
| 0 1 2
| 1 3 4
```

| Slicing a single column will produce a DataFrame with the columns  
| having only one value:

```
|
| >>> df_a = df[['a']]
| >>> df_a
| a
| 0 1
| 1 3
```

| So the columns can be squeezed down, resulting in a Series:

```
|
| >>> df_a.squeeze('columns')
| 0 1
| 1 3
| Name: a, dtype: int64
```

| Slicing a single row from a single column will produce a single  
| scalar DataFrame:

```

| >>> df_0a = df.loc[df.index < 1, ['a']]
| >>> df_0a
| a
| 0 1
|
| Squeezing the rows produces a single scalar Series:
|
| >>> df_0a.squeeze('rows')
| a 1
| Name: 0, dtype: int64
|
| Squeezing all axes will project directly into a scalar:
|
| >>> df_0a.squeeze()
| 1
|
| swapaxes(self, axis1, axis2, copy=True)
| Interchange axes and swap values axes appropriately.
|
| Returns
| -----
| y : same as input
|
| tail(self, n=5)
| Return the last `n` rows.
|
| This function returns last `n` rows from the object based on
| position. It is useful for quickly verifying data, for example,
| after sorting or appending rows.
|
| Parameters
| -----
| n : int, default 5
| Number of rows to select.
|
| Returns
| -----
| type of caller
| The last `n` rows of the caller object.
|
| See Also
| -----
| DataFrame.head : The first `n` rows of the caller object.
|

```

```

| Examples
| -----
| >>> df = pd.DataFrame({'animal':['alligator', 'bee', 'falcon',
'lion',
| ... 'monkey', 'parrot', 'shark', 'whale',
'zebra']})
|
| >>> df
| animal
| 0 alligator
| 1 bee
| 2 falcon
| 3 lion
| 4 monkey
| 5 parrot
| 6 shark
| 7 whale
| 8 zebra
|
| Viewing the last 5 lines
|
| >>> df.tail()
| animal
| 4 monkey
| 5 parrot
| 6 shark
| 7 whale
| 8 zebra
|
| Viewing the last `n` lines (three in this case)
|
| >>> df.tail(3)
| animal
| 6 shark
| 7 whale
| 8 zebra
|
| take(self, indices, axis=0, is_copy=True, **kwargs)
| Return the elements in the given *positional* indices along an axis.
|
| This means that we are not indexing according to actual values in
| the index attribute of the object. We are indexing according to the
| actual position of the element in the object.
|
| Parameters

```

```

| -----
| indices : array-like
| An array of ints indicating which positions to take.
| axis : {0 or 'index', 1 or 'columns', None}, default 0
| The axis on which to select elements. ``0`` means that we are
| selecting rows, ``1`` means that we are selecting columns.
| is_copy : bool, default True
| Whether to return a copy of the original object or not.
| **kwargs
| For compatibility with :meth:`numpy.take`. Has no effect on the
| output.
|
| Returns
| -----
| taken : same type as caller
| An array-like containing the elements taken from the object.
|
| See Also
| -----
| DataFrame.loc : Select a subset of a DataFrame by labels.
| DataFrame.iloc : Select a subset of a DataFrame by positions.
| numpy.take : Take elements from an array along an axis.
|
| Examples
| -----
| >>> df = pd.DataFrame([('falcon', 'bird', 389.0),
| ... ('parrot', 'bird', 24.0),
| ... ('lion', 'mammal', 80.5),
| ... ('monkey', 'mammal', np.nan)],
| ... columns=['name', 'class', 'max_speed'],
| ... index=[0, 2, 3, 1])
| >>> df
|
| name class max_speed
| 0 falcon bird 389.0
| 2 parrot bird 24.0
| 3 lion mammal 80.5
| 1 monkey mammal NaN
|
| Take elements at positions 0 and 3 along the axis 0 (default).
|
| Note how the actual indices selected (0 and 1) do not correspond to
| our selected indices 0 and 3. That's because we are selecting the 0th
| and 3rd rows, not rows whose indices equal 0 and 3.
|

```

```
| >>> df.take([0, 3])
| name class max_speed
| 0 falcon bird 389.0
| 1 monkey mammal NaN
```

| Take elements at indices 1 and 2 along the axis 1 (column selection).

```
| >>> df.take([1, 2], axis=1)
| class max_speed
| 0 bird 389.0
| 2 bird 24.0
| 3 mammal 80.5
| 1 mammal NaN
```

| We may take elements using negative integers for positive indices, starting from the end of the object, just like with Python lists.

```
| >>> df.take([-1, -2])
| name class max_speed
| 1 monkey mammal NaN
| 3 lion mammal 80.5
```

```
| to_clipboard(self, excel=True, sep=None, **kwargs)
| Copy object to the system clipboard.
```

| Write a text representation of object to the system clipboard.  
| This can be pasted into Excel, for example.

| Parameters

| -----

| excel : bool, default True

| - True, use the provided separator, writing in a csv format for allowing easy pasting into excel.

| - False, write a string representation of the object to the clipboard.

| sep : str, default ``'\t'``

| Field delimiter.

| \*\*kwargs

| These parameters will be passed to DataFrame.to\_csv.

| See Also

| -----

| DataFrame.to\_csv : Write a DataFrame to a comma-separated values

```

| (csv) file.
| read_clipboard : Read text from clipboard and pass to read_table.
|
| Notes
| -----
| Requirements for your platform.
|
| - Linux : `xclip`, or `xsel` (with `PyQt4` modules)
| - Windows : none
| - OS X : none
|
| Examples
| -----
| Copy the contents of a DataFrame to the clipboard.
|
| >>> df = pd.DataFrame([[1, 2, 3], [4, 5, 6]], columns=['A', 'B',
| 'C'])
| >>> df.to_clipboard(sep=',')
| ... # Wrote the following to the system clipboard:
| ... # ,A,B,C
| ... # 0,1,2,3
| ... # 1,4,5,6
|
| We can omit the the index by passing the keyword `index` and setting
| it to false.
|
| >>> df.to_clipboard(sep=',', index=False)
| ... # Wrote the following to the system clipboard:
| ... # A,B,C
| ... # 1,2,3
| ... # 4,5,6
|
| to_csv(self, path_or_buf=None, sep=',', na_rep='', float_format=None,
columns=None, header=True, index=True, index_label=None, mode='w',
encoding=None, compression='infer', quoting=None, quotechar='"',
line_terminator=None, chunksize=None, date_format=None, doublequote=True,
escapechar=None, decimal='.')
|
| Write object to a comma-separated values (csv) file.
|
| .. versionchanged:: 0.24.0
| The order of arguments for Series was changed.
|
| Parameters
| -----

```

```

| path_or_buf : str or file handle, default None
| File path or object, if None is provided the result is returned
as
| a string. If a file object is passed it should be opened with
| `newline=''`, disabling universal newlines.
|
| .. versionchanged:: 0.24.0
|
| Was previously named "path" for Series.
|
| sep : str, default ','
| String of length 1. Field delimiter for the output file.
| na_rep : str, default ''
| Missing data representation.
| float_format : str, default None
| Format string for floating point numbers.
| columns : sequence, optional
| Columns to write.
| header : bool or list of str, default True
| Write out the column names. If a list of strings is given it is
| assumed to be aliases for the column names.
|
| .. versionchanged:: 0.24.0
|
| Previously defaulted to False for Series.
|
| index : bool, default True
| Write row names (index).
| index_label : str or sequence, or False, default None
| Column label for index column(s) if desired. If None is given,
and
| `header` and `index` are True, then the index names are used. A
| sequence should be given if the object uses MultiIndex. If
| False do not print fields for index names. Use index_label=False
| for easier importing in R.
| mode : str
| Python write mode, default 'w'.
| encoding : str, optional
| A string representing the encoding to use in the output file,
| defaults to 'utf-8'.
| compression : str, default 'infer'
| Compression mode among the following possible values: {'infer',
| 'gzip', 'bz2', 'zip', 'xz', None}. If 'infer' and `path_or_buf`
| is path-like, then detect compression from the following

```



```

| extensions: '.gz', '.bz2', '.zip' or '.xz'. (otherwise no
| compression).
|
| .. versionchanged:: 0.24.0
|
| 'infer' option added and set to default.
|
| quoting : optional constant from csv module
| Defaults to csv.QUOTE_MINIMAL. If you have set a `float_format`
| then floats are converted to strings and thus
csv.QUOTE_NONNUMERIC
| will treat them as non-numeric.
| quotechar : str, default '"'
| String of length 1. Character used to quote fields.
| line_terminator : str, optional
| The newline character or character sequence to use in the output
| file. Defaults to `os.linesep`, which depends on the OS in which
| this method is called ('\n' for linux, '\r\n' for Windows, i.e.).
|
| .. versionchanged:: 0.24.0
| chunksize : int or None
| Rows to write at a time.
| date_format : str, default None
| Format string for datetime objects.
| doublequote : bool, default True
| Control quoting of `quotechar` inside a field.
| escapechar : str, default None
| String of length 1. Character used to escape `sep` and
`quotechar`
| when appropriate.
| decimal : str, default '.'
| Character recognized as decimal separator. E.g. use ',' for
| European data.
|
| Returns
| -----
| None or str
| If path_or_buf is None, returns the resulting csv format as a
| string. Otherwise returns None.
|
| See Also
| -----
| read_csv : Load a CSV file into a DataFrame.
| to_excel : Write DataFrame to an Excel file.

```

```

|
| Examples
| -----
| >>> df = pd.DataFrame({'name': ['Raphael', 'Donatello'],
| ... 'mask': ['red', 'purple'],
| ... 'weapon': ['sai', 'bo staff']})
| >>> df.to_csv(index=False)
| 'name,mask,weapon\nRaphael,red,sai\nDonatello,purple,bo staff\n'
|
| to_dense(self)
| Return dense representation of Series/DataFrame (as opposed to
sparse).
|
| .. deprecated:: 0.25.0
|
| Returns
| -----
| %(klass)s
| Dense %(klass)s.
|
| to_excel(self, excel_writer, sheet_name='Sheet1', na_rep='',
float_format=None, columns=None, header=True, index=True, index_label=None,
startrow=0, startcol=0, engine=None, merge_cells=True, encoding=None,
inf_rep='inf', verbose=True, freeze_panes=None)
| Write object to an Excel sheet.
|
| To write a single object to an Excel .xlsx file it is only necessary
to
| specify a target file name. To write to multiple sheets it is
necessary to
| create an `ExcelWriter` object with a target file name, and specify a
sheet
| in the file to write to.
|
| Multiple sheets may be written to by specifying unique `sheet_name`.
| With all data written to the file it is necessary to save the
changes.
| Note that creating an `ExcelWriter` object with a file name that
already
| exists will result in the contents of the existing file being erased.
|
| Parameters
| -----
| excel_writer : str or ExcelWriter object

```

```

| File path or existing ExcelWriter.
| sheet_name : str, default 'Sheet1'
| Name of sheet which will contain DataFrame.
| na_rep : str, default ''
| Missing data representation.
| float_format : str, optional
| Format string for floating point numbers. For example
| ``float_format="%.2f"`` will format 0.1234 to 0.12.
| columns : sequence or list of str, optional
| Columns to write.
| header : bool or list of str, default True
| Write out the column names. If a list of string is given it is
| assumed to be aliases for the column names.
| index : bool, default True
| Write row names (index).
| index_label : str or sequence, optional
| Column label for index column(s) if desired. If not specified,
and
| `header` and `index` are True, then the index names are used. A
| sequence should be given if the DataFrame uses MultiIndex.
| startrow : int, default 0
| Upper left cell row to dump data frame.
| startcol : int, default 0
| Upper left cell column to dump data frame.
| engine : str, optional
| Write engine to use, 'openpyxl' or 'xlsxwriter'. You can also set
this
| via the options ``io.excel.xlsx.writer``,
``io.excel.xls.writer``, and
| ``io.excel.xlsm.writer``.
| merge_cells : bool, default True
| Write MultiIndex and Hierarchical Rows as merged cells.
| encoding : str, optional
| Encoding of the resulting excel file. Only necessary for xlwt,
| other writers support unicode natively.
| inf_rep : str, default 'inf'
| Representation for infinity (there is no native representation
for
| infinity in Excel).
| verbose : bool, default True
| Display more information in the error logs.
| freeze_panes : tuple of int (length 2), optional
| Specifies the one-based bottommost row and rightmost column that
| is to be frozen.

```

.. versionadded:: 0.20.0.

#### See Also

-----  
to\_csv : Write DataFrame to a comma-separated values (csv) file.  
ExcelWriter : Class for writing DataFrame objects into excel sheets.  
read\_excel : Read an Excel file into a pandas DataFrame.  
read\_csv : Read a comma-separated values (csv) file into DataFrame.

#### Notes

-----  
For compatibility with :meth:`~DataFrame.to\_csv`,  
to\_excel serializes lists and dicts to strings before writing.

Once a workbook has been saved it is not possible write further data  
without rewriting the whole workbook.

#### Examples

-----  
Create, write to and save a workbook:

```
>>> df1 = pd.DataFrame([['a', 'b'], ['c', 'd']],
... index=['row 1', 'row 2'],
... columns=['col 1', 'col 2'])
>>> df1.to_excel("output.xlsx") # doctest: +SKIP
```

To specify the sheet name:

```
>>> df1.to_excel("output.xlsx",
... sheet_name='Sheet_name_1') # doctest: +SKIP
```

If you wish to write to more than one sheet in the workbook, it is  
necessary to specify an ExcelWriter object:

```
>>> df2 = df1.copy()
>>> with pd.ExcelWriter('output.xlsx') as writer: # doctest: +SKIP
... df1.to_excel(writer, sheet_name='Sheet_name_1')
... df2.to_excel(writer, sheet_name='Sheet_name_2')
```

To set the library that is used to write the Excel file,  
you can pass the `engine` keyword (the default engine is  
automatically chosen depending on the file extension):

```

|
| >>> df1.to_excel('output1.xlsx', engine='xlsxwriter') # doctest:
+SKIP
|
| to_hdf(self, path_or_buf, key, **kwargs)
| Write the contained data to an HDF5 file using HDFStore.
|
| Hierarchical Data Format (HDF) is self-describing, allowing an
| application to interpret the structure and contents of a file with
| no outside information. One HDF file can hold a mix of related
objects
| which can be accessed as a group or as individual objects.
|
| In order to add another DataFrame or Series to an existing HDF file
| please use append mode and a different a key.
|
| For more information see the :ref:`user guide <io.hdf5>`.
|
| Parameters
| -----
| path_or_buf : str or pandas.HDFStore
| File path or HDFStore object.
| key : str
| Identifier for the group in the store.
| mode : {'a', 'w', 'r+'}, default 'a'
| Mode to open file:
|
| - 'w': write, a new file is created (an existing file with
| the same name would be deleted).
| - 'a': append, an existing file is opened for reading and
| writing, and if the file does not exist it is created.
| - 'r+': similar to 'a', but the file must already exist.
| format : {'fixed', 'table'}, default 'fixed'
| Possible values:
|
| - 'fixed': Fixed format. Fast writing/reading. Not-appendable,
| nor searchable.
| - 'table': Table format. Write as a PyTables Table structure
| which may perform worse but allow more flexible operations
| like searching / selecting subsets of the data.
| append : bool, default False
| For Table formats, append the input data to the existing.
| data_columns : list of columns or True, optional
| List of columns to create as indexed data columns for on-disk

```

```

| queries, or True to use all columns. By default only the axes
| of the object are indexed. See :ref:`io.hdf5-query-data-columns`.
| Applicable only to format='table'.
| complevel : {0-9}, optional
| Specifies a compression level for data.
| A value of 0 disables compression.
| complib : {'zlib', 'lzo', 'bzip2', 'blosc'}, default 'zlib'
| Specifies the compression library to be used.
| As of v0.20.2 these additional compressors for Blosc are
supported
| (default if no compressor specified: 'blosc:blosclz'):
| {'blosc:blosclz', 'blosc:lz4', 'blosc:lz4hc', 'blosc:snappy',
| 'blosc:zlib', 'blosc:zstd'}.
| Specifying a compression library which is not available issues
| a ValueError.
| fletcher32 : bool, default False
| If applying compression use the fletcher32 checksum.
| dropna : bool, default False
| If true, ALL nan rows will not be written to store.
| errors : str, default 'strict'
| Specifies how encoding and decoding errors are to be handled.
| See the errors argument for :func:`open` for a full list
| of options.
|
| See Also
| -----
| DataFrame.read_hdf : Read from HDF file.
| DataFrame.to_parquet : Write a DataFrame to the binary parquet
format.
| DataFrame.to_sql : Write to a sql table.
| DataFrame.to_feather : Write out feather-format for DataFrames.
| DataFrame.to_csv : Write out to a csv file.
|
| Examples
| -----
| >>> df = pd.DataFrame({'A': [1, 2, 3], 'B': [4, 5, 6]},
| ... index=['a', 'b', 'c'])
| >>> df.to_hdf('data.h5', key='df', mode='w')
|
| We can add another object to the same file:
|
| >>> s = pd.Series([1, 2, 3, 4])
| >>> s.to_hdf('data.h5', key='s')
|

```

```

| Reading from HDF file:
|
| >>> pd.read_hdf('data.h5', 'df')
| A B
| a 1 4
| b 2 5
| c 3 6
| >>> pd.read_hdf('data.h5', 's')
| 0 1
| 1 2
| 2 3
| 3 4
| dtype: int64
|
| Deleting file with data:
|
| >>> import os
| >>> os.remove('data.h5')
|
| to_json(self, path_or_buf=None, orient=None, date_format=None,
double_precision=10, force_ascii=True, date_unit='ms', default_handler=None,
lines=False, compression='infer', index=True)
| Convert the object to a JSON string.
|
| Note NaN's and None will be converted to null and datetime objects
| will be converted to UNIX timestamps.
|
| Parameters
| -----
| path_or_buf : string or file handle, optional
| File path or object. If not specified, the result is returned as
| a string.
| orient : string
| Indication of expected JSON string format.
|
| * Series
|
| - default is 'index'
| - allowed values are: {'split', 'records', 'index', 'table'}
|
| * DataFrame
|
| - default is 'columns'
| - allowed values are:

```

```

| {'split', 'records', 'index', 'columns', 'values', 'table'}
|
| * The format of the JSON string
|
| - 'split' : dict like {'index' -> [index],
| 'columns' -> [columns], 'data' -> [values]}
| - 'records' : list like
| [{column -> value}, ... , {column -> value}]
| - 'index' : dict like {index -> {column -> value}}
| - 'columns' : dict like {column -> {index -> value}}
| - 'values' : just the values array
| - 'table' : dict like {'schema': {schema}, 'data': {data}}
| describing the data, and the data component is
| like ``orient='records'``.
|
| .. versionchanged:: 0.20.0
|
| date_format : {None, 'epoch', 'iso'}
| Type of date conversion. 'epoch' = epoch milliseconds,
| 'iso' = ISO8601. The default depends on the `orient`. For
| ``orient='table'``, the default is 'iso'. For all other orients,
| the default is 'epoch'.
| double_precision : int, default 10
| The number of decimal places to use when encoding
| floating point values.
| force_ascii : bool, default True
| Force encoded string to be ASCII.
| date_unit : string, default 'ms' (milliseconds)
| The time unit to encode to, governs timestamp and ISO8601
| precision. One of 's', 'ms', 'us', 'ns' for second, millisecond,
| microsecond, and nanosecond respectively.
| default_handler : callable, default None
| Handler to call if object cannot otherwise be converted to a
| suitable format for JSON. Should receive a single argument which
is
| the object to convert and return a serialisable object.
| lines : bool, default False
| If 'orient' is 'records' write out line delimited json format.
Will
| throw ValueError if incorrect 'orient' since others are not list
| like.
|
| .. versionadded:: 0.19.0
|

```



```

| compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}
|
| A string representing the compression to use in the output file,
| only used when the first argument is a filename. By default, the
| compression is inferred from the filename.
|
|
| .. versionadded:: 0.21.0
| .. versionchanged:: 0.24.0
| 'infer' option added and set to default
| index : bool, default True
| Whether to include the index values in the JSON string. Not
| including the index (``index=False``) is only supported when
| orient is 'split' or 'table'.
|
| .. versionadded:: 0.23.0
|
| Returns
| -----
| None or str
| If path_or_buf is None, returns the resulting json format as a
| string. Otherwise returns None.
|
| See Also
| -----
| read_json
|
| Examples
| -----
|
| >>> df = pd.DataFrame([['a', 'b'], ['c', 'd']],
| ... index=['row 1', 'row 2'],
| ... columns=['col 1', 'col 2'])
| >>> df.to_json(orient='split')
| '{"columns":["col 1","col 2"],'
| "index":["row 1","row 2"],'
| "data":[["a","b"],["c","d"]]}'
```

Encoding/decoding a Dataframe using ``'records'`` formatted JSON.  
Note that index labels are not preserved with this encoding.

```

| >>> df.to_json(orient='records')
| '[{"col 1":"a","col 2":"b"}, {"col 1":"c","col 2":"d"}]'
```

Encoding/decoding a Dataframe using ``'index'`` formatted JSON:

```

|
| >>> df.to_json(orient='index')
| '{"row 1":{"col 1":"a","col 2":"b"},"row 2":{"col 1":"c","col
2":"d"}}'
|
| Encoding/decoding a Dataframe using ``'columns'`` formatted JSON:
|
| >>> df.to_json(orient='columns')
| '{"col 1":{"row 1":"a","row 2":"c"},"col 2":{"row 1":"b","row
2":"d"}}'
|
| Encoding/decoding a Dataframe using ``'values'`` formatted JSON:
|
| >>> df.to_json(orient='values')
| '[[["a","b"],["c","d"]]]'
|
| Encoding with Table Schema
|
| >>> df.to_json(orient='table')
| '{"schema": {"fields": [{"name": "index", "type": "string"},
| {"name": "col 1", "type": "string"},
| {"name": "col 2", "type": "string"}],
| "primaryKey": "index",
| "pandas_version": "0.20.0"},
| "data": [{"index": "row 1", "col 1": "a", "col 2": "b"},
| {"index": "row 2", "col 1": "c", "col 2": "d"}]}'
|
| to_latex(self, buf=None, columns=None, col_space=None, header=True,
index=True, na_rep='NaN', formatters=None, float_format=None, sparsify=None,
index_names=True, bold_rows=False, column_format=None, longtable=None,
escape=None, encoding=None, decimal='.', multicolumn=None,
multicolumn_format=None, multirow=None)
|
| Render an object to a LaTeX tabular environment table.
|
| Render an object to a tabular environment table. You can splice
| this into a LaTeX document. Requires \usepackage{booktabs}.
|
| .. versionchanged:: 0.20.2
| Added to Series
|
| Parameters
| -----
| buf : file descriptor or None
| Buffer to write to. If None, the output is returned as a string.

```

```

| columns : list of label, optional
| The subset of columns to write. Writes all columns by default.
| col_space : int, optional
| The minimum width of each column.
| header : bool or list of str, default True
| Write out the column names. If a list of strings is given,
| it is assumed to be aliases for the column names.
| index : bool, default True
| Write row names (index).
| na_rep : str, default 'NaN'
| Missing data representation.
| formatters : list of functions or dict of {str: function}, optional
| Formatter functions to apply to columns' elements by position or
| name. The result of each function must be a unicode string.
| List must be of length equal to the number of columns.
| float_format : one-parameter function or str, optional, default None
| Formatter for floating point numbers. For example
| ``float_format="%.2f"`` and ``float_format="{:0.2f}".format``
will
| both result in 0.1234 being formatted as 0.12.
| sparsify : bool, optional
| Set to False for a DataFrame with a hierarchical index to print
| every multiindex key at each row. By default, the value will be
| read from the config module.
| index_names : bool, default True
| Prints the names of the indexes.
| bold_rows : bool, default False
| Make the row labels bold in the output.
| column_format : str, optional
| The columns format as specified in `LaTeX table format
| <https://en.wikibooks.org/wiki/LaTeX/Tables>`__ e.g. 'rcl' for 3
| columns. By default, 'l' will be used for all columns except
| columns of numbers, which default to 'r'.
| longtable : bool, optional
| By default, the value will be read from the pandas config
| module. Use a longtable environment instead of tabular. Requires
| adding a \usepackage{longtable} to your LaTeX preamble.
| escape : bool, optional
| By default, the value will be read from the pandas config
| module. When set to False prevents from escaping latex special
| characters in column names.
| encoding : str, optional
| A string representing the encoding to use in the output file,
| defaults to 'utf-8'.

```

```

| decimal : str, default '.'
| Character recognized as decimal separator, e.g. ',' in Europe.
|
| .. versionadded:: 0.18.0
| multicolumn : bool, default True
| Use \multicolumn to enhance MultiIndex columns.
| The default will be read from the config module.
|
| .. versionadded:: 0.20.0
| multicolumn_format : str, default 'l'
| The alignment for multicolumns, similar to `column_format`
| The default will be read from the config module.
|
| .. versionadded:: 0.20.0
| multirow : bool, default False
| Use \multirow to enhance MultiIndex rows. Requires adding a
| \usepackage{multirow} to your LaTeX preamble. Will print
| centered labels (instead of top-aligned) across the contained
| rows, separating groups via clines. The default will be read
| from the pandas config module.
|
| .. versionadded:: 0.20.0
|
| Returns
| -----
| str or None
| If buf is None, returns the resulting LaTeX format as a
| string. Otherwise returns None.
|
| See Also
| -----
| DataFrame.to_string : Render a DataFrame to a console-friendly
| tabular output.
| DataFrame.to_html : Render a DataFrame as an HTML table.
|
| Examples
| -----
| >>> df = pd.DataFrame({'name': ['Raphael', 'Donatello'],
| ... 'mask': ['red', 'purple'],
| ... 'weapon': ['sai', 'bo staff']})
| >>> df.to_latex(index=False) # doctest: +NORMALIZE_WHITESPACE
| '\\begin{tabular}{lll}\\n\\toprule\\n name & mask & weapon\\n
| \\\\n\\midrule\\n Raphael & red & sai \\\\n\\n Donatello &
| purple & bo staff \\\\n\\n\\bottomrule\\n\\end{tabular}\\n'

```

```

|
| to_msgpack(self, path_or_buf=None, encoding='utf-8', **kwargs)
| Serialize object to input file path using msgpack format.
|
| .. deprecated:: 0.25.0
|
| to_msgpack is deprecated and will be removed in a future version.
| It is recommended to use pyarrow for on-the-wire transmission of
| pandas objects.
|
| Parameters
| -----
| path : string File path, buffer-like, or None
| if None, return generated bytes
| append : bool whether to append to an existing msgpack
| (default is False)
| compress : type of compressor (zlib or blosc), default to None (no
| compression)
|
| Returns
| -----
| None or bytes
| If path_or_buf is None, returns the resulting msgpack format as a
| byte string. Otherwise returns None.
|
| to_pickle(self, path, compression='infer', protocol=4)
| Pickle (serialize) object to file.
|
| Parameters
| -----
| path : str
| File path where the pickled object will be stored.
| compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None},
| default 'infer'
| A string representing the compression to use in the output file.
By
| default, infers from the file extension in specified path.
|
| .. versionadded:: 0.20.0
| protocol : int
| Int which indicates which protocol should be used by the pickler,
| default HIGHEST_PROTOCOL (see [1]_ paragraph 12.1.2). The
possible
| values are 0, 1, 2, 3, 4. A negative value for the protocol

```

```

| parameter is equivalent to setting its value to HIGHEST_PROTOCOL.
|
| .. [1] https://docs.python.org/3/library/pickle.html
| .. versionadded:: 0.21.0
|
| See Also
| -----
| read_pickle : Load pickled pandas object (or any object) from file.
| DataFrame.to_hdf : Write DataFrame to an HDF5 file.
| DataFrame.to_sql : Write DataFrame to a SQL database.
| DataFrame.to_parquet : Write a DataFrame to the binary parquet
format.
|
| Examples
| -----
| >>> original_df = pd.DataFrame({"foo": range(5), "bar": range(5,
10)})
| >>> original_df
| foo bar
| 0 0 5
| 1 1 6
| 2 2 7
| 3 3 8
| 4 4 9
| >>> original_df.to_pickle("./dummy.pkl")
|
| >>> unpickled_df = pd.read_pickle("./dummy.pkl")
| >>> unpickled_df
| foo bar
| 0 0 5
| 1 1 6
| 2 2 7
| 3 3 8
| 4 4 9
|
| >>> import os
| >>> os.remove("./dummy.pkl")
|
| to_sql(self, name, con, schema=None, if_exists='fail', index=True,
index_label=None, chunksize=None, dtype=None, method=None)
| Write records stored in a DataFrame to a SQL database.
|
| Databases supported by SQLAlchemy [1]_ are supported. Tables can be
| newly created, appended to, or overwritten.

```

```

|
| Parameters
| -----
| name : string
| Name of SQL table.
| con : sqlalchemy.engine.Engine or sqlite3.Connection
| Using SQLAlchemy makes it possible to use any DB supported by
that
| library. Legacy support is provided for sqlite3.Connection
objects.
| schema : string, optional
| Specify the schema (if database flavor supports this). If None,
use
| default schema.
| if_exists : {'fail', 'replace', 'append'}, default 'fail'
| How to behave if the table already exists.
|
| * fail: Raise a ValueError.
| * replace: Drop the table before inserting new values.
| * append: Insert new values to the existing table.
|
| index : bool, default True
| Write DataFrame index as a column. Uses `index_label` as the
column
| name in the table.
| index_label : string or sequence, default None
| Column label for index column(s). If None is given (default) and
| `index` is True, then the index names are used.
| A sequence should be given if the DataFrame uses MultiIndex.
| chunksize : int, optional
| Rows will be written in batches of this size at a time. By
default,
| all rows will be written at once.
| dtype : dict, optional
| Specifying the datatype for columns. The keys should be the
column
| names and the values should be the SQLAlchemy types or strings
for
| the sqlite3 legacy mode.
| method : {None, 'multi', callable}, default None
| Controls the SQL insertion clause used:
|
| * None : Uses standard SQL ``INSERT`` clause (one per row).
| * 'multi': Pass multiple values in a single ``INSERT`` clause.

```

```

| * callable with signature `` (pd_table, conn, keys, data_iter)``.
|
| Details and a sample callable implementation can be found in the
| section :ref:`insert method <io.sql.method>`.
|
| .. versionadded:: 0.24.0
|
| Raises
| -----
| ValueError
| When the table already exists and `if_exists` is 'fail' (the
| default).
|
| See Also
| -----
| read_sql : Read a DataFrame from a table.
|
| Notes
| -----
| Timezone aware datetime columns will be written as
| ``Timestamp with timezone`` type with SQLAlchemy if supported by the
| database. Otherwise, the datetimes will be stored as timezone unaware
| timestamps local to the original timezone.
|
| .. versionadded:: 0.24.0
|
| References
| -----
| .. [1] http://docs.sqlalchemy.org
| .. [2] https://www.python.org/dev/peps/pep-0249/
|
| Examples
| -----
|
| Create an in-memory SQLite database.
|
| >>> from sqlalchemy import create_engine
| >>> engine = create_engine('sqlite://', echo=False)
|
| Create a table from scratch with 3 rows.
|
| >>> df = pd.DataFrame({'name' : ['User 1', 'User 2', 'User 3']})
| >>> df
| name

```



```

| 0 User 1
| 1 User 2
| 2 User 3
|
| >>> df.to_sql('users', con=engine)
| >>> engine.execute("SELECT * FROM users").fetchall()
| [(0, 'User 1'), (1, 'User 2'), (2, 'User 3')]
|
| >>> df1 = pd.DataFrame({'name' : ['User 4', 'User 5']})
| >>> df1.to_sql('users', con=engine, if_exists='append')
| >>> engine.execute("SELECT * FROM users").fetchall()
| [(0, 'User 1'), (1, 'User 2'), (2, 'User 3'),
| (0, 'User 4'), (1, 'User 5')]
|
| Overwrite the table with just ``df1``.
|
| >>> df1.to_sql('users', con=engine, if_exists='replace',
| ... index_label='id')
| >>> engine.execute("SELECT * FROM users").fetchall()
| [(0, 'User 4'), (1, 'User 5')]
|
| Specify the dtype (especially useful for integers with missing
values).
| Notice that while pandas is forced to store the data as floating
point,
| the database supports nullable integers. When fetching the data with
Python, we get back integer scalars.
|
| >>> df = pd.DataFrame({"A": [1, None, 2]})
| >>> df
| A
| 0 1.0
| 1 NaN
| 2 2.0
|
| >>> from sqlalchemy.types import Integer
| >>> df.to_sql('integers', con=engine, index=False,
| ... dtype={"A": Integer()})
|
| >>> engine.execute("SELECT * FROM integers").fetchall()
| [(1,), (None,), (2,)]
|
| to_xarray(self)
| Return an xarray object from the pandas object.

```

```

|
| Returns
| -----
| xarray.DataArray or xarray.Dataset
| Data in the pandas structure converted to Dataset if the object
is
| a DataFrame, or a DataArray if the object is a Series.
|
| See Also
| -----
| DataFrame.to_hdf : Write DataFrame to an HDF5 file.
| DataFrame.to_parquet : Write a DataFrame to the binary parquet
format.
|
| Notes
| -----
| See the `xarray docs <http://xarray.pydata.org/en/stable/>`
|
| Examples
| -----
| >>> df = pd.DataFrame([('falcon', 'bird', 389.0, 2),
| ... ('parrot', 'bird', 24.0, 2),
| ... ('lion', 'mammal', 80.5, 4),
| ... ('monkey', 'mammal', np.nan, 4)],
| ... columns=['name', 'class', 'max_speed',
| ... 'num_legs'])
| >>> df
| name class max_speed num_legs
| 0 falcon bird 389.0 2
| 1 parrot bird 24.0 2
| 2 lion mammal 80.5 4
| 3 monkey mammal NaN 4
|
| >>> df.to_xarray()
| <xarray.Dataset>
| Dimensions: (index: 4)
| Coordinates:
| * index (index) int64 0 1 2 3
| Data variables:
| name (index) object 'falcon' 'parrot' 'lion' 'monkey'
| class (index) object 'bird' 'bird' 'mammal' 'mammal'
| max_speed (index) float64 389.0 24.0 80.5 nan
| num_legs (index) int64 2 2 4 4
|

```

```

| >>> df['max_speed'].to_xarray()
| <xarray.DataArray 'max_speed' (index: 4)>
| array([389. , 24. , 80.5, nan])
| Coordinates:
| * index (index) int64 0 1 2 3
|
| >>> dates = pd.to_datetime(['2018-01-01', '2018-01-01',
| ... '2018-01-02', '2018-01-02'])
| >>> df_multiindex = pd.DataFrame({'date': dates,
| ... 'animal': ['falcon', 'parrot', 'falcon',
| ... 'parrot'],
| ... 'speed': [350, 18, 361,
15]}).set_index(['date',
| ... 'animal'])
| >>> df_multiindex
|
| speed
| date animal
| 2018-01-01 falcon 350
| parrot 18
| 2018-01-02 falcon 361
| parrot 15
|
| >>> df_multiindex.to_xarray()
| <xarray.Dataset>
| Dimensions: (animal: 2, date: 2)
| Coordinates:
| * date (date) datetime64[ns] 2018-01-01 2018-01-02
| * animal (animal) object 'falcon' 'parrot'
| Data variables:
| speed (date, animal) int64 350 18 361 15
|
| truncate(self, before=None, after=None, axis=None, copy=True)
| Truncate a Series or DataFrame before and after some index value.
|
| This is a useful shorthand for boolean indexing based on index
| values above or below certain thresholds.
|
| Parameters
| -----
| before : date, string, int
| Truncate all rows before this index value.
| after : date, string, int
| Truncate all rows after this index value.
| axis : {0 or 'index', 1 or 'columns'}, optional

```

```

| Axis to truncate. Truncates the index (rows) by default.
| copy : boolean, default is True,
| Return a copy of the truncated section.
|
| Returns
| -----
| type of caller
| The truncated Series or DataFrame.
|
| See Also
| -----
| DataFrame.loc : Select a subset of a DataFrame by label.
| DataFrame.iloc : Select a subset of a DataFrame by position.
|
| Notes
| -----
| If the index being truncated contains only datetime values,
| `before` and `after` may be specified as strings instead of
| Timestamps.
|
| Examples
| -----
| >>> df = pd.DataFrame({'A': ['a', 'b', 'c', 'd', 'e'],
| ... 'B': ['f', 'g', 'h', 'i', 'j'],
| ... 'C': ['k', 'l', 'm', 'n', 'o']},
| ... index=[1, 2, 3, 4, 5])
| >>> df
|
| A B C
| 1 a f k
| 2 b g l
| 3 c h m
| 4 d i n
| 5 e j o
|
| >>> df.truncate(before=2, after=4)
|
| A B C
| 2 b g l
| 3 c h m
| 4 d i n
|
| The columns of a DataFrame can be truncated.
|
| >>> df.truncate(before="A", after="B", axis="columns")
|
| A B

```

```

| 1 a f
| 2 b g
| 3 c h
| 4 d i
| 5 e j
|
| For Series, only rows can be truncated.
|
| >>> df['A'].truncate(before=2, after=4)
| 2 b
| 3 c
| 4 d
| Name: A, dtype: object
|
| The index values in ``truncate`` can be datetimes or string
| dates.
|
| >>> dates = pd.date_range('2016-01-01', '2016-02-01', freq='s')
| >>> df = pd.DataFrame(index=dates, data={'A': 1})
| >>> df.tail()
|
| A
| 2016-01-31 23:59:56 1
| 2016-01-31 23:59:57 1
| 2016-01-31 23:59:58 1
| 2016-01-31 23:59:59 1
| 2016-02-01 00:00:00 1
|
| >>> df.truncate(before=pd.Timestamp('2016-01-05'),
| ... after=pd.Timestamp('2016-01-10')).tail()
|
| A
| 2016-01-09 23:59:56 1
| 2016-01-09 23:59:57 1
| 2016-01-09 23:59:58 1
| 2016-01-09 23:59:59 1
| 2016-01-10 00:00:00 1
|
| Because the index is a DatetimeIndex containing only dates, we can
| specify `before` and `after` as strings. They will be coerced to
| Timestamps before truncation.
|
| >>> df.truncate('2016-01-05', '2016-01-10').tail()
|
| A
| 2016-01-09 23:59:56 1
| 2016-01-09 23:59:57 1

```

```

| 2016-01-09 23:59:58 1
| 2016-01-09 23:59:59 1
| 2016-01-10 00:00:00 1
|
| Note that ``truncate`` assumes a 0 value for any unspecified time
| component (midnight). This differs from partial string slicing, which
| returns any partially matching dates.
|
| >>> df.loc['2016-01-05':'2016-01-10', :].tail()
|
| A
| 2016-01-10 23:59:55 1
| 2016-01-10 23:59:56 1
| 2016-01-10 23:59:57 1
| 2016-01-10 23:59:58 1
| 2016-01-10 23:59:59 1
|
| tshift(self, periods=1, freq=None, axis=0)
| Shift the time index, using the index's frequency if available.
|
| Parameters
| -----
| periods : int
| Number of periods to move, can be positive or negative
| freq : DateOffset, timedelta, or time rule string, default None
| Increment to use from the tseries module or time rule (e.g.
'EOM')
| axis : int or basestring
| Corresponds to the axis that contains the Index
|
| Returns
| -----
| shifted : Series/DataFrame
|
| Notes
| -----
| If freq is not specified then tries to use the freq or inferred_freq
| attributes of the index. If neither of those attributes exist, a
| ValueError is thrown
|
| tz_convert(self, tz, axis=0, level=None, copy=True)
| Convert tz-aware axis to target time zone.
|
| Parameters
| -----

```

```

| tz : string or pytz.timezone object
| axis : the axis to convert
| level : int, str, default None
| If axis is a MultiIndex, convert a specific level. Otherwise
| must be None
| copy : boolean, default True
| Also make a copy of the underlying data
|
| Returns
| -----
| %(klass)s
| Object with time zone converted axis.
|
| Raises
| -----
| TypeError
| If the axis is tz-naive.
|
| tz_localize(self, tz, axis=0, level=None, copy=True, ambiguous='raise',
nonexistent='raise')
| Localize tz-naive index of a Series or DataFrame to target time zone.
|
| This operation localizes the Index. To localize the values in a
| timezone-naive Series, use :meth:`Series.dt.tz_localize`.
|
| Parameters
| -----
| tz : string or pytz.timezone object
| axis : the axis to localize
| level : int, str, default None
| If axis is a MultiIndex, localize a specific level. Otherwise
| must be None
| copy : boolean, default True
| Also make a copy of the underlying data
| ambiguous : 'infer', bool-ndarray, 'NaT', default 'raise'
| When clocks moved backward due to DST, ambiguous times may arise.
| For example in Central European Time (UTC+01), when going from
| 03:00 DST to 02:00 non-DST, 02:30:00 local time occurs both at
| 00:30:00 UTC and at 01:30:00 UTC. In such a situation, the
| `ambiguous` parameter dictates how ambiguous times should be
| handled.
|
| - 'infer' will attempt to infer fall dst-transition hours based
on

```

```

| order
| - bool-ndarray where True signifies a DST time, False designates
| a non-DST time (note that this flag is only applicable for
| ambiguous times)
| - 'NaT' will return NaT where there are ambiguous times
| - 'raise' will raise an AmbiguousTimeError if there are ambiguous
| times
| nonexistent : str, default 'raise'
| A nonexistent time does not exist in a particular timezone
| where clocks moved forward due to DST. Valid values are:
|
| - 'shift_forward' will shift the nonexistent time forward to the
| closest existing time
| - 'shift_backward' will shift the nonexistent time backward to
the
| closest existing time
| - 'NaT' will return NaT where there are nonexistent times
| - timedelta objects will shift nonexistent times by the timedelta
| - 'raise' will raise a NonExistentTimeError if there are
| nonexistent times
|
| .. versionadded:: 0.24.0
|
| Returns
| -----
| Series or DataFrame
| Same type as the input.
|
| Raises
| -----
| TypeError
| If the TimeSeries is tz-aware and tz is not None.
|
| Examples
| -----
|
| Localize local times:
|
| >>> s = pd.Series([1],
| ... index=pd.DatetimeIndex(['2018-09-15 01:30:00']))
| >>> s.tz_localize('CET')
| 2018-09-15 01:30:00+02:00 1
| dtype: int64
|

```



```

| Be careful with DST changes. When there is sequential data, pandas
| can infer the DST time:
|
| >>> s = pd.Series(range(7), index=pd.DatetimeIndex([
| ... '2018-10-28 01:30:00',
| ... '2018-10-28 02:00:00',
| ... '2018-10-28 02:30:00',
| ... '2018-10-28 02:00:00',
| ... '2018-10-28 02:30:00',
| ... '2018-10-28 03:00:00',
| ... '2018-10-28 03:30:00'])))
| >>> s.tz_localize('CET', ambiguous='infer')
| 2018-10-28 01:30:00+02:00 0
| 2018-10-28 02:00:00+02:00 1
| 2018-10-28 02:30:00+02:00 2
| 2018-10-28 02:00:00+01:00 3
| 2018-10-28 02:30:00+01:00 4
| 2018-10-28 03:00:00+01:00 5
| 2018-10-28 03:30:00+01:00 6
| dtype: int64
|
| In some cases, inferring the DST is impossible. In such cases, you
can
| pass an ndarray to the ambiguous parameter to set the DST explicitly
|
| >>> s = pd.Series(range(3), index=pd.DatetimeIndex([
| ... '2018-10-28 01:20:00',
| ... '2018-10-28 02:36:00',
| ... '2018-10-28 03:46:00'])))
| >>> s.tz_localize('CET', ambiguous=np.array([True, True, False]))
| 2018-10-28 01:20:00+02:00 0
| 2018-10-28 02:36:00+02:00 1
| 2018-10-28 03:46:00+01:00 2
| dtype: int64
|
| If the DST transition causes nonexistent times, you can shift these
| dates forward or backwards with a timedelta object or
`shift_forward`
| or `shift_backwards`.
| >>> s = pd.Series(range(2), index=pd.DatetimeIndex([
| ... '2015-03-29 02:30:00',
| ... '2015-03-29 03:30:00'])))
| >>> s.tz_localize('Europe/Warsaw', nonexistent='shift_forward')
| 2015-03-29 03:00:00+02:00 0

```

```

| 2015-03-29 03:30:00+02:00 1
| dtype: int64
| >>> s.tz_localize('Europe/Warsaw', nonexistent='shift_backward')
| 2015-03-29 01:59:59.999999999+01:00 0
| 2015-03-29 03:30:00+02:00 1
| dtype: int64
| >>> s.tz_localize('Europe/Warsaw', nonexistent=pd.Timedelta('1H'))
| 2015-03-29 03:30:00+02:00 0
| 2015-03-29 03:30:00+02:00 1
| dtype: int64
|
| where(self, cond, other=nan, inplace=False, axis=None, level=None,
errors='raise', try_cast=False)
| Replace values where the condition is False.
|
| Parameters
| -----
| cond : boolean Series/DataFrame, array-like, or callable
| Where `cond` is True, keep the original value. Where
| False, replace with corresponding value from `other`.
| If `cond` is callable, it is computed on the Series/DataFrame and
| should return boolean Series/DataFrame or array. The callable
must
| not change input Series/DataFrame (though pandas doesn't check
it).
|
| .. versionadded:: 0.18.1
| A callable can be used as cond.
|
| other : scalar, Series/DataFrame, or callable
| Entries where `cond` is False are replaced with
| corresponding value from `other`.
| If other is callable, it is computed on the Series/DataFrame and
| should return scalar or Series/DataFrame. The callable must not
| change input Series/DataFrame (though pandas doesn't check it).
|
| .. versionadded:: 0.18.1
| A callable can be used as other.
|
| inplace : bool, default False
| Whether to perform the operation in place on the data.
| axis : int, default None
| Alignment axis if needed.
| level : int, default None

```

```

| Alignment level if needed.
| errors : str, {'raise', 'ignore'}, default 'raise'
| Note that currently this parameter won't affect
| the results and will always coerce to a suitable dtype.
|
| - 'raise' : allow exceptions to be raised.
| - 'ignore' : suppress exceptions. On error return original
object.
|
| try_cast : bool, default False
| Try to cast the result back to the input type (if possible).
|
| Returns
| -----
| Same type as caller
|
| See Also
| -----
| :func:`DataFrame.mask` : Return an object of same shape as
| self.
|
| Notes
| -----
| The where method is an application of the if-then idiom. For each
| element in the calling DataFrame, if ``cond`` is ``True`` the
| element is used; otherwise the corresponding element from the
DataFrame
| ``other`` is used.
|
| The signature for :func:`DataFrame.where` differs from
| :func:`numpy.where`. Roughly ``df1.where(m, df2)`` is equivalent to
| ``np.where(m, df1, df2)``.
|
| For further details and examples see the ``where`` documentation in
| :ref:`indexing <indexing.where_mask>`.
|
| Examples
| -----
| >>> s = pd.Series(range(5))
| >>> s.where(s > 0)
| 0 NaN
| 1 1.0
| 2 2.0
| 3 3.0

```

```

| 4 4.0
| dtype: float64
|
| >>> s.mask(s > 0)
| 0 0.0
| 1 NaN
| 2 NaN
| 3 NaN
| 4 NaN
| dtype: float64
|
| >>> s.where(s > 1, 10)
| 0 10
| 1 10
| 2 2
| 3 3
| 4 4
| dtype: int64
|
| >>> df = pd.DataFrame(np.arange(10).reshape(-1, 2), columns=['A',
'B'])
| >>> df
| A B
| 0 0 1
| 1 2 3
| 2 4 5
| 3 6 7
| 4 8 9
| >>> m = df % 3 == 0
| >>> df.where(m, -df)
| A B
| 0 0 -1
| 1 -2 3
| 2 -4 -5
| 3 6 -7
| 4 -8 9
| >>> df.where(m, -df) == np.where(m, df, -df)
| A B
| 0 True True
| 1 True True
| 2 True True
| 3 True True
| 4 True True
| >>> df.where(m, -df) == df.mask(~m, -df)

```

```

| A B
| 0 True True
| 1 True True
| 2 True True
| 3 True True
| 4 True True
|
|
| xs(self, key, axis=0, level=None, drop_level=True)
| Return cross-section from the Series/DataFrame.
|
| This method takes a `key` argument to select data at a particular
| level of a MultiIndex.
|
| Parameters
| -----
| key : label or tuple of label
| Label contained in the index, or partially in a MultiIndex.
| axis : {0 or 'index', 1 or 'columns'}, default 0
| Axis to retrieve cross-section on.
| level : object, defaults to first n levels (n=1 or len(key))
| In case of a key partially contained in a MultiIndex, indicate
| which levels are used. Levels can be referred by label or
position.
| drop_level : bool, default True
| If False, returns object with same levels as self.
|
| Returns
| -----
| Series or DataFrame
| Cross-section from the original Series or DataFrame
| corresponding to the selected index levels.
|
| See Also
| -----
| DataFrame.loc : Access a group of rows and columns
| by label(s) or a boolean array.
| DataFrame.iloc : Purely integer-location based indexing
| for selection by position.
|
| Notes
| -----
| `xs` can not be used to set values.
|
| MultiIndex Slicers is a generic way to get/set values on

```

any level or levels.  
 It is a superset of `xs` functionality, see  
[:ref:`MultiIndex Slicers <advanced.mi\\_slicers>`](#).

#### Examples

```
>>> d = {'num_legs': [4, 4, 2, 2],
... 'num_wings': [0, 0, 2, 2],
... 'class': ['mammal', 'mammal', 'mammal', 'bird'],
... 'animal': ['cat', 'dog', 'bat', 'penguin'],
... 'locomotion': ['walks', 'walks', 'flies', 'walks']}
>>> df = pd.DataFrame(data=d)
>>> df = df.set_index(['class', 'animal', 'locomotion'])
>>> df
```

			num_legs	num_wings
class	animal	locomotion		
mammal	cat	walks	4	0
	dog	walks	4	0
	bat	flies	2	2
bird	penguin	walks	2	2

#### Get values at specified index

```
>>> df.xs('mammal')
 num_legs num_wings
animal locomotion
cat walks 4 0
dog walks 4 0
bat flies 2 2
```

#### Get values at several indexes

```
>>> df.xs(('mammal', 'dog'))
 num_legs num_wings
locomotion
walks 4 0
```

#### Get values at specified index and level

```
>>> df.xs('cat', level=1)
 num_legs num_wings
class locomotion
mammal walks 4 0
```

```

| Get values at several indexes and levels
|
| >>> df.xs(('bird', 'walks'),
| ... level=[0, 'locomotion'])
| num_legs num_wings
| animal
| penguin 2 2
|
| Get values at specified column and axis
|
| >>> df.xs('num_wings', axis=1)
| class animal locomotion
| mammal cat walks 0
| dog walks 0
| bat flies 2
| bird penguin walks 2
| Name: num_wings, dtype: int64
|
| -----
| Data descriptors inherited from pandas.core.generic.NDFrame:
|
| at
| Access a single value for a row/column label pair.
|
| Similar to ``loc``, in that both provide label-based lookups. Use
| ``at`` if you only need to get or set a single value in a DataFrame
| or Series.
|
| Raises
| -----
| KeyError
| When label does not exist in DataFrame
|
| See Also
| -----
| DataFrame.iat : Access a single value for a row/column pair by
integer
| position.
| DataFrame.loc : Access a group of rows and columns by label(s).
| Series.at : Access a single value using a label.
|
| Examples
| -----
| >>> df = pd.DataFrame([[0, 2, 3], [0, 4, 1], [10, 20, 30]],

```

```

| ... index=[4, 5, 6], columns=['A', 'B', 'C'])
|
| >>> df
|
| A B C
| 4 0 2 3
| 5 0 4 1
| 6 10 20 30
|
|
| Get value at specified row/column pair
|
| >>> df.at[4, 'B']
|
| 2
|
|
| Set value at specified row/column pair
|
| >>> df.at[4, 'B'] = 10
| >>> df.at[4, 'B']
|
| 10
|
|
| Get value within a Series
|
| >>> df.loc[5].at['B']
|
| 4
|
|
| blocks
| Internal property, property synonym for as_blocks().
|
| .. deprecated:: 0.21.0
|
|
| dtypes
| Return the dtypes in the DataFrame.
|
| This returns a Series with the data type of each column.
| The result's index is the original DataFrame's columns. Columns
| with mixed types are stored with the ``object`` dtype. See
| :ref:`the User Guide <basics.dtypes>` for more.
|
|
| Returns
| -----
| pandas.Series
| The data type of each column.
|
|
| See Also
| -----
| DataFrame.ftypes : Dtype and sparsity information.

```



```

|
| Examples
| -----
| >>> df = pd.DataFrame({'float': [1.0],
| ... 'int': [1],
| ... 'datetime': [pd.Timestamp('20180310')],
| ... 'string': ['foo']})
| >>> df.dtypes
| float float64
| int int64
| datetime datetime64[ns]
| string object
| dtype: object
|
| empty
| Indicator whether DataFrame is empty.
|
| True if DataFrame is entirely empty (no items), meaning any of the
| axes are of length 0.
|
| Returns
| -----
| bool
| If DataFrame is empty, return True, if not return False.
|
| See Also
| -----
| Series.dropna
| DataFrame.dropna
|
| Notes
| -----
| If DataFrame contains only NaNs, it is still not considered empty.
See
| the example below.
|
| Examples
| -----
| An example of an actual empty DataFrame. Notice the index is empty:
|
| >>> df_empty = pd.DataFrame({'A' : []})
| >>> df_empty
| Empty DataFrame
| Columns: [A]

```

```

| Index: []
| >>> df_empty.empty
| True
|
| If we only have NaNs in our DataFrame, it is not considered empty! We
| will need to drop the NaNs to make the DataFrame empty:
|
| >>> df = pd.DataFrame({'A' : [np.nan]})
| >>> df
| A
| 0 NaN
| >>> df.empty
| False
| >>> df.dropna().empty
| True
|
| ftypes
| Return the ftypes (indication of sparse/dense and dtype) in
DataFrame.
|
| .. deprecated:: 0.25.0
| Use :func:`dtypes` instead.
|
| This returns a Series with the data type of each column.
| The result's index is the original DataFrame's columns. Columns
| with mixed types are stored with the ``object`` dtype. See
| :ref:`the User Guide <basics.dtypes>` for more.
|
| Returns
| -----
| pandas.Series
| The data type and indication of sparse/dense of each column.
|
| See Also
| -----
| DataFrame.dtypes: Series with just dtype information.
| SparseDataFrame : Container for sparse tabular data.
|
| Notes
| -----
| Sparse data should have the same dtypes as its dense representation.
|
| Examples
| -----

```

```

| >>> arr = np.random.RandomState(0).randn(100, 4)
| >>> arr[arr < .8] = np.nan
| >>> pd.DataFrame(arr).ftypes
|
| 0 float64:dense
| 1 float64:dense
| 2 float64:dense
| 3 float64:dense
|
| dtype: object
|
|
| >>> pd.SparseDataFrame(arr).ftypes # doctest: +SKIP
|
| 0 float64:sparse
| 1 float64:sparse
| 2 float64:sparse
| 3 float64:sparse
|
| dtype: object
|
| iat
|
| Access a single value for a row/column pair by integer position.
|
|
| Similar to ``iloc``, in that both provide integer-based lookups. Use
| ``iat`` if you only need to get or set a single value in a DataFrame
| or Series.
|
|
| Raises
| -----
|
| IndexError
|
| When integer position is out of bounds
|
|
| See Also
| -----
|
| DataFrame.at : Access a single value for a row/column label pair.
| DataFrame.loc : Access a group of rows and columns by label(s).
| DataFrame.iloc : Access a group of rows and columns by integer
| position(s).
|
|
| Examples
| -----
|
| >>> df = pd.DataFrame([[0, 2, 3], [0, 4, 1], [10, 20, 30]],
| ... columns=['A', 'B', 'C'])
| >>> df
|
| A B C
| 0 0 2 3
| 1 0 4 1
| 2 10 20 30

```

```

|
| Get value at specified row/column pair
|
| >>> df.iat[1, 2]
| 1
|
| Set value at specified row/column pair
|
| >>> df.iat[1, 2] = 10
| >>> df.iat[1, 2]
| 10
|
| Get value within a series
|
| >>> df.loc[0].iat[1]
| 2
|
| iloc
| Purely integer-location based indexing for selection by position.
|
| ``.iloc[]`` is primarily integer position based (from ``0`` to
| ``length-1`` of the axis), but may also be used with a boolean
| array.
|
| Allowed inputs are:
|
| - An integer, e.g. ``5``.
| - A list or array of integers, e.g. ``[4, 3, 0]``.
| - A slice object with ints, e.g. ``1:7``.
| - A boolean array.
| - A ``callable`` function with one argument (the calling Series or
| DataFrame) and that returns valid output for indexing (one of the
| above).
|
| This is useful in method chains, when you don't have a reference to
the
| calling object, but would like to base your selection on some
value.
|
| ``.iloc`` will raise ``IndexError`` if a requested indexer is
| out-of-bounds, except *slice* indexers which allow out-of-bounds
| indexing (this conforms with python/numpy *slice* semantics).
|
| See more at :ref:`Selection by Position <indexing.integer>`.
|

```

```

| See Also
| -----
| DataFrame.iat : Fast integer location scalar accessor.
| DataFrame.loc : Purely label-location based indexer for selection by
label.
| Series.iloc : Purely integer-location based indexing for
| selection by position.
|
| Examples
| -----
|
| >>> mydict = [{'a': 1, 'b': 2, 'c': 3, 'd': 4},
| ... {'a': 100, 'b': 200, 'c': 300, 'd': 400},
| ... {'a': 1000, 'b': 2000, 'c': 3000, 'd': 4000 }]
| >>> df = pd.DataFrame(mydict)
| >>> df
|
| a b c d
| 0 1 2 3 4
| 1 100 200 300 400
| 2 1000 2000 3000 4000
|
| **Indexing just the rows**
|
| With a scalar integer.
|
| >>> type(df.iloc[0])
| <class 'pandas.core.series.Series'>
| >>> df.iloc[0]
|
| a 1
| b 2
| c 3
| d 4
| Name: 0, dtype: int64
|
| With a list of integers.
|
| >>> df.iloc[[0]]
|
| a b c d
| 0 1 2 3 4
| >>> type(df.iloc[[0]])
| <class 'pandas.core.frame.DataFrame'>
|
| >>> df.iloc[[0, 1]]
|
| a b c d

```

```
| 0 1 2 3 4
| 1 100 200 300 400
```

| With a ``slice`` object.

```
| >>> df.iloc[:3]
| a b c d
| 0 1 2 3 4
| 1 100 200 300 400
| 2 1000 2000 3000 4000
```

| With a boolean mask the same length as the index.

```
| >>> df.iloc[[True, False, True]]
| a b c d
| 0 1 2 3 4
| 2 1000 2000 3000 4000
```

| With a callable, useful in method chains. The ``x`` passed to the ``lambda`` is the DataFrame being sliced. This selects the rows whose index label even.

```
| >>> df.iloc[lambda x: x.index % 2 == 0]
| a b c d
| 0 1 2 3 4
| 2 1000 2000 3000 4000
```

| **\*\*Indexing both axes\*\***

| You can mix the indexer types for the index and columns. Use ``:`` to select the entire axis.

| With scalar integers.

```
| >>> df.iloc[0, 1]
| 2
```

| With lists of integers.

```
| >>> df.iloc[[0, 2], [1, 3]]
| b d
| 0 2 4
| 2 2000 4000
```

|       With `slice` objects.

|       >>> df.iloc[1:3, 0:3]

	a	b	c
1	100	200	300
2	1000	2000	3000

|       With a boolean array whose length matches the columns.

|       >>> df.iloc[:, [True, False, True, False]]

	a	c
0	1	3
1	100	300
2	1000	3000

|       With a callable function that expects the Series or DataFrame.

|       >>> df.iloc[:, lambda df: [0, 2]]

	a	c
0	1	3
1	100	300
2	1000	3000

|   is\_copy

|       Return the copy.

|   ix

|       A primarily label-location based indexer, with integer position  
|       fallback.

|       Warning: Starting in 0.20.0, the .ix indexer is deprecated, in  
|       favor of the more strict .iloc and .loc indexers.

|       ``.ix[]`` supports mixed integer and label based access. It is  
|       primarily label based, but will fall back to integer positional  
|       access unless the corresponding axis is of integer type.

|       ``.ix`` is the most general indexer and will support any of the  
|       inputs in ``.loc`` and ``.iloc``. ``.ix`` also supports floating  
|       point label schemes. ``.ix`` is exceptionally useful when dealing  
|       with mixed positional and label based hierarchical indexes.

|       However, when an axis is integer based, ONLY label based access  
|       and not positional access is supported. Thus, in such cases, it's

```

| usually better to be explicit and use ``.iloc`` or ``.loc``.
|
| See more at :ref:`Advanced Indexing <advanced>`.
|
| loc
| Access a group of rows and columns by label(s) or a boolean array.
|
| ``.loc[]`` is primarily label based, but may also be used with a
| boolean array.
|
| Allowed inputs are:
|
| - A single label, e.g. ``5`` or ``'a'``, (note that ``5`` is
| interpreted as a *label* of the index, and never as an
| integer position along the index).
| - A list or array of labels, e.g. ``['a', 'b', 'c']``.
| - A slice object with labels, e.g. ``'a':'f'``.
|
| .. warning:: Note that contrary to usual python slices, both
the
| start and the stop are included
|
| - A boolean array of the same length as the axis being sliced,
| e.g. ``[True, False, True]``.
| - A ``callable`` function with one argument (the calling Series or
| DataFrame) and that returns valid output for indexing (one of the
above)
|
| See more at :ref:`Selection by Label <indexing.label>`
|
| Raises
| -----
| KeyError:
| when any items are not found
|
| See Also
| -----
| DataFrame.at : Access a single value for a row/column label pair.
| DataFrame.iloc : Access group of rows and columns by integer
position(s).
| DataFrame.xs : Returns a cross-section (row(s) or column(s)) from the
| Series/DataFrame.
| Series.loc : Access group of values using labels.
|

```



|     Examples

|     -----

|     \*\*Getting values\*\*

|     >>> df = pd.DataFrame([[1, 2], [4, 5], [7, 8]],  
|     ...       index=['cobra', 'viper', 'sidewinder'],  
|     ...       columns=['max\_speed', 'shield'])  
|     >>> df

	max_speed	shield
cobra	1	2
viper	4	5
sidewinder	7	8

|     Single label. Note this returns the row as a Series.

|     >>> df.loc['viper']  
|     max\_speed     4  
|     shield       5  
|     Name: viper, dtype: int64

|     List of labels. Note using ``[['']]`` returns a DataFrame.

|     >>> df.loc[['viper', 'sidewinder']]  
|           max\_speed shield  
|     viper           4       5  
|     sidewinder       7       8

|     Single label for row and column

|     >>> df.loc['cobra', 'shield']  
|     2

|     Slice with labels for row and single label for column. As mentioned  
|     above, note that both the start and stop of the slice are included.

|     >>> df.loc['cobra':'viper', 'max\_speed']  
|     cobra     1  
|     viper     4  
|     Name: max\_speed, dtype: int64

|     Boolean list with the same length as the row axis

|     >>> df.loc[[False, False, True]]  
|           max\_speed shield

```

| sidewinder 7 8
|
| Conditional that returns a boolean Series
|
| >>> df.loc[df['shield'] > 6]
| max_speed shield
| sidewinder 7 8
|
| Conditional that returns a boolean Series with column labels
specified
|
| >>> df.loc[df['shield'] > 6, ['max_speed']]
| max_speed
| sidewinder 7
|
| Callable that returns a boolean Series
|
| >>> df.loc[lambda df: df['shield'] == 8]
| max_speed shield
| sidewinder 7 8
|
| **Setting values**
|
| Set value for all items matching the list of labels
|
| >>> df.loc[['viper', 'sidewinder'], ['shield']] = 50
| >>> df
| max_speed shield
| cobra 1 2
| viper 4 50
| sidewinder 7 50
|
| Set value for an entire row
|
| >>> df.loc['cobra'] = 10
| >>> df
| max_speed shield
| cobra 10 10
| viper 4 50
| sidewinder 7 50
|
| Set value for an entire column
|
| >>> df.loc[:, 'max_speed'] = 30

```

```

| >>> df
| max_speed shield
| cobra 30 10
| viper 30 50
| sidewinder 30 50
|
| Set value for rows matching callable condition
|
| >>> df.loc[df['shield'] > 35] = 0
| >>> df
| max_speed shield
| cobra 30 10
| viper 0 0
| sidewinder 0 0
|
| **Getting values on a DataFrame with an index that has integer
labels**
|
| Another example using integers for the index
|
| >>> df = pd.DataFrame([[1, 2], [4, 5], [7, 8]],
| ... index=[7, 8, 9], columns=['max_speed', 'shield'])
| >>> df
| max_speed shield
| 7 1 2
| 8 4 5
| 9 7 8
|
| Slice with integer labels for rows. As mentioned above, note that
both
| the start and stop of the slice are included.
|
| >>> df.loc[7:9]
| max_speed shield
| 7 1 2
| 8 4 5
| 9 7 8
|
| **Getting values with a MultiIndex**
|
| A number of examples using a DataFrame with a MultiIndex
|
| >>> tuples = [
| ... ('cobra', 'mark i'), ('cobra', 'mark ii'),

```

```

| ... ('sidewinder', 'mark i'), ('sidewinder', 'mark ii'),
| ... ('viper', 'mark ii'), ('viper', 'mark iii')
| ...]
| >>> index = pd.MultiIndex.from_tuples(tuples)
| >>> values = [[12, 2], [0, 4], [10, 20],
| ... [1, 4], [7, 1], [16, 36]]
| >>> df = pd.DataFrame(values, columns=['max_speed', 'shield'],
index=index)
| >>> df
|
| max_speed shield
| cobra mark i 12 2
| mark ii 0 4
| sidewinder mark i 10 20
| mark ii 1 4
| viper mark ii 7 1
| mark iii 16 36
|
| Single label. Note this returns a DataFrame with a single index.
|
| >>> df.loc['cobra']
| max_speed shield
| mark i 12 2
| mark ii 0 4
|
| Single index tuple. Note this returns a Series.
|
| >>> df.loc[('cobra', 'mark ii')]
| max_speed 0
| shield 4
| Name: (cobra, mark ii), dtype: int64
|
| Single label for row and column. Similar to passing in a tuple, this
| returns a Series.
|
| >>> df.loc['cobra', 'mark i']
| max_speed 12
| shield 2
| Name: (cobra, mark i), dtype: int64
|
| Single tuple. Note using ``[[]]`` returns a DataFrame.
|
| >>> df.loc[['cobra', 'mark ii']]
| max_speed shield
| cobra mark ii 0 4

```

Single tuple for the index with a single label for the column

```
>>> df.loc[('cobra', 'mark i'), 'shield']
2
```

Slice from index tuple to single label

```
>>> df.loc[('cobra', 'mark i'):'viper']
 max_speed shield
cobra mark i 12 2
 mark ii 0 4
sidewinder mark i 10 20
 mark ii 1 4
viper mark ii 7 1
 mark iii 16 36
```

Slice from index tuple to index tuple

```
>>> df.loc[('cobra', 'mark i'):(('viper', 'mark ii'))]
 max_speed shield
cobra mark i 12 2
 mark ii 0 4
sidewinder mark i 10 20
 mark ii 1 4
viper mark ii 7 1
```

`ndim`

Return an int representing the number of axes / array dimensions.

Return 1 if Series. Otherwise return 2 if DataFrame.

See Also

-----

`ndarray.ndim` : Number of array dimensions.

Examples

-----

```
>>> s = pd.Series({'a': 1, 'b': 2, 'c': 3})
>>> s.ndim
1
```

```
>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
>>> df.ndim
```

```

| 2
|
| size
| Return an int representing the number of elements in this object.
|
| Return the number of rows if Series. Otherwise return the number of
| rows times number of columns if DataFrame.
|
| See Also
| -----
| ndarray.size : Number of elements in the array.
|
| Examples
| -----
| >>> s = pd.Series({'a': 1, 'b': 2, 'c': 3})
| >>> s.size
| 3
|
| >>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
| >>> df.size
| 4
|
| values
| Return a Numpy representation of the DataFrame.
|
| .. warning::
|
| We recommend using :meth:`DataFrame.to_numpy` instead.
|
| Only the values in the DataFrame will be returned, the axes labels
| will be removed.
|
| Returns
| -----
| numpy.ndarray
| The values of the DataFrame.
|
| See Also
| -----
| DataFrame.to_numpy : Recommended alternative to this method.
| DataFrame.index : Retrieve the index labels.
| DataFrame.columns : Retrieving the column names.
|
| Notes

```

```

| -----
|
| The dtype will be a lower-common-denominator dtype (implicit
| upcasting); that is to say if the dtypes (even of numeric types)
| are mixed, the one that accommodates all will be chosen. Use this
| with care if you are not dealing with the blocks.
|
|
| e.g. If the dtypes are float16 and float32, dtype will be upcast to
| float32. If dtypes are int32 and uint8, dtype will be upcast to
| int32. By :func:`numpy.find_common_type` convention, mixing int64
| and uint64 will result in a float64 dtype.
|
|
| Examples
| -----
|
| A DataFrame where all columns are the same type (e.g., int64) results
| in an array of the same type.
|
| >>> df = pd.DataFrame({'age': [3, 29],
| ... 'height': [94, 170],
| ... 'weight': [31, 115]})
| >>> df
| age height weight
| 0 3 94 31
| 1 29 170 115
| >>> df.dtypes
| age int64
| height int64
| weight int64
| dtype: object
| >>> df.values
| array([[3, 94, 31],
| [29, 170, 115]], dtype=int64)
|
|
| A DataFrame with mixed type columns (e.g., str/object, int64, float32)
| results in an ndarray of the broadest type that accommodates these
| mixed types (e.g., object).
|
| >>> df2 = pd.DataFrame([('parrot', 24.0, 'second'),
| ... ('lion', 80.5, 1),
| ... ('monkey', np.nan, None)],
| ... columns=('name', 'max_speed', 'rank'))
| >>> df2.dtypes
| name object
| max_speed float64
| rank object

```

```

| dtype: object
| >>> df2.values
| array([['parrot', 24.0, 'second'],
| ['lion', 80.5, 1],
| ['monkey', nan, None]], dtype=object)
|
| -----
| Data and other attributes inherited from pandas.core.generic.NDFrame:
|
| __array_priority__ = 1000
|
| -----
| Methods inherited from pandas.core.base.PandasObject:
|
| __sizeof__(self)
| Generates the total memory usage for an object that returns
| either a value or Series of values
|
| -----
| Methods inherited from pandas.core.accessor.DirNamesMixin:
|
| __dir__(self)
| Provide method name lookup and completion
| Only provide 'public' methods.
|
| -----
| Data descriptors inherited from pandas.core.accessor.DirNamesMixin:
|
| __dict__
| dictionary for instance variables (if defined)
|
| __weakref__
| list of weak references to the object (if defined)

```

#### df.columns

```

Index(['user_id', 'Movie1', 'Movie2', 'Movie3', 'Movie4', 'Movie5', 'Movie6',
 'Movie7', 'Movie8', 'Movie9',
 ...,
 'Movie197', 'Movie198', 'Movie199', 'Movie200', 'Movie201',
'Movie202',
 'Movie203', 'Movie204', 'Movie205', 'Movie206'],
 dtype='object', length=207)

```



```
melt_df = df.melt(id_vars = df.columns[0], value_vars = df.columns[1:], var_name="movie_name",
value_name="rating")
melt_df
```

	user_id	movie_name	rating
0	A3R5OBKS7OM2IR	Movie1	5.0
1	AH3QC2PC1VTGP	Movie1	NaN
2	A3LKP6WPMP9UKX	Movie1	NaN
3	AVIY68KEPQ5ZD	Movie1	NaN
4	A1CV1WROP5KTTW	Movie1	NaN
...	...	...	...
998683	A1IMQ9WMFYKWH5	Movie206	5.0
998684	A1KLIKPUF5E88I	Movie206	5.0
998685	A5HG6WFZLO10D	Movie206	5.0
998686	A3UU690TWXCG1X	Movie206	5.0
998687	AI4J762YI6S06	Movie206	5.0

998688 rows x 3 columns

```
from surprise import Dataset
reader = Reader(rating_scale=(-1,10))

data = Dataset.load_from_df(melt_df.fillna(0), reader=reader)
trainset, testset = train_test_split(data, test_size=0.25)

from surprise import SVD
algo = SVD()
algo.fit(trainset)
```

```
<surprise.prediction_algorithms.matrix_factorization.SVD at 0x1d13ee50148>
```

```
predictions = algo.test(testset)
```

```
accuracy.rmse(predictions)
```

```
RMSE: 0.2751
```

```
0.2751140081006602
```

```
user_id = 'A3R5OBKS7OM2IR'
```

```
muvi_id = 'Movie1'
```

```
r_ui = 5.0
```

```
algo.predict(user_id, muvi_id, r_ui=r_ui, verbose=True)
```

```
user: A3R5OBKS7OM2IR item: Movie1 r_ui = 5.00 est = 0.08
```

```
{'was_impossible': False}
```

```
Prediction(uid='A3R5OBKS7OM2IR', iid='Movie1', r_ui=5.0,
```

```
est=0.08078859694827747, details={'was_impossible': False})
```

```
from surprise.model_selection import cross_validate
```

```
cross_validate(algo, data, measures=['RMSE', 'MAE'], cv=3, verbose=True)
```

```
Evaluating RMSE, MAE of algorithm SVD on 3 split(s).
```

	Fold 1	Fold 2	Fold 3	Mean	Std
RMSE (testset)	0.2814	0.2816	0.2823	0.2818	0.0004
MAE (testset)	0.0425	0.0432	0.0426	0.0428	0.0003
Fit time	54.30	54.58	55.27	54.72	0.40
Test time	3.94	4.14	4.14	4.08	0.10

```
{'test_rmse': array([0.28139167, 0.28157056, 0.28231805]),
 'test_mae': array([0.04251172, 0.04319265, 0.04258213]),
 'fit_time': (54.30270433425903, 54.58338952064514, 55.2669312953949),
 'test_time': (3.9379587173461914, 4.143815279006958, 4.143894910812378)}
```

```
def repeat(algo_type, frame, min_, max_):
```

```
 reader = Reader(rating_scale=(min_,max_))
```

```
 data = Dataset.load_from_df(frame, reader=reader)
```

```
 algo = algo_type
```

```
print(cross_validate(algo, data, measures=['RMSE', 'MAE'], cv=3, verbose=True))
```

```
print("#*25)
```

```
user_id = 'A3R5OBKS7OM2IR'
```

```
muvi_id = 'Movie1'
```

```
r_ui = 5.0
```

```
print(algo.predict(user_id, muvi_id, r_ui=r_ui, verbose=True))
```

```
print("#*25)
```

```
print()
```

```

user: A3R5OBKS7OM2IR item: Movie1 r_ui = 5.00 est = 0.05
{'was_impossible': False}
user: A3R5OBKS7OM2IR item: Movie1 r_ui = 5.00 est = 0.05
{'was_impossible': False}
#####

```

```

df = df.iloc[:1212, :50]
melt_df = df.melt(id_vars = df.columns[0], value_vars= df.columns[1:], var_name="movie_name",
value_name="rating")
repeat(SVD(), melt_df.fillna(0), -1, 10)
repeat(SVD(), melt_df.fillna(melt_df.mean()), -1, 10)
repeat(SVD(), melt_df.fillna(melt_df.median()), -1, 10)

```

Evaluating RMSE, MAE of algorithm SVD on 3 split(s).

	Fold 1	Fold 2	Fold 3	Mean	Std
RMSE (testset)	0.4397	0.4553	0.4535	0.4495	0.0070
MAE (testset)	0.1020	0.1044	0.1031	0.1032	0.0010
Fit time	3.13	3.14	3.18	3.15	0.02
Test time	0.19	0.19	0.19	0.19	0.00

```

{'test_rmse': array([0.43966833, 0.45531782, 0.45347166]), 'test_mae':
array([0.10202712, 0.10438337, 0.10305242]), 'fit_time': (3.132913827896118,
3.1423702239990234, 3.175093412399292), 'test_time': (0.18749499320983887,
0.18749594688415527, 0.18759727478027344)}
#####

```

Evaluating RMSE, MAE of algorithm SVD on 3 split(s).

	Fold 1	Fold 2	Fold 3	Mean	Std
RMSE (testset)	0.0881	0.0960	0.0875	0.0905	0.0039
MAE (testset)	0.0199	0.0203	0.0206	0.0203	0.0003
Fit time	3.11	3.17	3.13	3.14	0.03
Test time	0.42	0.19	0.19	0.27	0.11

```

{'test_rmse': array([0.08809229, 0.09595195, 0.0874913]), 'test_mae':
array([0.01994956, 0.02032273, 0.02064232]), 'fit_time': (3.1124162673950195,
3.1722023487091064, 3.128126382827759), 'test_time': (0.42202162742614746,
0.18749451637268066, 0.18763136863708496)}
#####

```

Evaluating RMSE, MAE of algorithm SVD on 3 split(s).

	Fold 1	Fold 2	Fold 3	Mean	Std
RMSE (testset)	0.0878	0.1031	0.1015	0.0975	0.0069
MAE (testset)	0.0192	0.0207	0.0193	0.0197	0.0007
Fit time	3.13	3.14	3.23	3.17	0.05
Test time	0.19	0.19	0.45	0.28	0.13

```
{ 'test_rmse': array([0.08782268, 0.10307901, 0.10150333]), 'test_mae':
array([0.01918205, 0.02071211, 0.01933695]), 'fit_time': (3.1252105236053467,
3.1438000202178955, 3.234687089920044), 'test_time': (0.18749427795410156,
0.18749451637268066, 0.4531123638153076)}
#####
```

```
from surprise.model_selection import GridSearchCV
param_grid = { 'n_epochs': [20, 30],
 'lr_all': [0.005, 0.01],
 'n_factors': [50, 100]}
```

```
gs = GridSearchCV(SVD, param_grid, measures=['rmse', 'mae'], cv=3)
gs.fit(data)
gs.best_score
{ 'rmse': 0.27760112584310725, 'mae': 0.040579519380100994}
```

<end of document>