

## Project 1 - Mercedes-Benz Greener Manufacturing

DESCRIPTION Reduce the time a Mercedes-Benz spends on the test bench.

### Problem Statement Scenario :

Since the first automobile, the Benz Patent Motor Car in 1886, Mercedes-Benz has stood for important automotive innovations. These include the passenger safety cell with the crumple zone, the airbag, and intelligent assistance systems. Mercedes-Benz applies for nearly 2000 patents per year, making the brand the European leader among premium carmakers. Mercedes-Benz cars are leaders in the premium car industry. With a huge selection of features and options, customers can choose the customized Mercedes-Benz of their dreams. To ensure the safety and reliability of every unique car configuration before they hit the road, Daimler's engineers have developed a robust testing system. As one of the world's biggest manufacturers of premium cars, safety and efficiency are paramount on Daimler's production lines. However, optimizing the speed of their testing system for many possible feature combinations is complex and time-consuming without a powerful algorithmic approach. You are required to reduce the time that cars spend on the test bench. Others will work with a dataset representing different permutations of features in a Mercedes-Benz car to predict the time it takes to pass testing. Optimal algorithms will contribute to faster testing, resulting in lower carbon dioxide emissions without reducing Daimler's standards.

### Following actions should be performed:

- If for any column(s), the variance is equal to zero, then you need to remove those variable(s).
- Check for null and unique values for test and train sets
- Apply label encoder.
- Perform dimensionality reduction.
- Predict your test\_df values using xgboost

## Start of Project

### Step 1: Importing the required libraries

```
In [25]: # for linear algebra
import numpy as np

# for data processing - CSV file I/O
import pandas as pd

# for dimensionality reduction
from sklearn.decomposition import PCA

# for ignoring warnings
import warnings
warnings.filterwarnings('ignore')
```

### Step 2: Reading the data from train.csv

```
In [26]: df_train = pd.read_csv('train.csv')

# for understanding the data
print('Size of training set: {} rows and {} columns'.format(*df_train.shape))

# for printing few rows and see how the data looks like
df_train.head()
```

Size of training set: 4209 rows and 378 columns

Out[26]:

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380	X382	X383
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0	0	0
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0	0	0
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0	1	0
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0	0	0

5 rows × 378 columns



### Step 3: Collecting the Y values into an array

```
In [27]: # for seperating y from the data as we will use this to learn as the prediction output
y_train = df_train['y'].values
```

### Step 4: Understanding the data types we have

```
In [28]: # for iterating through all the columns which has X in the name of the column

cols = [c for c in df_train.columns if 'X' in c]
print('Number of features: {}'.format(len(cols)))

print('Feature types:')
df_train[cols].dtypes.value_counts()
```

Number of features: 376

Feature types:

```
Out[28]: int64      368
object         8
dtype: int64
```

### Step 5: Counting the data in each of the columns

```
In [29]: counts = [[], [], []]
for c in cols:
    typ = df_train[c].dtype
    uniq = len(np.unique(df_train[c]))
    if uniq == 1:
        counts[0].append(c)
    elif uniq == 2 and typ == np.int64:
        counts[1].append(c)
    else:
        counts[2].append(c)

print('Constant features: {} Binary features: {} Categorical features: {}'.format(*
[ len(c) for c in counts]))

print('Constant features:', counts[0])

print('Categorical features:', counts[2])
```

Constant features: 12 Binary features: 356 Categorical features: 8

Constant features: ['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330', 'X347']

Categorical features: ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']

## Step 6: Reading the test.csv data

```
In [30]: df_test = pd.read_csv('test.csv')

# for removing columns ID and Y from the data as they are not used for learning
usable_columns = list(set(df_train.columns) - set(['ID', 'y']))
y_train = df_train['y'].values
id_test = df_test['ID'].values

x_train = df_train[usable_columns]
x_test = df_test[usable_columns]
```

## Step 7: Checking for null and unique values in the test and train datasets

```
In [31]: def check_missing_values(df):
    if df.isnull().any().any():
        print("\nThere are missing values in the dataframe")
    else:
        print("\nThere are no missing values in the dataframe")

check_missing_values(x_train)
check_missing_values(x_test)
```

There are no missing values in the dataframe

There are no missing values in the dataframe


## Step 8: Checking if for any column(s), the variance is equal to zero, then need to remove those variable(s).

```
In [32]: for column in usable_columns:
          cardinality = len(np.unique(x_train[column]))
          if cardinality == 1:
              x_train.drop(column, axis=1) # As column is with only one value hence drop it
              x_test.drop(column, axis=1)
          if cardinality > 2: # As column is categorical
              mapper = lambda x: sum([ord(digit) for digit in x])
              x_train[column] = x_train[column].apply(mapper)
              x_test[column] = x_test[column].apply(mapper)
          x_train.head()
```

Out[32]:

	X2	X272	X317	X113	X289	X196	X158	X281	X263	X195	...	X74	X262	X214	X49	X139	X79
0	213	0	0	0	0	0	0	0	1	0	...	1	1	0	0	0	(
1	215	0	0	0	0	0	0	0	1	0	...	1	0	0	0	0	(
2	110	1	0	0	0	0	1	0	0	0	...	1	0	0	0	0	(
3	110	1	0	0	0	0	0	0	0	0	...	1	0	0	0	0	(
4	110	1	0	0	0	0	1	0	0	0	...	1	0	0	0	0	(

5 rows × 376 columns



## Step 9: Making sure that data is now changed into numerical

```
In [33]: print('Feature types:')
          x_train[cols].dtypes.value_counts()
```

Feature types:

Out[33]: int64 376  
dtype: int64

## Step 10: Performing dimensionality reduction, Linear Dimensionality Reduction by using Singular Value Decomposition for the data to project it to a lower dimensional space.

```
In [34]: n_comp = 12
          pca = PCA(n_components=n_comp, random_state=420)
          pca2_results_train = pca.fit_transform(x_train)
          pca2_results_test = pca.transform(x_test)
```

## Step 11: Training using XGboost

```
In [35]: !pip install xgboost
```

Requirement already satisfied: xgboost in e:\programs\anaconda3\lib\site-packages (1.2.1)  
Requirement already satisfied: numpy in e:\programs\anaconda3\lib\site-packages (from xgboost) (1.18.5)  
Requirement already satisfied: scipy in e:\programs\anaconda3\lib\site-packages (from xgboost) (1.5.0)

```
In [36]: import xgboost as xgb
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split

x_train, x_valid, y_train, y_valid = train_test_split(
    pca2_results_train,
    y_train, test_size=0.3,
    random_state=422)

d_train = xgb.DMatrix(x_train, label=y_train)
d_valid = xgb.DMatrix(x_valid, label=y_valid)

# for d_test is used as xgb.DMatrix(x_test)
d_test = xgb.DMatrix(pca2_results_test)

params = {}
params['objective'] = 'reg:linear'
params['eta'] = 0.02
params['max_depth'] = 4

def xgb_r2_score(preds, dtrain):
    labels = dtrain.get_label()
    return 'r2', r2_score(labels, preds)

watchlist = [(d_train, 'train'), (d_valid, 'valid')]

clf = xgb.train(params, d_train,
                1000, watchlist, early_stopping_rounds=50,
                feval=xgb_r2_score, maximize=True, verbose_eval=10)
```

[16:45:41] WARNING: C:/Users/Administrator/workspace/xgboost-win64\_release\_1.2.0/src/objective/regression\_obj.cu:174: reg:linear is now deprecated in favor of reg:squarederror.

[0] train-rmse:98.90280 valid-rmse:99.15183 train-r2:-59.81084 valid-r2:-60.29545

Multiple eval metrics have been passed: 'valid-r2' will be used for early stopping.

Will train until valid-r2 hasn't improved in 50 rounds.

[10] train-rmse:81.06773 valid-rmse:81.34905 train-r2:-39.85638 valid-r2:-40.26023

[20] train-rmse:66.53464 valid-rmse:66.84795 train-r2:-26.52069 valid-r2:-26.86139

[30] train-rmse:54.70628 valid-rmse:55.05984 train-r2:-17.60537 valid-r2:-17.90152

[40] train-rmse:45.09587 valid-rmse:45.47532 train-r2:-11.64263 valid-r2:-11.89372

[50] train-rmse:37.30692 valid-rmse:37.71103 train-r2:-7.65252 valid-r2:-7.86673

[60] train-rmse:31.00880 valid-rmse:31.44166 train-r2:-4.97770 valid-r2:-5.16365

[70] train-rmse:25.94241 valid-rmse:26.40733 train-r2:-3.18393 valid-r2:-3.34786

[80] train-rmse:21.89302 valid-rmse:22.38741 train-r2:-1.97972 valid-r2:-2.12489

[90] train-rmse:18.68544 valid-rmse:19.22096 train-r2:-1.17055 valid-r2:-1.30344

[100] train-rmse:16.16211 valid-rmse:16.75008 train-r2:-0.62390 valid-r2:-0.74929

[110] train-rmse:14.21262 valid-rmse:14.85639 train-r2:-0.25577 valid-r2:-0.37611

[120] train-rmse:12.72305 valid-rmse:13.42720 train-r2:-0.00634 valid-r2:-0.12408

[130] train-rmse:11.60097 valid-rmse:12.37028 train-r2:0.16333 valid-r2:0.04592

[140] train-rmse:10.76386 valid-rmse:11.59724 train-r2:0.27972 valid-r2:0.16144

[150] train-rmse:10.14333 valid-rmse:11.03807 train-r2:0.36038 valid-r2:0.24035

[160] train-rmse:9.68290 valid-rmse:10.64420 train-r2:0.41713 valid-r2:0.29360

[170] train-rmse:9.36013 valid-rmse:10.35524 train-r2:0.45534 valid-r2:0.33143

[180] train-rmse:9.11822 valid-rmse:10.16152 train-r2:0.48313 valid-r2:0.35621

[190] train-rmse:8.94065 valid-rmse:10.02018 train-r2:0.50306 valid-r2:0.37400

[200] train-rmse:8.78864 valid-rmse:9.92615 train-r2:0.51982 valid-r2:0.38569

[210] train-rmse:8.66969 valid-rmse:9.85424 train-r2:0.53273 valid-r2:0.39456

[220] train-rmse:8.56730 valid-rmse:9.80281 train-r2:0.54370 valid-r2:0.40086

[230] train-rmse:8.48600 valid-rmse:9.76267 train-r2:0.55232 valid-r2:0.40576

[240] train-rmse:8.39960 valid-rmse:9.73654 train-r2:0.56139 valid-r2:0.40893

[250] train-rmse:8.33889 valid-rmse:9.71448 train-r2:0.56771 valid-r2:0.41161

[260] train-rmse:8.27284 valid-rmse:9.69895 train-r2:0.57453 valid-r2:0.41349

[270] train-rmse:8.23122 valid-rmse:9.68557 train-r2:0.57880 valid-r2:0.41511

[280] train-rmse:8.19004 valid-rmse:9.67438 train-r2:0.58300 valid-r2:0.41673

```

d-r2:0.41646
[290]  train-rmse:8.15480      valid-rmse:9.66742      train-r2:0.58658      vali
d-r2:0.41730
[300]  train-rmse:8.11006      valid-rmse:9.66298      train-r2:0.59110      vali
d-r2:0.41783
[310]  train-rmse:8.07228      valid-rmse:9.65484      train-r2:0.59491      vali
d-r2:0.41881
[320]  train-rmse:8.04500      valid-rmse:9.64865      train-r2:0.59764      vali
d-r2:0.41956
[330]  train-rmse:7.99633      valid-rmse:9.64532      train-r2:0.60249      vali
d-r2:0.41996
[340]  train-rmse:7.96786      valid-rmse:9.64302      train-r2:0.60532      vali
d-r2:0.42023
[350]  train-rmse:7.93808      valid-rmse:9.64156      train-r2:0.60826      vali
d-r2:0.42041
[360]  train-rmse:7.90765      valid-rmse:9.64072      train-r2:0.61126      vali
d-r2:0.42051
[370]  train-rmse:7.86297      valid-rmse:9.63646      train-r2:0.61564      vali
d-r2:0.42102
[380]  train-rmse:7.83090      valid-rmse:9.63157      train-r2:0.61877      vali
d-r2:0.42161
[390]  train-rmse:7.78861      valid-rmse:9.63284      train-r2:0.62288      vali
d-r2:0.42146
[400]  train-rmse:7.74889      valid-rmse:9.63509      train-r2:0.62671      vali
d-r2:0.42119
[410]  train-rmse:7.71412      valid-rmse:9.63245      train-r2:0.63006      vali
d-r2:0.42151
[420]  train-rmse:7.68438      valid-rmse:9.63487      train-r2:0.63290      vali
d-r2:0.42121
[430]  train-rmse:7.64837      valid-rmse:9.63755      train-r2:0.63633      vali
d-r2:0.42089
Stopping. Best iteration:
[382]  train-rmse:7.82060      valid-rmse:9.63136      train-r2:0.61977      vali
d-r2:0.42163

```

[16:45:43] WARNING: C:/Users/Administrator/workspace/xgboost-win64\_release\_1.2.0/src/objective/regression\_obj.cu:174: reg:linear is now deprecated in favor of reg:squarederror.

## Step 12: Predicting the test values using XGboost

```

In [37]: p_test = clf.predict(d_test)

sub = pd.DataFrame()
sub['ID'] = id_test
sub['y'] = p_test
sub.to_csv('xgb.csv', index=False)

sub.head()

```

Out[37]:

	ID	y
0	1	80.308640
1	2	93.417969
2	3	81.051048
3	4	76.639000
4	5	109.382736

**End of Project**